# OpenMP
# Application Programming
# Interface

**Version 4.5 November 2015**

This page intentionally left blank in published version.

# Contents

1 **CHAPTER 1**

2 # Introduction

3 The collection of compiler directives, library routines, and environment variables described in this
4 document collectively define the specification of the OpenMP Application Program Interface
5 (OpenMP API) for parallelism in C, C++ and Fortran programs.

6 This specification provides a model for parallel programming that is portable across architectures
7 from different vendors. Compilers from numerous vendors support the OpenMP API. More
8 information about the OpenMP API can be found at the following web site

9 **http://www.openmp.org**

10 The directives, library routines, and environment variables defined in this document allow users to
11 create and to manage parallel programs while permitting portability. The directives extend the C,
12 C++ and Fortran base languages with single program multiple data (SPMD) constructs, tasking
13 constructs, device constructs, worksharing constructs, and synchronization constructs, and they
14 provide support for sharing, mapping and privatizing data. The functionality to control the runtime
15 environment is provided by library routines and environment variables. Compilers that support the
16 OpenMP API often include a command line option to the compiler that activates and allows
17 interpretation of all OpenMP directives.

18 ## 1.1 Scope

19 The OpenMP API covers only user-directed parallelization, wherein the programmer explicitly
20 specifies the actions to be taken by the compiler and runtime system in order to execute the program
21 in parallel. OpenMP-compliant implementations are not required to check for data dependencies,
22 data conflicts, race conditions, or deadlocks, any of which may occur in conforming programs. In
23 addition, compliant implementations are not required to check for code sequences that cause a

| | |
|---|---|
| 1 | program to be classified as non-conforming. Application developers are responsible for correctly |
| 2 | using the OpenMP API to produce a conforming program. The OpenMP API does not cover |
| 3 | compiler-generated automatic parallelization and directives to the compiler to assist such |
| 4 | parallelization. |

## 5  1.2  Glossary

### 6  1.2.1  Threading Concepts

| | | |
|---|---|---|
| 7<br>8 | **thread** | An execution entity with a stack and associated static memory, called *threadprivate memory*. |
| 9 | **OpenMP thread** | A *thread* that is managed by the OpenMP runtime system. |
| 10<br>11 | **thread-safe routine** | A routine that performs the intended function even when executed concurrently (by more than one *thread*). |
| 12<br>13 | **processor** | Implementation defined hardware unit on which one or more *OpenMP threads* can execute. |
| 14 | **device** | An implementation defined logical execution engine. |
| 15 | | COMMENT: A *device* could have one or more *processors*. |
| 16 | **host device** | The *device* on which the *OpenMP program* begins execution. |
| 17 | **target device** | A device onto which code and data may be offloaded from the *host device*. |

### 18  1.2.2  OpenMP Language Terminology

| | | |
|---|---|---|
| 19 | **base language** | A programming language that serves as the foundation of the OpenMP specification. |
| 20<br>21 | | COMMENT: See Section 1.6 on page 21 for a listing of current *base languages* for the OpenMP API. |
| 22 | **base program** | A program written in a *base language*. |

| | | |
|---|---|---|
| 1<br>2 | **structured block** | For C/C++, an executable statement, possibly compound, with a single entry at the top and a single exit at the bottom, or an OpenMP *construct*. |
| 3<br>4 | | For Fortran, a block of executable statements with a single entry at the top and a single exit at the bottom, or an OpenMP *construct*. |

5        COMMENTS:

6        For all *base languages*:

7          • Access to the *structured block* must not be the result of a branch; and

8          • The point of exit cannot be a branch out of the *structured block*.

9        For C/C++:

10          • The point of entry must not be a call to **setjmp()**;

11          • **longjmp()** and **throw()** must not violate the entry/exit criteria;

12          • Calls to **exit()** are allowed in a *structured block*; and

13<br>14<br>15<br>16          • An expression statement, iteration statement, selection statement, or try block is considered to be a *structured block* if the corresponding compound statement obtained by enclosing it in **{** and **}** would be a *structured block*.

17        For Fortran:

18          • **STOP** statements are allowed in a *structured block*.

| | | |
|---|---|---|
| 19 | **enclosing context** | In C/C++, the innermost scope enclosing an OpenMP *directive*. |
| 20 | | In Fortran, the innermost scoping unit enclosing an OpenMP *directive*. |
| 21<br>22 | **directive** | In C/C++, a **#pragma**, and in Fortran, a comment, that specifies *OpenMP program* behavior. |

23<br>24        COMMENT: See Section 2.1 on page 26 for a description of OpenMP *directive* syntax.

| | | |
|---|---|---|
| 25 | **white space** | A non-empty sequence of space and/or horizontal tab characters. |
| 26<br>27 | **OpenMP program** | A program that consists of a *base program*, annotated with OpenMP *directives* and runtime library routines. |
| 28<br>29 | **conforming program** | An *OpenMP program* that follows all rules and restrictions of the OpenMP specification. |
| 30<br>31<br>32 | **declarative directive** | An OpenMP *directive* that may only be placed in a declarative context. A *declarative directive* results in one or more declarations only; it is not associated with the immediate execution of any user code. |

| | |
|---|---|
| **executable directive** | An OpenMP *directive* that is not declarative. That is, it may be placed in an executable context. |
| **stand-alone directive** | An OpenMP *executable directive* that has no associated executable user code. |
| **construct** | An OpenMP *executable directive* (and for Fortran, the paired **end** *directive*, if any) and the associated statement, loop or *structured block*, if any, not including the code in any called routines. That is, the lexical extent of an *executable directive*. |
| **combined construct** | A construct that is a shortcut for specifying one construct immediately nested inside another construct. A combined construct is semantically identical to that of explicitly specifying the first construct containing one instance of the second construct and no other statements. |
| **composite construct** | A construct that is composed of two constructs but does not have identical semantics to specifying one of the constructs immediately nested inside the other. A composite construct either adds semantics not included in the constructs from which it is composed or the nesting of the one construct inside the other is not conforming. |
| **region** | All code encountered during a specific instance of the execution of a given *construct* or of an OpenMP library routine. A *region* includes any code in called routines as well as any implicit code introduced by the OpenMP implementation. The generation of a *task* at the point where a *task generating construct* is encountered is a part of the *region* of the *encountering thread*, but an *explicit task region* associated with a *task generating construct* is not unless it is an *included task region*. The point where a **target** or **teams** directive is encountered is a part of the *region* of the *encountering thread*, but the *region* associated with the **target** or **teams** directive is not. |

> COMMENTS:
>
> > A *region* may also be thought of as the dynamic or runtime extent of a *construct* or of an OpenMP library routine.
>
> > During the execution of an *OpenMP program*, a *construct* may give rise to many *regions*.

| | |
|---|---|
| **active parallel region** | A **parallel** *region* that is executed by a *team* consisting of more than one *thread*. |
| **inactive parallel region** | A **parallel** *region* that is executed by a *team* of only one *thread*. |
| **sequential part** | All code encountered during the execution of an *initial task region* that is not part of a **parallel** *region* corresponding to a **parallel** *construct* or a **task** *region* corresponding to a **task** *construct*. |

> COMMENTS:
>
> > A *sequential part* is enclosed by an *implicit parallel region*.

| | | |
|---:|---:|---|
| 1<br>2<br>3 | | Executable statements in called routines may be in both a *sequential part* and any number of explicit **parallel** *regions* at different points in the program execution. |
| 4<br>5<br>6<br>7 | **master thread** | An *OpenMP thread* that has *thread* number 0. A *master thread* may be an *initial thread* or the *thread* that encounters a **parallel** *construct*, creates a *team*, generates a set of *implicit tasks*, and then executes one of those *tasks* as *thread* number 0. |
| 8<br>9<br>10<br>11 | **parent thread** | The *thread* that encountered the **parallel** *construct* and generated a **parallel** *region* is the *parent thread* of each of the *threads* in the *team* of that **parallel** *region*. The *master thread* of a **parallel** *region* is the same *thread* as its *parent thread* with respect to any resources associated with an *OpenMP thread*. |
| 12<br>13<br>14<br>15 | **child thread** | When a thread encounters a **parallel** construct, each of the threads in the generated **parallel** region's team are *child threads* of the encountering *thread*. The **target** or **teams** region's *initial thread* is not a *child thread* of the thread that encountered the **target** or **teams** construct. |
| 16 | **ancestor thread** | For a given *thread*, its *parent thread* or one of its *parent thread's ancestor threads*. |
| 17<br>18 | **descendent thread** | For a given *thread*, one of its *child threads* or one of its *child threads' descendent threads*. |
| 19 | **team** | A set of one or more *threads* participating in the execution of a **parallel** *region*. |
| 20 | | COMMENTS: |
| 21<br>22 | | For an *active parallel region*, the team comprises the *master thread* and at least one additional *thread*. |
| 23 | | For an *inactive parallel region*, the *team* comprises only the *master thread*. |
| 24 | **league** | The set of *thread teams* created by a **teams** construct. |
| 25 | **contention group** | An initial *thread* and its *descendent threads*. |
| 26<br>27<br>28 | **implicit parallel region** | An *inactive parallel region* that is not generated from a **parallel** *construct*. *Implicit parallel regions* surround the whole *OpenMP program*, all **target** *regions*, and all **teams** *regions*. |
| 29 | **initial thread** | A *thread* that executes an *implicit parallel region*. |
| 30 | **nested construct** | A *construct* (lexically) enclosed by another *construct*. |
| 31<br>32 | **closely nested construct** | A *construct* nested inside another *construct* with no other *construct* nested between them. |
| 33<br>34 | **nested region** | A *region* (dynamically) enclosed by another *region*. That is, a *region* encountered during the execution of another *region*. |

**closely nested region**    A *region nested* inside another *region* with no **parallel** *region nested* between them.

**strictly nested region**    A *region nested* inside another *region* with no other *region nested* between them.

**all threads**    All OpenMP *threads* participating in the *OpenMP program*.

**current team**    All *threads* in the *team* executing the innermost enclosing **parallel** *region*.

**encountering thread**    For a given *region*, the *thread* that encounters the corresponding *construct*.

**all tasks**    All *tasks* participating in the *OpenMP program*.

**current team tasks**    All *tasks* encountered by the corresponding *team*. The *implicit tasks* constituting the **parallel** *region* and any *descendent tasks* encountered during the execution of these *implicit tasks* are included in this set of tasks.

**generating task**    For a given *region*, the task for which execution by a *thread* generated the *region*.

**binding thread set**    The set of *threads* that are affected by, or provide the context for, the execution of a *region*.

The *binding thread* set for a given *region* can be *all threads* on a *device*, *all threads* in a *contention group*, all *master threads* executing an enclosing **teams** *region*, the *current team*, or the *encountering thread*.

COMMENT: The *binding thread set* for a particular *region* is described in its corresponding subsection of this specification.

**binding task set**    The set of *tasks* that are affected by, or provide the context for, the execution of a *region*.

The *binding task* set for a given *region* can be *all tasks*, the *current team tasks*, or the *generating task*.

COMMENT: The *binding task* set for a particular *region* (if applicable) is described in its corresponding subsection of this specification.

| | | |
|---|---|---|
| 1<br>2 | **binding region** | The enclosing *region* that determines the execution context and limits the scope of the effects of the bound *region* is called the *binding region*. |
| 3<br>4<br>5 | | *Binding region* is not defined for *regions* for which the *binding thread* set is *all threads* or the *encountering thread*, nor is it defined for *regions* for which the *binding task set* is *all tasks*. |

COMMENTS:

The *binding region* for an **ordered** *region* is the innermost enclosing *loop region*.

The *binding region* for a **taskwait** *region* is the innermost enclosing *task region*.

The *binding region* for a **cancel** *region* is the innermost enclosing *region* corresponding to the *construct-type-clause* of the **cancel** construct.

The *binding region* for a **cancellation point** *region* is the innermost enclosing *region* corresponding to the *construct-type-clause* of the **cancellation point** construct.

For all other *regions* for which the *binding thread set* is the *current team* or the *binding task set* is the *current team tasks*, the *binding region* is the innermost enclosing **parallel** *region*.

For *regions* for which the *binding task set* is the *generating task*, the *binding region* is the *region* of the *generating task*.

A **parallel** *region* need not be *active* nor explicit to be a *binding region*.

A *task region* need not be explicit to be a *binding region*.

A *region* never binds to any *region* outside of the innermost enclosing **parallel** *region*.

| | | |
|---|---|---|
| 27<br>28 | **orphaned construct** | A *construct* that gives rise to a *region* for which the *binding thread set* is the *current team*, but is not nested within another *construct* giving rise to the *binding region*. |
| 29<br>30 | **worksharing construct** | A *construct* that defines units of work, each of which is executed exactly once by one of the *threads* in the *team* executing the *construct*. |

For C/C++, *worksharing constructs* are **for**, **sections**, and **single**.

For Fortran, *worksharing constructs* are **do**, **sections**, **single** and **workshare**.

| | |
|---|---|
| **place** | Unordered set of *processors* on a device that is treated by the execution environment as a location unit when dealing with OpenMP thread affinity. |
| **place list** | The ordered list that describes all OpenMP *places* available to the execution environment. |
| **place partition** | An ordered list that corresponds to a contiguous interval in the OpenMP *place list*. It describes the *places* currently available to the execution environment for a given parallel *region*. |
| **place number** | A number that uniquely identifies a *place* in the *place list*, with zero identifying the first *place* in the *place list*, and each consecutive whole number identifying the next *place* in the *place list*. |
| **SIMD instruction** | A single machine instruction that can operate on multiple data elements. |
| **SIMD lane** | A software or hardware mechanism capable of processing one data element from a *SIMD instruction*. |
| **SIMD chunk** | A set of iterations executed concurrently, each by a *SIMD lane*, by a single *thread* by means of *SIMD instructions*. |

## 1.2.3 Loop Terminology

| | |
|---|---|
| **loop directive** | An OpenMP *executable* directive for which the associated user code must be a loop nest that is a *structured block*. |
| **associated loop(s)** | The loop(s) controlled by a *loop directive*. |
| | COMMENT: If the *loop directive* contains a **collapse** or an **ordered(***n***)** clause then it may have more than one *associated loop*. |
| **sequential loop** | A loop that is not associated with any OpenMP *loop directive*. |
| **SIMD loop** | A loop that includes at least one *SIMD chunk*. |
| **doacross loop nest** | A loop nest that has cross-iteration dependence. An iteration is dependent on one or more lexicographically earlier iterations. |
| | COMMENT: The **ordered** clause parameter on a loop directive identifies the loop(s) associated with the *doacross loop nest*. |

## 1.2.4  Synchronization Terminology

**barrier**   A point in the execution of a program encountered by a *team* of *threads*, beyond
              which no *thread* in the team may execute until all *threads* in the *team* have reached
              the barrier and all *explicit tasks* generated by the *team* have executed to completion.
              If *cancellation* has been requested, threads may proceed to the end of the canceled
              *region* even if some threads in the team have not reached the *barrier*.

**cancellation**   An action that cancels (that is, aborts) an OpenMP *region* and causes executing
                   *implicit* or *explicit* tasks to proceed to the end of the canceled *region*.

**cancellation point**   A point at which implicit and explicit tasks check if cancellation has been requested.
                         If cancellation has been observed, they perform the *cancellation*.

> COMMENT: For a list of cancellation points, see Section 2.14.1 on
> page 172

## 1.2.5  Tasking Terminology

**task**   A specific instance of executable code and its *data environment*, generated when a
           *thread* encounters a **task**, **taskloop**, **parallel**, **target**, or **teams** *construct*
           (or any *combined construct* that specifies any of these *constructs*).

**task region**   A *region* consisting of all code encountered during the execution of a *task*.

> COMMENT: A **parallel** *region* consists of one or more implicit *task*
> *regions*.

**explicit task**   A *task* generated when a **task** *construct* is encountered during execution.

**implicit task**   A *task* generated by an *implicit parallel region* or generated when a **parallel**
                    *construct* is encountered during execution.

**initial task**   An *implicit task* associated with an *implicit parallel region*.

**current task**   For a given *thread*, the *task* corresponding to the *task region* in which it is executing.

**child task**   A *task* is a *child task* of its generating *task region*. A *child task region* is not part of
                 its generating *task region*.

**sibling tasks**   *Tasks* that are *child tasks* of the same *task region*.

**descendent task**   A *task* that is the *child task* of a *task region* or of one of its *descendent task regions*.

| | |
|---|---|
| **task completion** | *Task completion* occurs when the end of the *structured block* associated with the *construct* that generated the *task* is reached. |
| | COMMENT: Completion of the *initial task* that is generated when the program begins occurs at program exit. |
| **task scheduling point** | A point during the execution of the current *task region* at which it can be suspended to be resumed later; or the point of *task completion*, after which the executing thread may switch to a different *task region*. |
| | COMMENT: For a list of *task scheduling points*, see Section 2.9.5 on page 94. |
| **task switching** | The act of a *thread* switching from the execution of one *task* to another *task*. |
| **tied task** | A *task* that, when its *task region* is suspended, can be resumed only by the same *thread* that suspended it. That is, the *task* is tied to that *thread*. |
| **untied task** | A *task* that, when its *task region* is suspended, can be resumed by any *thread* in the team. That is, the *task* is not tied to any *thread*. |
| **undeferred task** | A *task* for which execution is not deferred with respect to its generating *task region*. That is, its generating *task region* is suspended until execution of the *undeferred task* is completed. |
| **included task** | A *task* for which execution is sequentially included in the generating *task region*. That is, an *included task* is *undeferred* and executed immediately by the *encountering thread*. |
| **merged task** | A *task* for which the *data environment*, inclusive of ICVs, is the same as that of its generating *task region*. |
| **mergeable task** | A *task* that may be a *merged task* if it is an *undeferred task* or an *included task*. |
| **final task** | A *task* that forces all of its *child tasks* to become *final* and *included tasks*. |
| **task dependence** | An ordering relation between two *sibling tasks*: the *dependent task* and a previously generated *predecessor task*. The *task dependence* is fulfilled when the *predecessor task* has completed. |
| **dependent task** | A *task* that because of a *task dependence* cannot be executed until its *predecessor tasks* have completed. |
| **predecessor task** | A *task* that must complete before its *dependent tasks* can be executed. |
| **task synchronization construct** | A **taskwait**, **taskgroup**, or a **barrier** *construct*. |
| **task generating construct** | A **task** or a **taskloop** *construct*. |

| | | |
|---|---|---|
| 1 | **target task** | A *mergeable task* that is generated by a **target**, **target enter data**, |
| 2 | | **target exit data**, or **target update** *construct*. |
| 3 | **taskgroup set** | A set of tasks that are logically grouped by a **taskgroup** *region*. |

## 1.2.6  Data Terminology

| | | |
|---|---|---|
| 5 | **variable** | A named data storage block, for which the value can be defined and redefined during |
| 6 | | the execution of a program. |

▼                                                                                              ▼

7    Note – An array or structure element is a variable that is part of another variable.

▲                                                                                              ▲

| | | |
|---|---|---|
| 8 | **scalar variable** | For C/C++: A scalar variable, as defined by the base language. |
| 9 | | For Fortran: A scalar variable with intrinsic type, as defined by the base language, |
| 10 | | excluding character type. |
| 11 | **array section** | A designated subset of the elements of an array. |
| 12 | **array item** | An array, an array section, or an array element. |
| 13 | **structure** | A structure is a variable that contains one or more variables. |
| 14 | | For C/C++: Implemented using struct types. |
| 15 | | For C++: Implemented using class types. |
| 16 | | For Fortran: Implemented using derived types. |
| 17 | **private variable** | With respect to a given set of *task regions* or *SIMD lanes* that bind to the same |
| 18 | | **parallel** *region*, a *variable* for which the name provides access to a different |
| 19 | | block of storage for each *task region* or *SIMD lane*. |
| 20 | | A *variable* that is part of another variable (as an array or structure element) cannot be |
| 21 | | made private independently of other components. |
| 22 | **shared variable** | With respect to a given set of *task regions* that bind to the same **parallel** *region*, a |
| 23 | | *variable* for which the name provides access to the same block of storage for each |
| 24 | | *task region*. |
| 25 | | A *variable* that is part of another variable (as an array or structure element) cannot be |
| 26 | | *shared* independently of the other components, except for static data members of |
| 27 | | C++ classes. |

| | | |
|---|---|---|
| 1<br>2<br>3 | **threadprivate variable** | A *variable* that is replicated, one instance per *thread*, by the OpenMP implementation. Its name then provides access to a different block of storage for each *thread*. |
| 4<br>5<br>6 | | A *variable* that is part of another variable (as an array or structure element) cannot be made *threadprivate* independently of the other components, except for static data members of C++ classes. |
| 7 | **threadprivate memory** | The set of *threadprivate variables* associated with each *thread*. |
| 8 | **data environment** | The *variables* associated with the execution of a given *region*. |
| 9 | **device data environment** | The initial *data environment* associated with a device. |
| 10 | **device address** | An *implementation defined* reference to an address in a *device data environment*. |
| 11 | **device pointer** | A *variable* that contains a *device address*. |
| 12<br>13 | **mapped variable** | An original *variable* in a *data environment* with a corresponding *variable* in a device *data environment*. |
| 14 | | COMMENT: The original and corresponding *variables* may share storage. |
| 15<br>16<br>17 | **mappable type** | A type that is valid for a *mapped variable*. If a type is composed from other types (such as the type of an array or structure element) and any of the other types are not mappable then the type is not mappable. |
| 18<br>19 | | COMMENT: Pointer types are *mappable* but the memory block to which the pointer refers is not *mapped*. |
| 20 | | For C: The type must be a complete type. |
| 21 | | For C++: The type must be a complete type. |
| 22 | | In addition, for class types: |
| 23<br>24 | | • All member functions accessed in any **target** region must appear in a **declare target** directive. |
| 25 | | • All data members must be non-static. |
| 26 | | • A *mappable type* cannot contain virtual members. |
| 27 | | For Fortran: No restrictions on the type except that for derived types: |
| 28<br>29 | | • All type-bound procedures accessed in any target region must appear in a **declare target** directive. |
| | **defined** | |

| | | |
|---|---|---|
| 1 | | For *variables*, the property of having a valid value. |
| 2 | | For C: For the contents of *variables*, the property of having a valid value. |
| 3<br>4 | | For C++: For the contents of *variables* of POD (plain old data) type, the property of having a valid value. |
| 5<br>6 | | For *variables* of non-POD class type, the property of having been constructed but not subsequently destructed. |
| 7<br>8 | | For Fortran: For the contents of *variables*, the property of having a valid value. For the allocation or association status of *variables*, the property of having a valid status. |
| 9<br>10 | | COMMENT: Programs that rely upon *variables* that are not *defined* are *non-conforming programs*. |
| 11 | **class type** | For C++: *Variables* declared with one of the **class**, **struct**, or **union** keywords |
| 12 | **sequentially consistent atomic construct** | An **atomic** construct for which the **seq_cst** clause is specified. |
| 13 | **non-sequentially consistent atomic construct** | An **atomic** construct for which the **seq_cst** clause is not specified |

## 14  1.2.7  Implementation Terminology

| | | |
|---|---|---|
| 15<br>16 | **supporting *n* levels of parallelism** | Implies allowing an *active parallel region* to be enclosed by *n-1 active parallel regions*. |
| 17 | **supporting the OpenMP API** | Supporting at least one level of parallelism. |
| 18 | **supporting nested parallelism** | Supporting more than one level of parallelism. |
| 19<br>20 | **internal control variable** | A conceptual variable that specifies runtime behavior of a set of *threads* or *tasks* in an *OpenMP program*. |
| 21<br>22 | | COMMENT: The acronym ICV is used interchangeably with the term *internal control variable* in the remainder of this specification. |
| 23<br>24 | **compliant implementation** | An implementation of the OpenMP specification that compiles and executes any *conforming program* as defined by the specification. |
| 25<br>26 | | COMMENT: A *compliant implementation* may exhibit *unspecified behavior* when compiling or executing a *non-conforming program*. |

| 1<br>2 | **unspecified behavior** | A behavior or result that is not specified by the OpenMP specification or not known prior to the compilation or execution of an *OpenMP program*. |
|---|---|---|

1      **unspecified behavior**     A behavior or result that is not specified by the OpenMP specification or not known
2                                         prior to the compilation or execution of an *OpenMP program*.

3                                         Such *unspecified behavior* may result from:

4                                            • Issues documented by the OpenMP specification as having *unspecified behavior*.

5                                            • A *non-conforming program*.

6                                            • A *conforming program* exhibiting an *implementation defined* behavior.

7   **implementation defined**    Behavior that must be documented by the implementation, and is allowed to vary
8                                         among different *compliant implementations*. An implementation is allowed to define
9                                         this behavior as *unspecified*.

10                                         COMMENT: All features that have *implementation defined* behavior are
11                                         documented in Appendix C.

12                  **deprecated**     Implies a construct, clause or other feature is normative in the current specification
13                                         but is considered obsolescent and will be removed in the future.

## 14   1.3   Execution Model

15   The OpenMP API uses the fork-join model of parallel execution. Multiple threads of execution
16 perform tasks defined implicitly or explicitly by OpenMP directives. The OpenMP API is intended
17 to support programs that will execute correctly both as parallel programs (multiple threads of
18 execution and a full OpenMP support library) and as sequential programs (directives ignored and a
19 simple OpenMP stubs library). However, it is possible and permitted to develop a program that
20 executes correctly as a parallel program but not as a sequential program, or that produces different
21 results when executed as a parallel program compared to when it is executed as a sequential
22 program. Furthermore, using different numbers of threads may result in different numeric results
23 because of changes in the association of numeric operations. For example, a serial addition
24 reduction may have a different pattern of addition associations than a parallel reduction. These
25 different associations may change the results of floating-point addition.

26   An OpenMP program begins as a single thread of execution, called an initial thread. An initial
27 thread executes sequentially, as if enclosed in an implicit task region, called an initial task region,
28 that is defined by the implicit parallel region surrounding the whole program.

29   The thread that executes the implicit parallel region that surrounds the whole program executes on
30 the *host device*. An implementation may support other *target devices*. If supported, one or more
31 devices are available to the host device for offloading code and data. Each device has its own
32 threads that are distinct from threads that execute on another device. Threads cannot migrate from

one device to another device. The execution model is host-centric such that the host device offloads **target** regions to target devices.

When a **target** construct is encountered, a new *target task* is generated. The *target task* region encloses the **target** region. The *target task* is complete after the execution of the **target** region is complete.

When a *target task* executes, the enclosed **target** region is executed by an initial thread. The initial thread may execute on a *target device*. The initial thread executes sequentially, as if enclosed in an implicit task region, called an initial task region, that is defined by an implicit **parallel** region that surrounds the entire **target** region. If the target device does not exist or the implementation does not support the target device, all **target** regions associated with that device execute on the host device.

The implementation must ensure that the **target** region executes as if it were executed in the data environment of the target device unless an **if** clause is present and the **if** clause expression evaluates to *false*.

The **teams** construct creates a *league of thread teams* where the master thread of each team executes the region. Each of these master threads is an initial thread, and executes sequentially, as if enclosed in an implicit task region that is defined by an implicit parallel region that surrounds the entire **teams** region.

If a construct creates a data environment, the data environment is created at the time the construct is encountered. Whether a construct creates a data environment is defined in the description of the construct.

When any thread encounters a **parallel** construct, the thread creates a team of itself and zero or more additional threads and becomes the master of the new team. A set of implicit tasks, one per thread, is generated. The code for each task is defined by the code inside the **parallel** construct. Each task is assigned to a different thread in the team and becomes tied; that is, it is always executed by the thread to which it is initially assigned. The task region of the task being executed by the encountering thread is suspended, and each member of the new team executes its implicit task. There is an implicit barrier at the end of the **parallel** construct. Only the master thread resumes execution beyond the end of the **parallel** construct, resuming the task region that was suspended upon encountering the **parallel** construct. Any number of **parallel** constructs can be specified in a single program.

**parallel** regions may be arbitrarily nested inside each other. If nested parallelism is disabled, or is not supported by the OpenMP implementation, then the new team that is created by a thread encountering a **parallel** construct inside a **parallel** region will consist only of the encountering thread. However, if nested parallelism is supported and enabled, then the new team can consist of more than one thread. A **parallel** construct may include a **proc_bind** clause to specify the places to use for the threads in the team within the **parallel** region.

When any team encounters a worksharing construct, the work inside the construct is divided among the members of the team, and executed cooperatively instead of being executed by every thread.

There is a default barrier at the end of each worksharing construct unless the **nowait** clause is present. Redundant execution of code by every thread in the team resumes after the end of the worksharing construct.

When any thread encounters a **task** construct, a new explicit task is generated. Execution of explicitly generated tasks is assigned to one of the threads in the current team, subject to the thread's availability to execute work. Thus, execution of the new task could be immediate, or deferred until later according to task scheduling constraints and thread availability. Threads are allowed to suspend the current task region at a task scheduling point in order to execute a different task. If the suspended task region is for a tied task, the initially assigned thread later resumes execution of the suspended task region. If the suspended task region is for an untied task, then any thread may resume its execution. Completion of all explicit tasks bound to a given parallel region is guaranteed before the master thread leaves the implicit barrier at the end of the region. Completion of a subset of all explicit tasks bound to a given parallel region may be specified through the use of task synchronization constructs. Completion of all explicit tasks bound to the implicit parallel region is guaranteed by the time the program exits.

When any thread encounters a **simd** construct, the iterations of the loop associated with the construct may be executed concurrently using the SIMD lanes that are available to the thread.

The **cancel** construct can alter the previously described flow of execution in an OpenMP region. The effect of the **cancel** construct depends on its *construct-type-clause*. If a task encounters a **cancel** construct with a **taskgroup** *construct-type-clause*, then the task activates cancellation and continues execution at the end of its **task** region, which implies completion of that task. Any other task in that **taskgroup** that has begun executing completes execution unless it encounters a **cancellation point** construct, in which case it continues execution at the end of its **task** region, which implies its completion. Other tasks in that **taskgroup** region that have not begun execution are aborted, which implies their completion.

For all other *construct-type-clause* values, if a thread encounters a **cancel** construct, it activates cancellation of the innermost enclosing region of the type specified and the thread continues execution at the end of that region. Threads check if cancellation has been activated for their region at cancellation points and, if so, also resume execution at the end of the canceled region.

If cancellation has been activated regardless of *construct-type-clause*, threads that are waiting inside a barrier other than an implicit barrier at the end of the canceled region exit the barrier and resume execution at the end of the canceled region. This action can occur before the other threads reach that barrier.

Synchronization constructs and library routines are available in the OpenMP API to coordinate tasks and data access in **parallel** regions. In addition, library routines and environment variables are available to control or to query the runtime environment of OpenMP programs.

The OpenMP specification makes no guarantee that input or output to the same file is synchronous when executed in parallel. In this case, the programmer is responsible for synchronizing input and output statements (or routines) using the provided synchronization constructs or library routines.

1  For the case where each thread accesses a different file, no synchronization by the programmer is
2  necessary.

# 1.4  Memory Model

## 1.4.1  Structure of the OpenMP Memory Model

5   The OpenMP API provides a relaxed-consistency, shared-memory model. All OpenMP threads
6   have access to a place to store and to retrieve variables, called the *memory*. In addition, each thread
7   is allowed to have its own *temporary view* of the memory. The temporary view of memory for each
8   thread is not a required part of the OpenMP memory model, but can represent any kind of
9   intervening structure, such as machine registers, cache, or other local storage, between the thread
10  and the memory. The temporary view of memory allows the thread to cache variables and thereby
11  to avoid going to memory for every reference to a variable. Each thread also has access to another
12  type of memory that must not be accessed by other threads, called *threadprivate memory*.

13  A directive that accepts data-sharing attribute clauses determines two kinds of access to variables
14  used in the directive's associated structured block: shared and private. Each variable referenced in
15  the structured block has an original variable, which is the variable by the same name that exists in
16  the program immediately outside the construct. Each reference to a shared variable in the structured
17  block becomes a reference to the original variable. For each private variable referenced in the
18  structured block, a new version of the original variable (of the same type and size) is created in
19  memory for each task or SIMD lane that contains code associated with the directive. Creation of
20  the new version does not alter the value of the original variable. However, the impact of attempts to
21  access the original variable during the region associated with the directive is unspecified; see
22  Section 2.15.3.3 on page 192 for additional details. References to a private variable in the
23  structured block refer to the private version of the original variable for the current task or SIMD
24  lane. The relationship between the value of the original variable and the initial or final value of the
25  private version depends on the exact clause that specifies it. Details of this issue, as well as other
26  issues with privatization, are provided in Section 2.15 on page 178.

27  The minimum size at which a memory update may also read and write back adjacent variables that
28  are part of another variable (as array or structure elements) is implementation defined but is no
29  larger than required by the base language.

30  A single access to a variable may be implemented with multiple load or store instructions, and
31  hence is not guaranteed to be atomic with respect to other accesses to the same variable. Accesses
32  to variables smaller than the implementation defined minimum size or to C or C++ bit-fields may
33  be implemented by reading, modifying, and rewriting a larger unit of memory, and may thus
34  interfere with updates of variables or fields in the same unit of memory.

1   If multiple threads write without synchronization to the same memory unit, including cases due to
2   atomicity considerations as described above, then a data race occurs. Similarly, if at least one
3   thread reads from a memory unit and at least one thread writes without synchronization to that
4   same memory unit, including cases due to atomicity considerations as described above, then a data
5   race occurs. If a data race occurs then the result of the program is unspecified.

6   A private variable in a task region that eventually generates an inner nested **parallel** region is
7   permitted to be made shared by implicit tasks in the inner **parallel** region. A private variable in
8   a task region can be shared by an explicit **task** region generated during its execution. However, it
9   is the programmer's responsibility to ensure through synchronization that the lifetime of the
10  variable does not end before completion of the explicit **task** region sharing it. Any other access by
11  one task to the private variables of another task results in unspecified behavior.

## 1.4.2  Device Data Environments

13  When an OpenMP program begins, an implicit **target data** region for each device surrounds
14  the whole program. Each device has a device data environment that is defined by its implicit
15  **target data** region. Any **declare target** directives and the directives that accept
16  data-mapping attribute clauses determine how an original variable in a data environment is mapped
17  to a corresponding variable in a device data environment.

18  When an original variable is mapped to a device data environment and the associated
19  corresponding variable is not present in the device data environment, a new corresponding variable
20  (of the same type and size as the original variable) is created in the device data environment. The
21  initial value of the new corresponding variable is determined from the clauses and the data
22  environment of the encountering thread.

23  The corresponding variable in the device data environment may share storage with the original
24  variable. Writes to the corresponding variable may alter the value of the original variable. The
25  impact of this on memory consistency is discussed in Section 1.4.4 on page 20. When a task
26  executes in the context of a device data environment, references to the original variable refer to the
27  corresponding variable in the device data environment.

28  The relationship between the value of the original variable and the initial or final value of the
29  corresponding variable depends on the *map-type*. Details of this issue, as well as other issues with
30  mapping a variable, are provided in Section 2.15.5.1 on page 216.

31  The original variable in a data environment and the corresponding variable(s) in one or more device
32  data environments may share storage. Without intervening synchronization data races can occur.

## 1.4.3 The Flush Operation

The memory model has relaxed-consistency because a thread's temporary view of memory is not required to be consistent with memory at all times. A value written to a variable can remain in the thread's temporary view until it is forced to memory at a later time. Likewise, a read from a variable may retrieve the value from the thread's temporary view, unless it is forced to read from memory. The OpenMP flush operation enforces consistency between the temporary view and memory.

The flush operation is applied to a set of variables called the *flush-set*. The flush operation restricts reordering of memory operations that an implementation might otherwise do. Implementations must not reorder the code for a memory operation for a given variable, or the code for a flush operation for the variable, with respect to a flush operation that refers to the same variable.

If a thread has performed a write to its temporary view of a shared variable since its last flush of that variable, then when it executes another flush of the variable, the flush does not complete until the value of the variable has been written to the variable in memory. If a thread performs multiple writes to the same variable between two flushes of that variable, the flush ensures that the value of the last write is written to the variable in memory. A flush of a variable executed by a thread also causes its temporary view of the variable to be discarded, so that if its next memory operation for that variable is a read, then the thread will read from memory when it may again capture the value in the temporary view. When a thread executes a flush, no later memory operation by that thread for a variable involved in that flush is allowed to start until the flush completes. The completion of a flush of a set of variables executed by a thread is defined as the point at which all writes to those variables performed by the thread before the flush are visible in memory to all other threads and that thread's temporary view of all variables involved is discarded.

The flush operation provides a guarantee of consistency between a thread's temporary view and memory. Therefore, the flush operation can be used to guarantee that a value written to a variable by one thread may be read by a second thread. To accomplish this, the programmer must ensure that the second thread has not written to the variable since its last flush of the variable, and that the following sequence of events happens in the specified order:

1. The value is written to the variable by the first thread.

2. The variable is flushed by the first thread.

3. The variable is flushed by the second thread.

4. The value is read from the variable by the second thread.

Note – OpenMP synchronization operations, described in Section 2.13 on page 148 and in Section 3.3 on page 270, are recommended for enforcing this order. Synchronization through variables is possible but is not recommended because the proper timing of flushes is difficult.

## 1.4.4 OpenMP Memory Consistency

The restrictions in Section 1.4.3 on page 19 on reordering with respect to flush operations guarantee the following:

- If the intersection of the flush-sets of two flushes performed by two different threads is non-empty, then the two flushes must be completed as if in some sequential order, seen by all threads.

- If two operations performed by the same thread either access, modify, or flush the same variable, then they must be completed as if in that thread's program order, as seen by all threads.

- If the intersection of the flush-sets of two flushes is empty, the threads can observe these flushes in any order.

The flush operation can be specified using the **flush** directive, and is also implied at various locations in an OpenMP program: see Section 2.13.7 on page 162 for details.

▼                                                                                        ▼

Note – Since flush operations by themselves cannot prevent data races, explicit flush operations are only useful in combination with non-sequentially consistent atomic directives.

▲                                                                                        ▲

OpenMP programs that:

- do not use non-sequentially consistent atomic directives,

- do not rely on the accuracy of a *false* result from **omp_test_lock** and **omp_test_nest_lock**, and

- correctly avoid data races as required in Section 1.4.1 on page 17

behave as though operations on shared variables were simply interleaved in an order consistent with the order in which they are performed by each thread. The relaxed consistency model is invisible for such programs, and any explicit flush operations in such programs are redundant.

Implementations are allowed to relax the ordering imposed by implicit flush operations when the result is only visible to programs using non-sequentially consistent atomic directives.

# 1.5  OpenMP Compliance

An implementation of the OpenMP API is compliant if and only if it compiles and executes all
conforming programs according to the syntax and semantics laid out in Chapters 1, 2, 3 and 4.
Appendices A, B, C and D and sections designated as Notes (see Section 1.7 on page 23) are for
information purposes only and are not part of the specification.

The OpenMP API defines constructs that operate in the context of the base language that is
supported by an implementation. If the base language does not support a language construct that
appears in this document, a compliant OpenMP implementation is not required to support it, with
the exception that for Fortran, the implementation must allow case insensitivity for directive and
API routines names, and must allow identifiers of more than six characters

All library, intrinsic and built-in routines provided by the base language must be thread-safe in a
compliant implementation. In addition, the implementation of the base language must also be
thread-safe. For example, **ALLOCATE** and **DEALLOCATE** statements must be thread-safe in
Fortran. Unsynchronized concurrent use of such routines by different threads must produce correct
results (although not necessarily the same as serial execution results, as in the case of random
number generation routines).

Starting with Fortran 90, variables with explicit initialization have the **SAVE** attribute implicitly.
This is not the case in Fortran 77. However, a compliant OpenMP Fortran implementation must
give such a variable the **SAVE** attribute, regardless of the underlying base language version.

Appendix C lists certain aspects of the OpenMP API that are implementation defined. A compliant
implementation is required to define and document its behavior for each of the items in Appendix C.

# 1.6  Normative References

- ISO/IEC 9899:1990, *Information Technology - Programming Languages - C*.

  This OpenMP API specification refers to ISO/IEC 9899:1990 as C90.

- ISO/IEC 9899:1999, *Information Technology - Programming Languages - C*.

  This OpenMP API specification refers to ISO/IEC 9899:1999 as C99.

- ISO/IEC 14882:1998, *Information Technology - Programming Languages - C++*.

  This OpenMP API specification refers to ISO/IEC 14882:1998 as C++.

- ISO/IEC 1539:1980, *Information Technology - Programming Languages - Fortran*.

  This OpenMP API specification refers to ISO/IEC 1539:1980 as Fortran 77.

- ISO/IEC 1539:1991, *Information Technology - Programming Languages - Fortran*.

    This OpenMP API specification refers to ISO/IEC 1539:1991 as Fortran 90.

- ISO/IEC 1539-1:1997, *Information Technology - Programming Languages - Fortran*.

    This OpenMP API specification refers to ISO/IEC 1539-1:1997 as Fortran 95.

- ISO/IEC 1539-1:2004, *Information Technology - Programming Languages - Fortran*.

    This OpenMP API specification refers to ISO/IEC 1539-1:2004 as Fortran 2003. The following features are not supported:

    – IEEE Arithmetic issues covered in Fortran 2003 Section 14

    – Parameterized derived types

    – The **PASS** attribute

    – Procedures bound to a type as operators

    – Overriding a type-bound procedure

    – Polymorphic entities

    – **SELECT TYPE** construct

    – Deferred bindings and abstract types

    – Controlling IEEE underflow

    – Another IEEE class value

Where this OpenMP API specification refers to C, C++ or Fortran, reference is made to the base language supported by the implementation.

# 1.7  Organization of this Document

The remainder of this document is structured as follows:

- Chapter 2 "Directives"
- Chapter 3 "Runtime Library Routines"
- Chapter 4 "Environment Variables"
- Appendix A "Stubs for Runtime Library Routines"
- Appendix B "Interface Declarations"
- Appendix C "OpenMP Implementation-Defined Behaviors"
- Appendix D "Features History"

Some sections of this document only apply to programs written in a certain base language. Text that applies only to programs for which the base language is C or C++ is shown as follows:

———————————————————————————————— C / C++ ————————————————————————————————

C/C++ specific text...

———————————————————————————————— C / C++ ————————————————————————————————

Text that applies only to programs for which the base language is C only is shown as follows:

———————————————————————————————— C ————————————————————————————————

C specific text...

———————————————————————————————— C ————————————————————————————————

Text that applies only to programs for which the base language is C90 only is shown as follows:

———————————————————————————————— C90 ————————————————————————————————

C90 specific text...

———————————————————————————————— C90 ————————————————————————————————

Text that applies only to programs for which the base language is C99 only is shown as follows:

———————————————————————————————— C99 ————————————————————————————————

C99 specific text...

———————————————————————————————— C99 ————————————————————————————————

Text that applies only to programs for which the base language is C++ only is shown as follows:

————————————————————————— C++ —————————————————————————

C++ specific text...

————————————————————————— C++ —————————————————————————

Text that applies only to programs for which the base language is Fortran is shown as follows:

————————————————————————— Fortran —————————————————————————

Fortran specific text......

————————————————————————— Fortran —————————————————————————

Where an entire page consists of, for example, Fortran specific text, a marker is shown at the top of the page like this:

- - - - - - - - - - - - - - - - - - - - - - Fortran (cont.) - - - - - - - - - - - - - - - - - - - - - -

Some text is for information only, and is not part of the normative specification. Such text is designated as a note, like this:

Note – Non-normative text....
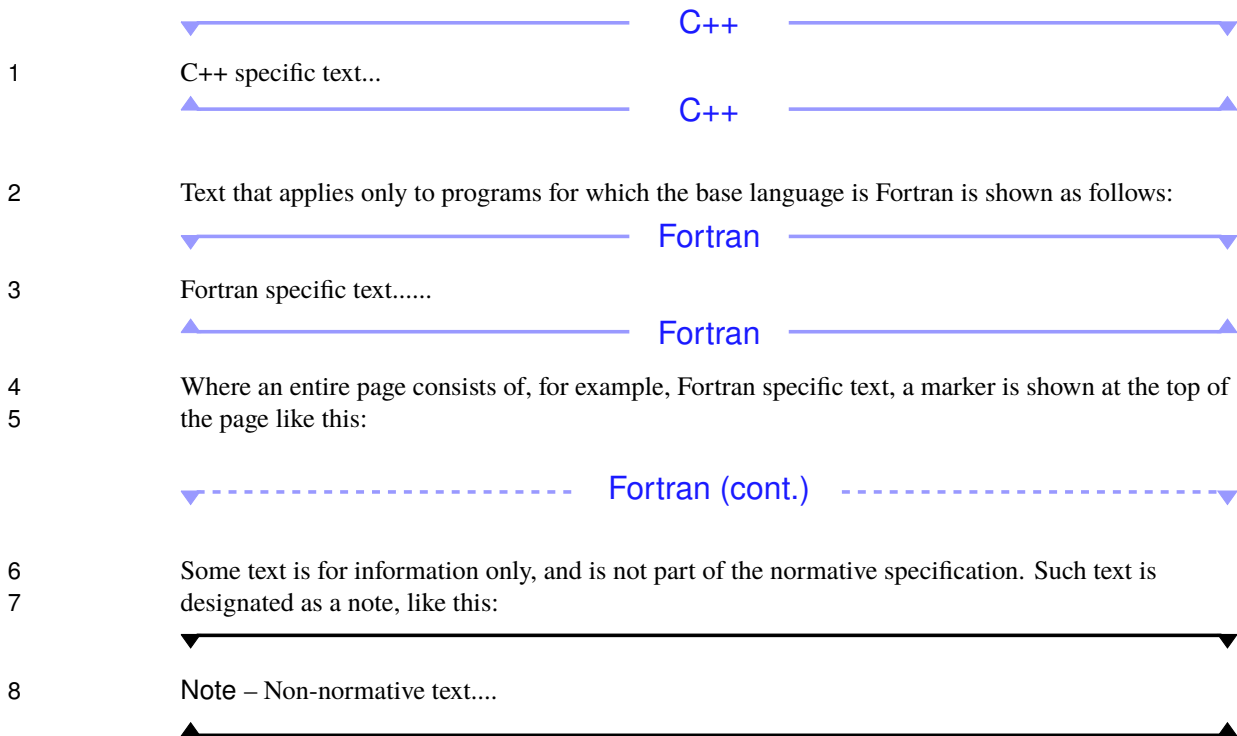
# <sub>2</sub> Directives

<sub>3</sub> This chapter describes the syntax and behavior of OpenMP directives, and is divided into the
<sub>4</sub> following sections:

<sub>5</sub>  • The language-specific directive format (Section 2.1 on page 26)

<sub>6</sub>  • Mechanisms to control conditional compilation (Section 2.2 on page 33)

<sub>7</sub>  • Control of OpenMP API ICVs (Section 2.3 on page 36)

<sub>8</sub>  • How to specify and to use array sections for all base languages (Section 2.4 on page 44)

<sub>9</sub>  • Details of each OpenMP directive (Section 2.5 on page 46 to Section 2.17 on page 227)

───────────────────────── C / C++ ─────────────────────────

<sub>10</sub> In C/C++, OpenMP directives are specified by using the **#pragma** mechanism provided by the C
<sub>11</sub> and C++ standards.

───────────────────────── C / C++ ─────────────────────────
───────────────────────── Fortran ─────────────────────────

<sub>12</sub> In Fortran, OpenMP directives are specified by using special comments that are identified by
<sub>13</sub> unique sentinels. Also, a special comment form is available for conditional compilation.

───────────────────────── Fortran ─────────────────────────

<sub>14</sub> Compilers can therefore ignore OpenMP directives and conditionally compiled code if support of
<sub>15</sub> the OpenMP API is not provided or enabled. A compliant implementation must provide an option
<sub>16</sub> or interface that ensures that underlying support of all OpenMP directives and OpenMP conditional
<sub>17</sub> compilation mechanisms is enabled. In the remainder of this document, the phrase *OpenMP*
<sub>18</sub> *compilation* is used to mean a compilation with these OpenMP features enabled.

---

***Fortran***

---

**Restrictions**

The following restriction applies to all OpenMP directives:

- OpenMP directives, except SIMD and **declare target** directives, may not appear in pure procedures.

---

***Fortran***

---

# 2.1 Directive Format

---

***C / C++***

---

OpenMP directives for C/C++ are specified with the **pragma** preprocessing directive. The syntax of an OpenMP directive is as follows:

> **#pragma omp** *directive-name [clause[ [ , ] clause] ... ] new-line*

Each directive starts with **#pragma omp**. The remainder of the directive follows the conventions of the C and C++ standards for compiler directives. In particular, white space can be used before and after the **#**, and sometimes white space must be used to separate the words in a directive. Preprocessing tokens following the **#pragma omp** are subject to macro replacement.

Some OpenMP directives may be composed of consecutive **#pragma** preprocessing directives if specified in their syntax.

Directives are case-sensitive.

An OpenMP executable directive applies to at most one succeeding statement, which must be a structured block.

---

***C / C++***

---

1    OpenMP directives for Fortran are specified as follows:

> *sentinel directive-name [clause[ [ , ] clause]...]*

2    All OpenMP compiler directives must begin with a directive *sentinel*. The format of a sentinel
3    differs between fixed and free-form source files, as described in Section 2.1.1 on page 28 and
4    Section 2.1.2 on page 29.

5    Directives are case insensitive. Directives cannot be embedded within continued statements, and
6    statements cannot be embedded within directives.

7    In order to simplify the presentation, free form is used for the syntax of OpenMP directives for
8    Fortran in the remainder of this document, except as noted.

9    Only one *directive-name* can be specified per directive (note that this includes combined directives,
10   see Section 2.11 on page 124). The order in which clauses appear on directives is not significant.
11   Clauses on directives may be repeated as needed, subject to the restrictions listed in the description
12   of each clause.

13   Some data-sharing attribute clauses (Section 2.15.3 on page 188), data copying clauses
14   (Section 2.15.4 on page 211), the **threadprivate** directive (Section 2.15.2 on page 183), the
15   **flush** directive (Section 2.13.7 on page 162), and the **link** clause of the **declare target**
16   directive (Section 2.10.6 on page 110) accept a *list*. The **to** clause of the **declare target**
17   directive (Section 2.10.6 on page 110) accepts an *extended-list*. A *list* consists of a
18   comma-separated collection of one or more *list items*. A *extended-list* consists of a
19   comma-separated collection of one or more *extended list items*.

20   A *list item* is a variable or array section. An *extended list item* is a *list item* or a function name.

21   A *list item* is a variable, array section or common block name (enclosed in slashes). An *extended*
22   *list item* is a *list item* or a procedure name.

23   For all base languages, a *list item* or an *extended list item* is subject to the restrictions specified in
24   Section 2.4 on page 44 and in each of the sections describing clauses and directives for which the
25   *list* or *extended-list* appears.

## 2.1.1 Fixed Source Form Directives

The following sentinels are recognized in fixed form source files:

```
!$omp | c$omp | *$omp
```

Sentinels must start in column 1 and appear as a single word with no intervening characters. Fortran fixed form line length, white space, continuation, and column rules apply to the directive line. Initial directive lines must have a space or zero in column 6, and continuation directive lines must have a character other than a space or a zero in column 6.

Comments may appear on the same line as a directive. The exclamation point initiates a comment when it appears after column 6. The comment extends to the end of the source line and is ignored. If the first non-blank character after the directive sentinel of an initial or continuation directive line is an exclamation point, the line is ignored.

Note – in the following example, the three formats for specifying the directive are equivalent (the first line represents the position of the first 9 columns):

```
c23456789
!$omp parallel do shared(a,b,c)

c$omp parallel do
c$omp+shared(a,b,c)

c$omp paralleldoshared(a,b,c)
```

## 2.1.2 Free Source Form Directives

The following sentinel is recognized in free form source files:

```
!$omp
```

The sentinel can appear in any column as long as it is preceded only by white space (spaces and tab characters). It must appear as a single word with no intervening character. Fortran free form line length, white space, and continuation rules apply to the directive line. Initial directive lines must have a space after the sentinel. Continued directive lines must have an ampersand (**&**) as the last non-blank character on the line, prior to any comment placed inside the directive. Continuation directive lines can have an ampersand after the directive sentinel with optional white space before and after the ampersand.

Comments may appear on the same line as a directive. The exclamation point (**!**) initiates a comment. The comment extends to the end of the source line and is ignored. If the first non-blank character after the directive sentinel is an exclamation point, the line is ignored.

One or more blanks or horizontal tabs must be used to separate adjacent keywords in directives in free source form, except in the following cases, where white space is optional between the given set of keywords:

```
declare reduction

declare simd

declare target

distribute parallel do

distribute parallel do simd

distribute simd

do simd

end atomic

end critical

end distribute

end distribute parallel do

end distribute parallel do simd
```

```
1          end distribute simd
2          end do
3          end do simd
4          end master
5          end ordered
6          end parallel
7          end parallel do
8          end parallel do simd
9          end parallel sections
10         end parallel workshare
11         end sections
12         end simd
13         end single
14         end target
15         end target data
16         end target parallel
17         end target parallel do
18         end target parallel do simd
19         end target simd
20         end target teams
21         end target teams distribute
22         end target teams distribute parallel do
23         end target teams distribute parallel do simd
24         end target teams distribute simd
25         end task
26         end taskgroup
27         end taskloop
```

```
1          end taskloop simd
2          end teams
3          end teams distribute
4          end teams distribute parallel do
5          end teams distribute parallel do simd
6          end teams distribute simd
7          end workshare
8          parallel do
9          parallel do simd
10         parallel sections
11         parallel workshare
12         target data
13         target enter data
14         target exit data
15         target parallel
16         target parallel do
17         target parallel do simd
18         target simd
19         target teams
20         target teams distribute
21         target teams distribute parallel do
22         target teams distribute parallel do simd
23         target teams distribute simd
24         target update
25         taskloop simd
26         teams distribute
27         teams distribute parallel do
```

```
teams distribute parallel do simd

teams distribute simd
```

▼                                                                        ▼

Note – in the following example the three formats for specifying the directive are equivalent (the first line represents the position of the first 9 columns):

```
!23456789
        !$omp parallel do &
                    !$omp shared(a,b,c)

        !$omp parallel &
        !$omp&do shared(a,b,c)

!$omp paralleldo shared(a,b,c)
```

▲                                                                        ▲

▲ ——————————————————————— Fortran ——————————————————— ▲

## 2.1.3  Stand-Alone Directives

**Summary**

Stand-alone directives are executable directives that have no associated user code.

**Description**

Stand-alone directives do not have any associated executable user code. Instead, they represent executable statements that typically do not have succinct equivalent statements in the base languages. There are some restrictions on the placement of a stand-alone directive within a program. A stand-alone directive may be placed only at a point where a base language executable statement is allowed.

**Restrictions**

——————————————————— C / C++ ———————————————————

2  For C/C++, a stand-alone directive may not be used in place of the statement following an **if**,
3  **while**, **do**, **switch**, or **label**.

——————————————————— C / C++ ———————————————————

——————————————————— Fortran ———————————————————

4  For Fortran, a stand-alone directive may not be used as the action statement in an **if** statement or
5  as the executable statement following a label if the label is referenced in the program.

——————————————————— Fortran ———————————————————

# 6  2.2  Conditional Compilation

7  In implementations that support a preprocessor, the **_OPENMP** macro name is defined to have the
8  decimal value *yyyymm* where *yyyy* and *mm* are the year and month designations of the version of
9  the OpenMP API that the implementation supports.

10  If this macro is the subject of a **#define** or a **#undef** preprocessing directive, the behavior is
11  unspecified.

——————————————————— Fortran ———————————————————

12  The OpenMP API requires Fortran lines to be compiled conditionally, as described in the following
13  sections.

## 2.2.1 Fixed Source Form Conditional Compilation Sentinels

The following conditional compilation sentinels are recognized in fixed form source files:

```
!$ | *$ | c$
```

To enable conditional compilation, a line with a conditional compilation sentinel must satisfy the following criteria:

- The sentinel must start in column 1 and appear as a single word with no intervening white space.

- After the sentinel is replaced with two spaces, initial lines must have a space or zero in column 6 and only white space and numbers in columns 1 through 5.

- After the sentinel is replaced with two spaces, continuation lines must have a character other than a space or zero in column 6 and only white space in columns 1 through 5.

If these criteria are met, the sentinel is replaced by two spaces. If these criteria are not met, the line is left unchanged.

Note – in the following example, the two forms for specifying conditional compilation in fixed source form are equivalent (the first line represents the position of the first 9 columns):

```
c23456789
!$ 10 iam = omp_get_thread_num() +
!$   &          index

#ifdef _OPENMP
   10 iam = omp_get_thread_num() +
      &              index
#endif
```

## 2.2.2 Free Source Form Conditional Compilation Sentinel

The following conditional compilation sentinel is recognized in free form source files:

```
!$
```

To enable conditional compilation, a line with a conditional compilation sentinel must satisfy the
following criteria:

- The sentinel can appear in any column but must be preceded only by white space.

- The sentinel must appear as a single word with no intervening white space.

- Initial lines must have a space after the sentinel.

- Continued lines must have an ampersand as the last non-blank character on the line, prior to any
  comment appearing on the conditionally compiled line. Continuation lines can have an
  ampersand after the sentinel, with optional white space before and after the ampersand.

If these criteria are met, the sentinel is replaced by two spaces. If these criteria are not met, the line
is left unchanged.

▼───────────────────────────────────────────────────────────────▼

Note – in the following example, the two forms for specifying conditional compilation in free
source form are equivalent (the first line represents the position of the first 9 columns):

```
c23456789
 !$ iam = omp_get_thread_num() +     &
 !$&    index

#ifdef _OPENMP
    iam = omp_get_thread_num() +     &
         index
#endif
```

▲───────────────────────────────────────────────────────────────▲

▲─────────────────────────── Fortran ───────────────────────────▲

# 2.3 Internal Control Variables

An OpenMP implementation must act as if there are internal control variables (ICVs) that control the behavior of an OpenMP program. These ICVs store information such as the number of threads to use for future **parallel** regions, the schedule to use for worksharing loops and whether nested parallelism is enabled or not. The ICVs are given values at various times (described below) during the execution of the program. They are initialized by the implementation itself and may be given values through OpenMP environment variables and through calls to OpenMP API routines. The program can retrieve the values of these ICVs only through OpenMP API routines.

For purposes of exposition, this document refers to the ICVs by certain names, but an implementation is not required to use these names or to offer any way to access the variables other than through the ways shown in Section 2.3.2 on page 37.

## 2.3.1 ICV Descriptions

The following ICVs store values that affect the operation of **parallel** regions.

- *dyn-var* - controls whether dynamic adjustment of the number of threads is enabled for encountered **parallel** regions. There is one copy of this ICV per data environment.

- *nest-var* - controls whether nested parallelism is enabled for encountered **parallel** regions. There is one copy of this ICV per data environment.

- *nthreads-var* - controls the number of threads requested for encountered **parallel** regions. There is one copy of this ICV per data environment.

- *thread-limit-var* - controls the maximum number of threads participating in the contention group. There is one copy of this ICV per data environment.

- *max-active-levels-var* - controls the maximum number of nested active **parallel** regions. There is one copy of this ICV per device.

- *place-partition-var* – controls the place partition available to the execution environment for encountered **parallel** regions. There is one copy of this ICV per implicit task.

- *active-levels-var* - the number of nested, active parallel regions enclosing the current task such that all of the **parallel** regions are enclosed by the outermost initial task region on the current device. There is one copy of this ICV per data environment.

- *levels-var* - the number of nested parallel regions enclosing the current task such that all of the **parallel** regions are enclosed by the outermost initial task region on the current device. There is one copy of this ICV per data environment.

- *bind-var* - controls the binding of OpenMP threads to places. When binding is requested, the variable indicates that the execution environment is advised not to move threads between places. The variable can also provide default thread affinity policies. There is one copy of this ICV per data environment.

The following ICVs store values that affect the operation of loop regions.

- *run-sched-var* - controls the schedule that the **runtime** schedule clause uses for loop regions. There is one copy of this ICV per data environment.
- *def-sched-var* - controls the implementation defined default scheduling of loop regions. There is one copy of this ICV per device.

The following ICVs store values that affect program execution.

- *stacksize-var* - controls the stack size for threads that the OpenMP implementation creates. There is one copy of this ICV per device.
- *wait-policy-var* - controls the desired behavior of waiting threads. There is one copy of this ICV per device.
- *cancel-var* - controls the desired behavior of the **cancel** construct and cancellation points. There is one copy of this ICV for the whole program.
- *default-device-var* - controls the default target device. There is one copy of this ICV per data environment.
- *max-task-priority-var* - controls the maximum priority value that can be specified in the **priority** clause of the **task** construct. There is one copy of this ICV for the whole program.

## 2.3.2 ICV Initialization

Table 2.1 shows the ICVs, associated environment variables, and initial values.

**TABLE 2.1:** ICV Initial Values

| ICV | Environment Variable | Initial value |
| --- | --- | --- |
| *dyn-var* | **OMP_DYNAMIC** | See description below |
| *nest-var* | **OMP_NESTED** | *false* |
| *nthreads-var* | **OMP_NUM_THREADS** | Implementation defined |

*table continued on next page*

| ICV | Environment Variable | Initial value |
|---|---|---|
| *run-sched-var* | **OMP_SCHEDULE** | Implementation defined |
| *def-sched-var* | (none) | Implementation defined |
| *bind-var* | **OMP_PROC_BIND** | Implementation defined |
| *stacksize-var* | **OMP_STACKSIZE** | Implementation defined |
| *wait-policy-var* | **OMP_WAIT_POLICY** | Implementation defined |
| *thread-limit-var* | **OMP_THREAD_LIMIT** | Implementation defined |
| *max-active-levels-var* | **OMP_MAX_ACTIVE_LEVELS** | See description below |
| *active-levels-var* | (none) | *zero* |
| *levels-var* | (none) | *zero* |
| *place-partition-var* | **OMP_PLACES** | Implementation defined |
| *cancel-var* | **OMP_CANCELLATION** | *false* |
| *default-device-var* | **OMP_DEFAULT_DEVICE** | Implementation defined |
| *max-task-priority-var* | **OMP_MAX_TASK_PRIORITY** | zero |

**Description**

- Each device has its own ICVs.

- The value of the *nthreads-var* ICV is a list.

- The value of the *bind-var* ICV is a list.

- The initial value of *dyn-var* is implementation defined if the implementation supports dynamic adjustment of the number of threads; otherwise, the initial value is *false*.

- The initial value of *max-active-levels-var* is the number of levels of parallelism that the implementation supports. See the definition of *supporting n levels of parallelism* in Section 1.2.7 on page 13 for further details.

The host and target device ICVs are initialized before any OpenMP API construct or OpenMP API routine executes. After the initial values are assigned, the values of any OpenMP environment variables that were set by the user are read and the associated ICVs for the host device are modified accordingly. The method for initializing a target device's ICVs is implementation defined.

**Cross References**

## 2.3.3 Modifying and Retrieving ICV Values

Table 2.2 shows the method for modifying and retrieving the values of ICVs through OpenMP API routines.

**TABLE 2.2:** Ways to Modify and to Retrieve ICV Values

| ICV | Ways to modify value | Ways to retrieve value |
|-----|----------------------|------------------------|
| *dyn-var* | `omp_set_dynamic()` | `omp_get_dynamic()` |
| *nest-var* | `omp_set_nested()` | `omp_get_nested()` |
| *nthreads-var* | `omp_set_num_threads()` | `omp_get_max_threads()` |
| *run-sched-var* | `omp_set_schedule()` | `omp_get_schedule()` |
| *def-sched-var* | (none) | (none) |

*table continued on next page*

| ICV | Ways to modify value | Ways to retrieve value |
|---|---|---|
| *bind-var* | (none) | **omp_get_proc_bind()** |
| *stacksize-var* | (none) | (none) |
| *wait-policy-var* | (none) | (none) |
| *thread-limit-var* | **thread_limit** clause | **omp_get_thread_limit()** |
| *max-active-levels-var* | **omp_set_max_active_levels()** | **omp_get_max_active_levels()** |
| *active-levels-var* | (none) | **omp_get_active_level()** |
| *levels-var* | (none) | **omp_get_level()** |
| *place-partition-var* | (none) | See description below |
| *cancel-var* | (none) | **omp_get_cancellation()** |
| *default-device-var* | **omp_set_default_device()** | **omp_get_default_device()** |
| *max-task-priority-var* | (none) | **omp_get_max_task_priority()** |

## Description

- The value of the *nthreads-var* ICV is a list. The runtime call **omp_set_num_threads()** sets the value of the first element of this list, and **omp_get_max_threads()** retrieves the value of the first element of this list.

- The value of the *bind-var* ICV is a list. The runtime call **omp_get_proc_bind()** retrieves the value of the first element of this list.

- Detailed values in the *place-partition-var* ICV are retrieved using the runtime calls **omp_get_partition_num_places()**, **omp_get_partition_place_nums()**, **omp_get_place_num_procs()**, and **omp_get_place_proc_ids()**.

## Cross References

- **thread_limit** clause of the **teams** construct, see Section 2.10.7 on page 114.

- **omp_set_num_threads** routine, see Section 3.2.1 on page 231.

- **omp_get_max_threads** routine, see Section 3.2.3 on page 233.

- **omp_set_dynamic** routine, see Section 3.2.7 on page 237.

- **omp_get_dynamic** routine, see Section 3.2.8 on page 239.

- **omp_get_cancellation** routine, see Section 3.2.9 on page 240.

- **omp_set_nested** routine, see Section 3.2.10 on page 240.

1  • **omp_get_nested** routine, see Section 3.2.11 on page 242.

2  • **omp_set_schedule** routine, see Section 3.2.12 on page 243.

3  • **omp_get_schedule** routine, see Section 3.2.13 on page 245.

4  • **omp_get_thread_limit** routine, see Section 3.2.14 on page 246.

5  • **omp_set_max_active_levels** routine, see Section 3.2.15 on page 246.

6  • **omp_get_max_active_levels** routine, see Section 3.2.16 on page 248.

7  • **omp_get_level** routine, see Section 3.2.17 on page 249.

8  • **omp_get_active_level** routine, see Section 3.2.20 on page 252.

9  • **omp_get_proc_bind** routine, see Section 3.2.22 on page 254.

10  • **omp_get_place_num_procs()** routine, see Section 3.2.24 on page 257.

11  • **omp_get_place_proc_ids()** routine, see Section 3.2.25 on page 258.

12  • **omp_get_partition_num_places()** routine, see Section 3.2.27 on page 260.

13  • **omp_get_partition_place_nums()** routine, see Section 3.2.28 on page 261.

14  • **omp_set_default_device** routine, see Section 3.2.29 on page 262.

15  • **omp_get_default_device** routine, see Section 3.2.30 on page 263.

16  • **omp_get_max_task_priority** routine, see Section 3.2.36 on page 268.

## 17  **2.3.4  How ICVs are Scoped**

18  Table 2.3 shows the ICVs and their scope.

19  **TABLE 2.3:** Scopes of ICVs

| ICV | Scope |
| --- | --- |
| *dyn-var* | data environment |
| *nest-var* | data environment |
| *nthreads-var* | data environment |
| *run-sched-var* | data environment |
| *def-sched-var* | device |

20

*table continued on next page*

| ICV | Scope |
|---|---|
| *bind-var* | data environment |
| *stacksize-var* | device |
| *wait-policy-var* | device |
| *thread-limit-var* | data environment |
| *max-active-levels-var* | device |
| *active-levels-var* | data environment |
| *levels-var* | data environment |
| *place-partition-var* | implicit task |
| *cancel-var* | global |
| *default-device-var* | data environment |
| *max-task-priority-var* | global |

**Description**

- There is one copy per device of each ICV with device scope

- Each data environment has its own copies of ICVs with data environment scope

- Each implicit task has its own copy of ICVs with implicit task scope

Calls to OpenMP API routines retrieve or modify data environment scoped ICVs in the data environment of their binding tasks.

## 2.3.4.1 How the Per-Data Environment ICVs Work

When a **task** construct or **parallel** construct is encountered, the generated task(s) inherit the values of the data environment scoped ICVs from the generating task's ICV values.

When a **task** construct is encountered, the generated task inherits the value of *nthreads-var* from the generating task's *nthreads-var* value. When a **parallel** construct is encountered, and the generating task's *nthreads-var* list contains a single element, the generated task(s) inherit that list as the value of *nthreads-var*. When a **parallel** construct is encountered, and the generating task's *nthreads-var* list contains multiple elements, the generated task(s) inherit the value of *nthreads-var* as the list obtained by deletion of the first element from the generating task's *nthreads-var* value. The *bind-var* ICV is handled in the same way as the *nthreads-var* ICV.

1 When a *target task* executes a **target** region, the generated initial task uses the values of the data
2 environment scoped ICVs from the device data environment ICV values of the device that will
3 execute the region.

4 If a **teams** construct with a **thread_limit** clause is encountered, the *thread-limit-var* ICV of
5 the construct's data environment is instead set to a value that is less than or equal to the value
6 specified in the clause.

7 When encountering a loop worksharing region with **schedule(runtime)**, all implicit task
8 regions that constitute the binding parallel region must have the same value for *run-sched-var* in
9 their data environments. Otherwise, the behavior is unspecified.

## 2.3.5 ICV Override Relationships

Table 2.4 shows the override relationships among construct clauses and ICVs.

**TABLE 2.4:** ICV Override Relationships

| ICV | construct clause, if used |
|---|---|
| *dyn-var* | (none) |
| *nest-var* | (none) |
| *nthreads-var* | **num_threads** |
| *run-sched-var* | **schedule** |
| *def-sched-var* | **schedule** |
| *bind-var* | **proc_bind** |
| *stacksize-var* | (none) |
| *wait-policy-var* | (none) |
| *thread-limit-var* | (none) |
| *max-active-levels-var* | (none) |
| *active-levels-var* | (none) |
| *levels-var* | (none) |

*table continued on next page*

| ICV | construct clause, if used |
|-----|---------------------------|
| *place-partition-var* | (none) |
| *cancel-var* | (none) |
| *default-device-var* | (none) |
| *max-task-priority-var* | (none) |

**Description**

- The **num_threads** clause overrides the value of the first element of the *nthreads-var* ICV.

- If *bind-var* is not set to *false* then the **proc_bind** clause overrides the value of the first element of the *bind-var* ICV; otherwise, the **proc_bind** clause has no effect.

**Cross References**

- **parallel** construct, see Section 2.5 on page 46.

- **proc_bind** clause, Section 2.5 on page 46.

- **num_threads** clause, see Section 2.5.1 on page 50.

- Loop construct, see Section 2.7.1 on page 56.

- **schedule** clause, see Section 2.7.1.1 on page 64.

# 2.4 Array Sections

An array section designates a subset of the elements in an array. An array section can appear only in clauses where it is explicitly allowed.

─────────────────── C / C++ ───────────────────

To specify an array section in an OpenMP construct, array subscript expressions are extended with the following syntax:

1         **[** *lower-bound* **:** *length* **]** or

2         **[** *lower-bound* **:** **]** or

3         **[** **:** *length* **]** or

4         **[ : ]**

5 The array section must be a subset of the original array.

6 Array sections are allowed on multidimensional arrays. Base language array subscript expressions
7 can be used to specify length-one dimensions of multidimensional array sections.

8 The *lower-bound* and *length* are integral type expressions. When evaluated they represent a set of
9 integer values as follows:

10 { *lower-bound*, *lower-bound* + 1, *lower-bound* + 2,... , *lower-bound* + *length* - 1 }

11 The *length* must evaluate to a non-negative integer.

12 When the size of the array dimension is not known, the *length* must be specified explicitly.

13 When the *length* is absent, it defaults to the size of the array dimension minus the *lower-bound*.

14 When the *lower-bound* is absent it defaults to 0.

15 Note – The following are examples of array sections:

16     `a[0:6]`

17     `a[:6]`

18     `a[1:10]`

19     `a[1:]`

20     `b[10][:][:0]`

21     `c[1:10][42][0:6]`

22 The first two examples are equivalent. If **a** is declared to be an eleven element array, the third and
23 fourth examples are equivalent. The fifth example is a zero-length array section. The last example
24 is not contiguous.

C / C++

1 Fortran has built-in support for array sections but the following restrictions apply for OpenMP
2 constructs:

3 • A stride expression may not be specified.

4 • The upper bound for the last dimension of an assumed-size dummy array must be specified.

5 **Restrictions**

6 Restrictions to array sections are as follows:

7 • An array section can appear only in clauses where it is explicitly allowed.

8 • An array section can only be specified for a base language identifier.

9 • The type of the variable appearing in an array section must be array or pointer.

10 • If the type of the variable appearing in an array section is a reference to a type *T* then the type
11 will be considered to be *T* for all purposes of the array section.

12 • An array section cannot be used in a C++ user-defined **[]**-operator.

13 # 2.5 `parallel` Construct

14 **Summary**

15 This fundamental construct starts parallel execution. See Section 1.3 on page 14 for a general
16 description of the OpenMP execution model.

1      **Syntax**

─────────────────────── C / C++ ───────────────────────

2      The syntax of the **parallel** construct is as follows:

```
#pragma omp parallel [clause[ [,] clause] ... ] new-line
    structured-block
```

3      where *clause* is one of the following:

4            **if(**[**parallel** **:**] *scalar-expression***)**

5            **num_threads(***integer-expression***)**

6            **default(shared** | **none)**

7            **private(***list***)**

8            **firstprivate(***list***)**

9            **shared(***list***)**

10           **copyin(***list***)**

11           **reduction(***reduction-identifier* **:** *list***)**

12           **proc_bind(master** | **close** | **spread)**

─────────────────────── C / C++ ───────────────────────
─────────────────────── Fortran ───────────────────────

13     The syntax of the **parallel** construct is as follows:

```
!$omp parallel [clause[ [,] clause] ... ]
    structured-block
!$omp end parallel
```

1     where *clause* is one of the following:

2         **if(**[**parallel** :] *scalar-logical-expression***)**

3         **num_threads(***scalar-integer-expression***)**

4         **default(private | firstprivate | shared | none)**

5         **private(***list***)**

6         **firstprivate(***list***)**

7         **shared(***list***)**

8         **copyin(***list***)**

9         **reduction(***reduction-identifier* : *list***)**

10         **proc_bind(master | close | spread)**

11     The **end parallel** directive denotes the end of the **parallel** construct.

— Fortran —

## Binding

13 The binding thread set for a **parallel** region is the encountering thread. The encountering thread
14 becomes the master thread of the new team.

## Description

16 When a thread encounters a **parallel** construct, a team of threads is created to execute the
17 **parallel** region (see Section 2.5.1 on page 50 for more information about how the number of
18 threads in the team is determined, including the evaluation of the **if** and **num_threads** clauses).
19 The thread that encountered the **parallel** construct becomes the master thread of the new team,
20 with a thread number of zero for the duration of the new **parallel** region. All threads in the new
21 team, including the master thread, execute the region. Once the team is created, the number of
22 threads in the team remains constant for the duration of that **parallel** region.

23 The optional **proc_bind** clause, described in Section 2.5.2 on page 52, specifies the mapping of
24 OpenMP threads to places within the current place partition, that is, within the places listed in the
25 *place-partition-var* ICV for the implicit task of the encountering thread.

26 Within a **parallel** region, thread numbers uniquely identify each thread. Thread numbers are
27 consecutive whole numbers ranging from zero for the master thread up to one less than the number
28 of threads in the team. A thread may obtain its own thread number by a call to the
29 **omp_get_thread_num** library routine.

30 A set of implicit tasks, equal in number to the number of threads in the team, is generated by the
31 encountering thread. The structured block of the **parallel** construct determines the code that

will be executed in each implicit task. Each task is assigned to a different thread in the team and becomes tied. The task region of the task being executed by the encountering thread is suspended and each thread in the team executes its implicit task. Each thread can execute a path of statements that is different from that of the other threads

The implementation may cause any thread to suspend execution of its implicit task at a task scheduling point, and switch to execute any explicit task generated by any of the threads in the team, before eventually resuming execution of the implicit task (for more details see Section 2.9 on page 83).

There is an implied barrier at the end of a **parallel** region. After the end of a **parallel** region, only the master thread of the team resumes execution of the enclosing task region.

If a thread in a team executing a **parallel** region encounters another **parallel** directive, it creates a new team, according to the rules in Section 2.5.1 on page 50, and it becomes the master of that new team.

If execution of a thread terminates while inside a **parallel** region, execution of all threads in all teams terminates. The order of termination of threads is unspecified. All work done by a team prior to any barrier that the team has passed in the program is guaranteed to be complete. The amount of work done by each thread after the last barrier that it passed and before it terminates is unspecified.

### Restrictions

Restrictions to the **parallel** construct are as follows:

- A program that branches into or out of a **parallel** region is non-conforming.

- A program must not depend on any ordering of the evaluations of the clauses of the **parallel** directive, or on any side effects of the evaluations of the clauses.

- At most one **if** clause can appear on the directive.

- At most one **proc_bind** clause can appear on the directive.

- At most one **num_threads** clause can appear on the directive. The **num_threads** expression must evaluate to a positive integer value.

———————————————— C / C++ ————————————————

A **throw** executed inside a **parallel** region must cause execution to resume within the same **parallel** region, and the same thread that threw the exception must catch it.

———————————————— C / C++ ————————————————
———————————————— Fortran ————————————————

Unsynchronized use of Fortran I/O statements by multiple threads on the same unit has unspecified behavior.

———————————————— Fortran ————————————————

**Cross References**

- **if** clause, see Section 2.12 on page 147.

- **default**, **shared**, **private**, **firstprivate**, and **reduction** clauses, see
  Section 2.15.3 on page 188.

- **copyin** clause, see Section 2.15.4 on page 211.

- **omp_get_thread_num** routine, see Section 3.2.4 on page 235.

## 2.5.1 Determining the Number of Threads for **a parallel** Region

When execution encounters a **parallel** directive, the value of the **if** clause or **num_threads**
clause (if any) on the directive, the current parallel context, and the values of the *nthreads-var*,
*dyn-var*, *thread-limit-var*, *max-active-levels-var*, and *nest-var* ICVs are used to determine the
number of threads to use in the region.

Using a variable in an **if** or **num_threads** clause expression of a **parallel** construct causes
an implicit reference to the variable in all enclosing constructs. The **if** clause expression and the
**num_threads** clause expression are evaluated in the context outside of the **parallel**
construct, and no ordering of those evaluations is specified. It is also unspecified whether, in what
order, or how many times any side effects of the evaluation of the **num_threads** or **if** clause
expressions occur.

When a thread encounters a **parallel** construct, the number of threads is determined according
to Algorithm 2.1.

---

## Algorithm 2.1

---

> **let** *ThreadsBusy* be the number of OpenMP threads currently executing in this
> contention group;
>
> **let** *ActiveParRegions* be the number of enclosing active parallel regions;
>
> **if** an **if** clause exists
>
> **then let** *IfClauseValue* be the value of the **if** clause expression;
>
> **else let** *IfClauseValue* = *true*;
>
> **if** a **num_threads** clause exists
>
> **then let** *ThreadsRequested* be the value of the **num_threads** clause expression;

1    **else let** *ThreadsRequested* = value of the first element of *nthreads-var*;

2    **let** *ThreadsAvailable* = (*thread-limit-var* - *ThreadsBusy* + 1);

3    **if** (*IfClauseValue* = *false*)

4    **then** number of threads = 1;

5    **else if** (*ActiveParRegions* >= 1) **and** (*nest-var* = *false*)

6    **then** number of threads = 1;

7    **else if** (*ActiveParRegions* = *max-active-levels-var*)

8    **then** number of threads = 1;

9    **else if** (*dyn-var* = *true*) **and** (*ThreadsRequested* <= *ThreadsAvailable*)

10   **then** number of threads = [ 1 : *ThreadsRequested* ];

11   **else if** (*dyn-var* = *true*) **and** (*ThreadsRequested* > *ThreadsAvailable*)

12   **then** number of threads = [ 1 : *ThreadsAvailable* ];

13   **else if** (*dyn-var* = *false*) **and** (*ThreadsRequested* <= *ThreadsAvailable*)

14   **then** number of threads = *ThreadsRequested*;

15   **else if** (*dyn-var* = *false*) **and** (*ThreadsRequested* > *ThreadsAvailable*)

16   **then** behavior is implementation defined;

---

Note – Since the initial value of the *dyn-var* ICV is implementation defined, programs that depend on a specific number of threads for correct execution should explicitly disable dynamic adjustment of the number of threads.

**Cross References**

- *nthreads-var*, *dyn-var*, *thread-limit-var*, *max-active-levels-var*, and *nest-var* ICVs, see Section 2.3 on page 36.

# 2.5.2 Controlling OpenMP Thread Affinity

When a thread encounters a **parallel** directive without a **proc_bind** clause, the *bind-var* ICV is used to determine the policy for assigning OpenMP threads to places within the current place partition, that is, the places listed in the *place-partition-var* ICV for the implicit task of the encountering thread. If the **parallel** directive has a **proc_bind** clause then the binding policy specified by the **proc_bind** clause overrides the policy specified by the first element of the *bind-var* ICV. Once a thread in the team is assigned to a place, the OpenMP implementation should not move it to another place.

The **master** thread affinity policy instructs the execution environment to assign every thread in the team to the same place as the master thread. The place partition is not changed by this policy, and each implicit task inherits the *place-partition-var* ICV of the parent implicit task.

The **close** thread affinity policy instructs the execution environment to assign the threads in the team to places close to the place of the parent thread. The place partition is not changed by this policy, and each implicit task inherits the *place-partition-var* ICV of the parent implicit task. If $T$ is the number of threads in the team, and $P$ is the number of places in the parent's place partition, then the assignment of threads in the team to places is as follows:

- $T \leq P$. The master thread executes on the place of the parent thread. The thread with the next smallest thread number executes on the next place in the place partition, and so on, with wrap around with respect to the place partition of the master thread.

- $T > P$. Each place $P$ will contain $S_p$ threads with consecutive thread numbers, where $\lfloor T/P \rfloor \leq Sp \leq \lceil T/P \rceil$. The first $S_0$ threads (including the master thread) are assigned to the place of the parent thread. The next $S_1$ threads are assigned to the next place in the place partition, and so on, with wrap around with respect to the place partition of the master thread. When $P$ does not divide $T$ evenly, the exact number of threads in a particular place is implementation defined.

The purpose of the **spread** thread affinity policy is to create a sparse distribution for a team of $T$ threads among the $P$ places of the parent's place partition. A sparse distribution is achieved by first subdividing the parent partition into $T$ subpartitions if $T \leq P$, or $P$ subpartitions if $T > P$. Then one thread ($T \leq P$) or a set of threads ($T > P$) is assigned to each subpartition. The *place-partition-var* ICV of each implicit task is set to its subpartition. The subpartitioning is not only a mechanism for achieving a sparse distribution, it also defines a subset of places for a thread to use when creating a nested **parallel** region. The assignment of threads to places is as follows:

- $T \leq P$. The parent thread's place partition is split into $T$ subpartitions, where each subpartition contains $\lfloor P/T \rfloor$ or $\lceil P/T \rceil$ consecutive places. A single thread is assigned to each subpartition. The master thread executes on the place of the parent thread and is assigned to the subpartition that includes that place. The thread with the next smallest thread number is assigned to the first place in the next subpartition, and so on, with wrap around with respect to the original place partition of the master thread.

- $T > P$. The parent thread's place partition is split into $P$ subpartitions, each consisting of a single place. Each subpartition is assigned $S_p$ threads with consecutive thread numbers, where $\lfloor T/P \rfloor \leq S_p \leq \lceil T/P \rceil$. The first $S_0$ threads (including the master thread) are assigned to the subpartition containing the place of the parent thread. The next $S_1$ threads are assigned to the next subpartition, and so on, with wrap around with respect to the original place partition of the master thread. When P does not divide $T$ evenly, the exact number of threads in a particular subpartition is implementation defined.

The determination of whether the affinity request can be fulfilled is implementation defined. If the affinity request cannot be fulfilled, then the affinity of threads in the team is implementation defined.

▼ ▼

Note - Wrap around is needed if the end of a place partition is reached before all thread assignments are done. For example, wrap around may be needed in the case of **close** and $T \leq P$, if the master thread is assigned to a place other than the first place in the place partition. In this case, thread 1 is assigned to the place after the place of the master place, thread 2 is assigned to the place after that, and so on. The end of the place partition may be reached before all threads are assigned. In this case, assignment of threads is resumed with the first place in the place partition.

▲ ▲

## 2.6 Canonical Loop Form

—————————————— C / C++ ——————————————

A loop has *canonical loop form* if it conforms to the following:

| **for** (*init-expr; test-expr; incr-expr*) *structured-block* | |
|---|---|
| *init-expr* | One of the following: <br> *var = lb* <br> *integer-type var = lb* <br> *random-access-iterator-type var = lb* <br> *pointer-type var = lb* |
| *test-expr* | One of the following: <br> *var relational-op b* <br> *b relational-op var* |

*continued on next page*

*continued from previous page*

| *incr-expr* | One of the following: |
| --- | --- |
| | ++*var* |
| | *var*++ |
| | - - *var* |
| | *var* - - |
| | *var* += *incr* |
| | *var* - = *incr* |
| | *var* = *var* + *incr* |
| | *var* = *incr* + *var* |
| | *var* = *var* - *incr* |

1

| *var* | One of the following: |
| --- | --- |
| | A variable of a signed or unsigned integer type. |
| | For C++, a variable of a random access iterator type. |
| | For C, a variable of a pointer type. |
| | If this variable would otherwise be shared, it is implicitly made private in the loop construct. This variable must not be modified during the execution of the *for-loop* other than in *incr-expr*. Unless the variable is specified **lastprivate** or **linear** on the loop construct, its value after the loop is unspecified. |
| *relational-op* | One of the following: |
| | **<** |
| | **<=** |
| | **>** |
| | **>=** |
| *lb* and *b* | Loop invariant expressions of a type compatible with the type of *var*. |
| *incr* | A loop invariant integer expression. |

2

3
4
5

The canonical form allows the iteration count of all associated loops to be computed before executing the outermost loop. The computation is performed for each loop in an integer type. This type is derived from the type of *var* as follows:

6

- If *var* is of an integer type, then the type is the type of *var*.

7
8

- For C++, if *var* is of a random access iterator type, then the type is the type that would be used by *std::distance* applied to variables of the type of *var*.

9

- For C, if *var* is of a pointer type, then the type is **ptrdiff_t**.

10

The behavior is unspecified if any intermediate result required to compute the iteration count

cannot be represented in the type determined above.

There is no implied synchronization during the evaluation of the *lb*, *b*, or *incr* expressions. It is unspecified whether, in what order, or how many times any side effects within the *lb*, *b*, or *incr* expressions occur.

▼                                                                                              ▼

Note – Random access iterators are required to support random access to elements in constant time. Other iterators are precluded by the restrictions since they can take linear time or offer limited functionality. It is therefore advisable to use tasks to parallelize those cases.

▲                                                                                              ▲


**Restrictions**

The following restrictions also apply:

- If *test-expr* is of the form *var relational-op b* and *relational-op* is < or <= then *incr-expr* must cause *var* to increase on each iteration of the loop. If *test-expr* is of the form *var relational-op b* and *relational-op* is > or >= then *incr-expr* must cause *var* to decrease on each iteration of the loop.

- If *test-expr* is of the form *b relational-op var* and *relational-op* is < or <= then *incr-expr* must cause *var* to decrease on each iteration of the loop. If *test-expr* is of the form *b relational-op var* and *relational-op* is > or >= then *incr-expr* must cause *var* to increase on each iteration of the loop.

- For C++, in the **simd** construct the only random access iterator types that are allowed for *var* are pointer types.

- The *b*, *lb* and *incr* expressions may not reference *var* of any of the associated loops.

▲ ────────────────────── C / C++ ────────────────────── ▲

# 2.7 Worksharing Constructs

A worksharing construct distributes the execution of the associated region among the members of the team that encounters it. Threads execute portions of the region in the context of the implicit tasks each one is executing. If the team consists of only one thread then the worksharing region is not executed in parallel.

A worksharing region has no barrier on entry; however, an implied barrier exists at the end of the worksharing region, unless a **nowait** clause is specified. If a **nowait** clause is present, an implementation may omit the barrier at the end of the worksharing region. In this case, threads that finish early may proceed straight to the instructions following the worksharing region without waiting for the other members of the team to finish the worksharing region, and without performing a flush operation.

The OpenMP API defines the following worksharing constructs, and these are described in the sections that follow:

- loop construct
- **sections** construct
- **single** construct
- **workshare** construct

## Restrictions

The following restrictions apply to worksharing constructs:

- Each worksharing region must be encountered by all threads in a team or by none at all, unless cancellation has been requested for the innermost enclosing parallel region.
- The sequence of worksharing regions and **barrier** regions encountered must be the same for every thread in a team

## 2.7.1 Loop Construct

### Summary

The loop construct specifies that the iterations of one or more associated loops will be executed in parallel by threads in the team in the context of their implicit tasks. The iterations are distributed across threads that already exist in the team executing the **parallel** region to which the loop region binds.

1 **Syntax**

--------------------------------- C / C++ ---------------------------------

2 The syntax of the loop construct is as follows:

```
#pragma omp for [clause[ [,] clause] ... ] new-line
    for-loops
```

3 where clause is one of the following:

4     **private(***list***)**

5     **firstprivate(***list***)**

6     **lastprivate(***list***)**

7     **linear(***list[ : linear-step]***)**

8     **reduction(***reduction-identifier : list***)**

9     **schedule(***[modifier [, modifier]:]kind[, chunk_size]***)**

10     **collapse(***n***)**

11     **ordered***[(***n***)]*

12     **nowait**

13 The **for** directive places restrictions on the structure of all associated *for-loops*. Specifically, all
14 associated *for-loops* must have *canonical loop form* (see Section 2.6 on page 53).

--------------------------------- C / C++ ---------------------------------
--------------------------------- Fortran ---------------------------------

15 The syntax of the loop construct is as follows:

```
!$omp do [clause[ [,] clause] ... ]
    do-loops
[!$omp end do [nowait]]
```

16 where *clause* is one of the following:

| | |
|---|---|
| 1 | **private(**_list_**)** |
| 2 | **firstprivate(**_list_**)** |
| 3 | **lastprivate(**_list_**)** |
| 4 | **linear(**_list[ : linear-step]_**)** |
| 5 | **reduction(**_reduction-identifier : list_**)** |
| 6 | **schedule(**_[modifier [, modifier]:]kind[, chunk_size]_**)** |
| 7 | **collapse(**_n_**)** |
| 8 | **ordered**_[(**n**)]_ |

9      If an **end do** directive is not specified, an **end do** directive is assumed at the end of the *do-loops*.

10      Any associated *do-loop* must be a *do-construct* or an *inner-shared-do-construct* as defined by the
11     Fortran standard. If an **end do** directive follows a *do-construct* in which several loop statements
12     share a **DO** termination statement, then the directive can only be specified for the outermost of these
13     **DO** statements.

14      If any of the loop iteration variables would otherwise be shared, they are implicitly made private on
15     the loop construct.

       ▲ ——————————————————————— Fortran ——————————————————————— ▲

16      **Binding**

17      The binding thread set for a loop region is the current team. A loop region binds to the innermost
18     enclosing **parallel** region. Only the threads of the team executing the binding **parallel**
19     region participate in the execution of the loop iterations and the implied barrier of the loop region if
20     the barrier is not eliminated by a **nowait** clause.

21      **Description**

22      The loop construct is associated with a loop nest consisting of one or more loops that follow the
23     directive.

24      There is an implicit barrier at the end of a loop construct unless a **nowait** clause is specified.

25      The **collapse** clause may be used to specify how many loops are associated with the loop
26     construct. The parameter of the **collapse** clause must be a constant positive integer expression.
27     If a **collapse** clause is specified with a parameter value greater than 1, then the iterations of the
28     associated loops to which the clause applies are collapsed into one larger iteration space that is then
29     divided according to the **schedule** clause. The sequential execution of the iterations in these
30     associated loops determines the order of the iterations in the collapsed iteration space. If no
31     **collapse** clause is present or its parameter is 1, the only loop that is associated with the loop

construct for the purposes of determining how the iteration space is divided according to the **schedule** clause is the one that immediately follows the loop directive.

The iteration count for each associated loop is computed before entry to the outermost loop. If execution of any associated loop changes any of the values used to compute any of the iteration counts, then the behavior is unspecified.

The integer type (or kind, for Fortran) used to compute the iteration count for the collapsed loop is implementation defined.

A worksharing loop has logical iterations numbered 0,1,...,N-1 where N is the number of loop iterations, and the logical numbering denotes the sequence in which the iterations would be executed if a set of associated loop(s) were executed sequentially. The **schedule** clause specifies how iterations of these associated loops are divided into contiguous non-empty subsets, called chunks, and how these chunks are distributed among threads of the team. Each thread executes its assigned chunk(s) in the context of its implicit task. The iterations of a given chunk are executed in sequential order by the assigned thread. The *chunk_size* expression is evaluated using the original list items of any variables that are made private in the loop construct. It is unspecified whether, in what order, or how many times, any side effects of the evaluation of this expression occur. The use of a variable in a **schedule** clause expression of a loop construct causes an implicit reference to the variable in all enclosing constructs.

Different loop regions with the same schedule and iteration count, even if they occur in the same parallel region, can distribute iterations among threads differently. The only exception is for the **static** schedule as specified in Table 2.5. Programs that depend on which thread executes a particular iteration under any other circumstances are non-conforming.

See Section 2.7.1.1 on page 64 for details of how the schedule for a worksharing loop is determined.

The schedule *kind* can be one of those specified in Table 2.5.

The schedule *modifier* can be one of those specified in Table 2.6. If the **static** schedule kind is specified or if the **ordered** clause is specified, and if no **monotonic** modifier is specified, the effect will be as if the **monotonic** modifier was specified.

◆───────────────────────────────────────────────────────────◆

Note – The next release of the OpenMP specification will include the following statement:

Otherwise, unless the **monotonic** modifier is specified, the effect will be as if the **nonmonotonic** modifier was specified.

◆───────────────────────────────────────────────────────────◆

The **ordered** clause with the parameter may also be used to specify how many loops are associated with the loop construct. The parameter of the **ordered** clause must be a constant positive integer expression if specified. The parameter of the **ordered** clause does not affect how the logical iteration space is then divided. If an **ordered** clause with the parameter is specified for the loop construct, then those associated loops form a *doacross loop nest*.

1　　If the value of the parameter in the **collapse** or **ordered** clause is larger than the number of
2　　nested loops following the construct, the behavior is unspecified.

3　　**TABLE 2.5: schedule** Clause *kind* Values

| **static** | When **schedule(static,** *chunk_size***)** is specified, iterations are divided into chunks of size *chunk_size*, and the chunks are assigned to the threads in the team in a round-robin fashion in the order of the thread number. |
| | When no *chunk_size* is specified, the iteration space is divided into chunks that are approximately equal in size, and at most one chunk is distributed to each thread. The size of the chunks is unspecified in this case. |
| | A compliant implementation of the **static** schedule must ensure that the same assignment of logical iteration numbers to threads will be used in two loop regions if the following conditions are satisfied: 1) both loop regions have the same number of loop iterations, 2) both loop regions have the same value of *chunk_size* specified, or both loop regions have no *chunk_size* specified, 3) both loop regions bind to the same parallel region, and 4) neither loop is associated with a SIMD construct. A data dependence between the same logical iterations in two such loops is guaranteed to be satisfied allowing safe use of the **nowait** clause. |
| **dynamic** | When **schedule(dynamic,** *chunk_size***)** is specified, the iterations are distributed to threads in the team in chunks. Each thread executes a chunk of iterations, then requests another chunk, until no chunks remain to be distributed. |
| | Each chunk contains *chunk_size* iterations, except for the chunk that contains the sequentially last iteration, which may have fewer iterations. |
| | When no *chunk_size* is specified, it defaults to 1. |
| **guided** | When **schedule(guided,** *chunk_size***)** is specified, the iterations are assigned to threads in the team in chunks. Each thread executes a chunk of iterations, then requests another chunk, until no chunks remain to be assigned. |

*table continued on next page*

4

For a *chunk_size* of 1, the size of each chunk is proportional to the number of unassigned iterations divided by the number of threads in the team, decreasing to 1. For a *chunk_size* with value $k$ (greater than 1), the size of each chunk is determined in the same way, with the restriction that the chunks do not contain fewer than $k$ iterations (except for the chunk that contains the sequentially last iteration, which may have fewer than $k$ iterations).

When no *chunk_size* is specified, it defaults to 1.

**auto**    When **schedule(auto)** is specified, the decision regarding scheduling is delegated to the compiler and/or runtime system. The programmer gives the implementation the freedom to choose any possible mapping of iterations to threads in the team.

**runtime**    When **schedule(runtime)** is specified, the decision regarding scheduling is deferred until run time, and the schedule and chunk size are taken from the *run-sched-var* ICV. If the ICV is set to **auto**, the schedule is implementation defined.

Note – For a team of $p$ threads and a loop of $n$ iterations, let $\lceil n/p \rceil$ be the integer $q$ that satisfies $n = p * q - r$, with $0 <= r < p$. One compliant implementation of the **static** schedule (with no specified *chunk_size*) would behave as though *chunk_size* had been specified with value $q$. Another compliant implementation would assign $q$ iterations to the first $p - r$ threads, and $q - 1$ iterations to the remaining $r$ threads. This illustrates why a conforming program must not rely on the details of a particular implementation.

A compliant implementation of the **guided** schedule with a *chunk_size* value of $k$ would assign $q = \lceil n/p \rceil$ iterations to the first available thread and set $n$ to the larger of $n - q$ and $p * k$. It would then repeat this process until $q$ is greater than or equal to the number of remaining iterations, at which time the remaining iterations form the final chunk. Another compliant implementation could use the same method, except with $q = \lceil n/(2p) \rceil$, and set $n$ to the larger of $n - q$ and $2 * p * k$.

**TABLE 2.6: schedule** Clause *modifier* Values

| | |
|---|---|
| **monotonic** | When the **monotonic** modifier is specified then each thread executes the chunks that it is assigned in increasing logical iteration order. |
| **nonmonotonic** | When the **nonmonotonic** modifier is specified then chunks are assigned to threads in any order and the behavior of an application that depends on any execution order of the chunks is unspecified. |
| **simd** | When the **simd** modifier is specified and the loop is associated with a SIMD construct, the *chunk_size* for all chunks except the first and last chunks is $new\_chunk\_size = \lceil chunk\_size/simd\_width \rceil * simd\_width$ where *simd_width* is an implementation-defined value. The first chunk will have at least *new_chunk_size* iterations except if it is also the last chunk. The last chunk may have fewer iterations than *new_chunk_size*. If the **simd** modifier is specified and the loop is not associated with a SIMD construct, the modifier is ignored. |

### Restrictions

Restrictions to the loop construct are as follows:

- All loops associated with the loop construct must be perfectly nested; that is, there must be no intervening code nor any OpenMP directive between any two loops.

- The values of the loop control expressions of the loops associated with the loop construct must be the same for all threads in the team.

- Only one **schedule** clause can appear on a loop directive.

- Only one **collapse** clause can appear on a loop directive.

- *chunk_size* must be a loop invariant integer expression with a positive value.

- The value of the *chunk_size* expression must be the same for all threads in the team.

- The value of the *run-sched-var* ICV must be the same for all threads in the team.

- When **schedule(runtime)** or **schedule(auto)** is specified, *chunk_size* must not be specified.

- A *modifier* may not be specified on a **linear** clause.

- Only one **ordered** clause can appear on a loop directive.

- The **ordered** clause must be present on the loop construct if any **ordered** region ever binds to a loop region arising from the loop construct.

- The **nonmonotonic** modifier can only be specified with **schedule(dynamic)** or **schedule(guided)**.

- The **nonmonotonic** modifier cannot be specified if an **ordered** clause is specified.

- Either the **monotonic** modifier or the **nonmonotonic** modifier can be specified but not both.

- The loop iteration variable may not appear in a **threadprivate** directive.

- If both the **collapse** and **ordered** clause with a parameter are specified, the parameter of the **ordered** clause must be greater than or equal to the parameter of the **collapse** clause.

- A **linear** clause or an **ordered** clause with a parameter can be specified on a loop directive but not both.

$\rule{center}{C / C++}$

- The associated *for-loops* must be structured blocks.

- Only an iteration of the innermost associated loop may be curtailed by a **continue** statement.

- No statement can branch to any associated **for** statement.

- Only one **nowait** clause can appear on a **for** directive.

- A throw executed inside a loop region must cause execution to resume within the same iteration of the loop region, and the same thread that threw the exception must catch it.

$\rule{center}{C / C++}$

$\rule{center}{Fortran}$

- The associated *do-loops* must be structured blocks.

- Only an iteration of the innermost associated loop may be curtailed by a **CYCLE** statement.

- No statement in the associated loops other than the **DO** statements can cause a branch out of the loops.

- The *do-loop* iteration variable must be of type integer.

- The *do-loop* cannot be a **DO WHILE** or a **DO** loop without loop control.

$\rule{center}{Fortran}$

**Cross References**

- **private**, **firstprivate**, **lastprivate**, **linear**, and **reduction** clauses, see Section 2.15.3 on page 188.

- **OMP_SCHEDULE** environment variable, see Section 4.1 on page 292.

- **ordered** construct, see Section 2.13.8 on page 166.

- **depend** clause, see Section 2.13.9 on page 169.

### 2.7.1.1 Determining the Schedule of a Worksharing Loop

When execution encounters a loop directive, the **schedule** clause (if any) on the directive, and the *run-sched-var* and *def-sched-var* ICVs are used to determine how loop iterations are assigned to threads. See Section 2.3 on page 36 for details of how the values of the ICVs are determined. If the loop directive does not have a **schedule** clause then the current value of the *def-sched-var* ICV determines the schedule. If the loop directive has a **schedule** clause that specifies the **runtime** schedule kind then the current value of the *run-sched-var* ICV determines the schedule. Otherwise, the value of the **schedule** clause determines the schedule. Figure 2.1 describes how the schedule for a worksharing loop is determined.

**Cross References**

- ICVs, see Section 2.3 on page 36
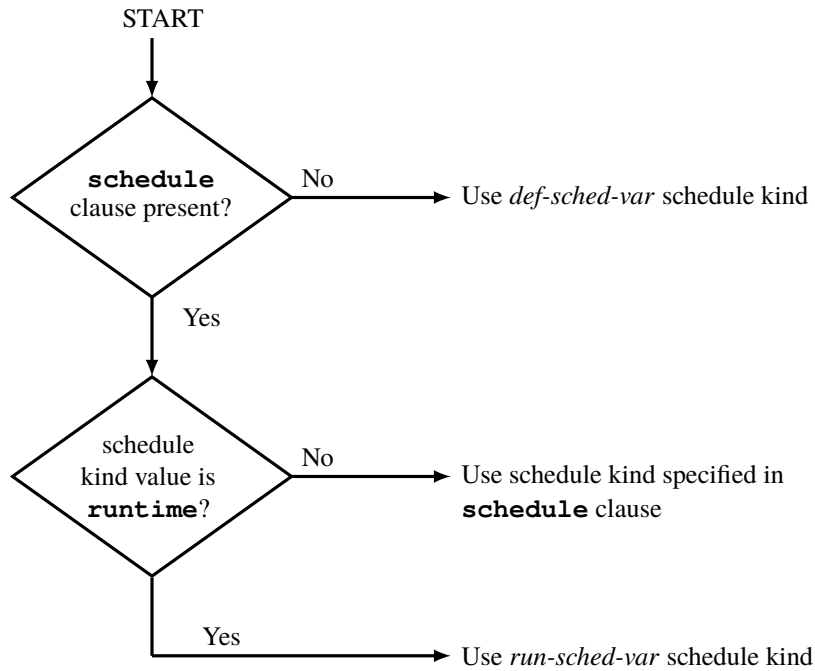


**FIGURE 2.1:** Determining the **schedule** for a Worksharing Loop

# 1 2.7.2 `sections` Construct

2 **Summary**

3 The `sections` construct is a non-iterative worksharing construct that contains a set of structured
4 blocks that are to be distributed among and executed by the threads in a team. Each structured
5 block is executed once by one of the threads in the team in the context of its implicit task.

6 **Syntax**

--- C / C++ ---

7 The syntax of the `sections` construct is as follows:

```
#pragma omp sections [clause[ [,] clause] ... ] new-line
    {
    [#pragma omp section new-line]
        structured-block
    [#pragma omp section new-line
        structured-block]
    ...
    }
```

8 where *clause* is one of the following:

9     `private(`*list*`)`

10     `firstprivate(`*list*`)`

11     `lastprivate(`*list*`)`

12     `reduction(`*reduction-identifier* `:` *list*`)`

13     `nowait`

--- C / C++ ---

1    The syntax of the **sections** construct is as follows:

```
!$omp sections [clause[ [,] clause] ... ]
   [!$omp section]
      structured-block
   [!$omp section
      structured-block]
   ...
!$omp end sections [nowait]
```

2    where *clause* is one of the following:

3        **private(***list***)**

4        **firstprivate(***list***)**

5        **lastprivate(***list***)**

6        **reduction(***reduction-identifier* **:** *list***)**

7    **Binding**

8    The binding thread set for a **sections** region is the current team. A **sections** region binds to
9    the innermost enclosing **parallel** region. Only the threads of the team executing the binding
10   **parallel** region participate in the execution of the structured blocks and the implied barrier of
11   the **sections** region if the barrier is not eliminated by a **nowait** clause.

12   **Description**

13   Each structured block in the **sections** construct is preceded by a **section** directive except
14   possibly the first block, for which a preceding **section** directive is optional.

15   The method of scheduling the structured blocks among the threads in the team is implementation
16   defined.

17   There is an implicit barrier at the end of a **sections** construct unless a **nowait** clause is
18   specified.

**Restrictions**

Restrictions to the **sections** construct are as follows:

- Orphaned **section** directives are prohibited. That is, the **section** directives must appear within the **sections** construct and must not be encountered elsewhere in the **sections** region.

- The code enclosed in a **sections** construct must be a structured block.

- Only a single **nowait** clause can appear on a **sections** directive.

—————————————————————  C++  —————————————————————

- A throw executed inside a **sections** region must cause execution to resume within the same section of the **sections** region, and the same thread that threw the exception must catch it.

—————————————————————  C++  —————————————————————

**Cross References**

- **private**, **firstprivate**, **lastprivate**, and **reduction** clauses, see Section 2.15.3 on page 188.

# 2.7.3  **single** Construct

**Summary**

The **single** construct specifies that the associated structured block is executed by only one of the threads in the team (not necessarily the master thread), in the context of its implicit task. The other threads in the team, which do not execute the block, wait at an implicit barrier at the end of the **single** construct unless a **nowait** clause is specified.

**Syntax**

—————————————————————  C / C++  —————————————————————

The syntax of the single construct is as follows:

```
#pragma omp single [clause[ [,] clause] ... ] new-line
    structured-block
```

1     where *clause* is one of the following:

2        **private(***list***)**

3        **firstprivate(***list***)**

4        **copyprivate(***list***)**

5        **nowait**

───────────────────────────── C / C++ ─────────────────────────────

───────────────────────────── Fortran ─────────────────────────────

6     The syntax of the **single** construct is as follows:

```
!$omp single [clause[ [,] clause] ... ]
   structured-block
!$omp end single [end_clause[ [,] end_clause] ... ]
```

7     where *clause* is one of the following:

8        **private(***list***)**

9        **firstprivate(***list***)**

10    and *end_clause* is one of the following:

11       **copyprivate(***list***)**

12       **nowait**

───────────────────────────── Fortran ─────────────────────────────

## Binding

14    The binding thread set for a **single** region is the current team. A **single** region binds to the
15    innermost enclosing **parallel** region. Only the threads of the team executing the binding
16    **parallel** region participate in the execution of the structured block and the implied barrier of the
17    **single** region if the barrier is not eliminated by a **nowait** clause.

## Description

19    The method of choosing a thread to execute the structured block is implementation defined. There
20    is an implicit barrier at the end of the **single** construct unless a **nowait** clause is specified.

**Restrictions**

Restrictions to the **single** construct are as follows:

- The **copyprivate** clause must not be used with the **nowait** clause.

- At most one **nowait** clause can appear on a **single** construct.

———————————————————————— C++ ————————————————————————

- A throw executed inside a **single** region must cause execution to resume within the same **single** region, and the same thread that threw the exception must catch it.

———————————————————————— C++ ————————————————————————

**Cross References**

- **private** and **firstprivate** clauses, see Section 2.15.3 on page 188.
- **copyprivate** clause, see Section 2.15.4.2 on page 213.

———————————————————————— Fortran ————————————————————————

# 2.7.4 **workshare** Construct

**Summary**

The **workshare** construct divides the execution of the enclosed structured block into separate units of work, and causes the threads of the team to share the work such that each unit is executed only once by one thread, in the context of its implicit task.

**Syntax**

The syntax of the **workshare** construct is as follows:

```
!$omp workshare
    structured-block
!$omp end workshare [nowait]
```

The enclosed structured block must consist of only the following:

- array assignments
- scalar assignments
- **FORALL** statements

1    • **FORALL** constructs

2    • **WHERE** statements

3    • **WHERE** constructs

4    • **atomic** constructs

5    • **critical** constructs

6    • **parallel** constructs

7    Statements contained in any enclosed **critical** construct are also subject to these restrictions.
8    Statements in any enclosed **parallel** construct are not restricted.

9    **Binding**

10    The binding thread set for a **workshare** region is the current team. A **workshare** region binds
11    to the innermost enclosing **parallel** region. Only the threads of the team executing the binding
12    **parallel** region participate in the execution of the units of work and the implied barrier of the
13    **workshare** region if the barrier is not eliminated by a **nowait** clause.

14    **Description**

15    There is an implicit barrier at the end of a **workshare** construct unless a **nowait** clause is
16    specified.

17    An implementation of the **workshare** construct must insert any synchronization that is required
18    to maintain standard Fortran semantics. For example, the effects of one statement within the
19    structured block must appear to occur before the execution of succeeding statements, and the
20    evaluation of the right hand side of an assignment must appear to complete prior to the effects of
21    assigning to the left hand side.

22    The statements in the **workshare** construct are divided into units of work as follows:

23    • For array expressions within each statement, including transformational array intrinsic functions
24      that compute scalar values from arrays:

25      – Evaluation of each element of the array expression, including any references to **ELEMENTAL**
26        functions, is a unit of work.

27      – Evaluation of transformational array intrinsic functions may be freely subdivided into any
28        number of units of work.

29    • For an array assignment statement, the assignment of each element is a unit of work.

30    • For a scalar assignment statement, the assignment operation is a unit of work.

- For a **WHERE** statement or construct, the evaluation of the mask expression and the masked assignments are each a unit of work.

- For a **FORALL** statement or construct, the evaluation of the mask expression, expressions occurring in the specification of the iteration space, and the masked assignments are each a unit of work

- For an **atomic** construct, the atomic operation on the storage location designated as $x$ is a unit of work.

- For a **critical** construct, the construct is a single unit of work.

- For a **parallel** construct, the construct is a unit of work with respect to the **workshare** construct. The statements contained in the **parallel** construct are executed by a new thread team.

- If none of the rules above apply to a portion of a statement in the structured block, then that portion is a unit of work.

The transformational array intrinsic functions are **MATMUL**, **DOT_PRODUCT**, **SUM**, **PRODUCT**, **MAXVAL**, **MINVAL**, **COUNT**, **ANY**, **ALL**, **SPREAD**, **PACK**, **UNPACK**, **RESHAPE**, **TRANSPOSE**, **EOSHIFT**, **CSHIFT**, **MINLOC**, and **MAXLOC**.

It is unspecified how the units of work are assigned to the threads executing a **workshare** region.

If an array expression in the block references the value, association status, or allocation status of private variables, the value of the expression is undefined, unless the same value would be computed by every thread.

If an array assignment, a scalar assignment, a masked array assignment, or a **FORALL** assignment assigns to a private variable in the block, the result is unspecified.

The **workshare** directive causes the sharing of work to occur only in the **workshare** construct, and not in the remainder of the **workshare** region.

**Restrictions**

The following restrictions apply to the **workshare** construct:

- All array assignments, scalar assignments, and masked array assignments must be intrinsic assignments.

- The construct must not contain any user defined function calls unless the function is **ELEMENTAL**.

Fortran

# 2.8 SIMD Constructs

## 2.8.1 `simd` Construct

**Summary**

The **simd** construct can be applied to a loop to indicate that the loop can be transformed into a SIMD loop (that is, multiple iterations of the loop can be executed concurrently using SIMD instructions).

**Syntax**

The syntax of the **simd** construct is as follows:

─────────────────── C / C++ ───────────────────

```
#pragma omp simd [clause[ [,] clause] ... ] new-line
    for-loops
```

where *clause* is one of the following:

**safelen**(*length*)

**simdlen**(*length*)

**linear**(*list[ : linear-step]*)

**aligned**(*list[ : alignment]*)

**private**(*list*)

**lastprivate**(*list*)

**reduction**(*reduction-identifier : list*)

**collapse**(*n*)

The **simd** directive places restrictions on the structure of the associated *for-loops*. Specifically, all associated *for-loops* must have *canonical loop form* (Section 2.6 on page 53).

─────────────────── C / C++ ───────────────────

```
!$omp simd [clause[ [,] clause ... ]
    do-loops
[!$omp end simd]
```

1    where *clause* is one of the following:

2        **safelen(***length***)**

3        **simdlen(***length***)**

4        **linear(***list[ : linear-step]***)**

5        **aligned(***list[ : alignment]***)**

6        **private(***list***)**

7        **lastprivate(***list***)**

8        **reduction(***reduction-identifier : list***)**

9        **collapse(***n***)**

10   If an **end simd** directive is not specified, an **end simd** directive is assumed at the end of the
11   *do-loops*.

12   Any associated *do-loop* must be a *do-construct* or an *inner-shared-do-construct* as defined by the
13   Fortran standard. If an **end simd** directive follows a *do-construct* in which several loop statements
14   share a **DO** termination statement, then the directive can only be specified for the outermost of these
15   **DO** statements.

16   **Binding**

17   A **simd** region binds to the current task region. The binding thread set of the **simd** region is the
18   current team.

## Description

The **simd** construct enables the execution of multiple iterations of the associated loops concurrently by means of SIMD instructions.

The **collapse** clause may be used to specify how many loops are associated with the construct. The parameter of the **collapse** clause must be a constant positive integer expression. If no **collapse** clause is present, the only loop that is associated with the loop construct is the one that immediately follows the directive.

If more than one loop is associated with the **simd** construct, then the iterations of all associated loops are collapsed into one larger iteration space that is then executed with SIMD instructions. The sequential execution of the iterations in all associated loops determines the order of the iterations in the collapsed iteration space.

The iteration count for each associated loop is computed before entry to the outermost loop. If execution of any associated loop changes any of the values used to compute any of the iteration counts, then the behavior is unspecified.

The integer type (or kind, for Fortran) used to compute the iteration count for the collapsed loop is implementation defined.

A SIMD loop has logical iterations numbered 0,1,...,N-1 where N is the number of loop iterations, and the logical numbering denotes the sequence in which the iterations would be executed if the associated loop(s) were executed with no SIMD instructions. If the **safelen** clause is used then no two iterations executed concurrently with SIMD instructions can have a greater distance in the logical iteration space than its value. The parameter of the **safelen** clause must be a constant positive integer expression. If used, the **simdlen** clause specifies the preferred number of iterations to be executed concurrently. The parameter of the **simdlen** clause must be a constant positive integer. The number of iterations that are executed concurrently at any given time is implementation defined. Each concurrent iteration will be executed by a different SIMD lane. Each set of concurrent iterations is a SIMD chunk. Lexical forward dependencies in the iterations of the original loop must be preserved within each SIMD chunk.

---------------------------------- C / C++ ----------------------------------

The **aligned** clause declares that the object to which each list item points is aligned to the number of bytes expressed in the optional parameter of the **aligned** clause

---------------------------------- C / C++ ----------------------------------
---------------------------------- Fortran ----------------------------------

The **aligned** clause declares that the location of each list item is aligned to the number of bytes expressed in the optional parameter of the **aligned** clause.

---------------------------------- Fortran ----------------------------------

The optional parameter of the **aligned** clause, *alignment*, must be a constant positive integer expression. If no optional parameter is specified, implementation-defined default alignments for SIMD instructions on the target platforms are assumed.

**Restrictions**

- All loops associated with the construct must be perfectly nested; that is, there must be no intervening code nor any OpenMP directive between any two loops.

- The associated loops must be structured blocks.

- A program that branches into or out of a **simd** region is non-conforming.

- Only one **collapse** clause can appear on a **simd** directive.

- A *list-item* cannot appear in more than one **aligned** clause.

- Only one **safelen** clause can appear on a **simd** directive.

- Only one **simdlen** clause can appear on a **simd** directive.

- If both **simdlen** and **safelen** clauses are specified, the value of the **simdlen** parameter must be less than or equal to the value of the **safelen** parameter.

- A *modifier* may not be specified on a **linear** clause.

- An **ordered** construct with the **simd** clause is the only OpenMP construct that can be encountered during execution of a **simd** region.

---------------------------------- C / C++ ----------------------------------

- The **simd** region cannot contain calls to the **longjmp** or **setjmp** functions.

---------------------------------- C / C++ ----------------------------------

------------------------------------- C -------------------------------------

- The type of list items appearing in the **aligned** clause must be array or pointer.

------------------------------------- C -------------------------------------

------------------------------------ C++ ------------------------------------

- The type of list items appearing in the **aligned** clause must be array, pointer, reference to array, or reference to pointer.

- No exception can be raised in the **simd** region.

------------------------------------ C++ ------------------------------------

1     • The *do-loop* iteration variable must be of type **integer**.

2     • The *do-loop* cannot be a **DO WHILE** or a **DO** loop without loop control.

3     • If a list item on the **aligned** clause has the **ALLOCATABLE** attribute, the allocation status must
4     be allocated.

5     • If a list item on the **aligned** clause has the **POINTER** attribute, the association status must be
6     associated.

7     • If the type of a list item on the **aligned** clause is either **C_PTR** or Cray pointer, the list item
8     must be defined.

•

———————— Fortran ————————

9   **Cross References**

10     • **private**, **lastprivate**, **linear** and **reduction** clauses, see Section 2.15.3 on page 188.

## 2.8.2  **declare simd** Construct

12   **Summary**

13   The **declare simd** construct can be applied to a function (C, C++ and Fortran) or a subroutine
14   (Fortran) to enable the creation of one or more versions that can process multiple arguments using
15   SIMD instructions from a single invocation in a SIMD loop. The **declare simd** directive is a
16   declarative directive. There may be multiple **declare simd** directives for a function (C, C++,
17   Fortran) or subroutine (Fortran).

18   **Syntax**

19   The syntax of the **declare simd** construct is as follows:

```
#pragma omp declare simd [clause[ [,] clause] ... ] new-line
[#pragma omp declare simd [clause[ [,] clause] ... ] new-line]
[ ... ]
    function definition or declaration
```

1    where *clause* is one of the following:

2        **simdlen(***length***)**

3        **linear(***linear-list[* : *linear-step]***)**

4        **aligned(***argument-list[* : *alignment]***)**

5        **uniform(***argument-list***)**

6        **inbranch**

7        **notinbranch**

```
!$omp declare simd [ (proc-name) ] [clause[ [,] clause] ... ]
```

8    where *clause* is one of the following:

9        **simdlen(***length***)**

10       **linear(***linear-list[* : *linear-step]***)**

11       **aligned(***argument-list[* : *alignment]***)**

12       **uniform(***argument-list***)**

13       **inbranch**

14       **notinbranch**

**Description**

─────────────────────────── C / C++ ───────────────────────────

The use of a **declare simd** construct on a function enables the creation of SIMD versions of the associated function that can be used to process multiple arguments from a single invocation in a SIMD loop concurrently.

The expressions appearing in the clauses of this directive are evaluated in the scope of the arguments of the function declaration or definition.

─────────────────────────── C / C++ ───────────────────────────
─────────────────────────── Fortran ───────────────────────────

The use of a **declare simd** construct enables the creation of SIMD versions of the specified subroutine or function that can be used to process multiple arguments from a single invocation in a SIMD loop concurrently.

─────────────────────────── Fortran ───────────────────────────

If a **declare simd** directive contains multiple SIMD declarations, each declaration enables the creation of SIMD versions.

If a SIMD version is created, the number of concurrent arguments for the function is determined by the **simdlen** clause. If the **simdlen** clause is used its value corresponds to the number of concurrent arguments of the function. The parameter of the **simdlen** clause must be a constant positive integer expression. Otherwise, the number of concurrent arguments for the function is implementation defined.

─────────────────────────── C++ ───────────────────────────

The special *this* pointer can be used as if was one of the arguments to the function in any of the **linear**, **aligned**, or **uniform** clauses.

─────────────────────────── C++ ───────────────────────────

The **uniform** clause declares one or more arguments to have an invariant value for all concurrent invocations of the function in the execution of a single SIMD loop.

─────────────────────────── C / C++ ───────────────────────────

The **aligned** clause declares that the object to which each list item points is aligned to the number of bytes expressed in the optional parameter of the **aligned** clause.

─────────────────────────── C / C++ ───────────────────────────

1 The **aligned** clause declares that the target of each list item is aligned to the number of bytes
2 expressed in the optional parameter of the **aligned** clause.

3 The optional parameter of the **aligned** clause, *alignment*, must be a constant positive integer
4 expression. If no optional parameter is specified, implementation-defined default alignments for
5 SIMD instructions on the target platforms are assumed.

6 The **inbranch** clause specifies that the SIMD version of the function will always be called from
7 inside a conditional statement of a SIMD loop. The **notinbranch** clause specifies that the SIMD
8 version of the function will never be called from inside a conditional statement of a SIMD loop. If
9 neither clause is specified, then the SIMD version of the function may or may not be called from
10 inside a conditional statement of a SIMD loop.

11 **Restrictions**

12 • Each argument can appear in at most one **uniform** or **linear** clause.

13 • At most one **simdlen** clause can appear in a **declare simd** directive.

14 • Either **inbranch** or **notinbranch** may be specified, but not both.

15 • When a *linear-step* expression is specified in a **linear** clause it must be either a constant integer
16 expression or an integer-typed parameter that is specified in a **uniform** clause on the directive.

17 • The function or subroutine body must be a structured block.

18 • The execution of the function or subroutine, when called from a SIMD loop, cannot result in the
19 execution of an OpenMP construct except for an **ordered** construct with the **simd** clause.

20 • The execution of the function or subroutine cannot have any side effects that would alter its
21 execution for concurrent iterations of a SIMD chunk.

22 • A program that branches into or out of the function is non-conforming.

23 • If the function has any declarations, then the **declare simd** construct for any declaration that
24 has one must be equivalent to the one specified for the definition. Otherwise, the result is
25 unspecified.

26 • The function cannot contain calls to the **longjmp** or **setjmp** functions.

1    • The type of list items appearing in the **aligned** clause must be array or pointer.

2    • The function cannot contain any calls to **throw**.

3    • The type of list items appearing in the **aligned** clause must be array, pointer, reference to
4      array, or reference to pointer.

5    • *proc-name* must not be a generic name, procedure pointer or entry name.

6    • If *proc-name* is omitted, the **declare simd** directive must appear in the specification part of a
7      subroutine subprogram or a function subprogram for which creation of the SIMD versions is
8      enabled.

9    • Any **declare simd** directive must appear in the specification part of a subroutine subprogram,
10     function subprogram or interface body to which it applies.

11   • If a **declare simd** directive is specified in an interface block for a procedure, it must match a
12     **declare simd** directive in the definition of the procedure.

13   • If a procedure is declared via a procedure declaration statement, the procedure *proc-name* should
14     appear in the same specification.

15   • If a **declare simd** directive is specified for a procedure name with explicit interface and a
16     **declare simd** directive is also specified for the definition of the procedure then the two
17     **declare simd** directives must match. Otherwise the result is unspecified.

18   • Procedure pointers may not be used to access versions created by the **declare simd** directive.

19   • The type of list items appearing in the **aligned** clause must be **C_PTR** or Cray pointer, or the
20     list item must have the **POINTER** or **ALLOCATABLE** attribute.

**Cross References**

- **reduction** clause, see Section 2.15.3.6 on page 201.

- **linear** clause, see Section 2.15.3.7 on page 207.

## 2.8.3   Loop SIMD Construct

**Summary**

The loop SIMD construct specifies that the iterations of one or more associated loops will be
distributed across threads that already exist in the team and that the iterations executed by each
thread can also be executed concurrently using SIMD instructions. The loop SIMD construct is a
composite construct.

**Syntax**

─────────────────── C / C++ ───────────────────

```
#pragma omp for simd [clause[ [,] clause] ... ] new-line
    for-loops
```

where *clause* can be any of the clauses accepted by the **for** or **simd** directives with identical
meanings and restrictions.

─────────────────── C / C++ ───────────────────
─────────────────── Fortran ───────────────────

```
!$omp do simd [clause[ [,] clause] ... ]
    do-loops
[!$omp end do simd [nowait] ]
```

where *clause* can be any of the clauses accepted by the **simd** or **do** directives, with identical
meanings and restrictions.

If an **end do simd** directive is not specified, an **end do simd** directive is assumed at the end of
the *do-loops*.

─────────────────── Fortran ───────────────────

**Description**

The loop SIMD construct will first distribute the iterations of the associated loop(s) across the implicit tasks of the parallel region in a manner consistent with any clauses that apply to the loop construct. The resulting chunks of iterations will then be converted to a SIMD loop in a manner consistent with any clauses that apply to the **simd** construct. The effect of any clause that applies to both constructs is as if it were applied to both constructs separately except the **collapse** clause, which is applied once.

**Restrictions**

All restrictions to the loop construct and the **simd** construct apply to the loop SIMD construct. In addition, the following restrictions apply:

- No **ordered** clause with a parameter can be specified.

- A list item may appear in a **linear** or **firstprivate** clause but not both.

**Cross References**

- loop construct, see Section 2.7.1 on page 56.

- **simd** construct, see Section 2.8.1 on page 72.

- Data attribute clauses, see Section 2.15.3 on page 188.

## 1 2.9  Tasking Constructs

## 2 2.9.1  `task` Construct

3 **Summary**

4 The **task** construct defines an explicit task.

5 **Syntax**

<div align="center">C / C++</div>

6 The syntax of the **task** construct is as follows:

```
#pragma omp task [clause[ [,] clause] ... ] new-line
    structured-block
```

7 where *clause* is one of the following:

8        **if(**[ **task** : ] *scalar-expression***)**

9        **final(***scalar-expression***)**

10        **untied**

11        **default(shared** | **none)**

12        **mergeable**

13        **private(***list***)**

14        **firstprivate(***list***)**

15        **shared(***list***)**

16        **depend(***dependence-type* : *list***)**

17        **priority(***priority-value***)**

<div align="center">C / C++</div>

1    The syntax of the **task** construct is as follows:

```
!$omp task [clause[ [,] clause] ... ]
    structured-block
!$omp end task
```

2    where *clause* is one of the following:

3    **if**(*[* **task** *:] scalar-logical-expression*)

4    **final**(*scalar-logical-expression*)

5    **untied**

6    **default**(**private** | **firstprivate** | **shared** | **none**)

7    **mergeable**

8    **private**(*list*)

9    **firstprivate**(*list*)

10   **shared**(*list*)

11   **depend**(*dependence-type* : *list*)

12   **priority**(*priority-value*)

13   **Binding**

14   The binding thread set of the **task** region is the current team. A **task** region binds to the
15   innermost enclosing **parallel** region.

## Description

When a thread encounters a **task** construct, a task is generated from the code for the associated structured block. The data environment of the task is created according to the data-sharing attribute clauses on the **task** construct, per-data environment ICVs, and any defaults that apply.

The encountering thread may immediately execute the task, or defer its execution. In the latter case, any thread in the team may be assigned the task. Completion of the task can be guaranteed using task synchronization constructs. If a **task** construct is encountered during execution of an outer task, the generated **task** region associated with this construct is not a part of the outer task region unless the generated task is an included task.

When an **if** clause is present on a **task** construct, and the **if** clause expression evaluates to *false*, an undeferred task is generated, and the encountering thread must suspend the current task region, for which execution cannot be resumed until the generated task is completed. The use of a variable in an **if** clause expression of a **task** construct causes an implicit reference to the variable in all enclosing constructs.

When a **final** clause is present on a **task** construct and the **final** clause expression evaluates to *true*, the generated task will be a final task. All **task** constructs encountered during execution of a final task will generate final and included tasks. Note that the use of a variable in a **final** clause expression of a **task** construct causes an implicit reference to the variable in all enclosing constructs.

The **if** clause expression and the **final** clause expression are evaluated in the context outside of the **task** construct, and no ordering of those evaluations is specified.

A thread that encounters a task scheduling point within the **task** region may temporarily suspend the **task** region. By default, a task is tied and its suspended **task** region can only be resumed by the thread that started its execution. If the **untied** clause is present on a **task** construct, any thread in the team can resume the **task** region after a suspension. The **untied** clause is ignored if a **final** clause is present on the same **task** construct and the **final** clause expression evaluates to *true*, or if a task is an included task.

The **task** construct includes a task scheduling point in the task region of its generating task, immediately following the generation of the explicit task. Each explicit **task** region includes a task scheduling point at its point of completion.

When the **mergeable** clause is present on a **task** construct, the generated task is a *mergeable task*.

The **priority** clause is a hint for the priority of the generated task. The *priority-value* is a non-negative numerical scalar expression that provides a hint for task execution order. Among all tasks ready to be executed, higher priority tasks (those with a higher numerical value in the **priority** clause expression) are recommended to execute before lower priority ones. The default *priority-value* when no **priority** clause is specified is zero (the lowest priority). If a value is specified in the **priority** clause that is higher than the *max-task-priority-var* ICV then the

implementation will use the value of that ICV. A program that relies on task execution order being
determined by this *priority-value* may have unspecified behavior.

Note – When storage is shared by an explicit **task** region, the programmer must ensure, by adding
proper synchronization, that the storage does not reach the end of its lifetime before the explicit
**task** region completes its execution.

## Restrictions

Restrictions to the **task** construct are as follows:

- A program that branches into or out of a **task** region is non-conforming.
- A program must not depend on any ordering of the evaluations of the clauses of the **task**
  directive, or on any side effects of the evaluations of the clauses.
- At most one **if** clause can appear on the directive.
- At most one **final** clause can appear on the directive.
- At most one **priority** clause can appear on the directive.

——————————————————— C / C++ ———————————————————

- A throw executed inside a **task** region must cause execution to resume within the same **task**
  region, and the same thread that threw the exception must catch it.

——————————————————— C / C++ ———————————————————
——————————————————— Fortran ———————————————————

- Unsynchronized use of Fortran I/O statements by multiple tasks on the same unit has unspecified
  behavior

——————————————————— Fortran ———————————————————

## Cross References

- Task scheduling constraints, see Section 2.9.5 on page 94.
- **depend** clause, see Section 2.13.9 on page 169.
- **if** Clause, see Section 2.12 on page 147.

# 1  2.9.2  `taskloop` **Construct**

2  **Summary**

3  The **taskloop** construct specifies that the iterations of one or more associated loops will be
4  executed in parallel using OpenMP tasks. The iterations are distributed across tasks created by the
5  construct and scheduled to be executed.

6  **Syntax**

───────────────────────  C / C++  ───────────────────────

7  The syntax of the **taskloop** construct is as follows:

```
#pragma omp taskloop [clause[[,] clause] ...] new-line
    for-loops
```

8  where *clause* is one of the following:

9  **if(**[ **taskloop** :] *scalar-expr***)**

10  **shared(***list***)**

11  **private(***list***)**

12  **firstprivate(***list***)**

13  **lastprivate(***list***)**

14  **default(shared | none)**

15  **grainsize(***grain-size***)**

16  **num_tasks(***num-tasks***)**

17  **collapse(***n***)**

18  **final(***scalar-expr***)**

19  **priority(***priority-value***)**

20  **untied**

21  **mergeable**

22  **nogroup**

23  The **taskloop** directive places restrictions on the structure of all associated *for-loops*.
24  Specifically, all associated *for-loops* must have canonical loop form (see Section 2.6 on page 53).

───────────────────────  C / C++  ───────────────────────

1    The syntax of the **taskloop** construct is as follows:

```
!$omp taskloop [clause[[,] clause] ...]
    do-loops
[!$omp end taskloop]
```

2    where *clause* is one of the following:

3        **if(**[ **taskloop** :] *scalar-logical-expr***)**

4        **shared(***list***)**

5        **private(***list***)**

6        **firstprivate(***list***)**

7        **lastprivate(***list***)**

8        **default(private | firstprivate | shared | none)**

9        **grainsize(***grain-size***)**

10       **num_tasks(***num-tasks***)**

11       **collapse(***n***)**

12       **final(***scalar-logical-expr***)**

13       **priority(***priority-value***)**

14       **untied**

15       **mergeable**

16       **nogroup**

17   If an **end taskloop** directive is not specified, an **end taskloop** directive is assumed at the end
18   of the *do-loops*.

19   Any associated *do-loop* must be *do-construct* or an *inner-shared-do-construct* as defined by the
20   Fortran standard. If an **end taskloop** directive follows a *do-construct* in which several loop
21   statements share a **DO** termination statement, then the directive can only be specified for the
22   outermost of these **DO** statements.

23   If any of the loop iteration variables would otherwise be shared, they are implicitly made private for
24   the loop-iteration tasks created by the **taskloop** construct. Unless the loop iteration variables are
25   specified in a **lastprivate** clause on the **taskloop** construct, their values after the loop are
26   unspecified.

## Binding

The binding thread set of the **taskloop** region is the current team. A **taskloop** region binds to the innermost enclosing **parallel** region.

## Description

When a thread encounters a **taskloop** construct, the construct partitions the associated loops into tasks for parallel execution of the loops' iterations. The data environment of the created tasks is created according to the data-sharing attribute clauses on the **taskloop** construct, per-data environment ICVs, and any defaults that apply. The order of the creation of the loop tasks is unspecified. Programs that rely on any execution order of the logical loop iterations are non-conforming.

If a **grainsize** clause is present on the **taskloop** construct, the number of logical loop iterations assigned to each created task is greater than or equal to the minimum of the value of the *grain-size* expression and the number of logical loop iterations, but less than two times the value of the *grain-size* expression. The parameter of the **grainsize** clause must be a positive integer expression. If **num_tasks** is specified, the **taskloop** construct creates as many tasks as the minimum of the *num-tasks* expression and the number of logical loop iterations. Each task must have at least one logical loop iteration. The parameter of the **num_tasks** clause must evaluate to a positive integer. If neither a **grainsize** nor **num_tasks** clause is present, the number of loop tasks created and the number of logical loop iterations assigned to these tasks is implementation defined.

The **collapse** clause may be used to specify how many loops are associated with the **taskloop** construct. The parameter of the **collapse** clause must be a constant positive integer expression. If no **collapse** clause is present, the only loop that is associated with the **taskloop** construct is the one that immediately follows the **taskloop** directive.

If more than one loop is associated with the **taskloop** construct, then the iterations of all associated loops are collapsed into one larger iteration space that is then divided according to the **grainsize** and **num_tasks** clauses. The sequential execution of the iterations in all associated loops determines the order of the iterations in the collapsed iteration space.

The iteration count for each associated loop is computed before entry to the outermost loop. If execution of any associated loop changes any of the values used to compute any of the iteration counts, then the behavior is unspecified.

The integer type (or kind, for Fortran) used to compute the iteration count for the collapsed loop is implementation defined.

When an **if** clause is present on a **taskloop** construct, and if the **if** clause expression evaluates to *false*, undeferred tasks are generated. The use of a variable in an **if** clause expression of a **taskloop** construct causes an implicit reference to the variable in all enclosing constructs.

When a **final** clause is present on a **taskloop** construct and the **final** clause expression evaluates to *true*, the generated tasks will be final tasks. The use of a variable in a **final** clause expression of a **taskloop** construct causes an implicit reference to the variable in all enclosing constructs.

When a **priority** clause is present on a **taskloop** construct, the generated tasks have the *priority-value* as if it was specified for each individual task. If the **priority** clause is not specified, tasks generated by the **taskloop** construct have the default task priority (zero).

If the **untied** clause is specified, all tasks created by the **taskloop** construct are untied tasks.

When the **mergeable** clause is present on a **taskloop** construct, each generated task is a *mergeable task*.

By default, the **taskloop** construct executes as if it was enclosed in a **taskgroup** construct with no statements or directives outside of the **taskloop** construct. Thus, the **taskloop** construct creates an implicit **taskgroup** region. If the **nogroup** clause is present, no implicit **taskgroup** region is created.

---
C++
---

For **firstprivate** variables of class type, the number of invocations of copy constructors to perform the initialization is implementation-defined.

---
C++
---

Note – When storage is shared by a **taskloop** region, the programmer must ensure, by adding proper synchronization, that the storage does not reach the end of its lifetime before the **taskloop** region and its descendant tasks complete their execution.

**Restrictions**

The restrictions of the **taskloop** construct are as follows:

- A program that branches into or out of a **taskloop** region is non-conforming.

- All loops associated with the **taskloop** construct must be perfectly nested; that is, there must be no intervening code nor any OpenMP directive between any two loops.

- At most one **grainsize** clause can appear on a **taskloop** directive.

- At most one **num_tasks** clause can appear on a **taskloop** directive.

- The **grainsize** clause and **num_tasks** clause are mutually exclusive and may not appear on the same **taskloop** directive.

- At most one **collapse** clause can appear on a **taskloop** directive.

- At most one **if** clause can appear on the directive.

- At most one **final** clause can appear on the directive.

- At most one **priority** clause can appear on the directive.

**Cross References**

- **task** construct, Section 2.9.1 on page 83.

- **taskgroup** construct, Section 2.13.5 on page 153.

- Data-sharing attribute clauses, Section 2.15.3 on page 188.

- **if** Clause, see Section 2.12 on page 147.


## 2.9.3 **taskloop simd** Construct

**Summary**

The **taskloop simd** construct specifies a loop that can be executed concurrently using SIMD instructions and that those iterations will also be executed in parallel using OpenMP tasks. The **taskloop simd** construct is a composite construct.

**Syntax**

$$\text{C / C++}$$

The syntax of the **taskloop simd** construct is as follows:

```
#pragma omp taskloop simd [clause[[,] clause] ...] new-line
    for-loops
```

where *clause* can be any of the clauses accepted by the **taskloop** or **simd** directives with identical meanings and restrictions.

$$\text{C / C++}$$

1    The syntax of the **taskloop simd** construct is as follows:

```
 !$omp taskloop simd [clause[[,] clause] ...]
     do-loops
[!$omp end taskloop simd]
```

2    where *clause* can be any of the clauses accepted by the **taskloop** or **simd** directives with
3    identical meanings and restrictions.

4    If an **end taskloop simd** directive is not specified, an **end taskloop simd** directive is
5    assumed at the end of the *do-loops*.

6    **Binding**

7    The binding thread set of the **taskloop simd** region is the current team. A **taskloop simd**
8    region binds to the innermost enclosing parallel region.

9    **Description**

10   The **taskloop simd** construct will first distribute the iterations of the associated loop(s) across
11   tasks in a manner consistent with any clauses that apply to the **taskloop** construct. The resulting
12   tasks will then be converted to a SIMD loop in a manner consistent with any clauses that apply to
13   the **simd** construct. The effect of any clause that applies to both constructs is as if it were applied
14   to both constructs separately except the **collapse** clause, which is applied once.

15   **Restrictions**

16   • The restrictions for the **taskloop** and **simd** constructs apply.

17   • No **reduction** clause can be specified.

18   **Cross References**

19   • **taskloop** construct, see Section 2.9.2 on page 87.

20   • **simd** construct, see Section 2.8.1 on page 72.

21   • Data-sharing attribute clauses, see Section 2.15.3 on page 188.

# 2.9.4 `taskyield` Construct

**Summary**

The **`taskyield`** construct specifies that the current task can be suspended in favor of execution of a different task. The **`taskyield`** construct is a stand-alone directive.

**Syntax**

────────────────────────── C / C++ ──────────────────────────

The syntax of the **`taskyield`** construct is as follows:

```
#pragma omp taskyield new-line
```

────────────────────────── C / C++ ──────────────────────────
────────────────────────── Fortran ──────────────────────────

The syntax of the **`taskyield`** construct is as follows:

```
!$omp taskyield
```

────────────────────────── Fortran ──────────────────────────

**Binding**

A **`taskyield`** region binds to the current task region. The binding thread set of the **`taskyield`** region is the current team.

**Description**

The **`taskyield`** region includes an explicit task scheduling point in the current task region.

**Cross References**

- Task scheduling, see Section 2.9.5 on page 94.

## 2.9.5 Task Scheduling

Whenever a thread reaches a task scheduling point, the implementation may cause it to perform a task switch, beginning or resuming execution of a different task bound to the current team. Task scheduling points are implied at the following locations:

- the point immediately following the generation of an explicit task;
- after the point of completion of a **task** region;
- in a **taskyield** region;
- in a **taskwait** region;
- at the end of a **taskgroup** region;
- in an implicit and explicit **barrier** region;
- the point immediately following the generation of a **target** region;
- at the beginning and end of a **target data** region;
- in a **target update** region;
- in a **target enter data** region;
- in a **target exit data** region;
- in the **omp_target_memcpy** routine;
- in the **omp_target_memcpy_rect** routine;

When a thread encounters a task scheduling point it may do one of the following, subject to the *Task Scheduling Constraints* (below):

- begin execution of a tied task bound to the current team
- resume any suspended task region, bound to the current team, to which it is tied
- begin execution of an untied task bound to the current team
- resume any suspended untied task region bound to the current team.

If more than one of the above choices is available, it is unspecified as to which will be chosen.

*Task Scheduling Constraints* are as follows:

1. An included task is executed immediately after generation of the task.

2. Scheduling of new tied tasks is constrained by the set of task regions that are currently tied to the thread, and that are not suspended in a **barrier** region. If this set is empty, any new tied task may be scheduled. Otherwise, a new tied task may be scheduled only if it is a descendent task of every task in the set.

3. A dependent task shall not be scheduled until its task dependences are fulfilled.

4. When an explicit task is generated by a construct containing an **if** clause for which the expression evaluated to *false*, and the previous constraints are already met, the task is executed immediately after generation of the task.

A program relying on any other assumption about task scheduling is non-conforming.

▼ ▼

Note – Task scheduling points dynamically divide task regions into parts. Each part is executed uninterrupted from start to end. Different parts of the same task region are executed in the order in which they are encountered. In the absence of task synchronization constructs, the order in which a thread executes parts of different schedulable tasks is unspecified.

A correct program must behave correctly and consistently with all conceivable scheduling sequences that are compatible with the rules above.

For example, if **threadprivate** storage is accessed (explicitly in the source code or implicitly in calls to library routines) in one part of a task region, its value cannot be assumed to be preserved into the next part of the same task region if another schedulable task exists that modifies it.

As another example, if a lock acquire and release happen in different parts of a task region, no attempt should be made to acquire the same lock in any part of another task that the executing thread may schedule. Otherwise, a deadlock is possible. A similar situation can occur when a **critical** region spans multiple parts of a task and another schedulable task contains a **critical** region with the same name.

The use of threadprivate variables and the use of locks or critical sections in an explicit task with an **if** clause must take into account that when the **if** clause evaluates to *false*, the task is executed immediately, without regard to *Task Scheduling Constraint* 2.

▲ ▲

# 2.10 Device Constructs

## 2.10.1 **target data** Construct

### Summary

Map variables to a device data environment for the extent of the region.

1   **Syntax**

2   The syntax of the **target data** construct is as follows:

```
#pragma omp target data clause[ [ [ , ] clause] ... ] new-line
    structured-block
```

3   where *clause* is one of the following:

4       **if(**[ **target data :**] *scalar-expression***)**

5       **device(***integer-expression***)**

6       **map(**[[*map-type-modifier[,]*] *map-type*: ] *list***)**

7       **use_device_ptr(***list***)**

8   The syntax of the **target data** construct is as follows:

```
!$omp target data clause[ [ [ , ] clause] ... ]
    structured-block
!$omp end target data
```

9   where *clause* is one of the following:

10      **if(**[ **target data :**] *scalar-logical-expression***)**

11      **device(***scalar-integer-expression***)**

12      **map(**[[*map-type-modifier[,]*] *map-type*: ] *list***)**

13      **use_device_ptr(***list***)**

14  The **end target data** directive denotes the end of the **target data** construct.

15  **Binding**

16  The binding task set for a **target data** region is the generating task. The **target data** region
17  binds to the region of the generating task.

**Description**

When a **target data** construct is encountered, the encountering task executes the region. If
there is no **device** clause, the default device is determined by the *default-device-var* ICV.
Variables are mapped for the extent of the region, according to any data-mapping clauses, from the
data environment of the encountering task to the device data environment. When an **if** clause is
present and the **if** clause expression evaluates to *false*, the device is the host.

List items that appear in a **use_device_ptr** clause are converted into device pointers to the
corresponding list item in the device data environment.

**Restrictions**

- A program must not depend on any ordering of the evaluations of the clauses of the
  **target data** directive, or on any side effects of the evaluations of the clauses.

- At most one **device** clause can appear on the directive. The **device** expression must evaluate
  to a non-negative integer value less than the value of **omp_get_num_devices()**.

- At most one **if** clause can appear on the directive.

- A *map-type* in a **map** clause must be **to**, **from**, **tofrom** or **alloc**.

- At least one **map** clause must appear on the directive.

- A list item in a **use_device_ptr** clause must have a corresponding list item in the device
  data environment.

- References in the construct to a list item that appears in a **use_device_ptr** clause must be to
  the address of the list item.

**Cross References**

- *default-device-var*, see Section 2.3 on page 36.

- **if** Clause, see Section 2.12 on page 147.

- **map** clause, see Section 2.15.5.1 on page 216.

## 2.10.2 **target enter data Construct**

**Summary**

The **target enter data** directive specifies that variables are mapped to a device data
environment. The **target enter data** directive is a stand-alone directive.

1    **Syntax**

2    The syntax of the **target enter data** construct is as follows:

```
#pragma omp target enter data [ clause[ [,] clause]...] new-line
```

3    where *clause* is one of the following:

4        **if(**[ **target enter data** :] *scalar-expression***)**

5        **device(***integer-expression***)**

6        **map(**[ [map-type-modifier[,]] map-type : ] list**)**

7        **depend(***dependence-type* : *list***)**

8        **nowait**

9    The syntax of the **target enter data** is as follows:

```
!$omp target enter data [ clause[ [,] clause]...]
```

10    where clause is one of the following:

11        **if(**[ **target enter data** :] *scalar-logical-expression***)**

12        **device(***scalar-integer-expression***)**

13        **map(**[ [map-type-modifier[,]] map-type : ] list**)**

14        **depend(***dependence-type* : *list***)**

15        **nowait**

16    **Binding**

17    The binding task set for a **target enter data** region is the generating task, which is the *target*
18    *task* generated by the **target enter data** construct. The **target enter data** region binds
19    to the corresponding *target task* region.

**Description**

When a `target enter data` construct is encountered, the list items are mapped to the device data environment according to the `map` clause semantics.

The `target enter data` construct is a task generating construct. The generated task is a *target task*. The generated task region encloses the `target enter data` region.

All clauses are evaluated when the `target enter data` construct is encountered. The data environment of the *target task* is created according to the data-sharing attribute clauses on the `target enter data` construct, per-data environment ICVs, and any default data-sharing attribute rules that apply to the `target enter data` construct. A variable that is mapped in the `target enter data` construct has a default data-sharing attribute of shared in the data environment of the *target task*.

Assignment operations associated with mapping a variable (see Section 2.15.5.1 on page 216) occur when the *target task* executes.

If the `nowait` clause is present, execution of the *target task* may be deferred. If the `nowait` clause is not present, the *target task* is an included task.

If a `depend` clause is present, it is associated with the *target task*.

If there is no `device` clause, the default device is determined by the *default-device-var* ICV.

When an `if` clause is present and the `if` clause expression evaluates to *false*, the device is the host.

**Restrictions**

- A program must not depend on any ordering of the evaluations of the clauses of the `target enter data` directive, or on any side effects of the evaluations of the clauses.

- At least one `map` clause must appear on the directive.

- At most one `device` clause can appear on the directive. The `device` expression must evaluate to a non-negative integer value.

- At most one `if` clause can appear on the directive.

- A *map-type* must be specified in all `map` clauses and must be either `to` or `alloc`.

**Cross References**

- *default-device-var*, see Section 2.3.1 on page 36.

- `task`, see Section 2.9.1 on page 83.

- `task scheduling constraints`, see Section 2.9.5 on page 94.

- `target data`, see Section 2.10.1 on page 95.

1　　　• **target exit data**, see Section 2.10.3 on page 100.

2　　　• **if** Clause, see Section 2.12 on page 147.

3　　　• **map** clause, see Section 2.15.5.1 on page 216.

4 ## 2.10.3 **target exit data** Construct

5 **Summary**

6 The **target exit data** directive specifies that list items are unmapped from a device data
7 environment. The **target exit data** directive is a stand-alone directive.

8 **Syntax**

$$\text{C / C++}$$

9 The syntax of the **target exit data** construct is as follows:

```
#pragma omp target exit data [ clause[ [,] clause]...] new-line
```

10 where *clause* is one of the following:

11　　**if(**[ **target exit data** :] *scalar-expression***)**

12　　**device(***integer-expression***)**

13　　**map(**[ [map-type-modifier[,]] map-type : ] list**)**

14　　**depend(***dependence-type* : *list***)**

15　　**nowait**

$$\text{C / C++}$$

1    The syntax of the **target exit data** is as follows:

```
!$omp target exit data [ clause[ [,] clause]...]
```

2    where clause is one of the following:

3        **if(**[ **target exit data** :] *scalar-logical-expression***)**

4        **device(***scalar-integer-expression***)**

5        **map(**[ *[map-type-modifier[,]] map-type* : *] list***)**

6        **depend(***dependence-type* : *list***)**

7        **nowait**

8    **Binding**

9    The binding task set for a **target exit data** region is the generating task, which is the *target*
10   *task* generated by the **target exit data** construct. The **target exit data** region binds to
11   the corresponding *target task* region.

12   **Description**

13   When a **target exit data** construct is encountered, the list items in the **map** clauses are
14   unmapped from the device data environment according to the **map** clause semantics.

15   The **target exit data** construct is a task generating construct. The generated task is a *target*
16   *task*. The generated task region encloses the **target exit data** region.

17   All clauses are evaluated when the **target exit data** construct is encountered. The data
18   environment of the *target task* is created according to the data-sharing attribute clauses on the
19   **target exit data** construct, per-data environment ICVs, and any default data-sharing attribute
20   rules that apply to the **target exit data** construct. A variable that is mapped in the
21   **target exit data** construct has a default data-sharing attribute of shared in the data
22   environment of the *target task*.

23   Assignment operations associated with mapping a variable (see Section 2.15.5.1 on page 216)
24   occur when the *target task* executes.

25   If the **nowait** clause is present, execution of the *target task* may be deferred. If the **nowait**
26   clause is not present, the *target task* is an included task.

27   If a **depend** clause is present, it is associated with the *target task*.

1    If there is no **device** clause, the default device is determined by the *default-device-var* ICV.

2    When an **if** clause is present and the **if** clause expression evaluates to *false*, the device is the host.

3    **Restrictions**

4    • A program must not depend on any ordering of the evaluations of the clauses of the
5    **target exit data** directive, or on any side effects of the evaluations of the clauses.

6    • At least one **map** clause must appear on the directive.

7    • At most one **device** clause can appear on the directive. The **device** expression must evaluate
8    to a non-negative integer value.

9    • At most one **if** clause can appear on the directive.

10   • A *map-type* must be specified in all **map** clauses and must be either **from**, **release**, or
11   **delete**.

12   **Cross References**

13   • *default-device-var*, see Section 2.3.1 on page 36.

14   • **task**, see Section 2.9.1 on page 83.

15   • **task scheduling constraints**, see Section 2.9.5 on page 94.

16   • **target data**, see Section 2.10.1 on page 95.

17   • **target enter data**, see Section 2.10.2 on page 97.

18   • **if** Clause, see Section 2.12 on page 147.

19   • **map** clause, see Section 2.15.5.1 on page 216.

# 1  2.10.4  `target` Construct

2  **Summary**

3  Map variables to a device data environment and execute the construct on that device.

4  **Syntax**

―――――――――――――― C / C++ ――――――――――――――

5  The syntax of the **target** construct is as follows:

```
#pragma omp target [clause[ [,] clause] ... ] new-line
    structured-block
```

6  where *clause* is one of the following:

7      **if(**[ **target** :] *scalar-expression***)**

8      **device(***integer-expression***)**

9      **private(***list***)**

10      **firstprivate(***list***)**

11      **map(**[[*map-type-modifier[,]*] *map-type*: ] *list***)**

12      **is_device_ptr(***list***)**

13      **defaultmap(tofrom:scalar)**

14      **nowait**

15      **depend(***dependence-type*: *list***)**

―――――――――――――― C / C++ ――――――――――――――

1    The syntax of the **target** construct is as follows:

```
!$omp target [clause[ [ , ] clause] ... ]
    structured-block
!$omp end target
```

2    where *clause* is one of the following:

3        **if(**[ **target** :] *scalar-logical-expression***)**

4        **device(***scalar-integer-expression***)**

5        **private(***list***)**

6        **firstprivate(***list***)**

7        **map(**[[*map-type-modifier[,]*] *map-type*: *] list***)**

8        **is_device_ptr(***list***)**

9        **defaultmap(tofrom:scalar)**

10       **nowait**

11       **depend (***dependence-type* : *list***)**

12   The **end target** directive denotes the end of the **target** construct

### Binding

14   The binding task set for a **target** region is the generating task, which is the *target task* generated
15   by the **target** construct. The **target** region binds to the corresponding *target task* region.

**Description**

The **target** construct provides a superset of the functionality provided by the **target data** directive, except for the **use_device_ptr** clause.

The functionality added to the **target** directive is the inclusion of an executable region to be executed by a device. That is, the **target** directive is an executable directive.

The **target** construct is a task generating construct. The generated task is a *target task*. The generated task region encloses the **target** region.

All clauses are evaluated when the **target** construct is encountered. The data environment of the *target task* is created according to the data-sharing attribute clauses on the **target** construct, per-data environment ICVs, and any default data-sharing attribute rules that apply to the **target** construct. A variable that is mapped in the **target** construct has a default data-sharing attribute of shared in the data environment of the *target task*.

Assignment operations associated with mapping a variable (see Section 2.15.5.1 on page 216) occur when the *target task* executes.

If the **nowait** clause is present, execution of the *target task* may be deferred. If the **nowait** clause is not present, the *target task* is an included task.

If a **depend** clause is present, it is associated with the *target task*.

When an **if** clause is present and the **if** clause expression evaluates to *false*, the **target** region is executed by the host device in the host data environment.

The **is_device_ptr** clause is used to indicate that a list item is a device pointer already in the device data environment and that it should be used directly. Support for device pointers created outside of OpenMP, specifically outside of the **omp_target_alloc** routine and the **use_device_ptr** clause, is implementation defined.

---

$\blacktriangledown$ ———————————— C / C++ ————————————— $\blacktriangledown$

If an array section is a list item in a **map** clause and the array section is derived from a variable for which the type is pointer then the data-sharing attribute for that variable in the construct is firstprivate. Prior to the execution of the construct, the private variable is initialized with the address of the storage location of the corresponding array section in the device data environment.

If a zero-length array section is a list item in a **map** clause, and the array section is derived from a variable for the which the type is pointer then that variable is initialized with the address of the corresponding storage location in the device data environment. If the corresponding storage location is not present in the device data environment then the private variable is initialized to NULL.

$\blacktriangle$ ———————————— C / C++ ————————————— $\blacktriangle$

**Restrictions**

- If a **target**, **target update**, **target data**, **target enter data**, or **target exit data** construct is encountered during execution of a **target** region, the behavior is unspecified.

- The result of an **omp_set_default_device**, **omp_get_default_device**, or **omp_get_num_devices** routine called within a **target** region is unspecified.

- The effect of an access to a **threadprivate** variable in a target region is unspecified.

- If a list item in a **map** clause is a structure element, any other element of that structure that is referenced in the **target** construct must also appear as a list item in a **map** clause.

- A variable referenced in a **target** region but not the **target** construct that is not declared in the **target** region must appear in a **declare target** directive.

- At most one **defaultmap** clause can appear on the directive.

- A *map-type* in a **map** clause must be **to**, **from**, **tofrom** or **alloc**.

- A list item that appears in an **is_device_ptr** clause must be a valid device pointer in the device data environment.

——————————— C ———————————

- A list item that appears in an **is_device_ptr** clause must have a type of pointer or array.

——————————— C ———————————

——————————— C++ ———————————

- A list item that appears in an **is_device_ptr** clause must have a type of pointer, array, reference to pointer or reference to array.

- A throw executed inside a **target** region must cause execution to resume within the same **target** region, and the same thread that threw the exception must catch it.

——————————— C++ ———————————

——————————— Fortran ———————————

- A list item that appears in an **is_device_ptr** clause must be a dummy argument.

- If a list item in a **map** clause is an array section, and the array section is derived from a variable with a **POINTER** or **ALLOCATABLE** attribute then the behavior is unspecified if the corresponding list item's variable is modified in the region.

——————————— Fortran ———————————

**Cross References**

- *default-device-var*, see Section 2.3 on page 36.

- **task** construct, see Section 2.9.1 on page 83.

- **task** scheduling constraints, see Section 2.9.5 on page 94

- **target data** construct, see Section 2.10.1 on page 95.

- **if** Clause, see Section 2.12 on page 147.

- **private** and **firstprivate** clauses, see Section 2.15.3 on page 188.

- Data-mapping Attribute Rules and Clauses, see Section 2.15.5 on page 215.

## 2.10.5 `target update` Construct

### Summary

The **target update** directive makes the corresponding list items in the device data environment consistent with their original list items, according to the specified motion clauses. The **target update** construct is a stand-alone directive.

### Syntax

$$\text{— C / C++ —}$$

The syntax of the **target update** construct is as follows:

```
#pragma omp target update clause[ [ [, ] clause] ... ] new-line
```

where *clause* is either *motion-clause* or one of the following:

    **if(**[ **target update** :] *scalar-expression***)**

    **device(***integer-expression***)**

    **nowait**

    **depend (***dependence-type* : *list***)**

and *motion-clause* is one of the following:

    **to(***list***)**

    **from(***list***)**

$$\text{— C / C++ —}$$

The syntax of the **target update** construct is as follows:

```
!$omp target update clause[ [ , ] clause] ... ]
```

where *clause* is either *motion-clause* or one of the following:

**if(**[**target update :**] *scalar-logical-expression***)**

**device(***scalar-integer-expression***)**

**nowait**

**depend (***dependence-type* **:** *list***)**

and *motion-clause* is one of the following:

**to(***list***)**

**from(***list***)**

**Binding**

The binding task set for a **target update** region is the generating task, which is the *target task* generated by the **target update** construct. The **target update** region binds to the corresponding *target task* region.

**Description**

For each list item in a **to** or **from** clause there is a corresponding list item and an original list item. If the corresponding list item is not present in the device data environment then no assignment occurs to or from the original list item. Otherwise, each corresponding list item in the device data environment has an original list item in the current task's data environment.

For each list item in a **from** clause the value of the corresponding list item is assigned to the original list item.

For each list item in a **to** clause the value of the original list item is assigned to the corresponding list item.

The list items that appear in the **to** or **from** clauses may include array sections.

The **target update** construct is a task generating construct. The generated task is a *target task*. The generated task region encloses the **target update** region.

All clauses are evaluated when the **target update** construct is encountered. The data environment of the *target task* is created according to the data-sharing attribute clauses on the **target update** construct, per-data environment ICVs, and any default data-sharing attribute rules that apply to the **target update** construct. A variable that is mapped in the **target update** construct has a default data-sharing attribute of shared in the data environment of the *target task*.

Assignment operations associated with mapping a variable (see Section 2.15.5.1 on page 216) occur when the *target task* executes.

If the **nowait** clause is present, execution of the *target task* may be deferred. If the **nowait** clause is not present, the *target task* is an included task.

If a **depend** clause is present, it is associated with the *target task*.

The device is specified in the **device** clause. If there is no **device** clause, the device is determined by the *default-device-var* ICV. When an **if** clause is present and the **if** clause expression evaluates to *false* then no assignments occur.

**Restrictions**

- A program must not depend on any ordering of the evaluations of the clauses of the **target update** directive, or on any side effects of the evaluations of the clauses.

- At least one *motion-clause* must be specified.

- If a list item is an array section it must specify contiguous storage.

- A list item can only appear in a **to** or **from** clause, but not both.

- A list item in a **to** or **from** clause must have a mappable type.

- At most one **device** clause can appear on the directive. The **device** expression must evaluate to a non-negative integer value less than the value of **omp_get_num_devices()**.

- At most one **if** clause can appear on the directive.

**Cross References**

## 2.10.6 `declare target` Directive

**Summary**

The `declare target` directive specifies that variables, functions (C, C++ and Fortran), and subroutines (Fortran) are mapped to a device. The `declare target` directive is a declarative directive.

**Syntax**

―――――――――――――――――――― C / C++ ――――――――――――――――――――

The syntax of the `declare target` directive takes either of the following forms:

```
#pragma omp declare target new-line
declaration-definition-seq
#pragma omp end declare target new-line
```

or

```
#pragma omp declare target (extended-list) new-line
```

or

```
#pragma omp declare target clause[ [,] clause ... ] new-line
```

where *clause* is one of the following:

`to(`*extended-list*`)`

`link(`*list*`)`

―――――――――――――――――――― C / C++ ――――――――――――――――――――

1    The syntax of the **declare target** directive is as follows:

```
!$omp declare target (extended-list)
```

2    or

```
!$omp declare target [clause[ [,] clause] ... ]
```

3    where *clause* is one of the following:

4        **to(***extended-list***)**

5        **link(***list***)**

6    **Description**

7    The **declare target** directive ensures that procedures and global variables can be executed or
8    accessed on a device. Variables are mapped for all device executions, or for specific device
9    executions through a **link** clause.

10   If an *extended-list* is present with no clause then the **to** clause is assumed.

11   If a list item of a **to** clause is a function (C, C++, Fortran) or subroutine (Fortran) then a
12   device-specific version of the routine is created that can be called from a **target** region.

13   If a list item of a **to** clause is a variable then the original variable is mapped to a corresponding
14   variable in the device data environment of all devices as if it had appeared in a **map** clause with the
15   *map-type* **to** on the implicit **target data** construct for each device. The list item is never
16   removed from those device data environments as if its reference count is initialized to positive
17   infinity.

18   The list items of a **link** clause are not mapped by the **declare target** directive. Instead, their
19   mapping is deferred until they are mapped by **target data** or **target** constructs. They are
20   mapped only for such regions.

1  The form of the **declare target** directive that has no clauses and requires a matching
2  **end declare target** directive defines an implicit *extended-list* to an implicit **to** clause. The
3  implicit *extended-list* consists of the variable names of any variable declarations at file or
4  namespace scope that appear between the two directives and of the function names of any function
5  declarations at file, namespace or class scope that appear between the two directives.

———————————————— C / C++ ————————————————
———————————————— Fortran ————————————————

6  If a **declare target** does not have any clauses then an implicit *extended-list* to an implicit **to**
7  clause of one item is formed from the name of the enclosing subroutine subprogram, function
8  subprogram or interface body to which it applies.

———————————————— Fortran ————————————————

9  **Restrictions**

10  • A threadprivate variable cannot appear in a **declare target** directive.

11  • A variable declared in a **declare target** directive must have a mappable type.

12  • The same list item must not appear multiple times in clauses on the same directive.

13  • The same list item must not appear in both a **to** clause on one **declare target** directive and
14  a **link** clause on another **declare target** directive.

———————————————— C / C++ ————————————————

15  • All declarations and definitions for a function must have a **declare target** directive if one is
16  specified for any of them. Otherwise, the result is unspecified.

17  • The *declaration-definition-seq* defined by a **declare target** directive and an
18  **end declare target** directive must not contain any **declare target** directives.

———————————————— C / C++ ————————————————
———————————————— C++ ————————————————

19  • The function names of overloaded functions or template functions may only be specified within
20  an implicit *extended-list*.

———————————————— C++ ————————————————

- If a list item is a procedure name, it must not be a generic name, procedure pointer or entry name.

- Any **declare target** directive with clauses must appear in a specification part of a subroutine subprogram, function subprogram, program or module.

- Any **declare target** directive without clauses must appear in a specification part of a subroutine subprogram, function subprogram or interface body to which it applies.

- If a **declare target** directive is specified in an interface block for a procedure, it must match a **declare target** directive in the definition of the procedure.

- If an external procedure is a type-bound procedure of a derived type and a **declare target** directive is specified in the definition of the external procedure, such a directive must appear in the interface block that is accessible to the derived type definition.

- If any procedure is declared via a procedure declaration statement that is not in the type-bound procedure part of a derived-type definition, any **declare target** with the procedure name must appear in the same specification part.

- A variable that is part of another variable (as an array or structure element) cannot appear in a **declare target** directive.

- The **declare target** directive must appear in the declaration section of a scoping unit in which the common block or variable is declared. Although variables in common blocks can be accessed by use association or host association, common block names cannot. This means that a common block name specified in a **declare target** directive must be declared to be a common block in the same scoping unit in which the **declare target** directive appears.

- If a **declare target** directive specifying a common block name appears in one program unit, then such a directive must also appear in every other program unit that contains a **COMMON** statement specifying the same name. It must appear after the last such **COMMON** statement in the program unit.

- If a list item is declared with the **BIND** attribute, the corresponding C entities must also be specified in a **declare target** directive in the C program.

- A blank common block cannot appear in a **declare target** directive.

- A variable can only appear in a **declare target** directive in the scope in which it is declared. It must not be an element of a common block or appear in an **EQUIVALENCE** statement.

- A variable that appears in a **declare target** directive must be declared in the Fortran scope of a module or have the **SAVE** attribute, either explicitly or implicitly.

# 2.10.7 `teams` Construct

**Summary**

The `teams` construct creates a league of thread teams and the master thread of each team executes the region.

**Syntax**

———————————————— C / C++ ————————————————

The syntax of the `teams` construct is as follows:

```
#pragma omp teams [clause[ [,] clause] ... ] new-line
    structured-block
```

where *clause* is one of the following:

      `num_teams(`*integer-expression*`)`

      `thread_limit(`*integer-expression*`)`

      `default(shared` | `none)`

      `private(`*list*`)`

      `firstprivate(`*list*`)`

      `shared(`*list*`)`

      `reduction(`*reduction-identifier* : *list*`)`

———————————————— C / C++ ————————————————

1    The syntax of the **teams** construct is as follows:

```
!$omp teams [clause[ [,] clause] ... ]
    structured-block
!$omp end teams
```

2    where *clause* is one of the following:

3        **num_teams(***scalar-integer-expression***)**

4        **thread_limit(***scalar-integer-expression***)**

5        **default(shared** | **firstprivate** | **private** | **none)**

6        **private(***list***)**

7        **firstprivate(***list***)**

8        **shared(***list***)**

9        **reduction(***reduction-identifier* **:** *list***)**

10   The **end teams** directive denotes the end of the **teams** construct.

### Binding

12   The binding thread set for a **teams** region is the encountering thread, which is the initial thread of
13   the **target** region.

### Description

15   When a thread encounters a **teams** construct, a league of thread teams is created and the master
16   thread of each thread team executes the **teams** region.

17   The number of teams created is implementation defined, but is less than or equal to the value
18   specified in the **num_teams** clause. A thread may obtain the number of teams by a call to the
19   **omp_get_num_teams** routine.

20   The maximum number of threads participating in the contention group that each team initiates is
21   implementation defined, but is less than or equal to the value specified in the **thread_limit**
22   clause.

23   On a combined or composite construct that includes **target** and **teams** constructs, the
24   expressions in **num_teams** and **thread_limit** clauses are evaluated on the host device on
25   entry to the **target** construct.

Once the teams are created, the number of teams remains constant for the duration of the **teams** region.

Within a **teams** region, team numbers uniquely identify each team. Team numbers are consecutive whole numbers ranging from zero to one less than the number of teams. A thread may obtain its own team number by a call to the **omp_get_team_num** library routine.

After the teams have completed execution of the **teams** region, the encountering thread resumes execution of the enclosing **target** region.

There is no implicit barrier at the end of a **teams** construct.

### Restrictions

Restrictions to the **teams** construct are as follows:

- A program that branches into or out of a **teams** region is non-conforming.

- A program must not depend on any ordering of the evaluations of the clauses of the **teams** directive, or on any side effects of the evaluation of the clauses.

- At most one **thread_limit** clause can appear on the directive. The **thread_limit** expression must evaluate to a positive integer value.

- At most one **num_teams** clause can appear on the directive. The **num_teams** expression must evaluate to a positive integer value.

- If specified, a **teams** construct must be contained within a **target** construct. That **target** construct must contain no statements, declarations or directives outside of the **teams** construct.

- **distribute**, **distribute simd**, distribute parallel loop, distribute parallel loop SIMD, and **parallel** regions, including any **parallel** regions arising from combined constructs, are the only OpenMP regions that may be strictly nested inside the **teams** region.

### Cross References

- **default**, **shared**, **private**, **firstprivate**, and **reduction** clauses, see Section 2.15.3 on page 188.

- **omp_get_num_teams** routine, see Section 3.2.32 on page 264.

- **omp_get_team_num** routine, see Section 3.2.33 on page 266.

# 1 2.10.8 **`distribute` Construct**

2 **Summary**

3 The **`distribute`** construct specifies that the iterations of one or more loops will be executed by
4 the thread teams in the context of their implicit tasks. The iterations are distributed across the
5 master threads of all teams that execute the **`teams`** region to which the **`distribute`** region binds.

6 **Syntax**

—————————————————————— C / C++ ——————————————————————

7 The syntax of the **`distribute`** construct is as follows:

```
#pragma omp distribute [clause[ [,] clause] ... ] new-line
    for-loops
```

8 Where *clause* is one of the following:

9       **`private(`***list***`)`**

10       **`firstprivate(`***list***`)`**

11       **`lastprivate(`***list***`)`**

12       **`collapse(`***n***`)`**

13       **`dist_schedule(`***kind[, chunk_size]***`)`**

14 All associated *for-loops* must have the canonical form described in Section 2.6 on page 53.

—————————————————————— C / C++ ——————————————————————

1　The syntax of the **distribute** construct is as follows:

```
!$omp distribute [clause[ [,] clause] ... ]
    do-loops
[!$omp end distribute]
```

2　Where *clause* is one of the following:

3　　　**private(***list***)**

4　　　**firstprivate(***list***)**

5　　　**lastprivate(***list***)**

6　　　**collapse(***n***)**

7　　　**dist_schedule(***kind[, chunk_size]***)**

8　If an **end distribute** directive is not specified, an **end distribute** directive is assumed at
9　the end of the *do-loops*.

10　Any associated *do-loop* must be a *do-construct* or an *inner-shared-do-construct* as defined by the
11　Fortran standard. If an **end distribute** directive follows a *do-construct* in which several loop
12　statements share a **DO** termination statement, then the directive can only be specified for the
13　outermost of these **DO** statements.

### 14　Binding

15　The binding thread set for a **distribute** region is the set of master threads executing an
16　enclosing **teams** region. A **distribute** region binds to this **teams** region. Only the threads
17　executing the binding **teams** region participate in the execution of the loop iterations.

### 18　Description

19　The **distribute** construct is associated with a loop nest consisting of one or more loops that
20　follow the directive.

21　There is no implicit barrier at the end of a **distribute** construct. To avoid data races the
22　original list items modified due to **lastprivate** or **linear** clauses should not be accessed
23　between the end of the **distribute** construct and the end of the **teams** region to which the
24　**distribute** binds.

25　The **collapse** clause may be used to specify how many loops are associated with the
26　**distribute** construct. The parameter of the **collapse** clause must be a constant positive

integer expression. If no **collapse** clause is present, the only loop that is associated with the **distribute** construct is the one that immediately follows the **distribute** construct.

If more than one loop is associated with the **distribute** construct, then the iteration of all associated loops are collapsed into one larger iteration space. The sequential execution of the iterations in all associated loops determines the order of the iterations in the collapsed iteration space.

The iteration count for each associated loop is computed before entry to the outermost loop. If execution of any associated loop changes any of the values used to compute any of the iteration counts, then the behavior is unspecified.

The integer type (or kind, for Fortran) used to compute the iteration count for the collapsed loop is implementation defined.

If **dist_schedule** is specified, *kind* must be **static**. If specified, iterations are divided into chunks of size *chunk_size*, chunks are assigned to the teams of the league in a round-robin fashion in the order of the team number. When no *chunk_size* is specified, the iteration space is divided into chunks that are approximately equal in size, and at most one chunk is distributed to each team of the league. The size of the chunks is unspecified in this case.

When no **dist_schedule** clause is specified, the schedule is implementation defined.

### Restrictions

Restrictions to the **distribute** construct are as follows:

- The **distribute** construct inherits the restrictions of the loop construct.
- The region associated with the **distribute** construct must be strictly nested inside a **teams** region.
- A list item may appear in a **firstprivate** or **lastprivate** clause but not both.

### Cross References

- loop construct, see Section 2.7.1 on page 56.
- **teams** construct, see Section 2.10.7 on page 114

## 2.10.9  **distribute simd** Construct

### Summary

The **distribute simd** construct specifies a loop that will be distributed across the master threads of the **teams** region and executed concurrently using SIMD instructions. The **distribute simd** construct is a composite construct.

**Syntax**

2 The syntax of the **distribute simd** construct is as follows:

───────── C / C++ ─────────

```
#pragma omp distribute simd [clause[ [,] clause] ... ] newline
    for-loops
```

3 where *clause* can be any of the clauses accepted by the **distribute** or **simd** directives with
4 identical meanings and restrictions.

───────── C / C++ ─────────

───────── Fortran ─────────

```
!$omp distribute simd [clause[ [,] clause] ... ]
    do-loops
[!$omp end distribute simd]
```

5 where *clause* can be any of the clauses accepted by the **distribute** or **simd** directives with
6 identical meanings and restrictions.

7 If an **end distribute simd** directive is not specified, an **end distribute simd** directive is
8 assumed at the end of the *do-loops*.

───────── Fortran ─────────

**Description**

10 The **distribute simd** construct will first distribute the iterations of the associated loop(s)
11 according to the semantics of the **distribute** construct and any clauses that apply to the
12 distribute construct. The resulting chunks of iterations will then be converted to a SIMD loop in a
13 manner consistent with any clauses that apply to the **simd** construct. The effect of any clause that
14 applies to both constructs is as if it were applied to both constructs separately except the
15 **collapse** clause, which is applied once.

**Restrictions**

17 • The restrictions for the **distribute** and **simd** constructs apply.

18 • A list item may not appear in a **linear** clause, unless it is the loop iteration variable.

2 • **simd** construct, see Section 2.8.1 on page 72.

3 • **distribute** construct, see Section 2.10.8 on page 117.

4 • Data attribute clauses, see Section 2.15.3 on page 188.

## 5   2.10.10   Distribute Parallel Loop Construct

6 **Summary**

7 The distribute parallel loop construct specifies a loop that can be executed in parallel by multiple
8 threads that are members of multiple teams. The distribute parallel loop construct is a composite
9 construct.

10 **Syntax**

11 The syntax of the distribute parallel loop construct is as follows:

--- C / C++ ---

```
#pragma omp distribute parallel for [clause[ [,] clause] ... ] newline
    for-loops
```

12 where *clause* can be any of the clauses accepted by the **distribute** or parallel loop directives
13 with identical meanings and restrictions.

--- C / C++ ---

--- Fortran ---

```
!$omp distribute parallel do [clause[ [,] clause] ... ]
    do-loops
[!$omp end distribute parallel do]
```

14 where *clause* can be any of the clauses accepted by the **distribute** or parallel loop directives
15 with identical meanings and restrictions.

16 If an **end distribute parallel do** directive is not specified, an
17 **end distribute parallel do** directive is assumed at the end of the *do-loops*.

--- Fortran ---

**Description**

The distribute parallel loop construct will first distribute the iterations of the associated loop(s) into chunks according to the semantics of the **distribute** construct and any clauses that apply to the **distribute** construct. Each of these chunks will form a loop. Each resulting loop will then be distributed across the threads within the teams region to which the **distribute** construct binds in a manner consistent with any clauses that apply to the parallel loop construct. The effect of any clause that applies to both constructs is as if it were applied to both constructs separately except the **collapse** clause, which is applied once.

**Restrictions**

- The restrictions for the **distribute** and parallel loop constructs apply.

- No **ordered** clause can be specified.

- No **linear** clause can be specified.

**Cross References**

- **distribute** construct, see Section 2.10.8 on page 117.

- Parallel loop construct, see Section 2.11.1 on page 124.

- Data attribute clauses, see Section 2.15.3 on page 188.

## 2.10.11 Distribute Parallel Loop SIMD Construct

**Summary**

The distribute parallel loop SIMD construct specifies a loop that can be executed concurrently using SIMD instructions in parallel by multiple threads that are members of multiple teams. The distribute parallel loop SIMD construct is a composite construct.

**Syntax**

$$\blacktriangledown \text{────────── C / C++ ──────────} \blacktriangledown$$

The syntax of the distribute parallel loop SIMD construct is as follows:

```
#pragma omp distribute parallel for simd [clause[ [,] clause] ... ] newline
    for-loops
```

where *clause* can be any of the clauses accepted by the **distribute** or parallel loop SIMD directives with identical meanings and restrictions

$$\blacktriangle \text{────────── C / C++ ──────────} \blacktriangle$$

1    The syntax of the distribute parallel loop SIMD construct is as follows:

```
!$omp distribute parallel do simd [clause[ [,] clause] ... ]
     do-loops
[!$omp end distribute parallel do simd]
```

2    where *clause* can be any of the clauses accepted by the **distribute** or parallel loop SIMD
3    directives with identical meanings and restrictions.

4    If an **end distribute parallel do simd** directive is not specified, an
5    **end distribute parallel do simd** directive is assumed at the end of the *do-loops*.

6    **Description**

7    The distribute parallel loop SIMD construct will first distribute the iterations of the associated
8    loop(s) according to the semantics of the **distribute** construct and any clauses that apply to the
9    **distribute** construct. The resulting loops will then be distributed across the threads contained
10   within the **teams** region to which the **distribute** construct binds in a manner consistent with
11   any clauses that apply to the parallel loop construct. The resulting chunks of iterations will then be
12   converted to a SIMD loop in a manner consistent with any clauses that apply to the **simd** construct.
13   The effect of any clause that applies to both constructs is as if it were applied to both constructs
14   separately except the **collapse** clause, which is applied once.

15   **Restrictions**

16   • The restrictions for the **distribute** and parallel loop SIMD constructs apply.

17   • No **ordered** clause can be specified.

18   • A list item may not appear in a **linear** clause, unless it is the loop iteration variable.

19   **Cross References**

20   • **distribute** construct, see Section 2.10.8 on page 117.

21   • Parallel loop SIMD construct, see Section 2.11.4 on page 128.

22   • Data attribute clauses, see Section 2.15.3 on page 188.

# 2.11 Combined Constructs

Combined constructs are shortcuts for specifying one construct immediately nested inside another construct. The semantics of the combined constructs are identical to that of explicitly specifying the first construct containing one instance of the second construct and no other statements.

Some combined constructs have clauses that are permitted on both constructs that were combined. Where specified, the effect is as if applying the clauses to one or both constructs. If not specified and applying the clause to one construct would result in different program behavior than applying the clause to the other construct then the program's behavior is unspecified.

## 2.11.1 Parallel Loop Construct

**Summary**

The parallel loop construct is a shortcut for specifying a **parallel** construct containing one loop constuct with one or more associated loops and no other statements.

**Syntax**

--- C / C++ ---

The syntax of the parallel loop construct is as follows:

```
#pragma omp parallel for [clause[ [,] clause] ... ] new-line
    for-loops
```

where *clause* can be any of the clauses accepted by the **parallel** or **for** directives, except the **nowait** clause, with identical meanings and restrictions.

--- C / C++ ---
--- Fortran ---

The syntax of the parallel loop construct is as follows:

```
!$omp parallel do [clause[ [,] clause] ... ]
    do-loops
[!$omp end parallel do]
```

where *clause* can be any of the clauses accepted by the **parallel** or **do** directives, with identical meanings and restrictions.

If an **end parallel do** directive is not specified, an **end parallel do** directive is assumed at the end of the *do-loops*. **nowait** may not be specified on an **end parallel do** directive.

--- Fortran ---

**Description**

The semantics are identical to explicitly specifying a **parallel** directive immediately followed by a loop directive.

**Restrictions**

- The restrictions for the **parallel** construct and the loop construct apply.

**Cross References**

- **parallel** construct, see Section 2.5 on page 46.
- loop SIMD construct, see Section 2.8.3 on page 81.
- Data attribute clauses, see Section 2.15.3 on page 188.

## 2.11.2 **parallel sections** Construct

**Summary**

The **parallel sections** construct is a shortcut for specifying a **parallel** construct containing one **sections** construct and no other statements.

**Syntax**

C / C++

The syntax of the **parallel sections** construct is as follows:

```
#pragma omp parallel sections [clause[ [, ] clause] ... ] new-line
    {
    [#pragma omp section new-line]
        structured-block
    [#pragma omp section new-line
        structured-block]
    ...
    }
```

where *clause* can be any of the clauses accepted by the **parallel** or **sections** directives, except the **nowait** clause, with identical meanings and restrictions.

C / C++

--- Fortran ---

1    The syntax of the **parallel sections** construct is as follows:

```
!$omp parallel sections [clause[ [,] clause] ... ]
    [!$omp section]
        structured-block
    [!$omp section
        structured-block]
    ...
!$omp end parallel sections
```

2    where *clause* can be any of the clauses accepted by the **parallel** or **sections** directives, with
3    identical meanings and restrictions.

4    The last section ends at the **end parallel sections** directive. **nowait** cannot be specified
5    on an **end parallel sections** directive.

--- Fortran ---

6    **Description**

--- C / C++ ---

7    The semantics are identical to explicitly specifying a **parallel** directive immediately followed
8    by a **sections** directive.

--- C / C++ ---
--- Fortran ---

9    The semantics are identical to explicitly specifying a **parallel** directive immediately followed
10   by a **sections** directive, and an **end sections** directive immediately followed by an
11   **end parallel** directive.

--- Fortran ---

12   **Restrictions**

13   The restrictions for the **parallel** construct and the **sections** construct apply.

- **parallel** construct, see Section 2.5 on page 46.

- **sections** construct, see Section 2.7.2 on page 65.

- Data attribute clauses, see Section 2.15.3 on page 188.

———————————— Fortran ————————————

## 2.11.3  **parallel workshare** Construct

**Summary**

The **parallel workshare** construct is a shortcut for specifying a **parallel** construct containing one **workshare** construct and no other statements.

**Syntax**

The syntax of the **parallel workshare** construct is as follows:

```
!$omp parallel workshare [clause[ [,] clause] ... ]
   structured-block
!$omp end parallel workshare
```

where *clause* can be any of the clauses accepted by the **parallel** directive, with identical meanings and restrictions. **nowait** may not be specified on an **end parallel workshare** directive.

**Description**

The semantics are identical to explicitly specifying a **parallel** directive immediately followed by a **workshare** directive, and an **end workshare** directive immediately followed by an **end parallel** directive.

**Restrictions**

The restrictions for the **parallel** construct and the **workshare** construct apply.

- **parallel** construct, see Section 2.5 on page 46.

- **workshare** construct, see Section 2.7.4 on page 69.

- Data attribute clauses, see Section 2.15.3 on page 188.

—————————————— Fortran ——————————————

## 2.11.4 Parallel Loop SIMD Construct

**Summary**

The parallel loop SIMD construct is a shortcut for specifying a **parallel** construct containing one loop SIMD construct and no other statement.

**Syntax**

—————————————— C / C++ ——————————————

The syntax of the parallel loop SIMD construct is as follows:

```
#pragma omp parallel for simd [clause[ [,] clause] ... ] new-line
    for-loops
```

where *clause* can be any of the clauses accepted by the **parallel** or **for simd** directives, except the **nowait** clause, with identical meanings and restrictions.

—————————————— C / C++ ——————————————
—————————————— Fortran ——————————————

The syntax of the parallel loop SIMD construct is as follows:

```
!$omp parallel do simd [clause[ [,] clause] ... ]
    do-loops
[!$omp end parallel do simd]
```

where *clause* can be any of the clauses accepted by the **parallel** or **do simd** directives, with identical meanings and restrictions.

If an **end parallel do simd** directive is not specified, an **end parallel do simd** directive is assumed at the end of the *do-loops*. **nowait** may not be specified on an **end parallel do simd** directive.

—————————————— Fortran ——————————————

**Description**

The semantics of the parallel loop SIMD construct are identical to explicitly specifying a **parallel** directive immediately followed by a loop SIMD directive. The effect of any clause that applies to both constructs is as if it were applied to the loop SIMD construct and not to the **parallel** construct.

**Restrictions**

The restrictions for the **parallel** construct and the loop SIMD construct apply.

**Cross References**

- **parallel** construct, see Section 2.5 on page 46.
- loop SIMD construct, see Section 2.8.3 on page 81.
- Data attribute clauses, see Section 2.15.3 on page 188.

## 2.11.5 `target parallel` Construct

**Summary**

The **target parallel** construct is a shortcut for specifying a **target** construct containing a **parallel** construct and no other statements.

**Syntax**

$\blacktriangledown$ —————————— C / C++ —————————— $\blacktriangledown$

The syntax of the **target parallel** construct is as follows:

```
#pragma omp target parallel [clause[ [,] clause] ... ] new-line
    structured-block
```

where *clause* can be any of the clauses accepted by the **target** or **parallel** directives, except for **copyin**, with identical meanings and restrictions.

$\blacktriangle$ —————————— C / C++ —————————— $\blacktriangle$

1    The syntax of the **target parallel** construct is as follows:

```
!$omp target parallel [clause[ [,] clause] ... ]
    structured-block
!$omp end target parallel
```

2    where *clause* can be any of the clauses accepted by the **target** or **parallel** directives, except
3    for **copyin**, with identical meanings and restrictions.

4    **Description**

5    The semantics are identical to explicitly specifying a **target** directive immediately followed by a
6    **parallel** directive.

7    **Restrictions**

8    The restrictions for the **target** and **parallel** constructs apply except for the following explicit
9    modifications:

10   • If any **if** clause on the directive includes a *directive-name-modifier* then all **if** clauses on the
11      directive must include a *directive-name-modifier*.

12   • At most one **if** clause without a *directive-name-modifier* can appear on the directive.

13   • At most one **if** clause with the **parallel** *directive-name-modifier* can appear on the directive.

14   • At most one **if** clause with the **target** *directive-name-modifier* can appear on the directive.

15   **Cross References**

16   • **parallel** construct, see Section 2.5 on page 46.

17   • **target** construct, see Section 2.10.4 on page 103.

18   • **if** Clause, see Section 2.12 on page 147.

19   • Data attribute clauses, see Section 2.15.3 on page 188.

# 2.11.6 Target Parallel Loop Construct

**Summary**

The target parallel loop construct is a shortcut for specifying a **target** construct containing a parallel loop construct and no other statements.

**Syntax**

———————————————— C / C++ ————————————————

The syntax of the target parallel loop construct is as follows:

```
#pragma omp target parallel for [clause[ [,] clause] ... ] new-line
    for-loops
```

where *clause* can be any of the clauses accepted by the **target** or **parallel for** directives, except for **copyin**, with identical meanings and restrictions.

———————————————— C / C++ ————————————————

———————————————— Fortran ————————————————

The syntax of the target parallel loop construct is as follows:

```
!$omp target parallel do [clause[ [,] clause] ... ]
    do-loops
[!$omp end target parallel do]
```

where *clause* can be any of the clauses accepted by the **target** or **parallel do** directives, except for **copyin**, with identical meanings and restrictions.

If an **end target parallel do** directive is not specified, an **end target parallel do** directive is assumed at the end of the *do-loops*.

———————————————— Fortran ————————————————

**Description**

The semantics are identical to explicitly specifying a **target** directive immediately followed by a parallel loop directive.

**Restrictions**

The restrictions for the **target** and parallel loop constructs apply except for the following explicit modifications:

- If any **if** clause on the directive includes a *directive-name-modifier* then all **if** clauses on the directive must include a *directive-name-modifier*.

- At most one **if** clause without a *directive-name-modifier* can appear on the directive.

- At most one **if** clause with the **parallel** *directive-name-modifier* can appear on the directive.

- At most one **if** clause with the **target** *directive-name-modifier* can appear on the directive.

**Cross References**

- **target** construct, see Section 2.10.4 on page 103.

- Parallel loop construct, see Section 2.11.1 on page 124.

- **if** Clause, see Section 2.12 on page 147.

- Data attribute clauses, see Section 2.15.3 on page 188.

## 2.11.7 Target Parallel Loop SIMD Construct

**Summary**

The target parallel loop SIMD construct is a shortcut for specifying a **target** construct containing a parallel loop SIMD construct and no other statements.

**Syntax**

―――――――――――――――― C / C++ ――――――――――――――――

The syntax of the target parallel loop SIMD construct is as follows:

```
#pragma omp target parallel for simd [clause[[,] clause] ... ] new-line
    for-loops
```

where *clause* can be any of the clauses accepted by the **target** or **parallel for simd** directives, except for **copyin**, with identical meanings and restrictions.

―――――――――――――――― C / C++ ――――――――――――――――

1 The syntax of the target parallel loop SIMD construct is as follows:

```
!$omp target parallel do simd [clause[ [ , ] clause] ... ]
    do-loops
[!$omp end target parallel do simd]
```

2 where *clause* can be any of the clauses accepted by the **target** or **parallel do simd**
3 directives, except for **copyin**, with identical meanings and restrictions.

4 If an **end target parallel do simd** directive is not specified, an
5 **end target parallel do simd** directive is assumed at the end of the *do-loops*.

6 **Description**

7 The semantics are identical to explicitly specifying a **target** directive immediately followed by a
8 parallel loop SIMD directive.

9 **Restrictions**

10 The restrictions for the **target** and parallel loop SIMD constructs apply except for the following
11 explicit modifications:

12 • If any **if** clause on the directive includes a *directive-name-modifier* then all **if** clauses on the
13 directive must include a *directive-name-modifier*.

14 • At most one **if** clause without a *directive-name-modifier* can appear on the directive.

15 • At most one **if** clause with the **parallel** *directive-name-modifier* can appear on the directive.

16 • At most one **if** clause with the **target** *directive-name-modifier* can appear on the directive.

17 **Cross References**

18 • **target** construct, see Section 2.10.4 on page 103.

19 • Parallel loop SIMD construct, see Section 2.11.4 on page 128.

20 • **if** Clause, see Section 2.12 on page 147.

21 • Data attribute clauses, see Section 2.15.3 on page 188.

## 2.11.8 `target simd` Construct

**Summary**

The **target simd** construct is a shortcut for specifying a **target** construct containing a **simd** construct and no other statements.

**Syntax**

─────────────────── C / C++ ───────────────────

The syntax of the **target simd** construct is as follows:

```
#pragma omp target simd [clause[ [,] clause] ... ] new-line
    for-loops
```

where *clause* can be any of the clauses accepted by the **target** or **simd** directives with identical meanings and restrictions.

─────────────────── C / C++ ───────────────────
─────────────────── Fortran ───────────────────

The syntax of the **target simd** construct is as follows:

```
!$omp target simd [clause[ [,] clause] ... ]
    do-loops
[!$omp end target simd]
```

where *clause* can be any of the clauses accepted by the **target** or **simd** directives with identical meanings and restrictions.

If an **end target simd** directive is not specified, an **end target simd** directive is assumed at the end of the *do-loops*.

─────────────────── Fortran ───────────────────

**Description**

The semantics are identical to explicitly specifying a **target** directive immediately followed by a **simd** directive.

**Restrictions**

The restrictions for the **target** and **simd** constructs apply.

**Cross References**

- **simd** construct, see Section 2.8.1 on page 72.

- **target** construct, see Section 2.10.4 on page 103.

- Data attribute clauses, see Section 2.15.3 on page 188.

## 2.11.9 **target teams** Construct

**Summary**

The **target teams** construct is a shortcut for specifying a **target** construct containing a **teams** construct and no other statements.

**Syntax**

─────────────────────────── C / C++ ───────────────────────────

The syntax of the **target teams** construct is as follows:

```
#pragma omp target teams [clause[ [,] clause] ... ] new-line
    structured-block
```

where *clause* can be any of the clauses accepted by the **target** or **teams** directives with identical meanings and restrictions.

─────────────────────────── C / C++ ───────────────────────────
─────────────────────────── Fortran ───────────────────────────

The syntax of the **target teams** construct is as follows:

```
!$omp target teams [clause[ [,] clause] ... ]
    structured-block
!$omp end target teams
```

where *clause* can be any of the clauses accepted by the **target** or **teams** directives with identical meanings and restrictions.

─────────────────────────── Fortran ───────────────────────────

**Description**

The semantics are identical to explicitly specifying a **target** directive immediately followed by a **teams** directive.

**Restrictions**

The restrictions for the **target** and **teams** constructs apply.

**Cross References**

- **target** construct, see Section 2.10.4 on page 103.
- **teams** construct, see Section 2.10.7 on page 114.
- Data attribute clauses, see Section 2.15.3 on page 188.

# 2.11.10 **teams distribute** Construct

**Summary**

The **teams distribute** construct is a shortcut for specifying a **teams** construct containing a **distribute** construct and no other statements.

**Syntax**

———————————— C / C++ ————————————

The syntax of the **teams distribute** construct is as follows:

```
#pragma omp teams distribute [clause[ [ , ] clause] ... ] new-line
    for-loops
```

where *clause* can be any of the clauses accepted by the **teams** or **distribute** directives with identical meanings and restrictions.

———————————— C / C++ ————————————

1  The syntax of the **teams distribute** construct is as follows:

```
!$omp teams distribute [clause[ [,] clause] ... ]
    do-loops
[!$omp end teams distribute]
```

2  where *clause* can be any of the clauses accepted by the **teams** or **distribute** directives with
3  identical meanings and restrictions.

4  If an **end teams distribute** directive is not specified, an **end teams distribute**
5  directive is assumed at the end of the *do-loops*.

6  **Description**

7  The semantics are identical to explicitly specifying a **teams** directive immediately followed by a
8  **distribute** directive. The effect of any clause that applies to both constructs is as if it were
9  applied to both constructs separately.

10  **Restrictions**

11  The restrictions for the **teams** and **distribute** constructs apply.

12  **Cross References**

13  • **teams** construct, see Section 2.10.7 on page 114.

14  • **distribute** construct, see Section 2.10.8 on page 117.

15  • Data attribute clauses, see Section 2.15.3 on page 188.

## 16  2.11.11  **teams distribute simd** Construct

17  **Summary**

18  The **teams distribute simd** construct is a shortcut for specifying a **teams** construct
19  containing a **distribute simd** construct and no other statements.

**Syntax**

--- C / C++ ---

The syntax of the **teams distribute simd** construct is as follows:

```
#pragma omp teams distribute simd [clause[ [, ] clause] ... ] new-line
    for-loops
```

where *clause* can be any of the clauses accepted by the **teams** or **distribute simd** directives with identical meanings and restrictions.

--- C / C++ ---

--- Fortran ---

The syntax of the **teams distribute simd** construct is as follows:

```
!$omp teams distribute simd [clause[ [, ] clause] ... ]
    do-loops
[!$omp end teams distribute simd]
```

where *clause* can be any of the clauses accepted by the **teams** or **distribute simd** directives with identical meanings and restrictions.

If an **end teams distribute simd** directive is not specified, an
**end teams distribute simd** directive is assumed at the end of the *do-loops*.

--- Fortran ---

**Description**

The semantics are identical to explicitly specifying a **teams** directive immediately followed by a **distribute simd** directive. The effect of any clause that applies to both constructs is as if it were applied to both constructs separately.

**Restrictions**

The restrictions for the **teams** and **distribute simd** constructs apply.

**Cross References**

- **teams** construct, see Section 2.10.7 on page 114.

- **distribute simd** construct, see Section 2.10.9 on page 119.

- Data attribute clauses, see Section 2.15.3 on page 188.

# 2.11.12 `target teams distribute` Construct

**Summary**

The **`target teams distribute`** construct is a shortcut for specifying a **`target`** construct containing a **`teams distribute`** construct and no other statements.

**Syntax**

--- C / C++ ---

The syntax of the **`target teams distribute`** construct is as follows:

```
#pragma omp target teams distribute [clause[ [,] clause] ... ] new-line
    for-loops
```

where *clause* can be any of the clauses accepted by the **`target`** or **`teams distribute`** directives with identical meanings and restrictions.

--- C / C++ ---

--- Fortran ---

The syntax of the **`target teams distribute`** construct is as follows:

```
!$omp target teams distribute [clause[ [,] clause] ... ]
    do-loops
[!$omp end target teams distribute]
```

where *clause* can be any of the clauses accepted by the **`target`** or **`teams distribute`** directives with identical meanings and restrictions.

If an **`end target teams distribute`** directive is not specified, an **`end target teams distribute`** directive is assumed at the end of the *do-loops*.

--- Fortran ---

**Description**

The semantics are identical to explicitly specifying a **`target`** directive immediately followed by a **`teams distribute`** directive.

**Restrictions**

The restrictions for the **`target`** and **`teams distribute`** constructs apply.

## 2.11.13 **target teams distribute simd** Construct

**Summary**

The **target teams distribute simd** construct is a shortcut for specifying a **target** construct containing a **teams distribute simd** construct and no other statements.

**Syntax**

--- C / C++ ---

The syntax of the **target teams distribute simd** construct is as follows:

```
#pragma omp target teams distribute simd [clause[ [,] clause] ... ] new-line
    for-loops
```

where *clause* can be any of the clauses accepted by the **target** or **teams distribute simd** directives with identical meanings and restrictions.

--- C / C++ ---

--- Fortran ---

The syntax of the **target teams distribute simd** construct is as follows:

```
!$omp target teams distribute simd [clause[ [,] clause] ... ]
    do-loops
[!$omp end target teams distribute simd]
```

where *clause* can be any of the clauses accepted by the **target** or **teams distribute simd** directives with identical meanings and restrictions.

If an **end target teams distribute simd** directive is not specified, an **end target teams distribute simd** directive is assumed at the end of the *do-loops*.

--- Fortran ---

**Description**

The semantics are identical to explicitly specifying a **target** directive immediately followed by a **teams distribute simd** directive.

**Restrictions**

The restrictions for the **target** and **teams distribute simd** constructs apply.

**Cross References**

- **target** construct, see Section 2.10.1 on page 95.
- **teams distribute simd** construct, see Section 2.11.11 on page 137.
- Data attribute clauses, see Section 2.15.3 on page 188.

## 2.11.14 Teams Distribute Parallel Loop Construct

**Summary**

The teams distribute parallel loop construct is a shortcut for specifying a **teams** construct containing a distribute parallel loop construct and no other statements.

**Syntax**

$\blacktriangledown$ ——————————————— C / C++ ———————————————— $\blacktriangledown$

The syntax of the teams distribute parallel loop construct is as follows:

```
#pragma omp teams distribute parallel for [clause[ [,] clause] ... ] new-line
    for-loops
```

where *clause* can be any of the clauses accepted by the **teams** or **distribute parallel for** directives with identical meanings and restrictions.

$\blacktriangle$ ——————————————— C / C++ ———————————————— $\blacktriangle$

1 The syntax of the teams distribute parallel loop construct is as follows:

```
!$omp teams distribute parallel do [clause[ [ , ] clause] ... ]
    do-loops
[ !$omp end teams distribute parallel do ]
```

2 where *clause* can be any of the clauses accepted by the **teams** or **distribute parallel do**
3 directives with identical meanings and restrictions.

4 If an **end teams distribute parallel do** directive is not specified, an
5 **end teams distribute parallel do** directive is assumed at the end of the *do-loops*.

### Description

7 The semantics are identical to explicitly specifying a **teams** directive immediately followed by a
8 distribute parallel loop directive. The effect of any clause that applies to both constructs is as if it
9 were applied to both constructs separately.

### Restrictions

11 The restrictions for the **teams** and distribute parallel loop constructs apply.

### Cross References

13 • **teams** construct, see Section 2.10.7 on page 114.

14 • Distribute parallel loop construct, see Section 2.10.10 on page 121.

15 • Data attribute clauses, see Section 2.15.3 on page 188.

## 2.11.15  Target Teams Distribute Parallel Loop Construct

### Summary

18 The target teams distribute parallel loop construct is a shortcut for specifying a **target** construct
19 containing a teams distribute parallel loop construct and no other statements.

## Syntax

The syntax of the target teams distribute parallel loop construct is as follows:

```
#pragma omp target teams distribute parallel for [clause[ [,] clause] ... ] new-line
    for-loops
```

where *clause* can be any of the clauses accepted by the **target** or
**teams distribute parallel for** directives with identical meanings and restrictions.

The syntax of the target teams distribute parallel loop construct is as follows:

```
!$omp target teams distribute parallel do [clause[ [,] clause] ... ]
    do-loops
[!$omp end target teams distribute parallel do]
```

where *clause* can be any of the clauses accepted by the **target** or
**teams distribute parallel do** directives with identical meanings and restrictions.

If an **end target teams distribute parallel do** directive is not specified, an
**end target teams distribute parallel do** directive is assumed at the end of the
*do-loops*.

## Description

The semantics are identical to explicitly specifying a **target** directive immediately followed by a
teams distribute parallel loop directive.

## Restrictions

The restrictions for the **target** and teams distribute parallel loop constructs apply except for the
following explicit modifications:

- If any **if** clause on the directive includes a *directive-name-modifier* then all **if** clauses on the
  directive must include a *directive-name-modifier*.

- At most one **if** clause without a *directive-name-modifier* can appear on the directive.

- At most one **if** clause with the **parallel** *directive-name-modifier* can appear on the directive.

- At most one **if** clause with the **target** *directive-name-modifier* can appear on the directive.

## 2.11.16 Teams Distribute Parallel Loop SIMD Construct

**Summary**

The teams distribute parallel loop SIMD construct is a shortcut for specifying a **teams** construct containing a distribute parallel loop SIMD construct and no other statements.

**Syntax**

$$\text{------} \quad \text{C / C++} \quad \text{------}$$

The syntax of the teams distribute parallel loop construct is as follows:

```
#pragma omp teams distribute parallel for simd [clause[ [,] clause] ... ] new-line
    for-loops
```

where *clause* can be any of the clauses accepted by the **teams** or **distribute parallel for simd** directives with identical meanings and restrictions.

$$\text{------} \quad \text{C / C++} \quad \text{------}$$
$$\text{------} \quad \text{Fortran} \quad \text{------}$$

The syntax of the teams distribute parallel loop construct is as follows:

```
!$omp teams distribute parallel do simd [clause[ [,] clause] ... ]
    do-loops
[!$omp end teams distribute parallel do simd]
```

where *clause* can be any of the clauses accepted by the **teams** or **distribute parallel do simd** directives with identical meanings and restrictions.

If an **end teams distribute parallel do simd** directive is not specified, an **end teams distribute parallel do simd** directive is assumed at the end of the *do-loops*.

$$\text{------} \quad \text{Fortran} \quad \text{------}$$

**Description**

The semantics are identical to explicitly specifying a **teams** directive immediately followed by a distribute parallel loop SIMD directive. The effect of any clause that applies to both constructs is as if it were applied to both constructs separately.

**Restrictions**

The restrictions for the **teams** and distribute parallel loop SIMD constructs apply.

**Cross References**

- **teams** construct, see Section 2.10.7 on page 114.
- Distribute parallel loop SIMD construct, see Section 2.10.11 on page 122.
- Data attribute clauses, see Section 2.15.3 on page 188.

## 2.11.17 Target Teams Distribute Parallel Loop SIMD Construct

**Summary**

The target teams distribute parallel loop SIMD construct is a shortcut for specifying a **target** construct containing a teams distribute parallel loop SIMD construct and no other statements.

**Syntax**

———————————————— C / C++ ————————————————

The syntax of the target teams distribute parallel loop SIMD construct is as follows:

```
#pragma omp target teams distribute parallel for simd \
            [clause[ [ , ] clause] ... ] new-line
    for-loops
```

where *clause* can be any of the clauses accepted by the **target** or
**teams distribute parallel for simd** directives with identical meanings and restrictions.

———————————————— C / C++ ————————————————

1 The syntax of the target teams distribute parallel loop SIMD construct is as follows:

```
!$omp target teams distribute parallel do simd [clause[ [,] clause] ... ]
    do-loops
[!$omp end target teams distribute parallel do simd]
```

2 where *clause* can be any of the clauses accepted by the **target** or
3 **teams distribute parallel do simd** directives with identical meanings and restrictions.

4 If an **end target teams distribute parallel do simd** directive is not specified, an
5 **end target teams distribute parallel do simd** directive is assumed at the end of the
6 *do-loops*.

### Description

8 The semantics are identical to explicitly specifying a **target** directive immediately followed by a
9 teams distribute parallel loop SIMD directive.

### Restrictions

11 The restrictions for the **target** and teams distribute parallel loop SIMD constructs apply except
12 for the following explicit modifications:

13 • If any **if** clause on the directive includes a *directive-name-modifier* then all **if** clauses on the
14   directive must include a *directive-name-modifier*.

15 • At most one **if** clause without a *directive-name-modifier* can appear on the directive.

16 • At most one **if** clause with the **parallel** *directive-name-modifier* can appear on the directive.

17 • At most one **if** clause with the **target** *directive-name-modifier* can appear on the directive.

### Cross References

19 • **target** construct, see Section 2.10.4 on page 103.

20 • Teams distribute parallel loop SIMD construct, see Section 2.11.16 on page 144.

21 • **if** Clause, see Section 2.12 on page 147.

22 • Data attribute clauses, see Section 2.15.3 on page 188.

## 1 2.12 `if` Clause

2 **Summary**

3 The semantics of an `if` clause are described in the section on the construct to which it applies. The
4 `if` clause *directive-name-modifier* names the associated construct to which an expression applies,
5 and is particularly useful for composite and combined constructs.

6 **Syntax**

— C / C++ —

7 The syntax of the `if` clause is as follows:

```
if([ directive-name-modifier : ] scalar-expression)
```

— C / C++ —

— Fortran —

8 The syntax of the `if` clause is as follows:

```
if([ directive-name-modifier : ] scalar-logical-expression)
```

— Fortran —

9 **Description**

10 The effect of the `if` clause depends on the construct to which it is applied. For combined or
11 composite constructs, the `if` clause only applies to the semantics of the construct named in the
12 *directive-name-modifier* if one is specified. If no *directive-name-modifier* is specified for a
13 combined or composite construct then the `if` clause applies to all constructs to which an `if` clause
14 can apply.

# 2.13 Master and Synchronization Constructs and Clauses

OpenMP provides the following synchronization constructs:

- the **master** construct;
- the **critical** construct;
- the **barrier** construct;
- the **taskwait** construct;
- the **taskgroup** construct;
- the **atomic** construct;
- the **flush** construct;
- the **ordered** construct.

## 2.13.1 **master** Construct

**Summary**

The **master** construct specifies a structured block that is executed by the master thread of the team.

**Syntax**

———————————————— C / C++ ————————————————

The syntax of the **master** construct is as follows:

```
#pragma omp master new-line
    structured-block
```

———————————————— C / C++ ————————————————
———————————————— Fortran ————————————————

The syntax of the **master** construct is as follows:

```
!$omp master
    structured-block
!$omp end master
```

———————————————— Fortran ————————————————

### Binding

The binding thread set for a **master** region is the current team. A **master** region binds to the innermost enclosing **parallel** region. Only the master thread of the team executing the binding **parallel** region participates in the execution of the structured block of the **master** region.

### Description

Other threads in the team do not execute the associated structured block. There is no implied barrier either on entry to, or exit from, the **master** construct.

### Restrictions

─────────────────────── C++ ───────────────────────

- A throw executed inside a **master** region must cause execution to resume within the same **master** region, and the same thread that threw the exception must catch it

─────────────────────── C++ ───────────────────────

## 2.13.2 `critical` Construct

### Summary

The **critical** construct restricts execution of the associated structured block to a single thread at a time.

### Syntax

─────────────────────── C / C++ ───────────────────────

The syntax of the **critical** construct is as follows:

```
#pragma omp critical [(name) [hint(hint-expression)] ] new-line
    structured-block
```

where *hint-expression* is an integer constant expression that evaluates to a valid lock hint (as described in Section 3.3.2 on page 273).

─────────────────────── C / C++ ───────────────────────

1 The syntax of the **critical** construct is as follows:

```
!$omp critical [(name) [hint(hint-expression)] ]
    structured-block
!$omp end critical [(name)]
```

2 where *hint-expression* is a constant expression that evaluates to a scalar value with kind
3 **omp_lock_hint_kind** and a value that is a valid lock hint (as described in Section 3.3.2 on
4 page 273).

## Binding

6 The binding thread set for a **critical** region is all threads in the contention group. The region is
7 executed as if only a single thread at a time among all threads in the contention group is entering
8 the region for execution, without regard to the team(s) to which the threads belong.

## Description

10 An optional *name* may be used to identify the **critical** construct. All **critical** constructs
11 without a name are considered to have the same unspecified name.

12 Identifiers used to identify a **critical** construct have external linkage and are in a name space
13 that is separate from the name spaces used by labels, tags, members, and ordinary identifiers.

14 The names of **critical** constructs are global entities of the program. If a name conflicts with
15 any other entity, the behavior of the program is unspecified.

16 The threads of a contention group execute the **critical** region as if only one thread of the
17 contention group is executing the **critical** region at a time. The **critical** construct enforces
18 these execution semantics with respect to all **critical** constructs with the same name in all
19 threads in the contention group, not just those threads in the current team.

20 The presence of a **hint** clause does not affect the isolation guarantees provided by the **critical**
21 construct. If no **hint** clause is specified, the effect is as if **hint(omp_lock_hint_none)** had
22 been specified.

**Restrictions**

2   • If the **hint** clause is specified, the **critical** construct must have a *name*.

3   • If the **hint** clause is specified, each of the **critical** constructs with the same *name* must
4     have a **hint** clause for which the *hint-expression* evaluates to the same value.

──────────────────────── C++ ────────────────────────

5   • A throw executed inside a **critical** region must cause execution to resume within the same
6     **critical** region, and the same thread that threw the exception must catch it.

──────────────────────── C++ ────────────────────────

──────────────────────── Fortran ────────────────────────

7   The following restrictions apply to the critical construct:

8   • If a *name* is specified on a **critical** directive, the same *name* must also be specified on the
9     **end critical** directive.

10  • If no *name* appears on the **critical** directive, no *name* can appear on the **end critical**
11    directive.

──────────────────────── Fortran ────────────────────────

**Cross References**

13  • **omp_init_lock_with_hint** and **omp_init_nest_lock_with_hint** routines, see
14    Section 3.3.2 on page 273.

## 2.13.3  **barrier** Construct

**Summary**

17  The **barrier** construct specifies an explicit barrier at the point at which the construct appears.
18  The **barrier** construct is a stand-alone directive.

**Syntax**

$\blacktriangledown$ ——————————— C / C++ ——————————— $\blacktriangledown$

The syntax of the **barrier** construct is as follows:

```
#pragma omp barrier new-line
```

$\blacktriangle$ ——————————— C / C++ ——————————— $\blacktriangle$
$\blacktriangledown$ ——————————— Fortran ——————————— $\blacktriangledown$

The syntax of the **barrier** construct is as follows:

```
!$omp barrier
```

$\blacktriangle$ ——————————— Fortran ——————————— $\blacktriangle$

**Binding**

The binding thread set for a **barrier** region is the current team. A **barrier** region binds to the innermost enclosing **parallel** region.

**Description**

All threads of the team executing the binding **parallel** region must execute the **barrier** region and complete execution of all explicit tasks bound to this **parallel** region before any are allowed to continue execution beyond the barrier.

The **barrier** region includes an implicit task scheduling point in the current task region.

**Restrictions**

The following restrictions apply to the **barrier** construct:

- Each **barrier** region must be encountered by all threads in a team or by none at all, unless cancellation has been requested for the innermost enclosing parallel region.

- The sequence of worksharing regions and **barrier** regions encountered must be the same for every thread in a team.

# 2.13.4 `taskwait` Construct

### Summary

The **taskwait** construct specifies a wait on the completion of child tasks of the current task. The **taskwait** construct is a stand-alone directive.

### Syntax

―――――――――――― C / C++ ――――――――――――

The syntax of the **taskwait** construct is as follows:

```
#pragma omp taskwait newline
```

―――――――――――― C / C++ ――――――――――――

―――――――――――― Fortran ――――――――――――

The syntax of the **taskwait** construct is as follows:

```
!$omp taskwait
```

―――――――――――― Fortran ――――――――――――

### Binding

The **taskwait** region binds to the current task region. The binding thread set of the **taskwait** region is the current team.

### Description

The **taskwait** region includes an implicit task scheduling point in the current task region. The current task region is suspended at the task scheduling point until all child tasks that it generated before the **taskwait** region complete execution.

# 2.13.5 `taskgroup` Construct

### Summary

The **taskgroup** construct specifies a wait on completion of child tasks of the current task and their descendent tasks.

**Syntax**

C / C++

The syntax of the **taskgroup** construct is as follows:

```
#pragma omp taskgroup new-line
    structured-block
```

C / C++

Fortran

The syntax of the **taskgroup** construct is as follows:

```
!$omp taskgroup
    structured-block
!$omp end taskgroup
```

Fortran

**Binding**

A **taskgroup** region binds to the current task region. A **taskgroup** region binds to the innermost enclosing **parallel** region.

**Description**

When a thread encounters a **taskgroup** construct, it starts executing the region. All child tasks generated in the **taskgroup** region and all of their descendants that bind to the same **parallel** region as the **taskgroup** region are part of the *taskgroup set* associated with the **taskgroup** region.

There is an implicit task scheduling point at the end of the **taskgroup** region. The current task is suspended at the task scheduling point until all tasks in the *taskgroup set* complete execution.

**Cross References**

- Task scheduling, see Section 2.9.5 on page 94.

**2.13.6** `atomic` **Construct**

2 **Summary**

3  The `atomic` construct ensures that a specific storage location is accessed atomically, rather than
4  exposing it to the possibility of multiple, simultaneous reading and writing threads that may result
5  in indeterminate values.

6 **Syntax**

7  In the following syntax, *atomic-clause* is a clause that indicates the semantics for which atomicity is
8  enforced and is one of the following:

9      `read`

10     `write`

11     `update`

12     `capture`

───────────────────── C / C++ ─────────────────────

13  The syntax of the `atomic` construct takes one of the following forms:

```
#pragma omp atomic [seq_cst[,]] atomic-clause [[,]seq_cst] new-line
    expression-stmt
```

14  or

```
#pragma omp atomic [seq_cst] new-line
    expression-stmt
```

15  or

```
#pragma omp atomic [seq_cst[,]] capture [[,]seq_cst] new-line
    structured-block
```

16  where *expression-stmt* is an expression statement with one of the following forms:

17  • If *atomic-clause* is `read`:
18    *v* = *x*;

19  • If *atomic-clause* is `write`:
20    *x* = *expr*;

- If *atomic-clause* is **update** or not present:
  *x*++;
  *x*--;
  ++*x*;
  --*x*;
  *x binop*= *expr*;
  *x* = *x binop expr*;
  *x* = *expr binop x*;

- If *atomic-clause* is **capture**:
  *v* = *x*++;
  *v* = *x*--;
  *v* = ++*x*;
  *v* = --*x*;
  *v* = *x binop*= *expr*;
  *v* = *x* = *x binop expr*;
  *v* = *x* = *expr binop x*;

  and where *structured-block* is a structured block with one of the following forms:

  {*v* = *x*; *x binop*= *expr*;}
  {*x binop*= *expr*; *v* = *x*;}
  {*v* = *x*; *x* = *x binop expr*;}
  {*v* = *x*; *x* = *expr binop x*;}
  {*x* = *x binop expr*; *v* = *x*;}
  {*x* = *expr binop x*; *v* = *x*;}
  {*v* = *x*; *x* = *expr*;}
  {*v* = *x*; *x*++;}
  {*v* = *x*; ++*x*;}
  {++*x*; *v* = *x*;}
  {*x*++; *v* = *x*;}
  {*v* = *x*; *x*--;}
  {*v* = *x*; --*x*;}
  {--*x*; *v* = *x*;}
  {*x*--; *v* = *x*;}

In the preceding expressions:

- *x* and *v* (as applicable) are both *l-value* expressions with scalar type.

- During the execution of an atomic region, multiple syntactic occurrences of *x* must designate the same storage location.

- Neither of *v* and *expr* (as applicable) may access the storage location designated by *x*.

1. • Neither of *x* and *expr* (as applicable) may access the storage location designated by *v*.

2. • *expr* is an expression with scalar type.

3. • *binop* is one of **+**, **\***, **−**, **/**, **&**, **^**, **|**, **<<**, or **>>**.

4. • *binop*, *binop***=**, **++**, and **−−** are not overloaded operators.

5. • The expression *x binop expr* must be numerically equivalent to *x binop (expr)*. This requirement
6.   is satisfied if the operators in *expr* have precedence greater than *binop*, or by using parentheses
7.   around *expr* or subexpressions of *expr*.

8. • The expression *expr binop x* must be numerically equivalent to *(expr) binop x*. This requirement
9.   is satisfied if the operators in *expr* have precedence equal to or greater than *binop*, or by using
10.   parentheses around *expr* or subexpressions of *expr*.

11. • For forms that allow multiple occurrences of *x*, the number of times that *x* is evaluated is
12.   unspecified.

━━━━━━━━━━━━━━━━━━━ C / C++ ━━━━━━━━━━━━━━━━━━━

━━━━━━━━━━━━━━━━━━━ Fortran ━━━━━━━━━━━━━━━━━━━

13. The syntax of the **atomic** construct takes any of the following forms:

```
!$omp atomic [seq_cst[,]] read [[,]seq_cst]
    capture-statement
[!$omp end atomic]
```

14. or

```
!$omp atomic [seq_cst[,]] write [[,]seq_cst]
    write-statement
[!$omp end atomic]
```

15. or

```
!$omp atomic [seq_cst[,]] update [[,]seq_cst]
    update-statement
[!$omp end atomic]
```

16. or

```
!$omp atomic [seq_cst]
    update-statement
[!$omp end atomic]
```

17. or

```
!$omp atomic [seq_cst[,]] capture [[,]seq_cst]
    update-statement
    capture-statement
!$omp end atomic
```

1    or

```
!$omp atomic [seq_cst[,]] capture [[,]seq_cst]
    capture-statement
    update-statement
!$omp end atomic
```

2    or

```
!$omp atomic [seq_cst[,]] capture [[,]seq_cst]
    capture-statement
    write-statement
!$omp end atomic
```

3    where *write-statement* has the following form (if *atomic-clause* is **capture** or **write**):

4        *x* **=** *expr*

5    where *capture-statement* has the following form (if *atomic-clause* is **capture** or **read**):

6        *v* **=** *x*

7    and where *update-statement* has one of the following forms (if *atomic-clause* is **update**,
8    **capture**, or not present):

9        *x* **=** *x operator expr*

10       *x* **=** *expr operator x*

11       *x* **=** *intrinsic_procedure_name* **(***x***,** *expr_list***)**

12       *x* **=** *intrinsic_procedure_name* **(***expr_list***,** *x***)**

13   In the preceding statements:

14   • *x* and *v* (as applicable) are both scalar variables of intrinsic type.

15   • *x* must not have the **ALLOCATABLE** attribute.

16   • During the execution of an atomic region, multiple syntactic occurrences of *x* must designate the
17     same storage location.

- None of *v*, *expr*, and *expr_list* (as applicable) may access the same storage location as *x*.

- None of *x*, *expr*, and *expr_list* (as applicable) may access the same storage location as *v*.

- *expr* is a scalar expression.

- *expr_list* is a comma-separated, non-empty list of scalar expressions. If *intrinsic_procedure_name* refers to **IAND**, **IOR**, or **IEOR**, exactly one expression must appear in *expr_list*.

- *intrinsic_procedure_name* is one of **MAX**, **MIN**, **IAND**, **IOR**, or **IEOR**.

- *operator* is one of **+**, **\***, **−**, **/**, **.AND.**, **.OR.**, **.EQV.**, or **.NEQV.**.

- The expression *x operator expr* must be numerically equivalent to *x operator (expr)*. This requirement is satisfied if the operators in *expr* have precedence greater than *operator*, or by using parentheses around *expr* or subexpressions of *expr*.

- The expression *expr operator x* must be numerically equivalent to *(expr) operator x*. This requirement is satisfied if the operators in *expr* have precedence equal to or greater than *operator*, or by using parentheses around *expr* or subexpressions of *expr*.

- *intrinsic_procedure_name* must refer to the intrinsic procedure name and not to other program entities.

- *operator* must refer to the intrinsic operator and not to a user-defined operator.

- All assignments must be intrinsic assignments.

- For forms that allow multiple occurrences of *x*, the number of times that *x* is evaluated is unspecified.

──────── Fortran ────────

**Binding**

If the size of *x* is 8, 16, 32, or 64 bits and *x* is aligned to a multiple of its size, the binding thread set for the **atomic** region is all threads on the device. Otherwise, the binding thread set for the **atomic** region is all threads in the contention group. **atomic** regions enforce exclusive access with respect to other **atomic** regions that access the same storage location *x* among all threads in the binding thread set without regard to the teams to which the threads belong.

**Description**

The **atomic** construct with the **read** clause forces an atomic read of the location designated by *x* regardless of the native machine word size.

The **atomic** construct with the **write** clause forces an atomic write of the location designated by *x* regardless of the native machine word size.

The **atomic** construct with the **update** clause forces an atomic update of the location designated by *x* using the designated operator or intrinsic. Note that when no clause is present, the semantics are equivalent to atomic update. Only the read and write of the location designated by *x* are performed mutually atomically. The evaluation of *expr* or *expr_list* need not be atomic with respect to the read or write of the location designated by *x*. No task scheduling points are allowed between the read and the write of the location designated by *x*.

The **atomic** construct with the **capture** clause forces an atomic update of the location designated by *x* using the designated operator or intrinsic while also capturing the original or final value of the location designated by *x* with respect to the atomic update. The original or final value of the location designated by *x* is written in the location designated by *v* depending on the form of the **atomic** construct structured block or statements following the usual language semantics. Only the read and write of the location designated by *x* are performed mutually atomically. Neither the evaluation of *expr* or *expr_list*, nor the write to the location designated by *v*, need be atomic with respect to the read or write of the location designated by *x*. No task scheduling points are allowed between the read and the write of the location designated by *x*.

Any **atomic** construct with a **seq_cst** clause forces the atomically performed operation to include an implicit flush operation without a list.

▼————————————————————————————————————————————————————————▼

Note – As with other implicit flush regions, Section 1.4.4 on page 20 reduces the ordering that must be enforced. The intent is that, when the analogous operation exists in C++11 or C11, a sequentially consistent **atomic** construct has the same semantics as a **memory_order_seq_cst** atomic operation in C++11/C11. Similarly, a non-sequentially consistent **atomic** construct has the same semantics as a **memory_order_relaxed** atomic operation in C++11/C11.

Unlike non-sequentially consistent **atomic** constructs, sequentially consistent **atomic** constructs preserve the interleaving (sequentially consistent) behavior of correct, data race free programs. However, they are not designed to replace the **flush** directive as a mechanism to enforce ordering for non-sequentially consistent **atomic** constructs, and attempts to do so require extreme caution. For example, a sequentially consistent **atomic write** construct may appear to be reordered with a subsequent non-sequentially consistent **atomic write** construct, since such reordering would not be observable by a correct program if the second write were outside an **atomic** directive.

▲————————————————————————————————————————————————————————▲

For all forms of the **atomic** construct, any combination of two or more of these **atomic** constructs enforces mutually exclusive access to the locations designated by *x* among threads in the binding thread set. To avoid race conditions, all accesses of the locations designated by *x* that could potentially occur in parallel must be protected with an **atomic** construct.

**atomic** regions do not guarantee exclusive access with respect to any accesses outside of **atomic** regions to the same storage location *x* even if those accesses occur during a **critical** or **ordered** region, while an OpenMP lock is owned by the executing task, or during the execution of a **reduction** clause.

However, other OpenMP synchronization can ensure the desired exclusive access. For example, a barrier following a series of atomic updates to *x* guarantees that subsequent accesses do not form a race with the atomic accesses.

A compliant implementation may enforce exclusive access between **atomic** regions that update different storage locations. The circumstances under which this occurs are implementation defined.

If the storage location designated by *x* is not size-aligned (that is, if the byte alignment of *x* is not a multiple of the size of *x*), then the behavior of the **atomic** region is implementation defined.

**Restrictions**

The following restrictions apply to the **atomic** construct:

- At most one **seq_cst** clause may appear on the construct.

———————————————————— C / C++ ————————————————————

- All atomic accesses to the storage locations designated by *x* throughout the program are required to have a compatible type.

———————————————————— C / C++ ————————————————————

———————————————————— Fortran ————————————————————

- All atomic accesses to the storage locations designated by *x* throughout the program are required to have the same type and type parameters.

———————————————————— Fortran ————————————————————

- OpenMP constructs may not be encountered during execution of an **atomic** region.

**Cross References**

- **critical** construct, see Section 2.13.2 on page 149.
- **barrier** construct, see Section 2.13.3 on page 151.
- **flush** construct, see Section 2.13.7 on page 162.
- **ordered** construct, see Section 2.13.8 on page 166.
- **reduction** clause, see Section 2.15.3.6 on page 201.
- lock routines, see Section 3.3 on page 270.

# 2.13.7 `flush` Construct

**Summary**

The **flush** construct executes the OpenMP flush operation. This operation makes a thread's temporary view of memory consistent with memory and enforces an order on the memory operations of the variables explicitly specified or implied. See the memory model description in Section 1.4 on page 17 for more details. The **flush** construct is a stand-alone directive.

**Syntax**

──────────────────────── C / C++ ────────────────────────

The syntax of the **flush** construct is as follows:

```
#pragma omp flush [(list)] new-line
```

──────────────────────── C / C++ ────────────────────────
──────────────────────── Fortran ────────────────────────

The syntax of the **flush** construct is as follows:

```
!$omp flush [(list)]
```

──────────────────────── Fortran ────────────────────────

**Binding**

The binding thread set for a **flush** region is the encountering thread. Execution of a **flush** region affects the memory and the temporary view of memory of only the thread that executes the region. It does not affect the temporary view of other threads. Other threads must themselves execute a flush operation in order to be guaranteed to observe the effects of the encountering thread's flush operation

**Description**

A `flush` construct without a list, executed on a given thread, operates as if the whole thread-visible data state of the program, as defined by the base language, is flushed. A `flush` construct with a list applies the flush operation to the items in the list, and does not return until the operation is complete for all specified list items. An implementation may implement a `flush` with a list by ignoring the list, and treating it the same as a `flush` without a list.

---- C / C++ ----

If a pointer is present in the list, the pointer itself is flushed, not the memory block to which the pointer refers.

---- C / C++ ----
---- Fortran ----

If the list item or a subobject of the list item has the `POINTER` attribute, the allocation or association status of the `POINTER` item is flushed, but the pointer target is not. If the list item is a Cray pointer, the pointer is flushed, but the object to which it points is not. If the list item is of type `C_PTR`, the variable is flushed, but the storage that corresponds to that address is not flushed. If the list item or the subobject of the list item has the `ALLOCATABLE` attribute and has an allocation status of currently allocated, the allocated variable is flushed; otherwise the allocation status is flushed.

---- Fortran ----

Note – Use of a `flush` construct with a list is extremely error prone and users are strongly discouraged from attempting it. The following examples illustrate the ordering properties of the flush operation. In the following incorrect pseudocode example, the programmer intends to prevent simultaneous execution of the protected section by the two threads, but the program does not work properly because it does not enforce the proper ordering of the operations on variables **a** and **b**. Any shared data accessed in the protected section is not guaranteed to be current or consistent during or after the protected section. The atomic notation in the pseudocode in the following two examples indicates that the accesses to **a** and **b** are `ATOMIC` writes and captures. Otherwise both examples would contain data races and automatically result in unspecified behavior.

The problem with this example is that operations on variables **a** and **b** are not ordered with respect to each other. For instance, nothing prevents the compiler from moving the flush of **b** on thread 1 or the flush of **a** on thread 2 to a position completely after the protected section (assuming that the protected section on thread 1 does not reference **b** and the protected section on thread 2 does not reference **a**). If either re-ordering happens, both threads can simultaneously execute the protected section.

The following pseudocode example correctly ensures that the protected section is executed by not more than one of the two threads at any one time. Execution of the protected section by neither thread is considered correct in this example. This occurs if both flushes complete prior to either thread executing its **if** statement.

1    The compiler is prohibited from moving the flush at all for either thread, ensuring that the
2    respective assignment is complete and the data is flushed before the **if** statement is executed.

3    A **flush** region without a list is implied at the following locations:

4    • During a barrier region.

5    • At entry to a **target update** region whose corresponding construct has a **to** clause.

6    • At exit from a **target update** region whose corresponding construct has a **from** clause.

7    • At entry to and exit from **parallel**, **critical**, **target** and **target data** regions.

8    • At entry to and exit from an **ordered** region, if a **threads** clause or a **depend** clause is
9      present, or if no clauses are present.

10   • At entry to a **target enter data** region.

11   • At exit from a **target exit data** region.

12   • At exit from worksharing regions unless a **nowait** is present.

13   • At entry to and exit from the **atomic** operation (read, write, update, or capture) performed in a
14     sequentially consistent atomic region.

15   • During **omp_set_lock** and **omp_unset_lock** regions.

16   • During **omp_test_lock**, **omp_set_nest_lock**, **omp_unset_nest_lock** and
17     **omp_test_nest_lock** regions, if the region causes the lock to be set or unset.

18   • Immediately before and immediately after every task scheduling point.

19   • During a **cancel** or **cancellation point** region, if the *cancel-var* ICV is *true* and
20     cancellation has been activated.

21   A **flush** region with a list is implied at the following locations:

22   • At entry to and exit from the **atomic** operation (read, write, update, or capture) performed in a
23     non-sequentially consistent **atomic** region, where the list contains only the storage location
24     designated as x according to the description of the syntax of the **atomic** construct in
25     Section 2.13.6 on page 155.

26   Note – A **flush** region is not implied at the following locations:

27   • At entry to worksharing regions.

28   • At entry to or exit from a **master** region.

# 2.13.8 `ordered` **Construct**

**Summary**

The **ordered** construct either specifies a structured block in a loop, **simd**, or loop SIMD region that will be executed in the order of the loop iterations, or it is a stand-alone directive that specifies cross-iteration dependences in a doacross loop nest. The **ordered** construct sequentializes and orders the execution of **ordered** regions while allowing code outside the region to run in parallel.

**Syntax**

———————————————————————————————— C / C++ ————————————————————————————————

The syntax of the **ordered** construct is as follows:

```
#pragma omp ordered [clause[ [,] clause] ] new-line
    structured-block
```

where *clause* is one of the following:

**threads**

**simd**

or

```
#pragma omp ordered clause [[[,] clause] ... ] new-line
```

where *clause* is one of the following:

**depend(source)**

**depend(sink : *vec*)**

———————————————————————————————— C / C++ ————————————————————————————————

1 The syntax of the **ordered** construct is as follows:

```
!$omp ordered [clause[ [,] clause] ]
    structured-block
!$omp end ordered
```

2 where *clause* is one of the following:

3    **threads**

4    **simd**

5 or

```
!$omp ordered clause [[[,] clause] ... ]
```

6 where *clause* is one of the following:

7    **depend(source)**

8    **depend(sink :** *vec***)**

9 If the **depend** clause is specified, the **ordered** construct is a stand-alone directive.

## Binding

11 The binding thread set for an **ordered** region is the current team. An **ordered** region binds to
12 the innermost enclosing **simd** or loop SIMD region if the **simd** clause is present, and otherwise it
13 binds to the innermost enclosing loop region. **ordered** regions that bind to different regions
14 execute independently of each other.

**Description**

If no clause is specified, the **ordered** construct behaves as if the **threads** clause had been specified. If the **threads** clause is specified, the threads in the team executing the loop region execute **ordered** regions sequentially in the order of the loop iterations. If any **depend** clauses are specified then those clauses specify the order in which the threads in the team execute **ordered** regions. If the **simd** clause is specified, the **ordered** regions encountered by any thread will use only a single SIMD lane to execute the **ordered** regions in the order of the loop iterations.

When the thread executing the first iteration of the loop encounters an **ordered** construct, it can enter the **ordered** region without waiting. When a thread executing any subsequent iteration encounters an **ordered** construct without a **depend** clause, it waits at the beginning of the **ordered** region until execution of all **ordered** regions belonging to all previous iterations has completed. When a thread executing any subsequent iteration encounters an **ordered** construct with one or more **depend(sink:***vec***)** clauses, it waits until its dependences on all valid iterations specified by the **depend** clauses are satisfied before it completes execution of the **ordered** region. A specific dependence is satisfied when a thread executing the corresponding iteration encounters an **ordered** construct with a **depend(source)** clause.

**Restrictions**

Restrictions to the **ordered** construct are as follows:

- At most one **threads** clause can appear on an **ordered** construct.

- At most one **simd** clause can appear on an **ordered** construct.

- At most one **depend(source)** clause can appear on an **ordered** construct.

- Either **depend(sink:***vec***)** clauses or **depend(source)** clauses may appear on an **ordered** construct, but not both.

- The loop or loop SIMD region to which an **ordered** region arising from an **ordered** construct without a **depend** clause binds must have an **ordered** clause without the parameter specified on the corresponding loop or loop SIMD directive.

- The loop region to which an **ordered** region arising from an **ordered** construct with any **depend** clauses binds must have an **ordered** clause with the parameter specified on the corresponding loop directive.

- An **ordered** construct with the **depend** clause specified must be closely nested inside a loop (or parallel loop) construct.

- An **ordered** region arising from an **ordered** construct with the **simd** clause specified must be closely nested inside a **simd** or loop SIMD region.

- An **ordered** region arising from an **ordered** construct with both the **simd** and **threads** clauses must be closely nested inside a loop SIMD region.

- During execution of an iteration of a loop or a loop nest within a loop, **simd**, or loop SIMD region, a thread must not execute more than one **ordered** region arising from an **ordered** construct without a **depend** clause.

---- C++ ----

- A throw executed inside a **ordered** region must cause execution to resume within the same **ordered** region, and the same thread that threw the exception must catch it.

---- C++ ----

### Cross References

- loop construct, see Section 2.7.1 on page 56.

- **simd** construct, see Section 2.8.1 on page 72.

- parallel loop construct, see Section 2.11.1 on page 124.

- **depend** Clause, see Section 2.13.9 on page 169

## 2.13.9  depend **Clause**

### Summary

The **depend** clause enforces additional constraints on the scheduling of tasks or loop iterations. These constraints establish dependences only between sibling tasks or between loop iterations.

### Syntax

The syntax of the **depend** clause is as follows:

> **depend(**_dependence-type_ **:** _list_**)**

where _dependence-type_ is one of the following:

 **in**

 **out**

 **inout**

or

```
depend(dependence-type)
```

1     where *dependence-type* is:

2         **source**

3     or

```
depend(dependence-type : vec)
```

4     where *dependence-type* is:

5         **sink**

6     and where *vec* is the iteration vector, which has the form:

7     $x_1 [\pm d_1], x_2 [\pm d_2], \ldots, x_n [\pm d_n]$

8     where *n* is the value specified by the **ordered** clause in the loop directive, $x_i$ denotes the loop
9     iteration variable of the *i*-th nested loop associated with the loop directive, and $d_i$ is a constant
10    non-negative integer.

## Description

12    Task dependences are derived from the *dependence-type* of a **depend** clause and its list items
13    when *dependence-type* is **in**, **out**, or **inout**.

14    For the **in** *dependence-type*, if the storage location of at least one of the list items is the same as the
15    storage location of a list item appearing in an **out** or **inout** *dependence-type* list of a **task**
16    construct from which a sibling task was previously generated, then the generated task will be a
17    dependent task of that sibling task.

18    For the **out** and **inout** *dependence-types*, if the storage location of at least one of the list items is
19    the same as the storage location of a list item appearing in an **in**, **out**, or **inout** *dependence-type*
20    list of a **task** construct from which a sibling task was previously generated, then the generated task
21    will be a dependent task of that sibling task.

─────────────────────  Fortran  ─────────────────────

22    If a list item has the **ALLOCATABLE** attribute and its allocation status is "not currently allocated",
23    the behavior is unspecified. If a list item has the **POINTER** attribute and its association status is
24    disassociated or undefined, the behavior is unspecified.

─────────────────────  Fortran  ─────────────────────

25    The list items that appear in the **depend** clause may include array sections.

The **source** *dependence-type* specifies the satisfaction of cross-iteration dependences that arise
from the current iteration.

The **sink** *dependence-type* specifies a cross-iteration dependence, where the iteration vector *vec*
indicates the iteration that satisfies the dependence.

If the iteration vector *vec* does not occur in the iteration space, the **depend** clause is ignored. If all
**depend** clauses on an **ordered** construct are ignored then the construct is ignored.

**Restrictions**

Restrictions to the **depend** clause are as follows:

- List items used in **depend** clauses of the same task or sibling tasks must indicate identical
  storage locations or disjoint storage locations.

- List items used in **depend** clauses cannot be zero-length array sections.

- A variable that is part of another variable (such as an element of a structure) but is not an array
  element or an array section cannot appear in a **depend** clause.

- For a *vec* element of **sink** *dependence-type* of the form $x_i + d_i$ or $x_i - d_i$ if the loop iteration
  variable $x_i$ has an integral or pointer type, the expression $x_i + d_i$ or $x_i - d_i$ for any value of the
  loop iteration variable $x_i$ that can encounter the **ordered** construct must be computable in the
  loop iteration variable's type without overflow.

--- C++ ---

- For a *vec* element of **sink** *dependence-type* of the form $x_i + d_i$ or $x_i - d_i$ if the loop iteration
  variable $x_i$ is of a random access iterator type other than pointer type, the expression $( x_i - lb_i )$
  $+ d_i$ or $( x_i - lb_i ) - d_i$ for any value of the loop iteration variable $x_i$ that can encounter the
  **ordered** construct must be computable in the type that would be used by *std::distance* applied
  to variables of the type of $x_i$ without overflow.

--- C++ ---

**Cross References**

- Array sections, see Section 2.4 on page 44.
- **task** construct, see Section 2.9.1 on page 83.
- Task scheduling constraints, see Section 2.9.5 on page 94.
- **ordered** construct, see Section 2.13.8 on page 166.

# 2.14 Cancellation Constructs

## 2.14.1 **cancel** Construct

**Summary**

The **cancel** construct activates cancellation of the innermost enclosing region of the type specified. The **cancel** construct is a stand-alone directive.

**Syntax**

C / C++

The syntax of the **cancel** construct is as follows:

```
#pragma omp cancel construct-type-clause [ [ , ] if-clause] new-line
```

where *construct-type-clause* is one of the following:

    **parallel**

    **sections**

    **for**

    **taskgroup**

and *if-clause* is

    **if (***scalar-expression***)**

C / C++

1 The syntax of the **cancel** construct is as follows:

```
!$omp cancel construct-type-clause [ [ , ] if-clause]
```

2 where *construct-type-clause* is one of the following:

3     **parallel**

4     **sections**

5     **do**

6     **taskgroup**

7 and *if-clause* is

8     **if (***scalar-logical-expression***)**

## Binding

10 The binding thread set of the **cancel** region is the current team. The binding region of the
11 **cancel** region is the innermost enclosing region of the type corresponding to the
12 *construct-type-clause* specified in the directive (that is, the innermost **parallel**, **sections**,
13 loop, or **taskgroup** region).

## Description

15 The **cancel** construct activates cancellation of the binding region only if the *cancel-var* ICV is
16 *true*, in which case the **cancel** construct causes the encountering task to continue execution at the
17 end of the binding region if *construct-type-clause* is **parallel**, **for**, **do**, or **sections**. If the
18 *cancel-var* ICV is *true* and *construct-type-clause* is **taskgroup**, the encountering task continues
19 execution at the end of the current task region. If the *cancel-var* ICV is *false*, the **cancel**
20 construct is ignored.

21 Threads check for active cancellation only at cancellation points that are implied at the following
22 locations:

23 • **cancel** regions;

24 • **cancellation point** regions;

25 • **barrier** regions;

26 • implicit barriers regions.

1    When a thread reaches one of the above cancellation points and if the *cancel-var* ICV is *true*, then:

2    • If the thread is at a **cancel** or **cancellation point** region and *construct-type-clause* is
3      **parallel**, **for**, **do**, or **sections**, the thread continues execution at the end of the canceled
4      region if cancellation has been activated for the innermost enclosing region of the type specified.

5    • If the thread is at a **cancel** or **cancellation point** region and *construct-type-clause* is
6      **taskgroup**, the encountering task checks for active cancellation of all of the *taskgroup sets* to
7      which the encountering task belongs, and continues execution at the end of the current task
8      region if cancellation has been activated for any of the *taskgroup sets*.

9    • If the encountering task is at a barrier region, the encountering task checks for active cancellation
10     of the innermost enclosing **parallel** region. If cancellation has been activated, then the
11     encountering task continues execution at the end of the canceled region.

12   Note – If one thread activates cancellation and another thread encounters a cancellation point, the
13   order of execution between the two threads is non-deterministic. Whether the thread that
14   encounters a cancellation point detects the activated cancellation depends on the underlying
15   hardware and operating system.

16   When cancellation of tasks is activated through the **cancel taskgroup** construct, the tasks that
17   belong to the *taskgroup set* of the innermost enclosing **taskgroup** region will be canceled. The
18   task that encountered the **cancel taskgroup** construct continues execution at the end of its
19   **task** region, which implies completion of that task. Any task that belongs to the innermost
20   enclosing **taskgroup** and has already begun execution must run to completion or until a
21   cancellation point is reached. Upon reaching a cancellation point and if cancellation is active, the
22   task continues execution at the end of its **task** region, which implies the task's completion. Any
23   task that belongs to the innermost enclosing **taskgroup** and that has not begun execution may be
24   discarded, which implies its completion.

25   When cancellation is active for a **parallel**, **sections**, or loop region, each thread of the
26   binding thread set resumes execution at the end of the canceled region if a cancellation point is
27   encountered. If the canceled region is a **parallel** region, any tasks that have been created by a
28   **task** construct and their descendent tasks are canceled according to the above **taskgroup**
29   cancellation semantics. If the canceled region is a **sections**, or loop region, no task cancellation
30   occurs.

─────────────────── C++ ───────────────────

31   The usual C++ rules for object destruction are followed when cancellation is performed.

─────────────────── C++ ───────────────────

All private objects or subobjects with **ALLOCATABLE** attribute that are allocated inside the canceled construct are deallocated.

If the canceled construct contains a **reduction** or **lastprivate** clause, the final value of the **reduction** or **lastprivate** variable is undefined.

When an **if** clause is present on a **cancel** construct and the **if** expression evaluates to *false*, the **cancel** construct does not activate cancellation. The cancellation point associated with the **cancel** construct is always encountered regardless of the value of the **if** expression.

Note – The programmer is responsible for releasing locks and other synchronization data structures that might cause a deadlock when a **cancel** construct is encountered and blocked threads cannot be canceled. The programmer is also responsible for ensuring proper synchronizations to avoid deadlocks that might arise from cancellation of OpenMP regions that contain OpenMP synchronization constructs.

**Restrictions**

The restrictions to the **cancel** construct are as follows:

- The behavior for concurrent cancellation of a region and a region nested within it is unspecified.

- If *construct-type-clause* is **taskgroup**, the **cancel** construct must be closely nested inside a **task** construct and the **cancel** region must be closely nested inside a **taskgroup** region. If *construct-type-clause* is **sections**, the **cancel** construct must be closely nested inside a **sections** or **section** construct. Otherwise, the **cancel** construct must be closely nested inside an OpenMP construct that matches the type specified in *construct-type-clause* of the **cancel** construct.

- A worksharing construct that is canceled must not have a **nowait** clause.

- A loop construct that is canceled must not have an **ordered** clause.

- During execution of a construct that may be subject to cancellation, a thread must not encounter an orphaned cancellation point. That is, a cancellation point must only be encountered within that construct and must not be encountered elsewhere in its region.

**Cross References**

- *cancel-var* ICV, see Section .

- **cancellation point** construct, see Section .

- **omp_get_cancellation** routine, see Section .

## 2.14.2 **cancellation point Construct**

**Summary**

The **cancellation point** construct introduces a user-defined cancellation point at which implicit or explicit tasks check if cancellation of the innermost enclosing region of the type specified has been activated. The **cancellation point** construct is a stand-alone directive.

**Syntax**

———————————————— C / C++ ————————————————

The syntax of the **cancellation point** construct is as follows:

```
#pragma omp cancellation point construct-type-clause new-line
```

where *construct-type-clause* is one of the following:

**parallel**

**sections**

**for**

**taskgroup**

———————————————— C / C++ ————————————————

1    The syntax of the **cancellation point** construct is as follows:

```
!$omp cancellation point construct-type-clause
```

2    where *construct-type-clause* is one of the following:

3         **parallel**

4         **sections**

5         **do**

6         **taskgroup**

7    **Binding**

8    The binding thread set of the **cancellation point** construct is the current team. The binding
9    region of the **cancellation point** region is the innermost enclosing region of the type
10   corresponding to the *construct-type-clause* specified in the directive (that is, the innermost
11   **parallel**, **sections**, loop, or **taskgroup** region).

12   **Description**

13   This directive introduces a user-defined cancellation point at which an implicit or explicit task must
14   check if cancellation of the innermost enclosing region of the type specified in the clause has been
15   requested. This construct does not implement any synchronization between threads or tasks.

16   When an implicit or explicit task reaches a user-defined cancellation point and if the *cancel-var*
17   ICV is *true*, then:

18   • If the *construct-type-clause* of the encountered **cancellation point** construct is
19     **parallel**, **for**, **do**, or **sections**, the thread continues execution at the end of the canceled
20     region if cancellation has been activated for the innermost enclosing region of the type specified.

21   • If the *construct-type-clause* of the encountered **cancellation point** construct is
22     **taskgroup**, the encountering task checks for active cancellation of all *taskgroup sets* to which
23     the encountering task belongs and continues execution at the end of the current task region if
24     cancellation has been activated for any of them.

**Restrictions**

- A **cancellation point** construct for which *construct-type-clause* is **taskgroup** must be closely nested inside a **task** construct, and the **cancellation point** region must be closely nested inside a **taskgroup** region. A **cancellation point** construct for which *construct-type-clause* is **sections** must be closely nested inside a **sections** or **section** construct. Otherwise, a **cancellation point** construct must be closely nested inside an OpenMP construct that matches the type specified in *construct-type-clause*.

**Cross References**

- *cancel-var* ICV, see Section 2.3.1 on page 36.
- **cancel** construct, see Section 2.14.1 on page 172.
- **omp_get_cancellation** routine, see Section 3.2.9 on page 240.

# 2.15 Data Environment

This section presents a directive and several clauses for controlling the data environment during the execution of **target**, **teams**, **parallel**, **simd**, task generating, and worksharing regions.

- Section 2.15.1 on page 179 describes how the data-sharing attributes of variables referenced in **target**, **teams**, **parallel**, **simd**, task generating, and worksharing regions are determined.
- The **threadprivate** directive, which is provided to create threadprivate memory, is described in Section 2.15.2 on page 183.
- Clauses that may be specified on directives to control the data-sharing attributes of variables referenced in **target**, **teams**, **parallel**, **simd**, task generating, or worksharing constructs are described in Section 2.15.3 on page 188
- Clauses that may be specified on directives to copy data values from private or threadprivate variables on one thread to the corresponding variables on other threads in the team are described in Section 2.15.4 on page 211.
- Clauses that may be specified on directives to control the data-mapping of variables to a device data environment are described in Section 2.15.5.1 on page 216.

# 2.15.1 Data-sharing Attribute Rules

This section describes how the data-sharing attributes of variables referenced in **target**, **parallel**, **task**, **taskloop**, **simd**, and worksharing regions are determined. The following two cases are described separately:

- Section 2.15.1.1 on page 179 describes the data-sharing attribute rules for variables referenced in a construct.

- Section 2.15.1.2 on page 183 describes the data-sharing attribute rules for variables referenced in a region, but outside any construct.

## 2.15.1.1 Data-sharing Attribute Rules for Variables Referenced in a Construct

The data-sharing attributes of variables that are referenced in a construct can be *predetermined*, *explicitly determined*, or *implicitly determined*, according to the rules outlined in this section.

Specifying a variable on a **firstprivate**, **lastprivate**, **linear**, **reduction**, or **copyprivate** clause of an enclosed construct causes an implicit reference to the variable in the enclosing construct. Specifying a variable on a **map** clause of an enclosed construct may cause an implicit reference to the variable in the enclosing construct. Such implicit references are also subject to the data-sharing attribute rules outlined in this section.

Certain variables and objects have *predetermined* data-sharing attributes as follows:

———————————————————— C / C++ ————————————————————

- Variables appearing in **threadprivate** directives are threadprivate.

- Variables with automatic storage duration that are declared in a scope inside the construct are private.

- Objects with dynamic storage duration are shared.

- Static data members are shared.

- The loop iteration variable(s) in the associated *for-loop(s)* of a **for**, **parallel for**, **taskloop**, or **distribute** construct is (are) private.

- The loop iteration variable in the associated *for-loop* of a **simd** construct with just one associated *for-loop* is linear with a *linear-step* that is the increment of the associated *for-loop*.

- The loop iteration variables in the associated *for-loops* of a **simd** construct with multiple associated *for-loops* are lastprivate.

- Variables with static storage duration that are declared in a scope inside the construct are shared.

- If an array section is a list item in a **map** clause on the **target** construct and the array section is derived from a variable for which the type is pointer then that variable is firstprivate.

---------------------------------- C / C++ ----------------------------------
---------------------------------- Fortran ----------------------------------

- Variables and common blocks appearing in **threadprivate** directives are threadprivate.

- The loop iteration variable(s) in the associated *do-loop(s)* of a **do**, **parallel do**, **taskloop**, or **distribute** construct is (are) private.

- The loop iteration variable in the associated *do-loop* of a **simd** construct with just one associated *do-loop* is linear with a *linear-step* that is the increment of the associated *do-loop*.

- The loop iteration variables in the associated *do-loops* of a **simd** construct with multiple associated *do-loops* are lastprivate.

- A loop iteration variable for a sequential loop in a **parallel** or task generating construct is private in the innermost such construct that encloses the loop.

- Implied-do indices and **forall** indices are private.

- Cray pointees have the same the data-sharing attribute as the storage with which their Cray pointers are associated.

- Assumed-size arrays are shared.

- An associate name preserves the association with the selector established at the **ASSOCIATE** statement.

---------------------------------- Fortran ----------------------------------

Variables with predetermined data-sharing attributes may not be listed in data-sharing attribute clauses, except for the cases listed below. For these exceptions only, listing a predetermined variable in a data-sharing attribute clause is allowed and overrides the variable's predetermined data-sharing attributes.

- The loop iteration variable(s) in the associated *for-loop(s)* of a **for**, **parallel for**, **taskloop**, or **distribute** construct may be listed in a **private** or **lastprivate** clause.

- The loop iteration variable in the associated *for-loop* of a **simd** construct with just one associated *for-loop* may be listed in a **linear** clause with a *linear-step* that is the increment of the associated *for-loop*.

- The loop iteration variables in the associated *for-loops* of a **simd** construct with multiple associated *for-loops* may be listed in a **lastprivate** clause.

- Variables with **const**-qualified type having no mutable member may be listed in a **firstprivate** clause, even if they are static data members.

- The loop iteration variable(s) in the associated *do-loop(s)* of a **do**, **parallel do**, **taskloop**, or **distribute** construct may be listed in a **private** or **lastprivate** clause.

- The loop iteration variable in the associated *d*o-loop of a **simd** construct with just one associated *do-loop* may be listed in a **linear** clause with a *linear-step* that is the increment of the associated loop.

- The loop iteration variables in the associated *do-loops* of a **simd** construct with multiple associated *do-loops* may be listed in a **lastprivate** clause.

- Variables used as loop iteration variables in sequential loops in a **parallel** or task generating construct may be listed in data-sharing clauses on the construct itself, and on enclosed constructs, subject to other restrictions.

- Assumed-size arrays may be listed in a **shared** clause.

Additional restrictions on the variables that may appear in individual clauses are described with each clause in Section 2.15.3 on page 188.

Variables with *explicitly determined* data-sharing attributes are those that are referenced in a given construct and are listed in a data-sharing attribute clause on the construct.

Variables with *implicitly determined* data-sharing attributes are those that are referenced in a given construct, do not have predetermined data-sharing attributes, and are not listed in a data-sharing attribute clause on the construct.

Rules for variables with *implicitly determined* data-sharing attributes are as follows:

- In a **parallel**, **teams**, or task generating construct, the data-sharing attributes of these variables are determined by the **default** clause, if present (see Section 2.15.3.1 on page 189).

- In a **parallel** construct, if no **default** clause is present, these variables are shared.

- For constructs other than task generating constructs or **target** constructs, if no **default** clause is present, these variables reference the variables with the same names that exist in the enclosing context.

- In a **target** construct, variables that are not mapped after applying data-mapping attribute rules (see Section 2.15.5 on page 215) are firstprivate.

---
C++
---

- In an orphaned task generating construct, if no **default** clause is present, formal arguments passed by reference are firstprivate.

---
C++
---

---
Fortran
---

- In an orphaned task generating construct, if no **default** clause is present, dummy arguments are firstprivate.

---
Fortran
---

- In a task generating construct, if no **default** clause is present, a variable for which the data-sharing attribute is not determined by the rules above and that in the enclosing context is determined to be shared by all implicit tasks bound to the current team is shared.

- In a task generating construct, if no **default** clause is present, a variable for which the data-sharing attribute is not determined by the rules above is firstprivate.

Additional restrictions on the variables for which data-sharing attributes cannot be implicitly determined in a task generating construct are described in Section 2.15.3.4 on page 196.

### 2.15.1.2 Data-sharing Attribute Rules for Variables Referenced in a Region but not in a Construct

The data-sharing attributes of variables that are referenced in a region, but not in a construct, are determined as follows:

---
**C / C++**
---

- Variables with static storage duration that are declared in called routines in the region are shared.

- File-scope or namespace-scope variables referenced in called routines in the region are shared unless they appear in a **threadprivate** directive.

- Objects with dynamic storage duration are shared.

- Static data members are shared unless they appear in a **threadprivate** directive.

- In C++, formal arguments of called routines in the region that are passed by reference have the same data-sharing attributes as the associated actual arguments.

- Other variables declared in called routines in the region are private.

---
**C / C++**
---
**Fortran**
---

- Local variables declared in called routines in the region and that have the **save** attribute, or that are data initialized, are shared unless they appear in a **threadprivate** directive.

- Variables belonging to common blocks, or accessed by host or use association, and referenced in called routines in the region are shared unless they appear in a **threadprivate** directive.

- Dummy arguments of called routines in the region that are passed by reference have the same data-sharing attributes as the associated actual arguments.

- Cray pointees have the same data-sharing attribute as the storage with which their Cray pointers are associated.

- Implied-do indices, **forall** indices, and other local variables declared in called routines in the region are private.

---
**Fortran**
---

## 2.15.2 `threadprivate` Directive

### Summary

The **threadprivate** directive specifies that variables are replicated, with each thread having its own copy. The **threadprivate** directive is a declarative directive.

**Syntax**

────────────────────── C / C++ ──────────────────────

The syntax of the **threadprivate** directive is as follows:

```
#pragma omp threadprivate(list) new-line
```

where *list* is a comma-separated list of file-scope, namespace-scope, or static block-scope variables that do not have incomplete types.

────────────────────── C / C++ ──────────────────────
────────────────────── Fortran ──────────────────────

The syntax of the **threadprivate** directive is as follows:

```
!$omp threadprivate(list)
```

where *list* is a comma-separated list of named variables and named common blocks. Common block names must appear between slashes.

────────────────────── Fortran ──────────────────────

**Description**

Each copy of a threadprivate variable is initialized once, in the manner specified by the program, but at an unspecified point in the program prior to the first reference to that copy. The storage of all copies of a threadprivate variable is freed according to how static variables are handled in the base language, but at an unspecified point in the program.

A program in which a thread references another thread's copy of a threadprivate variable is non-conforming.

The content of a threadprivate variable can change across a task scheduling point if the executing thread switches to another task that modifies the variable. For more details on task scheduling, see Section 1.3 on page 14 and Section 2.9 on page 83.

In **parallel** regions, references by the master thread will be to the copy of the variable in the thread that encountered the **parallel** region.

During a sequential part references will be to the initial thread's copy of the variable. The values of data in the initial thread's copy of a threadprivate variable are guaranteed to persist between any two consecutive references to the variable in the program.

The values of data in the threadprivate variables of non-initial threads are guaranteed to persist between two consecutive active **parallel** regions only if all of the following conditions hold:

- Neither **parallel** region is nested inside another explicit **parallel** region.

- The number of threads used to execute both **parallel** regions is the same.

- The thread affinity policies used to execute both **parallel** regions are the same.

- The value of the *dyn-var* internal control variable in the enclosing task region is *false* at entry to both **parallel** regions.

If these conditions all hold, and if a threadprivate variable is referenced in both regions, then threads with the same thread number in their respective regions will reference the same copy of that variable.

$\blacktriangledown$ ——————————— C / C++ ——————————— $\blacktriangledown$

If the above conditions hold, the storage duration, lifetime, and value of a thread's copy of a threadprivate variable that does not appear in any **copyin** clause on the second region will be retained. Otherwise, the storage duration, lifetime, and value of a thread's copy of the variable in the second region is unspecified.

If the value of a variable referenced in an explicit initializer of a threadprivate variable is modified prior to the first reference to any instance of the threadprivate variable, then the behavior is unspecified.

$\blacktriangle$ ——————————— C / C++ ——————————— $\blacktriangle$

$\blacktriangledown$ ——————————— C++ ——————————— $\blacktriangledown$

The order in which any constructors for different threadprivate variables of class type are called is unspecified. The order in which any destructors for different threadprivate variables of class type are called is unspecified.

$\blacktriangle$ ——————————— C++ ——————————— $\blacktriangle$

$\blacktriangledown$ ——————————— Fortran ——————————— $\blacktriangledown$

A variable is affected by a **copyin** clause if the variable appears in the **copyin** clause or it is in a common block that appears in the **copyin** clause.

If the above conditions hold, the definition, association, or allocation status of a thread's copy of a threadprivate variable or a variable in a threadprivate common block, that is not affected by any **copyin** clause that appears on the second region, will be retained. Otherwise, the definition and association status of a thread's copy of the variable in the second region are undefined, and the allocation status of an allocatable variable will be implementation defined.

If a threadprivate variable or a variable in a threadprivate common block is not affected by any **copyin** clause that appears on the first **parallel** region in which it is referenced, the variable or any subobject of the variable is initially defined or undefined according to the following rules:

- If it has the **ALLOCATABLE** attribute, each copy created will have an initial allocation status of not currently allocated.

- If it has the **POINTER** attribute:

  – if it has an initial association status of disassociated, either through explicit initialization or default initialization, each copy created will have an association status of disassociated;

  – otherwise, each copy created will have an association status of undefined.

- If it does not have either the **POINTER** or the **ALLOCATABLE** attribute:

  – if it is initially defined, either through explicit initialization or default initialization, each copy created is so defined;

  – otherwise, each copy created is undefined.

———————————————————— Fortran ————————————————————

**Restrictions**

The restrictions to the **threadprivate** directive are as follows:

- A threadprivate variable must not appear in any clause except the **copyin**, **copyprivate**, **schedule**, **num_threads**, **thread_limit**, and **if** clauses.

- A program in which an untied task accesses threadprivate storage is non-conforming.

———————————————————— C / C++ ————————————————————

- A variable that is part of another variable (as an array or structure element) cannot appear in a **threadprivate** clause unless it is a static data member of a C++ class.

- A **threadprivate** directive for file-scope variables must appear outside any definition or declaration, and must lexically precede all references to any of the variables in its list.

- A **threadprivate** directive for namespace-scope variables must appear outside any definition or declaration other than the namespace definition itself, and must lexically precede all references to any of the variables in its list.

- Each variable in the list of a **threadprivate** directive at file, namespace, or class scope must refer to a variable declaration at file, namespace, or class scope that lexically precedes the directive.

- A **threadprivate** directive for static block-scope variables must appear in the scope of the variable and not in a nested scope. The directive must lexically precede all references to any of the variables in its list.

- Each variable in the list of a **threadprivate** directive in block scope must refer to a variable declaration in the same scope that lexically precedes the directive. The variable declaration must use the static storage-class specifier.

- If a variable is specified in a **threadprivate** directive in one translation unit, it must be specified in a **threadprivate** directive in every translation unit in which it is declared.

- The address of a threadprivate variable is not an address constant.

---------------------------------- C / C++ ----------------------------------
---------------------------------- C++ ----------------------------------

- A **threadprivate** directive for static class member variables must appear in the class definition, in the same scope in which the member variables are declared, and must lexically precede all references to any of the variables in its list.

- A threadprivate variable must not have an incomplete type or a reference type.

- A threadprivate variable with class type must have:

  - an accessible, unambiguous default constructor in case of default initialization without a given initializer;

  - an accessible, unambiguous constructor accepting the given argument in case of direct initialization;

  - an accessible, unambiguous copy constructor in case of copy initialization with an explicit initializer

---------------------------------- C++ ----------------------------------

---------------------------------- Fortran ----------------------------------

- A variable that is part of another variable (as an array or structure element) cannot appear in a **threadprivate** clause.

- The **threadprivate** directive must appear in the declaration section of a scoping unit in which the common block or variable is declared. Although variables in common blocks can be accessed by use association or host association, common block names cannot. This means that a common block name specified in a **threadprivate** directive must be declared to be a common block in the same scoping unit in which the **threadprivate** directive appears.

- If a **threadprivate** directive specifying a common block name appears in one program unit, then such a directive must also appear in every other program unit that contains a **COMMON** statement specifying the same name. It must appear after the last such **COMMON** statement in the program unit.

- If a threadprivate variable or a threadprivate common block is declared with the **BIND** attribute, the corresponding C entities must also be specified in a **threadprivate** directive in the C program.

- A blank common block cannot appear in a **threadprivate** directive.

- A variable can only appear in a **threadprivate** directive in the scope in which it is declared. It must not be an element of a common block or appear in an **EQUIVALENCE** statement.

1  • A variable that appears in a **threadprivate** directive must be declared in the scope of a
2    module or have the **SAVE** attribute, either explicitly or implicitly.

<div align="center">Fortran</div>

3  **Cross References**

4  • *dyn-var* ICV, see Section 2.3 on page 36.

5  • Number of threads used to execute a **parallel** region, see Section 2.5.1 on page 50.

6  • **copyin** clause, see Section 2.15.4.1 on page 211.

## 2.15.3  Data-Sharing Attribute Clauses

8  Several constructs accept clauses that allow a user to control the data-sharing attributes of variables
9  referenced in the construct. Data-sharing attribute clauses apply only to variables for which the
10 names are visible in the construct on which the clause appears.

11 Not all of the clauses listed in this section are valid on all directives. The set of clauses that is valid
12 on a particular directive is described with the directive.

13 Most of the clauses accept a comma-separated list of list items (see Section 2.1 on page 26). All list
14 items appearing in a clause must be visible, according to the scoping rules of the base language.
15 With the exception of the **default** clause, clauses may be repeated as needed. A list item that
16 specifies a given variable may not appear in more than one clause on the same directive, except that
17 a variable may be specified in both **firstprivate** and **lastprivate** clauses.

<div align="center">C++</div>

18 If a variable referenced in a data-sharing attribute clause has a type derived from a template, and
19 there are no other references to that variable in the program, then any behavior related to that
20 variable is unspecified.

<div align="center">C++</div>

A named common block may be specified in a list by enclosing the name in slashes. When a named common block appears in a list, it has the same meaning as if every explicit member of the common block appeared in the list. An explicit member of a common block is a variable that is named in a **COMMON** statement that specifies the common block name and is declared in the same scoping unit in which the clause appears.

Although variables in common blocks can be accessed by use association or host association, common block names cannot. As a result, a common block name specified in a data-sharing attribute clause must be declared to be a common block in the same scoping unit in which the data-sharing attribute clause appears.

When a named common block appears in a **private**, **firstprivate**, **lastprivate**, or **shared** clause of a directive, none of its members may be declared in another data-sharing attribute clause in that directive. When individual members of a common block appear in a **private**, **firstprivate**, **lastprivate**, **reduction**, or **linear** clause of a directive, the storage of the specified variables is no longer Fortran associated with the storage of the common block itself.

Fortran

## 2.15.3.1 **default** Clause

### Summary

The **default** clause explicitly determines the data-sharing attributes of variables that are referenced in a **parallel**, **teams**, or task generating construct and would otherwise be implicitly determined (see Section 2.15.1.1 on page 179).

### Syntax

C / C++

The syntax of the **default** clause is as follows:

```
default(shared | none)
```

C / C++

Fortran

The syntax of the **default** clause is as follows:

```
default(private | firstprivate | shared | none)
```

Fortran

**Description**

The **default(shared)** clause causes all variables referenced in the construct that have implicitly determined data-sharing attributes to be shared.

————————————————— Fortran —————————————————

The **default(firstprivate)** clause causes all variables in the construct that have implicitly determined data-sharing attributes to be firstprivate.

The **default(private)** clause causes all variables referenced in the construct that have implicitly determined data-sharing attributes to be private.

————————————————— Fortran —————————————————

The **default(none)** clause requires that each variable that is referenced in the construct, and that does not have a predetermined data-sharing attribute, must have its data-sharing attribute explicitly determined by being listed in a data-sharing attribute clause.

**Restrictions**

The restrictions to the **default** clause are as follows:

- Only a single **default** clause may be specified on a **parallel**, **task**, **taskloop** or **teams** directive.

### 2.15.3.2 **shared** Clause

**Summary**

The **shared** clause declares one or more list items to be shared by tasks generated by a **parallel**, **teams**, or task generating construct.

**Syntax**

The syntax of the **shared** clause is as follows:

```
shared(list)
```

**Description**

All references to a list item within a task refer to the storage area of the original variable at the point the directive was encountered.

The programmer must ensure, by adding proper synchronization, that storage shared by an explicit **task** region does not reach the end of its lifetime before the explicit **task** region completes its execution.

———————————— Fortran ————————————

The association status of a shared pointer becomes undefined upon entry to and on exit from the **parallel**, **teams**, or task generating construct if it is associated with a target or a subobject of a target that is in a **private**, **firstprivate**, **lastprivate**, or **reduction** clause in the construct.

Under certain conditions, passing a shared variable to a non-intrinsic procedure may result in the value of the shared variable being copied into temporary storage before the procedure reference, and back out of the temporary storage into the actual argument storage after the procedure reference. When this situation occurs is implementation defined.

Note – Use of intervening temporary storage may occur when the following three conditions hold regarding an actual argument in a reference to a non-intrinsic procedure:

1. The actual argument is one of the following:

    • A shared variable.

    • A subobject of a shared variable.

    • An object associated with a shared variable.

    • An object associated with a subobject of a shared variable.

2. The actual argument is also one of the following:

    • An array section.

    • An array section with a vector subscript.

    • An assumed-shape array.

    • A pointer array.

3. The associated dummy argument for this actual argument is an explicit-shape array or an assumed-size array.

1　These conditions effectively result in references to, and definitions of, the temporary storage during
2　the procedure reference. Any references to (or definitions of) the shared storage that is associated
3　with the dummy argument by any other task must be synchronized with the procedure reference to
4　avoid possible race conditions.

Fortran

5　**Restrictions**

6　The restrictions for the **shared** clause are as follows:

C

7　• A variable that is part of another variable (as an array or structure element) cannot appear in a
8　shared clause.

C

C++

9　• A variable that is part of another variable (as an array or structure element) cannot appear in a
10　**shared** clause except if the **shared** clause is associated with a construct within a class
11　non-static member function and the variable is an accessible data member of the object for which
12　the non-static member function is invoked.

C++

Fortran

13　• A variable that is part of another variable (as an array or structure element) cannot appear in a
14　shared clause.

Fortran

15　## 2.15.3.3 **private** Clause

16　**Summary**

17　The **private** clause declares one or more list items to be private to a task or to a SIMD lane.

18　**Syntax**

19　The syntax of the private clause is as follows:

```
private(list)
```

**Description**

Each task that references a list item that appears in a **private** clause in any statement in the construct receives a new list item. Each SIMD lane used in a **simd** construct that references a list item that appears in a private clause in any statement in the construct receives a new list item. Language-specific attributes for new list items are derived from the corresponding original list item. Inside the construct, all references to the original list item are replaced by references to the new list item. In the rest of the region, it is unspecified whether references are to the new list item or the original list item.

───────────── C++ ─────────────

If the construct is contained in a member function, it is unspecified anywhere in the region if accesses through the implicit **this** pointer refer to the new list item or the original list item.

───────────── C++ ─────────────

Therefore, if an attempt is made to reference the original item, its value after the region is also unspecified. If a SIMD construct or a task does not reference a list item that appears in a **private** clause, it is unspecified whether SIMD lanes or the task receive a new list item.

The value and/or allocation status of the original list item will change only:

- if accessed and modified via pointer,

- if possibly accessed in the region but outside of the construct,

- as a side effect of directives or clauses, or

───────────── Fortran ─────────────

- if accessed and modified via construct association.

───────────── Fortran ─────────────

List items that appear in a **private**, **firstprivate**, or **reduction** clause in a **parallel** construct may also appear in a **private** clause in an enclosed **parallel**, worksharing, **task**, **taskloop**, **simd**, or **target** construct.

List items that appear in a **private** or **firstprivate** clause in a **task** or **taskloop** construct may also appear in a **private** clause in an enclosed **parallel**, **task**, **taskloop**, **simd**, or **target** construct.

List items that appear in a **private**, **firstprivate**, **lastprivate**, or **reduction** clause in a worksharing construct may also appear in a **private** clause in an enclosed **parallel**, **task**, **simd**, or **target** construct.

1 A new list item of the same type, with automatic storage duration, is allocated for the construct.
2 The storage and thus lifetime of these list items lasts until the block in which they are created exits.
3 The size and alignment of the new list item are determined by the type of the variable. This
4 allocation occurs once for each task generated by the construct and once for each SIMD lane used
5 by the construct.

6 The new list item is initialized, or has an undefined initial value, as if it had been locally declared
7 without an initializer.

8 If the type of a list item is a reference to a type $T$ then the type will be considered to be $T$ for all
9 purposes of this clause.

10 The order in which any default constructors for different private variables of class type are called is
11 unspecified. The order in which any destructors for different private variables of class type are
12 called is unspecified.

13 If any statement of the construct references a list item, a new list item of the same type and type
14 parameters is allocated. This allocation occurs once for each task generated by the construct and
15 once for each SIMD lane used by the construct. The initial value of the new list item is undefined.
16 The initial status of a private pointer is undefined.

17 For a list item or the subobject of a list item with the **ALLOCATABLE** attribute:

18 • if the allocation status is "not currently allocated", the new list item or the subobject of the new
19   list item will have an initial allocation status of "not currently allocated".

20 • if the allocation status is "currently allocated", the new list item or the subobject of the new list
21   item will have an initial allocation status of "currently allocated".

22 • If the new list item or the subobject of the new list item is an array, its bounds will be the same as
23   those of the original list item or the subobject of the original list item.

24 A list item that appears in a **private** clause may be storage-associated with other variables when
25 the **private** clause is encountered. Storage association may exist because of constructs such as
26 **EQUIVALENCE** or **COMMON**. If $A$ is a variable appearing in a **private** clause on a construct and
27 $B$ is a variable that is storage-associated with $A$, then:

28 • The contents, allocation, and association status of $B$ are undefined on entry to the region.

29 • Any definition of $A$, or of its allocation or association status, causes the contents, allocation, and
30   association status of $B$ to become undefined.

- Any definition of *B*, or of its allocation or association status, causes the contents, allocation, and association status of *A* to become undefined.

A list item that appears in a **private** clause may be a selector of an **ASSOCIATE** construct. If the construct association is established prior to a **parallel** region, the association between the associate name and the original list item will be retained in the region.

Finalization of a list item of a finalizable type or subojects of a list item of a finalizable type occurs at the end of the region. The order in which any final subroutines for different variables of a finalizable type are called is unspecified.

———————————————— Fortran ————————————————

## Restrictions

The restrictions to the **private** clause are as follows:

———————————————— C ————————————————

- A variable that is part of another variable (as an array or structure element) cannot appear in a **private** clause.

———————————————— C ————————————————

———————————————— C++ ————————————————

- A variable that is part of another variable (as an array or structure element) cannot appear in a **private** clause except if the **private** clause is associated with a construct within a class non-static member function and the variable is an accessible data member of the object for which the non-static member function is invoked.

- A variable of class type (or array thereof) that appears in a **private** clause requires an accessible, unambiguous default constructor for the class type.

———————————————— C++ ————————————————

———————————————— C / C++ ————————————————

- A variable that appears in a **private** clause must not have a **const**-qualified type unless it is of class type with a **mutable** member. This restriction does not apply to the **firstprivate** clause.

- A variable that appears in a **private** clause must not have an incomplete type or be a reference to an incomplete type.

———————————————— C / C++ ————————————————

1  • A variable that is part of another variable (as an array or structure element) cannot appear in a
2  **private** clause.

3  • A variable that appears in a **private** clause must either be definable, or an allocatable variable.
4  This restriction does not apply to the **firstprivate** clause.

5  • Variables that appear in namelist statements, in variable format expressions, and in expressions
6  for statement function definitions, may not appear in a **private** clause.

7  • Pointers with the **INTENT(IN)** attribute may not appear in a **private** clause. This restriction
8  does not apply to the **firstprivate** clause.

## 9  2.15.3.4  **firstprivate** Clause

### 10  Summary

11  The **firstprivate** clause declares one or more list items to be private to a task, and initializes
12  each of them with the value that the corresponding original item has when the construct is
13  encountered.

### 14  Syntax

15  The syntax of the **firstprivate** clause is as follows:

```
firstprivate(list)
```

**Description**

The **firstprivate** clause provides a superset of the functionality provided by the **private** clause.

A list item that appears in a **firstprivate** clause is subject to the **private** clause semantics described in Section 2.15.3.3 on page 192, except as noted. In addition, the new list item is initialized from the original list item existing before the construct. The initialization of the new list item is done once for each task that references the list item in any statement in the construct. The initialization is done prior to the execution of the construct.

For a **firstprivate** clause on a **parallel**, **task**, **taskloop**, **target**, or **teams** construct, the initial value of the new list item is the value of the original list item that exists immediately prior to the construct in the task region where the construct is encountered. For a **firstprivate** clause on a worksharing construct, the initial value of the new list item for each implicit task of the threads that execute the worksharing construct is the value of the original list item that exists in the implicit task immediately prior to the point in time that the worksharing construct is encountered.

To avoid race conditions, concurrent updates of the original list item must be synchronized with the read of the original list item that occurs as a result of the **firstprivate** clause.

If a list item appears in both **firstprivate** and **lastprivate** clauses, the update required for **lastprivate** occurs after all the initializations for **firstprivate**.

---------------------------------- C / C++ ----------------------------------

For variables of non-array type, the initialization occurs by copy assignment. For an array of elements of non-array type, each element is initialized as if by assignment from an element of the original array to the corresponding element of the new array.

---------------------------------- C / C++ ----------------------------------
------------------------------------- C++ -------------------------------------

For variables of class type, a copy constructor is invoked to perform the initialization. The order in which copy constructors for different variables of class type are called is unspecified.

------------------------------------- C++ -------------------------------------

----------------------------------- Fortran -----------------------------------

If the original list item does not have the **POINTER** attribute, initialization of the new list items occurs as if by intrinsic assignment, unless the original list item has the allocation status of not currently allocated, in which case the new list items will have the same status.

If the original list item has the **POINTER** attribute, the new list items receive the same association status of the original list item as if by pointer assignment.

----------------------------------- Fortran -----------------------------------

**Restrictions**

The restrictions to the **firstprivate** clause are as follows:

- A list item that is private within a **parallel** region must not appear in a **firstprivate** clause on a worksharing construct if any of the worksharing regions arising from the worksharing construct ever bind to any of the **parallel** regions arising from the **parallel** construct.

- A list item that is private within a **teams** region must not appear in a **firstprivate** clause on a **distribute** construct if any of the **distribute** regions arising from the **distribute** construct ever bind to any of the **teams** regions arising from the **teams** construct.

- A list item that appears in a **reduction** clause of a **parallel** construct must not appear in a **firstprivate** clause on a worksharing, **task**, or **taskloop** construct if any of the worksharing or task regions arising from the worksharing, **task**, or **taskloop** construct ever bind to any of the **parallel** regions arising from the **parallel** construct.

- A list item that appears in a **reduction** clause of a **teams** construct must not appear in a **firstprivate** clause on a **distribute** construct if any of the **distribute** regions arising from the **distribute** construct ever bind to any of the **teams** regions arising from the **teams** construct.

- A list item that appears in a **reduction** clause of a worksharing construct must not appear in a **firstprivate** clause in a **task** construct encountered during execution of any of the worksharing regions arising from the worksharing construct.

---------------------------------- C++ ----------------------------------

- A variable of class type (or array thereof) that appears in a **firstprivate** clause requires an accessible, unambiguous copy constructor for the class type.

---------------------------------- C++ ----------------------------------

---------------------------------- C / C++ ----------------------------------

- A variable that appears in a **firstprivate** clause must not have an incomplete C/C++ type or be a reference to an incomplete type.

- If a list item in a **firstprivate** clause on a worksharing construct has a reference type then it must bind to the same object for all threads of the team.

---------------------------------- C / C++ ----------------------------------

---------------------------------- Fortran ----------------------------------

- Variables that appear in namelist statements, in variable format expressions, or in expressions for statement function definitions, may not appear in a **firstprivate** clause.

---------------------------------- Fortran ----------------------------------

1 **2.15.3.5 lastprivate Clause**

**Summary**

3
4 The **lastprivate** clause declares one or more list items to be private to an implicit task or to a
SIMD lane, and causes the corresponding original list item to be updated after the end of the region.

5 **Syntax**

6 The syntax of the **lastprivate** clause is as follows:

```
lastprivate(list)
```

7 **Description**

8
9 The **lastprivate** clause provides a superset of the functionality provided by the **private**
clause.

10
11
12
13
14 A list item that appears in a **lastprivate** clause is subject to the **private** clause semantics
described in Section 2.15.3.3 on page 192. In addition, when a **lastprivate** clause appears on
the directive that identifies a worksharing construct or a SIMD construct, the value of each new list
item from the sequentially last iteration of the associated loops, or the lexically last **section**
construct, is assigned to the original list item.

—————————————— C / C++ ——————————————

15
16 For an array of elements of non-array type, each element is assigned to the corresponding element
of the original array.

—————————————— C / C++ ——————————————

—————————————— Fortran ——————————————

17
18 If the original list item does not have the **POINTER** attribute, its update occurs as if by intrinsic
assignment.

If the original list item has the **POINTER** attribute, its update occurs as if by pointer assignment.

———————————————— Fortran ————————————————

List items that are not assigned a value by the sequentially last iteration of the loops, or by the lexically last **section** construct, have unspecified values after the construct. Unassigned subcomponents also have unspecified values after the construct.

The original list item becomes defined at the end of the construct if there is an implicit barrier at that point. To avoid race conditions, concurrent reads or updates of the original list item must be synchronized with the update of the original list item that occurs as a result of the **lastprivate** clause.

If the **lastprivate** clause is used on a construct to which **nowait** is applied, accesses to the original list item may create a data race. To avoid this, synchronization must be inserted to ensure that the sequentially last iteration or lexically last section construct has stored and flushed that list item.

If the **lastprivate** clause is used on a **distribute simd**, distribute parallel loop, or distribute parallel loop SIMD, accesses to the original list item may create a data race. To avoid this, synchronization must be inserted to ensure that the sequentially last iteration has stored and flushed that list item.

If a list item appears in both **firstprivate** and **lastprivate** clauses, the update required for **lastprivate** occurs after all initializations for **firstprivate**.

### Restrictions

The restrictions to the **lastprivate** clause are as follows:

- A list item that is private within a **parallel** region, or that appears in the **reduction** clause of a **parallel** construct, must not appear in a **lastprivate** clause on a worksharing construct if any of the corresponding worksharing regions ever binds to any of the corresponding **parallel** regions.

———————————————— C++ ————————————————

- A variable of class type (or array thereof) that appears in a **lastprivate** clause requires an accessible, unambiguous default constructor for the class type, unless the list item is also specified in a **firstprivate** clause.

- A variable of class type (or array thereof) that appears in a **lastprivate** clause requires an accessible, unambiguous copy assignment operator for the class type. The order in which copy assignment operators for different variables of class type are called is unspecified.

———————————————— C++ ————————————————

- A variable that appears in a **lastprivate** clause must not have a **const**-qualified type unless it is of class type with a **mutable** member.

- A variable that appears in a **lastprivate** clause must not have an incomplete C/C++ type or be a reference to an incomplete type.

- If a list item in a **lastprivate** clause on a worksharing construct has a reference type then it must bind to the same object for all threads of the team.

- A variable that appears in a **lastprivate** clause must be definable.

- If the original list item has the **ALLOCATABLE** attribute, the corresponding list item in the sequentially last iteration or lexically last section must have an allocation status of allocated upon exit from that iteration or section.

- Variables that appear in namelist statements, in variable format expressions, or in expressions for statement function definitions, may not appear in a **lastprivate** clause.

## 2.15.3.6 **reduction Clause**

### Summary

The **reduction** clause specifies a *reduction-identifier* and one or more list items. For each list item, a private copy is created in each implicit task or SIMD lane, and is initialized with the initializer value of the *reduction-identifier*. After the end of the region, the original list item is updated with the values of the private copies using the combiner associated with the *reduction-identifier*.

**Syntax**

─────────────── C / C++ ───────────────

2 The syntax of the **reduction** clause is as follows:

**reduction(***reduction-identifier* **:** *list***)**

3 where:

─────────────── C ───────────────

4 *reduction-identifier* is either an *identifier* or one of the following operators: **+**, **−**, **\***, **&**, **|**, ^, **&&** and
5 **||**

─────────────── C ───────────────

─────────────── C++ ───────────────

6 *reduction-identifier* is either an *id-expression* or one of the following operators: **+**, **−**, **\***, **&**, **|**, ^, **&&**
7 and **||**

─────────────── C++ ───────────────

8 Table 2.7 lists each *reduction-identifier* that is implicitly declared at every scope for arithmetic
9 types and its semantic initializer value. The actual initializer value is that value as expressed in the
10 data type of the reduction list item.

**TABLE 2.7:** Implicitly Declared C/C++ *reduction-identifiers*

| Identifier | Initializer | Combiner |
|---|---|---|
| **+** | **omp_priv = 0** | **omp_out += omp_in** |
| **\*** | **omp_priv = 1** | **omp_out \*= omp_in** |
| **−** | **omp_priv = 0** | **omp_out += omp_in** |
| **&** | **omp_priv =  0** | **omp_out &= omp_in** |
| **|** | **omp_priv = 0** | **omp_out |= omp_in** |
| ^ | **omp_priv = 0** | **omp_out ^= omp_in** |
| **&&** | **omp_priv = 1** | **omp_out = omp_in && omp_out** |
| **||** | **omp_priv = 0** | **omp_out = omp_in || omp_out** |

*table continued on next page*

| Identifier | Initializer | Combiner |
|---|---|---|
| `max` | `omp_priv =` *Least representable number in the reduction list item type* | `omp_out = omp_in > omp_out ?` `omp_in :  omp_out` |
| `min` | `omp_priv =` *Largest representable number in the reduction list item type* | `omp_out = omp_in < omp_out ?` `omp_in :  omp_out` |

◄───────────────── C / C++ ─────────────────►

▼───────────────── Fortran ─────────────────▼

The syntax of the **reduction** clause is as follows:

---

**reduction(***reduction-identifier* **:** *list***)**

---

where *reduction-identifier* is either a base language identifier, or a user-defined operator, or one of the following operators: **+**, **−**, **∗**, **.and.**, **.or.**, **.eqv.**, **.neqv.**, or one of the following intrinsic procedure names: **max**, **min**, **iand**, **ior**, **ieor**.

Table 2.8 lists each *reduction-identifier* that is implicitly declared for numeric and logical types and its semantic initializer value. The actual initializer value is that value as expressed in the data type of the reduction list item.

**TABLE 2.8:** Implicitly Declared Fortran *reduction-identifiers*

| Identifier | Initializer | Combiner |
|---|---|---|
| `+` | `omp_priv = 0` | `omp_out = omp_in + omp_out` |
| `*` | `omp_priv = 1` | `omp_out = omp_in * omp_out` |
| `−` | `omp_priv = 0` | `omp_out = omp_in + omp_out` |
| `.and.` | `omp_priv = .true.` | `omp_out = omp_in .and. omp_out` |
| `.or.` | `omp_priv = .false.` | `omp_out = omp_in .or. omp_out` |
| `.eqv.` | `omp_priv = .true.` | `omp_out = omp_in .eqv. omp_out` |

| Identifier | Initializer | Combiner |
|---|---|---|
| `.neqv.` | `omp_priv = .false.` | `omp_out = omp_in .neqv. omp_out` |
| `max` | `omp_priv =` *Least representable number in the reduction list item type* | `omp_out = max(omp_in, omp_out)` |
| `min` | `omp_priv =` *Largest representable number in the reduction list item type* | `omp_out = min(omp_in, omp_out)` |
| `iand` | `omp_priv =` *All bits on* | `omp_out = iand(omp_in, omp_out)` |
| `ior` | `omp_priv = 0` | `omp_out = ior(omp_in, omp_out)` |
| `ieor` | `omp_priv = 0` | `omp_out = ieor(omp_in, omp_out)` |

▲ ———————————————————— Fortran ———————————————————— ▲

In the above tables, `omp_in` and `omp_out` correspond to two identifiers that refer to storage of the type of the list item. `omp_out` holds the final value of the combiner operation.

Any *reduction-identifier* that is defined with the `declare reduction` directive is also valid. In that case, the initializer and combiner of the *reduction-identifier* are specified by the *initializer-clause* and the *combiner* in the `declare reduction` directive.

**Description**

The reduction clause can be used to perform some forms of recurrence calculations (involving mathematically associative and commutative operators) in parallel.

For `parallel` and worksharing constructs, a private copy of each list item is created, one for each implicit task, as if the `private` clause had been used. For the `simd` construct, a private copy of each list item is created, one for each SIMD lane as if the `private` clause had been used. For the `teams` construct, a private copy of each list item is created, one for each team in the league as if the `private` clause had been used. The private copy is then initialized as specified above. At the end of the region for which the `reduction` clause was specified, the original list item is updated by combining its original value with the final value of each of the private copies, using the combiner of the specified *reduction-identifier*.

If the original list item has the **POINTER** attribute, the private copy of the list item is associated with a private target.

———————————————————————— Fortran ————————————————————————

The *reduction-identifier* specified in the **reduction** clause must match a previously declared *reduction-identifier* of the same name and type for each of the list items. This match is done by means of a name lookup in the base language.

———————————————————————— C / C++ ————————————————————————

The list items that appear in the **reduction** clause may include array sections.

———————————————————————— C / C++ ————————————————————————
———————————————————————— C++ ————————————————————————

If the type is a derived class, then any *reduction-identifier* that matches its base classes is also a match, if there is no specific match for the type.

If the *reduction-identifier* is not an *id-expression*, then it is implicitly converted to one by prepending the keyword operator (for example, **+** becomes *operator***+**).

If the *reduction-identifier* is qualified then a qualified name lookup is used to find the declaration.

If the *reduction-identifier* is unqualified then an *argument-dependent name lookup* must be performed using the type of each list item.

———————————————————————— C++ ————————————————————————

If the list item is an array or array section, it will be treated as if a **reduction** clause would be applied to each separate element of the array section. The elements of each private array section will be allocated contiguously.

If **nowait** is not used, the reduction computation will be complete at the end of the construct; however, if the reduction clause is used on a construct to which **nowait** is also applied, accesses to the original list item will create a race and, thus, have unspecified effect unless synchronization ensures that they occur after all threads have executed all of their iterations or **section** constructs, and the reduction computation has completed and stored the computed value of that list item. This can most simply be ensured through a barrier synchronization.

The location in the OpenMP program at which the values are combined and the order in which the values are combined are unspecified. Therefore, when comparing sequential and parallel runs, or when comparing one parallel run to another (even if the number of threads used is the same), there is no guarantee that bit-identical results will be obtained or that side effects (such as floating-point exceptions) will be identical or take place at the same location in the OpenMP program.

To avoid race conditions, concurrent reads or updates of the original list item must be synchronized with the update of the original list item that occurs as a result of the **reduction** computation.

<sup>1</sup> **Restrictions**

The restrictions to the **reduction** clause are as follows:

- A list item that appears in a **reduction** clause of a worksharing construct must be shared in the **parallel** regions to which any of the worksharing regions arising from the worksharing construct bind.

- A list item that appears in a **reduction** clause of the innermost enclosing worksharing or **parallel** construct may not be accessed in an explicit task.

- Any number of **reduction** clauses can be specified on the directive, but a list item can appear only once in the **reduction** clauses for that directive.

- For a *reduction-identifier* declared with the **declare reduction** construct, the directive must appear before its use in a **reduction** clause.

- If a list item is an array section, it must specify contiguous storage and it cannot be a zero-length array section.

- If a list item is an array section, accesses to the elements of the array outside the specified array section result in unspecified behavior.

--- C / C++ ---

- The type of a list item that appears in a **reduction** clause must be valid for the *reduction-identifier*. For a **max** or **min** reduction in C, the type of the list item must be an allowed arithmetic data type: **char**, **int**, **float**, **double**, or **_Bool**, possibly modified with **long**, **short**, **signed**, or **unsigned**. For a **max** or **min** reduction in C++, the type of the list item must be an allowed arithmetic data type: **char**, **wchar_t**, **int**, **float**, **double**, or **bool**, possibly modified with **long**, **short**, **signed**, or **unsigned**.

- A list item that appears in a **reduction** clause must not be **const**-qualified.

- If a list item in a **reduction** clause on a worksharing construct has a reference type then it must bind to the same object for all threads of the team.

- The *reduction-identifier* for any list item must be unambiguous and accessible.

--- C / C++ ---

- The type and the rank of a list item that appears in a **reduction** clause must be valid for the *combiner* and *initializer*.

- A list item that appears in a **reduction** clause must be definable.

- A procedure pointer may not appear in a **reduction** clause.

- A pointer with the **INTENT(IN)** attribute may not appear in the **reduction** clause.

- An original list item with the **POINTER** attribute or any pointer component of an original list item that is referenced in the *combiner* must be associated at entry to the construct that contains the **reduction** clause. Additionally, the list item or the pointer component of the list item must not be deallocated, allocated, or pointer assigned within the region.

- An original list item with the **ALLOCATABLE** attribute or any allocatable component of an original list item that is referenced in the *combiner* must be in the allocated state at entry to the construct that contains the **reduction** clause. Additionally, the list item or the allocatable component of the list item must be neither deallocated nor allocated within the region.

- If the *reduction-identifier* is defined in a **declare reduction** directive, the **declare reduction** directive must be in the same subprogram, or accessible by host or use association.

- If the *reduction-identifier* is a user-defined operator, the same explicit interface for that operator must be accessible as at the **declare reduction** directive.

- If the *reduction-identifier* is defined in a **declare reduction** directive, any subroutine or function referenced in the initializer clause or combiner expression must be an intrinsic function, or must have an explicit interface where the same explicit interface is accessible as at the **declare reduction** directive.

### 2.15.3.7  **linear** Clause

#### Summary

The **linear** clause declares one or more list items to be private to a SIMD lane and to have a linear relationship with respect to the iteration space of a loop.

1    **Syntax**

────────────────────── C ──────────────────────

2    The syntax of the **linear** clause is as follows:

    **linear(***linear-list[ : linear-step]***)**

3    where *linear-list* is one of the following

4         *list*
5         *modifier* **(***list***)**

6    where *modifier* is one of the following:

7         **val**

────────────────────── C ──────────────────────

────────────────────── C++ ──────────────────────

8    The syntax of the **linear** clause is as follows:

    **linear(***linear-list[ : linear-step]***)**

9    where *linear-list* is one of the following

10         *list*
11         *modifier* **(***list***)**

12    where *modifier* is one of the following:

13         **ref**
14         **val**
15         **uval**

────────────────────── C++ ──────────────────────

1    The syntax of the **linear** clause is as follows:

**linear(***linear-list[* **:** *linear-step]***)**

2    where *linear-list* is one of the following

3        *list*

4        *modifier* **(***list***)**

5    where *modifier* is one of the following:

6        **ref**

7        **val**

8        **uval**

## 9    **Description**

10   The **linear** clause provides a superset of the functionality provided by the **private** clause. A
11   list item that appears in a **linear** clause is subject to the **private** clause semantics described in
12   Section 2.15.3.3 on page 192 except as noted. If *linear-step* is not specified, it is assumed to be 1.

13   When a **linear** clause is specified on a construct, the value of the new list item on each iteration
14   of the associated loop(s) corresponds to the value of the original list item before entering the
15   construct plus the logical number of the iteration times *linear-step*. The value corresponding to the
16   sequentially last iteration of the associated loop(s) is assigned to the original list item.

17   When a **linear** clause is specified on a declarative directive, all list items must be formal
18   parameters (or, in Fortran, dummy arguments) of a function that will be invoked concurrently on
19   each SIMD lane. If no *modifier* is specified or the **val** or **uval** modifier is specified, the value of
20   each list item on each lane corresponds to the value of the list item upon entry to the function plus
21   the logical number of the lane times *linear-step*. If the **uval** modifier is specified, each invocation
22   uses the same storage location for each SIMD lane; this storage location is updated with the final
23   value of the logically last lane. If the **ref** modifier is specified, the storage location of each list
24   item on each lane corresponds to an array at the storage location upon entry to the function indexed
25   by the logical number of the lane times *linear-step*.

**Restrictions**

- The *linear-step* expression must be invariant during the execution of the region associated with the construct. Otherwise, the execution results in unspecified behavior.

- A *list-item* cannot appear in more than one **linear** clause.

- A *list-item* that appears in a **linear** clause cannot appear in any other data-sharing attribute clause.

———————————————————— C ————————————————————

- A *list-item* that appears in a **linear** clause must be of integral or pointer type.

———————————————————— C ————————————————————

———————————————————— C++ ————————————————————

- A *list-item* that appears in a **linear** clause without the **ref** modifier must be of integral or pointer type, or must be a reference to an integral or pointer type.

- The **ref** or **uval** modifier can only be used if the *list-item* is of a reference type.

- If a list item in a **linear** clause on a worksharing construct has a reference type then it must bind to the same object for all threads of the team.

- If the list item is of a reference type and the **ref** modifier is not specified and if any write to the list item occurs before any read of the list item then the result is unspecified.

———————————————————— C++ ————————————————————

———————————————————— Fortran ————————————————————

- A *list-item* that appears in a **linear** clause without the **ref** modifier must be of type **integer**.

- The **ref** or **uval** modifier can only be used if the *list-item* is a dummy argument without the **VALUE** attribute.

- Variables that have the **POINTER** attribute and Cray pointers may not appear in a linear clause.

- The list item with the **ALLOCATABLE** attribute in the sequentially last iteration must have an allocation status of allocated upon exit from that iteration.

- If the list item is a dummy argument without the **VALUE** attribute and the **ref** modifier is not specified and if any write to the list item occurs before any read of the list item then the result is unspecified.

———————————————————— Fortran ————————————————————

# 1  2.15.4  Data Copying Clauses

2   This section describes the **copyin** clause (allowed on the **parallel** directive and combined
3   parallel worksharing directives) and the **copyprivate** clause (allowed on the **single** directive).

4   These clauses support the copying of data values from private or threadprivate variables on one
5   implicit task or thread to the corresponding variables on other implicit tasks or threads in the team.

6   The clauses accept a comma-separated list of list items (see Section 2.1 on page 26). All list items
7   appearing in a clause must be visible, according to the scoping rules of the base language. Clauses
8   may be repeated as needed, but a list item that specifies a given variable may not appear in more
9   than one clause on the same directive.

---------- Fortran ----------

10   An associate name preserves the association with the selector established at the **ASSOCIATE**
11   statement. A list item that appears in a data copying clause may be a selector of an **ASSOCIATE**
12   construct. If the construct association is established prior to a parallel region, the association
13   between the associate name and the original list item will be retained in the region.

---------- Fortran ----------

## 14  2.15.4.1  `copyin` Clause

### 15  Summary

16   The **copyin** clause provides a mechanism to copy the value of the master thread's threadprivate
17   variable to the threadprivate variable of each other member of the team executing the **parallel**
18   region.

### 19  Syntax

20   The syntax of the **copyin** clause is as follows:

```
copyin(list)
```

<sup>1</sup> **Description**

————————————— C / C++ —————————————

<sup>2</sup> The copy is done after the team is formed and prior to the start of execution of the associated
<sup>3</sup> structured block. For variables of non-array type, the copy occurs by copy assignment. For an array
<sup>4</sup> of elements of non-array type, each element is copied as if by assignment from an element of the
<sup>5</sup> master thread's array to the corresponding element of the other thread's array.

————————————— C / C++ —————————————
————————————— C++ —————————————

<sup>6</sup> For class types, the copy assignment operator is invoked. The order in which copy assignment
<sup>7</sup> operators for different variables of class type are called is unspecified.

————————————— C++ —————————————

————————————— Fortran —————————————

<sup>8</sup> The copy is done, as if by assignment, after the team is formed and prior to the start of execution of
<sup>9</sup> the associated structured block.

<sup>10</sup> On entry to any **parallel** region, each thread's copy of a variable that is affected by a **copyin**
<sup>11</sup> clause for the **parallel** region will acquire the allocation, association, and definition status of the
<sup>12</sup> master thread's copy, according to the following rules:

<sup>13</sup> • If the original list item has the **POINTER** attribute, each copy receives the same association
<sup>14</sup>   status of the master thread's copy as if by pointer assignment.

<sup>15</sup> • If the original list item does not have the **POINTER** attribute, each copy becomes defined with
<sup>16</sup>   the value of the master thread's copy as if by intrinsic assignment, unless it has the allocation
<sup>17</sup>   status of not currently allocated, in which case each copy will have the same status.

————————————— Fortran —————————————

## Restrictions

The restrictions to the **copyin** clause are as follows:

---------------------------------- C / C++ ----------------------------------

- A list item that appears in a **copyin** clause must be threadprivate.

- A variable of class type (or array thereof) that appears in a **copyin** clause requires an accessible, unambiguous copy assignment operator for the class type.

---------------------------------- C / C++ ----------------------------------
---------------------------------- Fortran ----------------------------------

- A list item that appears in a **copyin** clause must be threadprivate. Named variables appearing in a threadprivate common block may be specified: it is not necessary to specify the whole common block.

- A common block name that appears in a **copyin** clause must be declared to be a common block in the same scoping unit in which the **copyin** clause appears.

---------------------------------- Fortran ----------------------------------


## 2.15.4.2 **copyprivate** Clause

### Summary

The **copyprivate** clause provides a mechanism to use a private variable to broadcast a value from the data environment of one implicit task to the data environments of the other implicit tasks belonging to the **parallel** region.

To avoid race conditions, concurrent reads or updates of the list item must be synchronized with the update of the list item that occurs as a result of the **copyprivate** clause.

### Syntax

The syntax of the **copyprivate** clause is as follows:

```
copyprivate(list)
```

**Description**

The effect of the `copyprivate` clause on the specified list items occurs after the execution of the structured block associated with the `single` construct (see Section 2.7.3 on page 67), and before any of the threads in the team have left the barrier at the end of the construct.

——————————————————— C / C++ ———————————————————

In all other implicit tasks belonging to the `parallel` region, each specified list item becomes defined with the value of the corresponding list item in the implicit task associated with the thread that executed the structured block. For variables of non-array type, the definition occurs by copy assignment. For an array of elements of non-array type, each element is copied by copy assignment from an element of the array in the data environment of the implicit task associated with the thread that executed the structured block to the corresponding element of the array in the data environment of the other implicit tasks

——————————————————— C / C++ ———————————————————
——————————————————— C++ ———————————————————

For class types, a copy assignment operator is invoked. The order in which copy assignment operators for different variables of class type are called is unspecified.

——————————————————— C++ ———————————————————

——————————————————— Fortran ———————————————————

If a list item does not have the `POINTER` attribute, then in all other implicit tasks belonging to the `parallel` region, the list item becomes defined as if by intrinsic assignment with the value of the corresponding list item in the implicit task associated with the thread that executed the structured block.

If the list item has the `POINTER` attribute, then, in all other implicit tasks belonging to the `parallel` region, the list item receives, as if by pointer assignment, the same association status of the corresponding list item in the implicit task associated with the thread that executed the structured block.

The order in which any final subroutines for different variables of a finalizable type are called is unspecified.

——————————————————— Fortran ———————————————————

Note – The `copyprivate` clause is an alternative to using a shared variable for the value when providing such a shared variable would be difficult (for example, in a recursion requiring a different variable at each level).

**Restrictions**

The restrictions to the **copyprivate** clause are as follows:

- All list items that appear in the **copyprivate** clause must be either threadprivate or private in the enclosing context.

- A list item that appears in a **copyprivate** clause may not appear in a **private** or **firstprivate** clause on the **single** construct.

———————————————————— C++ ————————————————————

- A variable of class type (or array thereof) that appears in a **copyprivate** clause requires an accessible unambiguous copy assignment operator for the class type.

———————————————————— C++ ————————————————————

———————————————————— Fortran ————————————————————

- A common block that appears in a **copyprivate** clause must be threadprivate.

- Pointers with the **INTENT(IN)** attribute may not appear in the **copyprivate** clause.

- The list item with the **ALLOCATABLE** attribute must have the allocation status of allocated when the intrinsic assignment is performed.

———————————————————— Fortran ————————————————————


# 2.15.5  Data-mapping Attribute Rules and Clauses

This section describes how the data-mapping attributes of any variable referenced in a **target** region are determined. When specified, explicit **map** clauses on **target data** and **target** directives determine these attributes. Otherwise, the following data-mapping rules apply for variables referenced in a **target** construct that are not declared in the construct and do not appear in data-sharing attribute or **map** clauses:

Certain variables and objects have predetermined data-mapping attributes as follows:

- If a variable appears in a **to** or **link** clause on a **declare target** directive then it is treated as if it had appeared in a **map** clause with a *map-type* of **tofrom**.

———————————————————— C / C++ ————————————————————

- A variable that is of type pointer is treated as if it had appeared in a **map** clause as a zero-length array section.

———————————————————— C / C++ ————————————————————

- A variable that is of type reference to pointer is treated as if it had appeared in a **map** clause as a zero-length array section.

Otherwise, the following implicit data-mapping attribute rules apply:

- If a **defaultmap(tofrom:scalar)** clause is not present then a scalar variable is not mapped, but instead has an implicit data-sharing attribute of firstprivate (see Section 2.15.1.1 on page 179).

- If a **defaultmap(tofrom:scalar)** clause is present then a scalar variable is treated as if it had appeared in a **map** clause with a *map-type* of **tofrom**.

- If a variable is not a scalar then it is treated as if it had appeared in a **map** clause with a *map-type* of **tofrom**.

### 2.15.5.1  **map** Clause

**Summary**

The **map** clause specifies how an original list item is mapped from the current task's data environment to a corresponding list item in the device data environment of the device identified by the construct.

**Syntax**

The syntax of the map clause is as follows:

```
map([ [map-type-modifier[,]] map-type  :  ] list)
```

where *map-type* is one of the following:

      **to**

      **from**

      **tofrom**

      **alloc**

      **release**

      **delete**

and *map-type-modifier* is **always**.

**Description**

The list items that appear in a **map** clause may include array sections and structure elements.

The *map-type* and *map-type-modifier* specify the effect of the **map** clause, as described below.

The original and corresponding list items may share storage such that writes to either item by one task followed by a read or write of the other item by another task without intervening synchronization can result in data races.

If the **map** clause appears on a **target**, **target data**, or **target enter data** construct then on entry to the region the following sequence of steps occurs:

1. If a corresponding list item of the original list item is not present in the device data environment, then:

    a) A new list item with language-specific attributes is derived from the original list item and created in the device data environment.

    b) The new list item becomes the corresponding list item to the original list item in the device data environment.

    c) The corresponding list item has a reference count that is initialized to zero.

2. The corresponding list item's reference count is incremented by one.

3. If the corresponding list item's reference count is one or the **always** *map-type-modifier* is present, then:

    a) If the *map-type* is **to** or **tofrom**, then the corresponding list item is assigned the value of the original list item.

4. If the corresponding list item's reference count is one, then:

    a) If the *map-type* is **from** or **alloc**, the value of the corresponding list item is undefined.

If the **map** clause appears on a **target**, **target data**, or **target exit data** construct then on exit from the region the following sequence of steps occurs:

1. If a corresponding list item of the original list item is not present in the device data environment, then the list item is ignored.

2. If a corresponding list item of the original list item is present in the device data environment, then:

    a) If the corresponding list item's reference count is greater than zero, then:

        i. If the *map-type* is **tofrom**, **from** or **release**, then the corresponding list item's reference count is decremented by one.

        ii. If the *map-type* is **delete**, then the corresponding list item's reference count is set to zero.

b) If the corresponding list item's reference count is zero or the **always** *map-type-modifier* is present, then:

    i. If the *map-type* is **from** or **tofrom**, then the original list item is assigned the value of the corresponding list item.

c) If the corresponding list item's reference count is zero, then the corresponding list item is removed from the device data environment

---
**C / C++**

---

If a new list item is created then a new list item of the same type, with automatic storage duration, is allocated for the construct. The size and alignment of the new list item are determined by the type of the variable. This allocation occurs if the region references the list item in any statement.

**C / C++**

---
**Fortran**

---

If a new list item is created then a new list item of the same type, type parameter, and rank is allocated.

**Fortran**

---

The *map-type* determines how the new list item is initialized.

If a *map-type* is not specified, the *map-type* defaults to **tofrom**.

## Restrictions

- A list item cannot appear in both a **map** clause and a data-sharing attribute clause on the same construct.

- If a list item is an array section, it must specify contiguous storage.

- At most one list item can be an array item derived from a given variable in **map** clauses of the same construct.

- List items of **map** clauses in the same construct must not share original storage.

- If any part of the original storage of a list item has corresponding storage in the device data environment, all of the original storage must have corresponding storage in the device data environment.

- If a list item is an element of a structure, and a different element of the structure has a corresponding list item in the device data environment prior to a task encountering the construct associated with the **map** clause, then the list item must also have a correspnding list item in the device data environment prior to the task encountering the construct.

- If a list item is an element of a structure, only the rightmost symbol of the variable reference can be an array section.

1 • If variables that share storage are mapped, the behavior is unspecified.

2 • A list item must have a mappable type.

3 • **threadprivate** variables cannot appear in a **map** clause.

––––––––––––––––––––––––– C++ –––––––––––––––––––––––––

4 • If the type of a list item is a reference to a type *T* then the type will be considered to be *T* for all
5 purposes of this clause.

––––––––––––––––––––––––– C++ –––––––––––––––––––––––––

––––––––––––––––––––––––– C / C++ –––––––––––––––––––––––––

6 • Initialization and assignment are through bitwise copy.

7 • A variable for which the type is pointer and an array section derived from that variable must not
8 appear as list items of **map** clauses of the same construct.

9 • A list item cannot be a variable that is a member of a structure with a union type.

10 • A bit-field cannot appear in a **map** clause.

––––––––––––––––––––––––– C / C++ –––––––––––––––––––––––––

––––––––––––––––––––––––– Fortran –––––––––––––––––––––––––

11 • The value of the new list item becomes that of the original list item in the map initialization and
12 assignment.

13 • A list item must not contain any components that have the **ALLOCATABLE** attribute.

14 • If the allocation status of a list item with the **ALLOCATABLE** attribute is unallocated upon entry
15 to a **target** region, the list item must be unallocated upon exit from the region.

16 • If the allocation status of a list item with the **ALLOCATABLE** attribute is allocated upon entry to
17 a **target** region, the allocation status of the corresponding list item must not be changed and
18 must not be reshaped in the region.

19 • If an array section of an allocatable array is mapped and the size of the section is smaller than
20 that of the whole array, the **target** region must not have any reference to the whole array.

––––––––––––––––––––––––– Fortran –––––––––––––––––––––––––

21 ## 2.15.5.2 **defaultmap** Clause

22 ### Summary

23 The **defaultmap** clause explicitly determines the data-mapping attributes of variables that are
24 referenced in a **target** construct and would otherwise be implicitly determined.

**Syntax**

$$\blacktriangledown \underline{\hspace{4cm}} \text{C / C++} \underline{\hspace{4cm}} \blacktriangledown$$

2    The syntax of the **defaultmap** clause is as follows:

```
defaultmap(tofrom:scalar)
```

$$\blacktriangle \underline{\hspace{4cm}} \text{C / C++} \underline{\hspace{4cm}} \blacktriangle$$

$$\blacktriangledown \underline{\hspace{4cm}} \text{Fortran} \underline{\hspace{4cm}} \blacktriangledown$$

3    The syntax of the **defaultmap** clause is as follows:

```
defaultmap(tofrom:scalar)
```

$$\blacktriangle \underline{\hspace{4cm}} \text{Fortran} \underline{\hspace{4cm}} \blacktriangle$$

**Description**

5    The **defaultmap(tofrom:scalar)** clause causes all scalar variables referenced in the
6    construct that have implicitly determined data-mapping attributes to have the **tofrom** *map-type*.

# 2.16 `declare reduction` **Directive**

**Summary**

9    The following section describes the directive for declaring user-defined reductions. The
10   **declare reduction** directive declares a *reduction-identifier* that can be used in a
11   **reduction** clause. The **declare reduction** directive is a declarative directive.

1 **Syntax**

———————————————— C ————————————————

```
#pragma omp declare reduction(reduction-identifier : typename-list :
combiner ) [initializer-clause] new-line
```

2 where:

3   • *reduction-identifier* is either a base language identifier or one of the following operators: **+**, **−**, **\***,
4     **&**, **|**, ˆ, **&&** and **||**

5   • *typename-list* is a list of type names

6   • *combiner* is an expression

7   • *initializer-clause* is **initializer(***initializer-expr***)** where *initializer-expr* is
8     **omp_priv =** *initializer* or *function-name* **(***argument-list***)**

———————————————— C ————————————————

——————————————— C++ ———————————————

```
#pragma omp declare reduction(reduction-identifier : typename-list :
combiner) [initializer-clause] new-line
```

9 where:

10   • *reduction-identifier* is either an *id-expression* or one of the following operators: **+**, **−**, **\***, **&**, **|**, ˆ,
11     **&&** and **||**

12   • *typename-list* is a list of type names

13   • *combiner* is an expression

14   • *initializer-clause* is **initializer(***initializer-expr***)** where *initializer-expr* is
15     **omp_priv** *initializer* or *function-name* **(***argument-list***)**

——————————————— C++ ———————————————

```
!$omp declare reduction(reduction-identifier : type-list : combiner)
[initializer-clause]
```

1   where:

2   • *reduction-identifier* is either a base language identifier, or a user-defined operator, or one of the
3   following operators: **+**, **−**, **\***, **.and.**, **.or.**, **.eqv.**, **.neqv.**, or one of the following intrinsic
4   procedure names: **max**, **min**, **iand**, **ior**, **ieor**.

5   • *type-list* is a list of type specifiers

6   • *combiner* is either an assignment statement or a subroutine name followed by an argument list

7   • *initializer-clause* is **initializer(***initializer-expr***)**, where *initializer-expr* is
8   **omp_priv =** *expression* or *subroutine-name***(***argument-list***)**

9   **Description**

10   Custom reductions can be defined using the **declare reduction** directive; the
11   *reduction-identifier* and the type identify the **declare reduction** directive. The
12   *reduction-identifier* can later be used in a **reduction** clause using variables of the type or types
13   specified in the **declare reduction** directive. If the directive applies to several types then it is
14   considered as if there were multiple **declare reduction** directives, one for each type.

15   If a type with deferred or assumed length type parameter is specified in a **declare reduction**
16   directive, the *reduction-identifier* of that directive can be used in a **reduction** clause with any
17   variable of the same type and the same kind parameter, regardless of the length type Fortran
18   parameters with which the variable is declared.

19   The visibility and accessibility of this declaration are the same as those of a variable declared at the
20   same point in the program. The enclosing context of the *combiner* and of the *initializer-expr* will be
21   that of the **declare reduction** directive. The *combiner* and the *initializer-expr* must be correct
22   in the base language as if they were the body of a function defined at the same point in the program.

If the *reduction-identifier* is the same as the name of a user-defined operator or an extended operator, or the same as a generic name that is one of the allowed intrinsic procedures, and if the operator or procedure name appears in an accessibility statement in the same module, the accessibility of the corresponding **declare reduction** directive is determined by the accessibility attribute of the statement.

If the *reduction-identifier* is the same as a generic name that is one of the allowed intrinsic procedures and is accessible, and if it has the same name as a derived type in the same module, the accessibility of the corresponding **declare reduction** directive is determined by the accessibility of the generic name according to the base language.

Fortran

C++

The **declare reduction** directive can also appear at points in the program at which a static data member could be declared. In this case, the visibility and accessibility of the declaration are the same as those of a static data member declared at the same point in the program.

C++

The *combiner* specifies how partial results can be combined into a single value. The *combiner* can use the special variable identifiers **omp_in** and **omp_out** that are of the type of the variables being reduced with this *reduction-identifier*. Each of them will denote one of the values to be combined before executing the *combiner*. It is assumed that the special **omp_out** identifier will refer to the storage that holds the resulting combined value after executing the *combiner*.

The number of times the *combiner* is executed, and the order of these executions, for any **reduction** clause is unspecified.

Fortran

If the *combiner* is a subroutine name with an argument list, the *combiner* is evaluated by calling the subroutine with the specified argument list.

If the *combiner* is an assignment statement, the *combiner* is evaluated by executing the assignment statement.

Fortran

As the *initializer-expr* value of a user-defined reduction is not known *a priori* the *initializer-clause* can be used to specify one. Then the contents of the *initializer-clause* will be used as the initializer for private copies of reduction list items where the **omp_priv** identifier will refer to the storage to be initialized. The special identifier **omp_orig** can also appear in the *initializer-clause* and it will refer to the storage of the original variable to be reduced.

The number of times that the *initializer-expr* is evaluated, and the order of these evaluations, is unspecified.

---
C / C++
---

1 If the *initializer-expr* is a function name with an argument list, the *initializer-expr* is evaluated by
2 calling the function with the specified argument list. Otherwise, the *initializer-expr* specifies how
3 **omp_priv** is declared and initialized.

---
C / C++
---

---
C
---

4 If no *initializer-clause* is specified, the private variables will be initialized following the rules for
5 initialization of objects with static storage duration.

---
C
---

---
C++
---

6 If no *initializer-expr* is specified, the private variables will be initialized following the rules for
7 *default-initialization*.

---
C++
---

---
Fortran
---

8 If the *initializer-expr* is a subroutine name with an argument list, the *initializer-expr* is evaluated by
9 calling the subroutine with the specified argument list.

10 If the *initializer-expr* is an assignment statement, the *initializer-expr* is evaluated by executing the
11 assignment statement.

12 If no *initializer-clause* is specified, the private variables will be initialized as follows:

13 • For **complex**, **real**, or **integer** types, the value 0 will be used.

14 • For **logical** types, the value **.false.** will be used.

15 • For derived types for which default initialization is specified, default initialization will be used.

16 • Otherwise, not specifying an *initializer-clause* results in unspecified behavior.

---
Fortran
---

---
C / C++
---

17 If *reduction-identifier* is used in a **target** region then a **declare target** construct must be
18 specified for any function that can be accessed through the *combiner* and *initializer-expr*.

---
C / C++
---

If *reduction-identifier* is used in a **target** region then a **declare target** construct must be
specified for any function or subroutine that can be accessed through the *combiner* and
*initializer-expr*.

**Restrictions**

- Only the variables **omp_in** and **omp_out** are allowed in the *combiner*.

- Only the variables **omp_priv** and **omp_orig** are allowed in the *initializer-clause*.

- If the variable **omp_orig** is modified in the *initializer-clause*, the behavior is unspecified.

- If execution of the *combiner* or the *initializer-expr* results in the execution of an OpenMP
  construct or an OpenMP API call, then the behavior is unspecified.

- A *reduction-identifier* may not be re-declared in the current scope for the same type or for a type
  that is compatible according to the base language rules.

- At most one *initializer-clause* can be specified.

- A type name in a **declare reduction** directive cannot be a function type, an array type, a
  reference type, or a type qualified with **const**, **volatile** or **restrict**.

- If the *initializer-expr* is a function name with an argument list, then one of the arguments must be
  the address of **omp_priv**.

- If the *initializer-expr* is a function name with an argument list, then one of the arguments must be
  **omp_priv** or the address of **omp_priv**.

- If the *initializer-expr* is a subroutine name with an argument list, then one of the arguments must be **omp_priv**.

- If the **declare reduction** directive appears in the specification part of a module and the corresponding reduction clause does not appear in the same module, the *reduction-identifier* must be the same as the name of a user-defined operator, one of the allowed operators that is extended or a generic name that is the same as the name of one of the allowed intrinsic procedures.

- If the **declare reduction** directive appears in the specification of a module, if the corresponding **reduction** clause does not appear in the same module, and if the *reduction-identifier* is the same as the name of a user-defined operator or an extended operator, or the same as a generic name that is the same as one of the allowed intrinsic procedures then the interface for that operator or the generic name must be defined in the specification of the same module, or must be accessible by use association.

- Any subroutine or function used in the **initializer** clause or *combiner* expression must be an intrinsic function, or must have an accessible interface.

- Any user-defined operator or extended operator used in the **initializer** clause or *combiner* expression must have an accessible interface.

- If any subroutine, function, user-defined operator, or extended operator is used in the **initializer** clause or *combiner* expression, it must be accessible to the subprogram in which the corresponding **reduction** clause is specified.

- If the length type parameter is specified for a character type, it must be a constant, a colon or an **\***.

- If a character type with deferred or assumed length parameter is specified in a **declare reduction** directive, no other **declare reduction** directive with Fortran character type of the same kind parameter and the same *reduction-identifier* is allowed in the same scope.

- Any subroutine used in the **initializer** clause or *combiner* expression must not have any alternate returns appear in the argument list.

**Cross References**

- **reduction** clause, Section 2.15.3.6 on page 201.

# 2.17  Nesting of Regions

This section describes a set of restrictions on the nesting of regions. The restrictions on nesting are as follows:

- A worksharing region may not be closely nested inside a worksharing, explicit **task**, **taskloop**, **critical**, **ordered**, **atomic**, or **master** region.

- A **barrier** region may not be closely nested inside a worksharing, explicit **task**, **taskloop**, **critical**, **ordered**, **atomic**, or **master** region.

- A **master** region may not be closely nested inside a worksharing, **atomic**, explicit **task**, or **taskloop** region.

- An **ordered** region arising from an **ordered** construct without any clause or with the **threads** or **depend** clause may not be closely nested inside a **critical**, **ordered**, **atomic**, explicit **task**, or **taskloop** region.

- An **ordered** region arising from an **ordered** construct without any clause or with the **threads** or **depend** clause must be closely nested inside a loop region (or parallel loop region) with an **ordered** clause.

- An **ordered** region arising from an **ordered** construct with the **simd** clause must be closely nested inside a **simd** (or loop SIMD) region.

- An **ordered** region arising from an **ordered** construct with both the **simd** and **threads** clauses must be closely nested inside a loop SIMD region.

- A **critical** region may not be nested (closely or otherwise) inside a **critical** region with the same name. This restriction is not sufficient to prevent deadlock.

- OpenMP constructs may not be encountered during execution of an **atomic** region.

- An ordered construct with the **simd** clause is the only OpenMP construct that can be encountered during execution of a **simd** region.

- If a **target**, **target update**, **target data**, **target enter data**, or **target exit data** construct is encountered during execution of a **target** region, the behavior is unspecified.

- If specified, a **teams** construct must be contained within a **target** construct. That **target** construct must not contain any statements or directives outside of the **teams** construct.

- **distribute**, **distribute simd**, distribute parallel loop, distribute parallel loop SIMD, and **parallel** regions, including any **parallel** regions arising from combined constructs, are the only OpenMP regions that may be strictly nested inside the **teams** region.

- The region associated with the **distribute** construct must be strictly nested inside a **teams** region.

- If *construct-type-clause* is **taskgroup**, the **cancel** construct must be closely nested inside a **task** construct and the **cancel** region must be closely nested inside a **taskgroup** region. If *construct-type-clause* is **sections**, the **cancel** construct must be closely nested inside a **sections** or **section** construct. Otherwise, the **cancel** construct must be closely nested inside an OpenMP construct that matches the type specified in *construct-type-clause* of the **cancel** construct.

- A **cancellation point** construct for which *construct-type-clause* is **taskgroup** must be closely nested inside a **task** construct, and the **cancellation point** region must be closely nested inside a **taskgroup** region. A **cancellation point** construct for which *construct-type-clause* is **sections** must be closely nested inside a **sections** or **section** construct. Otherwise, a **cancellation point** construct must be closely nested inside an OpenMP construct that matches the type specified in *construct-type-clause*.

# 2  Runtime Library Routines

3  This chapter describes the OpenMP API runtime library routines and is divided into the following
4  sections:

5  • Runtime library definitions (Section 3.1 on page 230).

6  • Execution environment routines that can be used to control and to query the parallel execution
7  environment (Section 3.2 on page 231).

8  • Lock routines that can be used to synchronize access to data (Section 3.3 on page 270).

9  • Portable timer routines (Section 3.4 on page 279).

10  • Device memory routines that can be used to allocate memory and to manage pointers on target
11  devices (Section 3.5 on page 282).

12  Throughout this chapter, *true* and *false* are used as generic terms to simplify the description of the
13  routines.

—————————————————— C / C++ ——————————————————

14  *true* means a nonzero integer value and *false* means an integer value of zero.

—————————————————— C / C++ ——————————————————

—————————————————— Fortran ——————————————————

15  *true* means a logical value of `.TRUE.` and *false* means a logical value of `.FALSE.`.

—————————————————— Fortran ——————————————————

—————————————————— Fortran ——————————————————

16  **Restrictions**

17  The following restriction applies to all OpenMP runtime library routines:

18  • OpenMP runtime library routines may not be called from `PURE` or `ELEMENTAL` procedures.

—————————————————— Fortran ——————————————————

# 3.1 Runtime Library Definitions

For each base language, a compliant implementation must supply a set of definitions for the OpenMP API runtime library routines and the special data types of their parameters. The set of definitions must contain a declaration for each OpenMP API runtime library routine and a declaration for the *simple lock*, *nestable lock*, *schedule*, and *thread affinity policy* data types. In addition, each set of definitions may specify other implementation specific values.

---
C / C++
---

The library routines are external functions with "C" linkage.

Prototypes for the C/C++ runtime library routines described in this chapter shall be provided in a header file named **omp.h**. This file defines the following:

- The prototypes of all the routines in the chapter.
- The type **omp_lock_t**.
- The type **omp_nest_lock_t**.
- The type **omp_lock_hint_t**.
- The type **omp_sched_t**.
- The type **omp_proc_bind_t**.

See Section Section B.1 on page 327 for an example of this file.

---
C / C++
---
Fortran
---

The OpenMP Fortran API runtime library routines are external procedures. The return values of these routines are of default kind, unless otherwise specified.

Interface declarations for the OpenMP Fortran runtime library routines described in this chapter shall be provided in the form of a Fortran **include** file named **omp_lib.h** or a Fortran 90 **module** named **omp_lib**. It is implementation defined whether the **include** file or the **module** file (or both) is provided.

These files define the following:

- The interfaces of all of the routines in this chapter.
- The **integer parameter omp_lock_kind**.
- The **integer parameter omp_nest_lock_kind**.
- The **integer parameter omp_lock_hint_kind**.
- The **integer parameter omp_sched_kind**.
- The **integer parameter omp_proc_bind_kind**.

1  • The **integer parameter openmp_version** with a value *yyyymm* where *yyyy* and *mm* are
2     the year and month designations of the version of the OpenMP Fortran API that the
3     implementation supports. This value matches that of the C preprocessor macro **_OPENMP**, when
4     a macro preprocessor is supported (see Section 2.2 on page 33).

5  See Section B.1 on page 331 and Section B.3 on page 335 for examples of these files.

6  It is implementation defined whether any of the OpenMP runtime library routines that take an
7  argument are extended with a generic interface so arguments of different **KIND** type can be
8  accommodated. See Appendix B.4 for an example of such an extension.

———————————————— Fortran ————————————————

# 9 3.2 Execution Environment Routines

10  This section describes routines that affect and monitor threads, processors, and the parallel
11  environment.

## 12 3.2.1 omp_set_num_threads

### 13 Summary

14  The **omp_set_num_threads** routine affects the number of threads to be used for subsequent
15  parallel regions that do not specify a **num_threads** clause, by setting the value of the first
16  element of the *nthreads-var* ICV of the current task.

### 17 Format

———————————————— C / C++ ————————————————

```
void omp_set_num_threads(int num_threads);
```

———————————————— C / C++ ————————————————
———————————————— Fortran ————————————————

```
subroutine omp_set_num_threads(num_threads)
integer num_threads
```

———————————————— Fortran ————————————————

**Constraints on Arguments**

2    The value of the argument passed to this routine must evaluate to a positive integer, or else the
3    behavior of this routine is implementation defined.

4    **Binding**

5    The binding task set for an **omp_set_num_threads** region is the generating task.

6    **Effect**

7    The effect of this routine is to set the value of the first element of the *nthreads-var* ICV of the
8    current task to the value specified in the argument.

9    **Cross References**

10    • *nthreads-var* ICV, see Section 2.3 on page 36.

11    • **parallel** construct and **num_threads** clause, see Section 2.5 on page 46.

12    • Determining the number of threads for a **parallel** region, see Section 2.5.1 on page 50.

13    • **omp_get_max_threads** routine, see Section 3.2.3 on page 233.

14    • **OMP_NUM_THREADS** environment variable, see Section 4.2 on page 293.

## 15    3.2.2   **omp_get_num_threads**

16    **Summary**

17    The **omp_get_num_threads** routine returns the number of threads in the current team.

18    **Format**

────────────────────────────  C / C++  ────────────────────────────

```
int omp_get_num_threads(void);
```

────────────────────────────  C / C++  ────────────────────────────

```
integer function omp_get_num_threads()
```

1    **Binding**

2    The binding region for an **omp_get_num_threads** region is the innermost enclosing
3    **parallel** region.

4    **Effect**

5    The **omp_get_num_threads** routine returns the number of threads in the team executing the
6    **parallel** region to which the routine region binds. If called from the sequential part of a
7    program, this routine returns 1.

8    **Cross References**

9    • **parallel** construct, see Section 2.5 on page 46.

10    • Determining the number of threads for a **parallel** region, see Section 2.5.1 on page 50.

11    • **omp_set_num_threads** routine, see Section 3.2.1 on page 231.

12    • **OMP_NUM_THREADS** environment variable, see Section 4.2 on page 293.

13    ## 3.2.3 omp_get_max_threads

14    **Summary**

15    The **omp_get_max_threads** routine returns an upper bound on the number of threads that
16    could be used to form a new team if a **parallel** construct without a **num_threads** clause were
17    encountered after execution returns from this routine.

<sup>1</sup> **Format**

---

C / C++

```
int omp_get_max_threads(void);
```

C / C++

---

Fortran

```
integer function omp_get_max_threads()
```

Fortran

---

<sup>2</sup> **Binding**

<sup>3</sup> The binding task set for an **omp_get_max_threads** region is the generating task.


<sup>4</sup> **Effect**

<sup>5</sup> The value returned by **omp_get_max_threads** is the value of the first element of the
<sup>6</sup> *nthreads-var* ICV of the current task. This value is also an upper bound on the number of threads
<sup>7</sup> that could be used to form a new team if a parallel region without a **num_threads** clause were
<sup>8</sup> encountered after execution returns from this routine.

---

<sup>9</sup> Note – The return value of the **omp_get_max_threads** routine can be used to dynamically
<sup>10</sup> allocate sufficient storage for all threads in the team formed at the subsequent active **parallel**
<sup>11</sup> region.

---

<sup>12</sup> **Cross References**

<sup>13</sup> • *nthreads-var* ICV, see Section 2.3 on page 36.

<sup>14</sup> • **parallel** construct, see Section 2.5 on page 46.

<sup>15</sup> • **num_threads** clause, see Section 2.5 on page 46.

<sup>16</sup> • Determining the number of threads for a **parallel** region, see Section 2.5.1 on page 50.

<sup>17</sup> • **omp_set_num_threads** routine, see Section 3.2.1 on page 231.

<sup>18</sup> • **OMP_NUM_THREADS** environment variable, see Section 4.2 on page 293.

<sup>1</sup> **3.2.4 omp_get_thread_num**

<sup>2</sup> **Summary**

<sup>3</sup> The **omp_get_thread_num** routine returns the thread number, within the current team, of the
<sup>4</sup> calling thread.

<sup>5</sup> **Format**

---

C / C++

```
int omp_get_thread_num(void);
```

C / C++

---

Fortran

```
integer function omp_get_thread_num()
```

Fortran

---

<sup>6</sup> **Binding**

<sup>7</sup> The binding thread set for an **omp_get_thread_num** region is the current team. The binding
<sup>8</sup> region for an **omp_get_thread_num** region is the innermost enclosing **parallel** region.

<sup>9</sup> **Effect**

<sup>10</sup> The **omp_get_thread_num** routine returns the thread number of the calling thread, within the
<sup>11</sup> team executing the **parallel** region to which the routine region binds. The thread number is an
<sup>12</sup> integer between 0 and one less than the value returned by **omp_get_num_threads**, inclusive.
<sup>13</sup> The thread number of the master thread of the team is 0. The routine returns 0 if it is called from
<sup>14</sup> the sequential part of a program.

---

<sup>15</sup> Note – The thread number may change during the execution of an untied task. The value returned
<sup>16</sup> by **omp_get_thread_num** is not generally useful during the execution of such a task region.

---

<sup>17</sup> **Cross References**
<sup>18</sup> • **omp_get_num_threads** routine, see Section 3.2.2 on page 232.

## 3.2.5 `omp_get_num_procs`

**Summary**

The **`omp_get_num_procs`** routine returns the number of processors available to the device.

**Format**

——————————————— C / C++ ———————————————

```
int omp_get_num_procs(void);
```

——————————————— C / C++ ———————————————
——————————————— Fortran ———————————————

```
integer function omp_get_num_procs()
```

——————————————— Fortran ———————————————
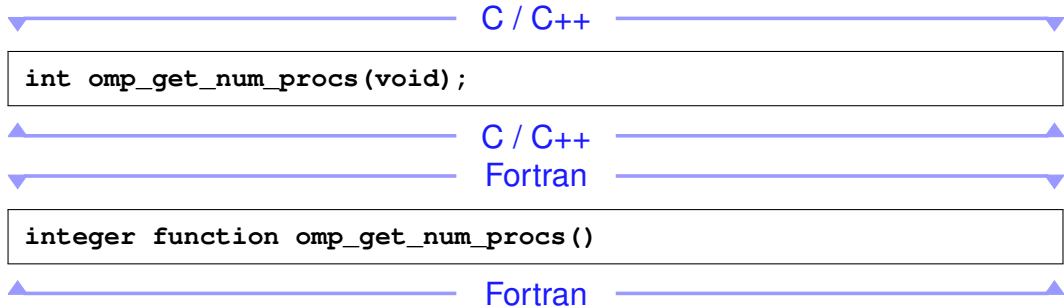
**Binding**

The binding thread set for an **`omp_get_num_procs`** region is all threads on a device. The effect
of executing this routine is not related to any specific region corresponding to any construct or API
routine.

**Effect**

The **`omp_get_num_procs`** routine returns the number of processors that are available to the
device at the time the routine is called. This value may change between the time that it is
determined by the **`omp_get_num_procs`** routine and the time that it is read in the calling
context due to system actions outside the control of the OpenMP implementation.

**Cross References**

None.

## 3.2.6 `omp_in_parallel`

**Summary**

The **`omp_in_parallel`** routine returns *true* if the *active-levels-var* ICV is greater than zero;
otherwise, it returns *false*.

1 **Format**

―――――――――――――― C / C++ ――――――――――――――

```
int omp_in_parallel(void);
```

―――――――――――――― C / C++ ――――――――――――――
―――――――――――――― Fortran ――――――――――――――

```
logical function omp_in_parallel()
```

―――――――――――――― Fortran ――――――――――――――

2 **Binding**

3 The binding task set for an **omp_in_parallel** region is the generating task.

4 **Effect**

5 The effect of the **omp_in_parallel** routine is to return *true* if the current task is enclosed by an
6 active **parallel** region, and the **parallel** region is enclosed by the outermost initial task
7 region on the device; otherwise it returns *false*.

8 **Cross References**

9 • *active-levels-var*, see Section 2.3 on page 36.

10 • **parallel** construct, see Section 2.5 on page 46.

11 • **omp_get_active_level** routine, see Section 3.2.20 on page 252.

12 ## 3.2.7 **omp_set_dynamic**

13 **Summary**

14 The **omp_set_dynamic** routine enables or disables dynamic adjustment of the number of
15 threads available for the execution of subsequent **parallel** regions by setting the value of the
16 *dyn-var* ICV.

**Format**

───────────────── C / C++ ─────────────────

```
void omp_set_dynamic(int dynamic_threads);
```

───────────────── C / C++ ─────────────────

───────────────── Fortran ─────────────────

```
subroutine omp_set_dynamic(dynamic_threads)
logical dynamic_threads
```

───────────────── Fortran ─────────────────

2   **Binding**

3   The binding task set for an **omp_set_dynamic** region is the generating task.

4   **Effect**

5   For implementations that support dynamic adjustment of the number of threads, if the argument to
6   **omp_set_dynamic** evaluates to *true*, dynamic adjustment is enabled for the current task;
7   otherwise, dynamic adjustment is disabled for the current task. For implementations that do not
8   support dynamic adjustment of the number of threads this routine has no effect: the value of
9   *dyn-var* remains *false*.

10  **Cross References**

11  • *dyn-var* ICV, see Section 2.3 on page 36.

12  • Determining the number of threads for a **parallel** region, see Section 2.5.1 on page 50.

13  • **omp_get_num_threads** routine, see Section 3.2.2 on page 232.

14  • **omp_get_dynamic** routine, see Section 3.2.8 on page 239.

15  • **OMP_DYNAMIC** environment variable, see Section 4.3 on page 294.

## 3.2.8 `omp_get_dynamic`

**Summary**

The **`omp_get_dynamic`** routine returns the value of the *dyn-var* ICV, which determines whether dynamic adjustment of the number of threads is enabled or disabled.

**Format**

─────────────────── C / C++ ───────────────────

```
int omp_get_dynamic(void);
```

─────────────────── C / C++ ───────────────────

─────────────────── Fortran ───────────────────

```
logical function omp_get_dynamic()
```

─────────────────── Fortran ───────────────────

**Binding**

The binding task set for an **`omp_get_dynamic`** region is the generating task.

**Effect**

This routine returns *true* if dynamic adjustment of the number of threads is enabled for the current task; it returns *false*, otherwise. If an implementation does not support dynamic adjustment of the number of threads, then this routine always returns *false*.

**Cross References**

- *dyn-var* ICV, see Section 2.3 on page 36.
- Determining the number of threads for a **`parallel`** region, see Section 2.5.1 on page 50.
- **`omp_set_dynamic`** routine, see Section 3.2.7 on page 237.
- **`OMP_DYNAMIC`** environment variable, see Section 4.3 on page 294.

## 3.2.9 `omp_get_cancellation`

**Summary**

The `omp_get_cancellation` routine returns the value of the *cancel-var* ICV, which determines if cancellation is enabled or disabled.

**Format**

―――――――――――――――――― C / C++ ――――――――――――――――――

```
int omp_get_cancellation(void);
```

―――――――――――――――――― C / C++ ――――――――――――――――――

―――――――――――――――――― Fortran ――――――――――――――――――

```
logical function omp_get_cancellation()
```

―――――――――――――――――― Fortran ――――――――――――――――――
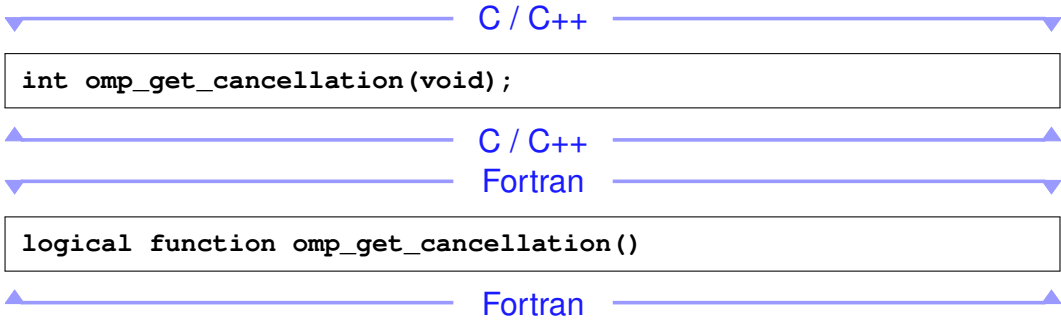
**Binding**

The binding task set for an `omp_get_cancellation` region is the whole program.

**Effect**

This routine returns *true* if cancellation is enabled. It returns *false* otherwise.

**Cross References**

- *cancel-var* ICV, see Section 2.3.1 on page 36.
- `cancel` construct, see Section 2.14.1 on page 172
- `OMP_CANCELLATION` environment variable, see Section 4.11 on page 300

## 3.2.10 `omp_set_nested`

**Summary**

The `omp_set_nested` routine enables or disables nested parallelism, by setting the *nest-var* ICV.

<sup>1</sup> **Format**

─────────────── C / C++ ───────────────

```
void omp_set_nested(int nested);
```

─────────────── C / C++ ───────────────
─────────────── Fortran ───────────────

```
subroutine omp_set_nested(nested)
logical nested
```

─────────────── Fortran ───────────────

<sup>2</sup> **Binding**

<sup>3</sup> The binding task set for an **omp_set_nested** region is the generating task.

<sup>4</sup> **Effect**

<sup>5</sup> For implementations that support nested parallelism, if the argument to **omp_set_nested**
<sup>6</sup> evaluates to *true*, nested parallelism is enabled for the current task; otherwise, nested parallelism is
<sup>7</sup> disabled for the current task. For implementations that do not support nested parallelism, this
<sup>8</sup> routine has no effect: the value of *nest-var* remains *false*.

<sup>9</sup> **Cross References**

<sup>10</sup>   &bull; *nest-var* ICV, see Section 2.3 on page 36.

<sup>11</sup>   &bull; Determining the number of threads for a **parallel** region, see Section 2.5.1 on page 50.

<sup>12</sup>   &bull; **omp_set_max_active_levels** routine, see Section 3.2.15 on page 246.

<sup>13</sup>   &bull; **omp_get_max_active_levels** routine, see Section 3.2.16 on page 248.

<sup>14</sup>   &bull; **omp_get_nested** routine, see Section 3.2.11 on page 242.

<sup>15</sup>   &bull; **OMP_NESTED** environment variable, see Section 4.6 on page 297.

# 3.2.11  `omp_get_nested`

**Summary**

The **`omp_get_nested`** routine returns the value of the *nest-var* ICV, which determines if nested parallelism is enabled or disabled.

**Format**

―――――――――――――――――――― C / C++ ――――――――――――――――――――

```
int omp_get_nested(void);
```

―――――――――――――――――――― C / C++ ――――――――――――――――――――

―――――――――――――――――――――― Fortran ――――――――――――――――――――――

```
logical function omp_get_nested()
```

―――――――――――――――――――――― Fortran ――――――――――――――――――――――

**Binding**

The binding task set for an **`omp_get_nested`** region is the generating task.

**Effect**

This routine returns *true* if nested parallelism is enabled for the current task; it returns *false*, otherwise. If an implementation does not support nested parallelism, this routine always returns *false*.

**Cross References**

- *nest-var* ICV, see Section 2.3 on page 36.
- Determining the number of threads for a **`parallel`** region, see Section 2.5.1 on page 50.
- **`omp_set_nested`** routine, see Section 3.2.10 on page 240.
- **`OMP_NESTED`** environment variable, see Section 4.6 on page 297.

**3.2.12** `omp_set_schedule`

2 **Summary**

3 The **omp_set_schedule** routine affects the schedule that is applied when **runtime** is used as
4 schedule kind, by setting the value of the *run-sched-var* ICV.

5 **Format**

―――――――――――――――― C / C++ ――――――――――――――――

```
void omp_set_schedule(omp_sched_t kind, int chunk_size);
```

―――――――――――――――― C / C++ ――――――――――――――――
―――――――――――――――― Fortran ――――――――――――――――

```
subroutine omp_set_schedule(kind, chunk_size)
integer (kind=omp_sched_kind) kind
integer chunk_size
```

―――――――――――――――― Fortran ――――――――――――――――

6 **Constraints on Arguments**

7 The first argument passed to this routine can be one of the valid OpenMP schedule kinds (except for
8 **runtime**) or any implementation specific schedule. The C/C++ header file (**omp.h**) and the
9 Fortran include file (**omp_lib.h**) and/or Fortran 90 module file (**omp_lib**) define the valid
10 constants. The valid constants must include the following, which can be extended with
11 implementation specific values:

```
typedef enum omp_sched_t {
    omp_sched_static = 1,
    omp_sched_dynamic = 2,
    omp_sched_guided = 3,
    omp_sched_auto = 4
} omp_sched_t;
```

C / C++

Fortran

```
integer(kind=omp_sched_kind), parameter :: omp_sched_static = 1
integer(kind=omp_sched_kind), parameter :: omp_sched_dynamic = 2
integer(kind=omp_sched_kind), parameter :: omp_sched_guided = 3
integer(kind=omp_sched_kind), parameter :: omp_sched_auto = 4
```

Fortran

## Binding

The binding task set for an **omp_set_schedule** region is the generating task.

## Effect

The effect of this routine is to set the value of the *run-sched-var* ICV of the current task to the values specified in the two arguments. The schedule is set to the schedule type specified by the first argument *kind*. It can be any of the standard schedule types or any other implementation specific one. For the schedule types **static**, **dynamic**, and **guided** the *chunk_size* is set to the value of the second argument, or to the default *chunk_size* if the value of the second argument is less than 1; for the schedule type **auto** the second argument has no meaning; for implementation specific schedule types, the values and associated meanings of the second argument are implementation defined.

## Cross References

- *run-sched-var* ICV, see Section 2.3 on page 36.

- Determining the schedule of a worksharing loop, see Section 2.7.1.1 on page 64.

- **omp_get_schedule** routine, see Section 3.2.13 on page 245.

- **OMP_SCHEDULE** environment variable, see Section 4.1 on page 292.

<sup>1</sup> ### 3.2.13 `omp_get_schedule`

<sup>2</sup> **Summary**

<sup>3</sup> The **omp_get_schedule** routine returns the schedule that is applied when the runtime schedule
<sup>4</sup> is used.

<sup>5</sup> **Format**

---

C / C++

---

```
void omp_get_schedule(omp_sched_t * kind, int * chunk_size);
```

---

C / C++

---

Fortran

---

```
subroutine omp_get_schedule(kind, chunk_size)
integer (kind=omp_sched_kind) kind
integer chunk_size
```

---

Fortran

---

<sup>6</sup> **Binding**

<sup>7</sup> The binding task set for an **omp_get_schedule** region is the generating task.

<sup>8</sup> **Effect**

<sup>9</sup> This routine returns the *run-sched-var* ICV in the task to which the routine binds. The first
<sup>10</sup> argument *kind* returns the schedule to be used. It can be any of the standard schedule types as
<sup>11</sup> defined in Section 3.2.12 on page 243, or any implementation specific schedule type. The second
<sup>12</sup> argument is interpreted as in the **omp_set_schedule** call, defined in Section 3.2.12 on
<sup>13</sup> page 243.

<sup>14</sup> **Cross References**

<sup>15</sup> - *run-sched-var* ICV, see Section 2.3 on page 36.

<sup>16</sup> - Determining the schedule of a worksharing loop, see Section 2.7.1.1 on page 64.

<sup>17</sup> - **omp_set_schedule** routine, see Section 3.2.12 on page 243.

<sup>18</sup> - **OMP_SCHEDULE** environment variable, see Section 4.1 on page 292.

## 3.2.14 `omp_get_thread_limit`

**Summary**

The `omp_get_thread_limit` routine returns the maximum number of OpenMP threads available to participate in the current contention group.

**Format**

---------------------- C / C++ ----------------------

```
int omp_get_thread_limit(void);
```

---------------------- C / C++ ----------------------
----------------------- Fortran -----------------------

```
integer function omp_get_thread_limit()
```

----------------------- Fortran -----------------------

**Binding**

The binding thread set for an `omp_get_thread_limit` region is all threads on the device. The effect of executing this routine is not related to any specific region corresponding to any construct or API routine.

**Effect**

The `omp_get_thread_limit` routine returns the value of the *thread-limit-var* ICV.

**Cross References**

- *thread-limit-var* ICV, see Section 2.3 on page 36.
- `OMP_THREAD_LIMIT` environment variable, see Section 4.10 on page 300.

## 3.2.15 `omp_set_max_active_levels`

**Summary**

The `omp_set_max_active_levels` routine limits the number of nested active parallel regions on the device, by setting the *max-active-levels-var* ICV

1 **Format**

─────────────── C / C++ ───────────────

```
void omp_set_max_active_levels(int max_levels);
```

─────────────── C / C++ ───────────────

─────────────── Fortran ───────────────

```
subroutine omp_set_max_active_levels(max_levels)
integer max_levels
```

─────────────── Fortran ───────────────

2 **Constraints on Arguments**

3 The value of the argument passed to this routine must evaluate to a non-negative integer, otherwise
4 the behavior of this routine is implementation defined.

5 **Binding**

6 When called from a sequential part of the program, the binding thread set for an
7 **omp_set_max_active_levels** region is the encountering thread. When called from within
8 any explicit parallel region, the binding thread set (and binding region, if required) for the
9 **omp_set_max_active_levels** region is implementation defined.

10 **Effect**

11 The effect of this routine is to set the value of the *max-active-levels-var* ICV to the value specified
12 in the argument.

13 If the number of parallel levels requested exceeds the number of levels of parallelism supported by
14 the implementation, the value of the *max-active-levels-var* ICV will be set to the number of parallel
15 levels supported by the implementation.

16 This routine has the described effect only when called from a sequential part of the program. When
17 called from within an explicit **parallel** region, the effect of this routine is implementation
18 defined.

19 **Cross References**

20 • *max-active-levels-var* ICV, see Section 2.3 on page 36.

21 • **omp_get_max_active_levels** routine, see Section 3.2.16 on page 248.

22 • **OMP_MAX_ACTIVE_LEVELS** environment variable, see Section 4.9 on page 300.

## 3.2.16 `omp_get_max_active_levels`

**Summary**

The `omp_get_max_active_levels` routine returns the value of the *max-active-levels-var*
ICV, which determines the maximum number of nested active parallel regions on the device.

**Format**

---------------------------------- C / C++ ----------------------------------

```
int omp_get_max_active_levels(void);
```

---------------------------------- C / C++ ----------------------------------
---------------------------------- Fortran ----------------------------------

```
integer function omp_get_max_active_levels()
```

---------------------------------- Fortran ----------------------------------

**Binding**

When called from a sequential part of the program, the binding thread set for an
`omp_get_max_active_levels` region is the encountering thread. When called from within
any explicit parallel region, the binding thread set (and binding region, if required) for the
`omp_get_max_active_levels` region is implementation defined.

**Effect**

The `omp_get_max_active_levels` routine returns the value of the *max-active-levels-var*
ICV, which determines the maximum number of nested active parallel regions on the device.

**Cross References**

- *max-active-levels-var* ICV, see Section 2.3 on page 36.
- `omp_set_max_active_levels` routine, see Section 3.2.15 on page 246.
- `OMP_MAX_ACTIVE_LEVELS` environment variable, see Section 4.9 on page 300.

## 3.2.17 `omp_get_level`

**Summary**

The `omp_get_level` routine returns the value of the *levels-var* ICV.

**Format**

------- C / C++ -------

```
int omp_get_level(void);
```

------- C / C++ -------

------- Fortran -------

```
integer function omp_get_level()
```

------- Fortran -------

**Binding**

The binding task set for an `omp_get_level` region is the generating task.

**Effect**

The effect of the `omp_get_level` routine is to return the number of nested `parallel` regions (whether active or inactive) enclosing the current task such that all of the `parallel` regions are enclosed by the outermost initial task region on the current device.

**Cross References**

- *levels-var* ICV, see Section 2.3 on page 36.
- `omp_get_active_level` routine, see Section 3.2.20 on page 252.
- `OMP_MAX_ACTIVE_LEVELS` environment variable, see Section 4.9 on page 300.

## 3.2.18 `omp_get_ancestor_thread_num`

**Summary**

The `omp_get_ancestor_thread_num` routine returns, for a given nested level of the current thread, the thread number of the ancestor of the current thread.

**Format**

--- C / C++ ---

```
int omp_get_ancestor_thread_num(int level);
```

--- C / C++ ---

--- Fortran ---

```
integer function omp_get_ancestor_thread_num(level)
integer level
```

--- Fortran ---
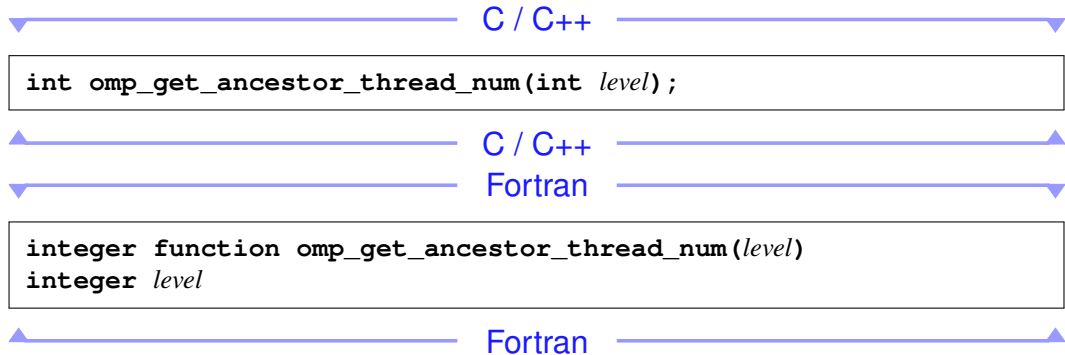
**Binding**

The binding thread set for an `omp_get_ancestor_thread_num` region is the encountering thread. The binding region for an `omp_get_ancestor_thread_num` region is the innermost enclosing **parallel** region.

**Effect**

The `omp_get_ancestor_thread_num` routine returns the thread number of the ancestor at a given nest level of the current thread or the thread number of the current thread. If the requested nest level is outside the range of 0 and the nest level of the current thread, as returned by the `omp_get_level` routine, the routine returns -1.

Note – When the `omp_get_ancestor_thread_num` routine is called with a value of **level**=0, the routine always returns 0. If **level**=`omp_get_level()`, the routine has the same effect as the `omp_get_thread_num` routine.

## 3.2.19  **omp_get_team_size**

### Summary

The **omp_get_team_size** routine returns, for a given nested level of the current thread, the size
of the thread team to which the ancestor or the current thread belongs.

### Format

-------------------------------------------------- C / C++ --------------------------------------------------

```
int omp_get_team_size(int level);
```

-------------------------------------------------- C / C++ --------------------------------------------------

-------------------------------------------------- Fortran --------------------------------------------------

```
integer function omp_get_team_size(level)
integer level
```

-------------------------------------------------- Fortran --------------------------------------------------

### Binding

The binding thread set for an **omp_get_team_size** region is the encountering thread. The
binding region for an **omp_get_team_size** region is the innermost enclosing **parallel**
region.

**Effect**

2  The **omp_get_team_size** routine returns the size of the thread team to which the ancestor or
3  the current thread belongs. If the requested nested level is outside the range of 0 and the nested
4  level of the current thread, as returned by the **omp_get_level** routine, the routine returns -1.
5  Inactive parallel regions are regarded like active parallel regions executed with one thread.

6  Note – When the **omp_get_team_size** routine is called with a value of **level**=0, the routine
7  always returns 1. If **level**=**omp_get_level()**, the routine has the same effect as the
8  **omp_get_num_threads** routine.

9  **Cross References**

10  • **omp_get_num_threads** routine, see Section 3.2.2 on page 232.

11  • **omp_get_level** routine, see Section 3.2.17 on page 249.

12  • **omp_get_ancestor_thread_num** routine, see Section 3.2.18 on page 250.

13  ## 3.2.20 omp_get_active_level

14  **Summary**

15  The **omp_get_active_level** routine returns the value of the *active-level-vars* ICV..

16  **Format**

C / C++

```
int omp_get_active_level(void);
```

C / C++

```
integer function omp_get_active_level()
```

1    **Binding**

2    The binding task set for the an **omp_get_active_level** region is the generating task.

3    **Effect**

4    The effect of the **omp_get_active_level** routine is to return the number of nested, active
5    **parallel** regions enclosing the current task such that all of the **parallel** regions are enclosed
6    by the outermost initial task region on the current device.

7    **Cross References**

8    • *active-levels-var* ICV, see Section 2.3 on page 36.

9    • **omp_get_level** routine, see Section 3.2.17 on page 249.

10   ## 3.2.21  omp_in_final

11   **Summary**

12   The **omp_in_final** routine returns *true* if the routine is executed in a final task region;
13   otherwise, it returns *false*.

14   **Format**

C / C++

```
int omp_in_final(void);
```

C / C++

Fortran

```
logical function omp_in_final()
```

Fortran

**Binding**

The binding task set for an **omp_in_final** region is the generating task.

**Effect**

**omp_in_final** returns *true* if the enclosing task region is final. Otherwise, it returns *false*.

**Cross References**

- **task** construct, see Section 2.9.1 on page 83.

## 3.2.22 `omp_get_proc_bind`

**Summary**

The **omp_get_proc_bind** routine returns the thread affinity policy to be used for the subsequent nested **parallel** regions that do not specify a **proc_bind** clause.

**Format**

———————————————— C / C++ ————————————————

```
omp_proc_bind_t omp_get_proc_bind(void);
```

———————————————— C / C++ ————————————————
———————————————— Fortran ————————————————

```
integer (kind=omp_proc_bind_kind) function omp_get_proc_bind()
```

———————————————— Fortran ————————————————

**Constraints on Arguments**

The value returned by this routine must be one of the valid affinity policy kinds. The C/ C++ header file (**omp.h**) and the Fortran include file (**omp_lib.h**) and/or Fortran 90 module file (**omp_lib**) define the valid constants. The valid constants must include the following:

─────────────────────── C / C++ ───────────────────────

```
typedef enum omp_proc_bind_t {
  omp_proc_bind_false = 0,
  omp_proc_bind_true = 1,
  omp_proc_bind_master = 2,
  omp_proc_bind_close = 3,
  omp_proc_bind_spread = 4
} omp_proc_bind_t;
```

─────────────────────── C / C++ ───────────────────────

─────────────────────── Fortran ───────────────────────

```
integer (kind=omp_proc_bind_kind), &
                parameter :: omp_proc_bind_false = 0
integer (kind=omp_proc_bind_kind), &
                parameter :: omp_proc_bind_true = 1
integer (kind=omp_proc_bind_kind), &
                parameter :: omp_proc_bind_master = 2
integer (kind=omp_proc_bind_kind), &
                parameter :: omp_proc_bind_close = 3
integer (kind=omp_proc_bind_kind), &
                parameter :: omp_proc_bind_spread = 4
```

─────────────────────── Fortran ───────────────────────

**Binding**

The binding task set for an **omp_get_proc_bind** region is the generating task

**Effect**

The effect of this routine is to return the value of the first element of the *bind-var* ICV of the current task. See Section 2.5.2 on page 52 for the rules governing the thread affinity policy.

**Cross References**

- *bind-var* ICV, see Section 2.3 on page 36.

- Controlling OpenMP thread affinity, see Section 2.5.2 on page 52.

- **OMP_PROC_BIND** environment variable, see Section 4.4 on page 294.

## 3.2.23 `omp_get_num_places`

**Summary**

The **omp_get_num_places** routine returns the number of places available to the execution environment in the place list.

**Format**

———————————————— C / C++ ————————————————

```
int omp_get_num_places(void);
```

———————————————— C / C++ ————————————————

———————————————— Fortran ————————————————

```
integer function omp_get_num_places()
```

———————————————— Fortran ————————————————

**Binding**

The binding thread set for an **omp_get_num_places** region is all threads on a device. The effect of executing this routine is not related to any specific region corresponding to any construct or API routine.

**Effect**

The **omp_get_num_places** routine returns the number of places in the place list. This value is equivalent to the number of places in the *place-partition-var* ICV in the execution environment of the initial task.

**Cross References**

- *place-partition-var* ICV, see Section 2.3 on page 36.
- **OMP_PLACES** environment variable, see Section 4.5 on page 295.

## 3.2.24 omp_get_place_num_procs

**Summary**

The **omp_get_place_num_procs** routine returns the number of processors available to the execution environment in the specified place.

**Format**

$\qquad$ C / C++ $\qquad$

```
int omp_get_place_num_procs(int place_num);
```

$\qquad$ C / C++ $\qquad$

$\qquad$ Fortran $\qquad$

```
integer function omp_get_place_num_procs(place_num)
integer place_num
```

$\qquad$ Fortran $\qquad$

**Binding**

The binding thread set for an **omp_get_place_num_procs** region is all threads on a device. The effect of executing this routine is not related to any specific region corresponding to any construct or API routine.

**Effect**

The **omp_get_place_num_procs** routine returns the number of processors associated with the place numbered *place_num*. The routine returns zero when *place_num* is negative, or is equal to or larger than the value returned by **omp_get_num_places()**.

3  ## 3.2.25  `omp_get_place_proc_ids`

4    **Summary**

5    The **omp_get_place_proc_ids** routine returns the numerical identifiers of the processors
6    available to the execution environment in the specified place.

7    **Format**

---
C / C++
---

```
void omp_get_place_proc_ids(int place_num, int *ids);
```

---
C / C++
---

---
Fortran
---

```
subroutine omp_get_place_proc_ids(place_num, ids)
integer place_num
integer ids(*)
```

---
Fortran
---

8    **Binding**

9    The binding thread set for an **omp_get_place_proc_ids** region is all threads on a device.
10   The effect of executing this routine is not related to any specific region corresponding to any
11   construct or API routine.

12   **Effect**

13   The **omp_get_place_proc_ids** routine returns the numerical identifiers of each processor
14   associated with the place numbered *place_num*. The numerical identifiers are non-negative, and
15   their meaning is implementation defined. The numerical identifiers are returned in the array *ids* and
16   their order in the array is implementation defined. The array must be sufficiently large to contain
17   **omp_get_place_num_procs(***place_num***)** integers; otherwise, the behavior is unspecified.
18   The routine has no effect when *place_num* has a negative value, or a value equal or larger than
19   **omp_get_num_places()**.

## 3.2.26 omp_get_place_num

### Summary

The **omp_get_place_num** routine returns the place number of the place to which the encountering thread is bound.

### Format

$C / C_{++}$

```
int omp_get_place_num(void);
```

$C / C_{++}$

Fortran

```
integer function omp_get_place_num()
```

Fortran

### Binding

The binding thread set for an **omp_get_place_num** region is the encountering thread.

### Effect

When the encountering thread is bound to a place, the **omp_get_place_num** routine returns the place number associated with the thread. The returned value is between 0 and one less than the value returned by **omp_get_num_places()**, inclusive. When the encountering thread is not bound to a place, the routine returns -1.

## 3.2.27 `omp_get_partition_num_places`

**Summary**

The **omp_get_partition_num_places** routine returns the number of places in the place partition of the innermost implicit task.

**Format**

------------------------------ C / C++ ------------------------------

```
int omp_get_partition_num_places(void);
```

------------------------------ C / C++ ------------------------------

------------------------------ Fortran ------------------------------

```
integer function omp_get_partition_num_places()
```

------------------------------ Fortran ------------------------------

**Binding**

The binding task set for an **omp_get_partition_num_places** region is the encountering implicit task.

**Effect**

The **omp_get_partition_num_places** routine returns the number of places in the *place-partition-var* ICV.

**Cross References**

• *place-partition-var* ICV, see Section 2.3 on page 36.

• Controlling OpenMP thread affinity, see Section 2.5.2 on page 52.

• **OMP_PLACES** environment variable, see Section 4.5 on page 295.

## 3.2.28 `omp_get_partition_place_nums`

### Summary

The **omp_get_partition_place_nums** routine returns the list of place numbers
corresponding to the places in the *place-partition-var* ICV of the innermost implicit task.

### Format

$C / C_{++}$

```
void omp_get_partition_place_nums(int *place_nums);
```

$C / C_{++}$

Fortran

```
subroutine omp_get_partition_place_nums(place_nums)
integer place_nums(*)
```

Fortran

### Binding

The binding task set for an **omp_get_partition_place_nums** region is the encountering
implicit task.

### Effect

The **omp_get_partition_place_nums** routine returns the list of place numbers
corresponding to the places in the *place-partition-var* ICV of the innermost implicit task. The array
must be sufficiently large to contain **omp_get_partition_num_places()** integers;
otherwise, the behavior is unspecified.

**Cross References**

- *place-partition-var* ICV, see Section 2.3 on page 36.

- Controlling OpenMP thread affinity, see Section 2.5.2 on page 52.

- **omp_get_partition_num_places** routine, see Section 3.2.27 on page 260.

- **OMP_PLACES** environment variable, see Section 4.5 on page 295.


## 3.2.29  omp_set_default_device

**Summary**

The **omp_set_default_device** routine controls the default target device by assigning the
value of the *default-device-var* ICV.

**Format**

──────────────────── C / C++ ────────────────────

```
void omp_set_default_device(int device_num);
```

──────────────────── C / C++ ────────────────────
──────────────────── Fortran ────────────────────

```
subroutine omp_set_default_device(device_num)
integer device_num
```

──────────────────── Fortran ────────────────────

**Binding**

The binding task set for an **omp_set_default_device** region is the generating task.

**Effect**

The effect of this routine is to set the value of the *default-device-var* ICV of the current task to the
value specified in the argument. When called from within a **target** region the effect of this
routine is unspecified.

**Cross References**

- *default-device-var*, see Section 2.3 on page 36.
- **omp_get_default_device**, see Section 3.2.30 on page 263.
- **OMP_DEFAULT_DEVICE** environment variable, see Section 4.13 on page 302

# 3.2.30 `omp_get_default_device`

**Summary**

The **omp_get_default_device** routine returns the default target device.

**Format**

---

C / C++

---

```
int omp_get_default_device(void);
```

---

C / C++

---

---

Fortran

---

```
integer function omp_get_default_device()
```

---

Fortran

---
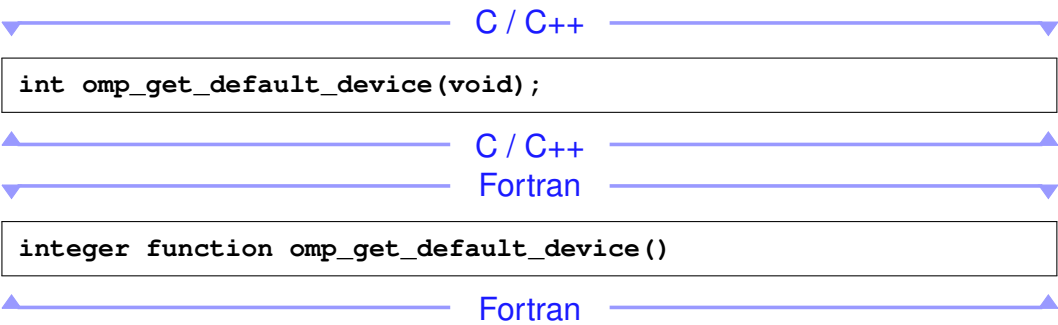
**Binding**

The binding task set for an **omp_get_default_device** region is the generating task.

**Effect**

The **omp_get_default_device** routine returns the value of the *default-device-var* ICV of the current task. When called from within a **target** region the effect of this routine is unspecified.

**Cross References**

- *default-device-var*, see Section 2.3 on page 36.
- **omp_set_default_device**, see Section 3.2.29 on page 262.
- **OMP_DEFAULT_DEVICE** environment variable, see Section 4.13 on page 302.

# 3.2.31 `omp_get_num_devices`

**Summary**

The `omp_get_num_devices` routine returns the number of target devices.

**Format**

---------------- C / C++ ----------------

```
int omp_get_num_devices(void);
```

---------------- C / C++ ----------------

---------------- Fortran ----------------

```
integer function omp_get_num_devices()
```

---------------- Fortran ----------------

**Binding**

The binding task set for an `omp_get_num_devices` region is the generating task.

**Effect**

The `omp_get_num_devices` routine returns the number of available target devices. When called from within a `target` region the effect of this routine is unspecified.

**Cross References**

None.

# 3.2.32 `omp_get_num_teams`

**Summary**

The `omp_get_num_teams` routine returns the number of teams in the current `teams` region.

1 **Format**

---------- C / C++ ----------

```
int omp_get_num_teams(void);
```

---------- C / C++ ----------

---------- Fortran ----------

```
integer function omp_get_num_teams()
```

---------- Fortran ----------

2 **Binding**

3 The binding task set for an **omp_get_num_teams** region is the generating task

4 **Effect**

5 The effect of this routine is to return the number of teams in the current **teams** region. The routine
6 returns 1 if it is called from outside of a **teams** region.

7 **Cross References**

8 • **teams** construct, see Section 2.10.7 on page 114.

## 3.2.33 `omp_get_team_num`

**Summary**

The `omp_get_team_num` routine returns the team number of the calling thread.

**Format**

— C / C++ —

```
int omp_get_team_num(void);
```

— C / C++ —

— Fortran —

```
integer function omp_get_team_num()
```

— Fortran —

**Binding**

The binding task set for an `omp_get_team_num` region is the generating task.

**Effect**

The `omp_get_team_num` routine returns the team number of the calling thread. The team number is an integer between 0 and one less than the value returned by `omp_get_num_teams()`, inclusive. The routine returns 0 if it is called outside of a `teams` region.

**Cross References**

- `teams` construct, see Section 2.10.7 on page 114.
- `omp_get_num_teams` routine, see Section 3.2.32 on page 264.

## 3.2.34 `omp_is_initial_device`

**Summary**

The `omp_is_initial_device` routine returns *true* if the current task is executing on the host device; otherwise, it returns *false*.

**Format**

$$\text{\color{blue}\blacktriangledown\text{————————}\quad C / C{+}{+} \quad\text{————————}\blacktriangledown}$$

```
int omp_is_initial_device(void);
```

$$\text{\color{blue}\blacktriangle\text{————————}\quad C / C{+}{+} \quad\text{————————}\blacktriangle}$$

$$\text{\color{blue}\blacktriangledown\text{————————}\quad Fortran \quad\text{————————}\blacktriangledown}$$

```
logical function omp_is_initial_device()
```

$$\text{\color{blue}\blacktriangle\text{————————}\quad Fortran \quad\text{————————}\blacktriangle}$$
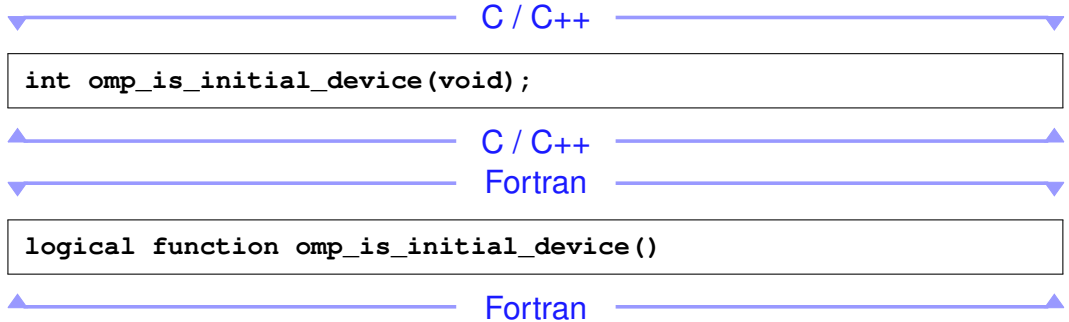
**Binding**

The binding task set for an `omp_is_initial_device` region is the generating task.

**Effect**

The effect of this routine is to return *true* if the current task is executing on the host device; otherwise, it returns *false*.

**Cross References**

- `target` construct, see Section 2.10.4 on page 103

## 3.2.35 `omp_get_initial_device`

**Summary**

The `omp_get_initial_device` routine returns a device number representing the host device.

**Format**

---
C / C++
---

```
int omp_get_initial_device(void);
```

---
C / C++
---

---
Fortran
---

```
integer function omp_get_initial_device()
```

---
Fortran
---

**Binding**

The binding task set for an **omp_get_initial_device** region is the generating task.

**Effect**

The effect of this routine is to return the device number of the host device. The value of the device number is implementation defined. If it is between 0 and one less than **omp_get_num_devices()** then it is valid for use with all device constructs and routines; if it is outside that range, then it is only valid for use with the device memory routines and not in the **device** clause. When called from within a **target** region the effect of this routine is unspecified.

**Cross References**

## 3.2.36  omp_get_max_task_priority

**Summary**

The **omp_get_max_task_priority** routine returns the maximum value that can be specified in the **priority** clause.

<sup>1</sup> **Format**

---
 C / C++
---

```
int omp_get_max_task_priority(void);
```

---
 C / C++
---

---
 Fortran
---

```
integer function omp_get_max_task_priority()
```

---
 Fortran
---

<sup>2</sup> **Binding**

<sup>3</sup> The binding thread set for an **omp_get_max_task_priority** region is all threads on the
<sup>4</sup> device. The effect of executing this routine is not related to any specific region corresponding to
<sup>5</sup> any construct or API routine.

<sup>6</sup> **Effect**

<sup>7</sup> The **omp_get_max_task_priority** routine returns the value of the *max-task-priority-var*
<sup>8</sup> ICV, which determines the maximum value that can be specified in the **priority** clause.

<sup>9</sup> **Cross References**

<sup>10</sup>   • *max-task-priority-var*, see Section 2.3 on page 36.

<sup>11</sup>   • **task** construct, see Section 2.9.1 on page 83.

# 3.3 Lock Routines

The OpenMP runtime library includes a set of general-purpose lock routines that can be used for synchronization. These general-purpose lock routines operate on OpenMP locks that are represented by OpenMP lock variables. OpenMP lock variables must be accessed only through the routines described in this section; programs that otherwise access OpenMP lock variables are non-conforming.

An OpenMP lock can be in one of the following states: *uninitialized*, *unlocked*, or *locked*. If a lock is in the *unlocked* state, a task can *set* the lock, which changes its state to *locked*. The task that sets the lock is then said to *own* the lock. A task that owns a lock can *unset* that lock, returning it to the *unlocked* state. A program in which a task unsets a lock that is owned by another task is non-conforming.

Two types of locks are supported: *simple locks* and *nestable locks*. A *nestable lock* can be set multiple times by the same task before being unset; a *simple lock* cannot be set if it is already owned by the task trying to set it. *Simple lock* variables are associated with *simple locks* and can only be passed to *simple lock* routines. *Nestable lock* variables are associated with *nestable locks* and can only be passed to *nestable lock* routines.

Each type of lock can also have a *lock hint* that contains information about the intended usage of the lock by the application code. The effect of the lock hint is implementation defined. An OpenMP implementation can use this hint to select a usage-specific lock, but lock hints do not change the mutual exclusion semantics of locks. A conforming implementation can safely ignore the lock hint.

Constraints on the state and ownership of the lock accessed by each of the lock routines are described with the routine. If these constraints are not met, the behavior of the routine is unspecified.

The OpenMP lock routines access a lock variable such that they always read and update the most current value of the lock variable. It is not necessary for an OpenMP program to include explicit **flush** directives to ensure that the lock variable's value is consistent among different tasks.

## Binding

The binding thread set for all lock routine regions is all threads in the contention group. As a consequence, for each OpenMP lock, the lock routine effects relate to all tasks that call the routines, without regard to which teams the threads in the contention group executing the tasks belong.

## Simple Lock Routines

$$\text{— C / C++ —}$$

The type **omp_lock_t** represents a simple lock. For the following routines, a simple lock variable must be of **omp_lock_t** type. All simple lock routines require an argument that is a pointer to a variable of type **omp_lock_t**.

$$\text{— C / C++ —}$$

1　For the following routines, a simple lock variable must be an integer variable of
2　**kind=omp_lock_kind**.

3　The simple lock routines are as follows:

4　　• The **omp_init_lock** routine initializes a simple lock.

5　　• The **omp_init_lock_with_hint** routine initializes a simple lock and attaches a hint to it.

6　　• The **omp_destroy_lock** routine uninitializes a simple lock.

7　　• The **omp_set_lock** routine waits until a simple lock is available, and then sets it.

8　　• The **omp_unset_lock** routine unsets a simple lock.

9　　• The **omp_test_lock** routine tests a simple lock, and sets it if it is available.

10　**Nestable Lock Routines**

11　The type **omp_nest_lock_t** represents a nestable lock. For the following routines, a nestable
12　lock variable must be of **omp_nest_lock_t** type. All nestable lock routines require an
13　argument that is a pointer to a variable of type **omp_nest_lock_t**.

14　For the following routines, a nestable lock variable must be an integer variable of
15　**kind=omp_nest_lock_kind**.

16　The nestable lock routines are as follows:

17　　• The **omp_init_nest_lock** routine initializes a nestable lock.

18　　• The **omp_init_nest_lock_with_hint** routine initializes a nestable lock and attaches a
19　　hint to it.

20　　• The **omp_destroy_nest_lock** routine uninitializes a nestable lock.

21　　• The **omp_set_nest_lock** routine waits until a nestable lock is available, and then sets it.

22　　• The **omp_unset_nest_lock** routine unsets a nestable lock.

23　　• The **omp_test_nest_lock** routine tests a nestable lock, and sets it if it is available

**Restrictions**

OpenMP lock routines have the following restrictions:

- The use of the same OpenMP lock in different contention groups results in unspecified behavior.

# 3.3.1 `omp_init_lock` and `omp_init_nest_lock`

**Summary**

These routines initialize an OpenMP lock without a hint.

**Format**

$$\text{C / C++}$$

```
void omp_init_lock(omp_lock_t *lock);
void omp_init_nest_lock(omp_nest_lock_t *lock);
```

$$\text{C / C++}$$

$$\text{Fortran}$$

```
subroutine omp_init_lock(svar)
integer (kind=omp_lock_kind) svar

subroutine omp_init_nest_lock(nvar)
integer (kind=omp_nest_lock_kind) nvar
```

$$\text{Fortran}$$

**Constraints on Arguments**

A program that accesses a lock that is not in the uninitialized state through either routine is non-conforming.

**Effect**

The effect of these routines is to initialize the lock to the unlocked state; that is, no task owns the lock. In addition, the nesting count for a nestable lock is set to zero.

**3.3.2  `omp_init_lock_with_hint` and**
**`omp_init_nest_lock_with_hint`**

3   **Summary**

4   These routines initialize an OpenMP lock with a hint. The effect of the hint is
5   implementation-defined. The OpenMP implementation can ignore the hint without changing
6   program semantics.

7   **Format**

$\qquad\qquad$ C / C++ $\qquad\qquad$

```
void omp_init_lock_with_hint(omp_lock_t *lock,
                               omp_lock_hint_t hint);
void omp_init_nest_lock_with_hint(omp_nest_lock_t *lock,
                                    omp_lock_hint_t hint);
```

$\qquad\qquad$ C / C++ $\qquad\qquad$
$\qquad\qquad$ Fortran $\qquad\qquad$

```
subroutine omp_init_lock_with_hint(svar, hint)
integer (kind=omp_lock_kind) svar
integer (kind=omp_lock_hint_kind) hint

subroutine omp_init_nest_lock_with_hint(nvar, hint)
integer (kind=omp_nest_lock_kind) nvar
integer (kind=omp_lock_hint_kind) hint
```

$\qquad\qquad$ Fortran $\qquad\qquad$

**Constraints on Arguments**

A program that accesses a lock that is not in the uninitialized state through either routine is non-conforming.

The second argument passed to this routine (*hint*) can be one of the valid OpenMP lock hints below or any implementation-defined hint. The C/C++ header file (**omp.h**) and the Fortran include file (**omp_lib.h**) and/or Fortran 90 module file (**omp_lib**) define the valid lock hint constants. The valid constants must include the following, which can be extended with implementation-defined values:

───────────────────────────── C / C++ ─────────────────────────────

```
typedef enum omp_lock_hint_t {
  omp_lock_hint_none = 0,
  omp_lock_hint_uncontended = 1,
  omp_lock_hint_contended = 2,
  omp_lock_hint_nonspeculative = 4,
  omp_lock_hint_speculative = 8
} omp_lock_hint_t;
```

───────────────────────────── C / C++ ─────────────────────────────

───────────────────────────── Fortran ─────────────────────────────

```
integer (kind=omp_lock_hint_kind), &
        parameter :: omp_lock_hint_none = 0
integer (kind=omp_lock_hint_kind), &
        parameter :: omp_lock_hint_uncontended = 1
integer (kind=omp_lock_hint_kind), &
        parameter :: omp_lock_hint_contended = 2
integer (kind=omp_lock_hint_kind), &
        parameter :: omp_lock_hint_nonspeculative = 4
integer (kind=omp_lock_hint_kind), &
        parameter :: omp_lock_hint_speculative = 8
```

───────────────────────────── Fortran ─────────────────────────────

The hints can be combined by using the **+** or **|** operators in C/C++ or the **+** operator in Fortran. The effect of the combined hint is implementation defined and can be ignored by the implementation. Combining **omp_lock_hint_none** with any other hint is equivalent to specifying the other hint. The following restrictions apply to combined hints; violating these restrictions results in unspecified behavior:

- the hints **omp_lock_hint_uncontended** and **omp_lock_hint_contended** cannot be combined,

- the hints **omp_lock_hint_nonspeculative** and **omp_lock_hint_speculative** cannot be combined.

6   **Effect**

7   The effect of these routines is to initialize the lock to the unlocked state and, optionally, to choose a
8   specific lock implementation based on the hint. After initialization no task owns the lock. In
9   addition, the nesting count for a nestable lock is set to zero.

10  ### 3.3.3  `omp_destroy_lock` and
11  `omp_destroy_nest_lock`

12  **Summary**

13  These routines ensure that the OpenMP lock is uninitialized.

14  **Format**

C / C++

```
void omp_destroy_lock(omp_lock_t *lock);
void omp_destroy_nest_lock(omp_nest_lock_t *lock);
```

C / C++

Fortran

```
subroutine omp_destroy_lock(svar)
integer (kind=omp_lock_kind) svar

subroutine omp_destroy_nest_lock(nvar)
integer (kind=omp_nest_lock_kind) nvar
```

Fortran

**Constraints on Arguments**

A program that accesses a lock that is not in the unlocked state through either routine is
non-conforming.

**Effect**

The effect of these routines is to change the state of the lock to uninitialized.

### 3.3.4  `omp_set_lock` and `omp_set_nest_lock`

**Summary**

These routines provide a means of setting an OpenMP lock. The calling task region behaves as if it
was suspended until the lock can be set by this task.

**Format**

--- C / C++ ---

```
void omp_set_lock(omp_lock_t *lock);
void omp_set_nest_lock(omp_nest_lock_t *lock);
```

--- C / C++ ---

--- Fortran ---

```
subroutine omp_set_lock(svar)
integer (kind=omp_lock_kind) svar

subroutine omp_set_nest_lock(nvar)
integer (kind=omp_nest_lock_kind) nvar
```

--- Fortran ---

**Constraints on Arguments**

A program that accesses a lock that is in the uninitialized state through either routine is
non-conforming. A simple lock accessed by **`omp_set_lock`** that is in the locked state must not
be owned by the task that contains the call or deadlock will result.

**Effect**

Each of these routines has an effect equivalent to suspension of the task executing the routine until the specified lock is available.

▼ ▼

Note – The semantics of these routines is specified *as if* they serialize execution of the region guarded by the lock. However, implementations may implement them in other ways provided that the isolation properties are respected so that the actual execution delivers a result that could arise from some serialization.

▲ ▲

A simple lock is available if it is unlocked. Ownership of the lock is granted to the task executing the routine.

A nestable lock is available if it is unlocked or if it is already owned by the task executing the routine. The task executing the routine is granted, or retains, ownership of the lock, and the nesting count for the lock is incremented.

## 3.3.5 `omp_unset_lock` and `omp_unset_nest_lock`

**Summary**

These routines provide the means of unsetting an OpenMP lock.

**Format**

――――――― C / C++ ―――――――

```
void omp_unset_lock(omp_lock_t *lock);
void omp_unset_nest_lock(omp_nest_lock_t *lock);
```

――――――― C / C++ ―――――――
――――――― Fortran ―――――――

```
subroutine omp_unset_lock(svar)
integer (kind=omp_lock_kind) svar

subroutine omp_unset_nest_lock(nvar)
integer (kind=omp_nest_lock_kind) nvar
```

――――――― Fortran ―――――――

**Constraints on Arguments**

2
3
A program that accesses a lock that is not in the locked state or that is not owned by the task that contains the call through either routine is non-conforming.

4
**Effect**

5
For a simple lock, the **omp_unset_lock** routine causes the lock to become unlocked.

6
7
For a nestable lock, the **omp_unset_nest_lock** routine decrements the nesting count, and causes the lock to become unlocked if the resulting nesting count is zero.

8
9
10
For either routine, if the lock becomes unlocked, and if one or more task regions were effectively suspended because the lock was unavailable, the effect is that one task is chosen and given ownership of the lock.

11
## 3.3.6 `omp_test_lock` and `omp_test_nest_lock`

12
**Summary**

13
14
These routines attempt to set an OpenMP lock but do not suspend execution of the task executing the routine.

15
**Format**

$\qquad$ C / C++ $\qquad$

```
int omp_test_lock(omp_lock_t *lock);
int omp_test_nest_lock(omp_nest_lock_t *lock);
```

$\qquad$ C / C++ $\qquad$

$\qquad$ Fortran $\qquad$

```
logical function omp_test_lock(svar)
integer (kind=omp_lock_kind) svar
integer function omp_test_nest_lock(nvar)
integer (kind=omp_nest_lock_kind) nvar
```

$\qquad$ Fortran $\qquad$

**Constraints on Arguments**

A program that accesses a lock that is in the uninitialized state through either routine is non-conforming. The behavior is unspecified if a simple lock accessed by **omp_test_lock** is in the locked state and is owned by the task that contains the call.

**Effect**

These routines attempt to set a lock in the same manner as **omp_set_lock** and **omp_set_nest_lock**, except that they do not suspend execution of the task executing the routine.

For a simple lock, the **omp_test_lock** routine returns *true* if the lock is successfully set; otherwise, it returns *false*.

For a nestable lock, the **omp_test_nest_lock** routine returns the new nesting count if the lock is successfully set; otherwise, it returns zero.

# 3.4 Timing Routines

This section describes routines that support a portable wall clock timer.

## 3.4.1 omp_get_wtime

**Summary**

The **omp_get_wtime** routine returns elapsed wall clock time in seconds.

**Format**

———————————————— C / C++ ————————————————

```
double omp_get_wtime(void);
```

———————————————— C / C++ ————————————————
———————————————— Fortran ————————————————

```
double precision function omp_get_wtime()
```

———————————————— Fortran ————————————————

<sup>1</sup> **Binding**

The binding thread set for an `omp_get_wtime` region is the encountering thread. The routine's return value is not guaranteed to be consistent across any set of threads.

**Effect**

The `omp_get_wtime` routine returns a value equal to the elapsed wall clock time in seconds since some "time in the past". The actual "time in the past" is arbitrary, but it is guaranteed not to change during the execution of the application program. The time returned is a "per-thread time", so it is not required to be globally consistent across all threads participating in an application.

▼ ──────────────────────────────────────────────────────── ▼

Note – It is anticipated that the routine will be used to measure elapsed times as shown in the following example:

───────────────── C / C++ ─────────────────

```
double start;
double end;
start = omp_get_wtime();
... work to be timed ...
end = omp_get_wtime();
printf("Work took %f seconds\n", end - start);
```

───────────────── C / C++ ─────────────────
───────────────── Fortran ─────────────────

```
DOUBLE PRECISION START, END
START = omp_get_wtime()
... work to be timed ...
END = omp_get_wtime()
PRINT *, "Work took", END - START, "seconds"
```

───────────────── Fortran ─────────────────

## 1 3.4.2 `omp_get_wtick`

2 **Summary**

3 The `omp_get_wtick` routine returns the precision of the timer used by `omp_get_wtime`.

4 **Format**

------------------------------ C / C++ ------------------------------

```
double omp_get_wtick(void);
```

------------------------------ C / C++ ------------------------------

------------------------------ Fortran ------------------------------

```
double precision function omp_get_wtick()
```

------------------------------ Fortran ------------------------------
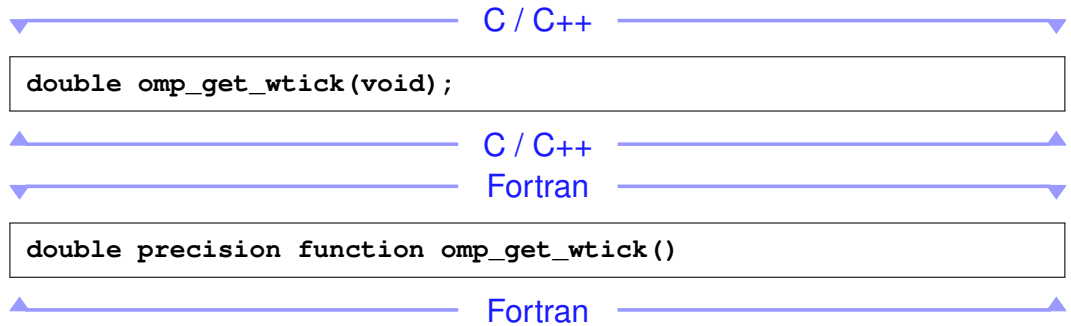
5 **Binding**

6 The binding thread set for an `omp_get_wtick` region is the encountering thread. The routine's
7 return value is not guaranteed to be consistent across any set of threads.

8 **Effect**

9 The `omp_get_wtick` routine returns a value equal to the number of seconds between successive
10 clock ticks of the timer used by `omp_get_wtime`.

# 3.5 Device Memory Routines

This section describes routines that support allocation of memory and management of pointers in the data environments of target devices.

## 3.5.1 omp_target_alloc

**Summary**

The **omp_target_alloc** routine allocates memory in a device data environment.

**Format**

```
void* omp_target_alloc(size_t size, int device_num);
```

**Effect**

The **omp_target_alloc** routine returns the device address of a storage location of *size* bytes. The storage location is dynamically allocated in the device data environment of the device specified by *device_num*, which must be greater than or equal to zero and less than the result of **omp_get_num_devices()** or the result of a call to **omp_get_initial_device()**. When called from within a **target** region the effect of this routine is unspecified.

The **omp_target_alloc** routine returns **NULL** if it cannot dynamically allocate the memory in the device data environment.

The device address returned by **omp_target_alloc** can be used in an **is_device_ptr** clause, Section 2.10.4 on page 103.

Pointer arithmetic is not supported on the device address returned by **omp_target_alloc**.

Freeing the storage returned by **omp_target_alloc** with any routine other than **omp_target_free** results in unspecified behavior.

**Cross References**

- **target** construct, see Section 2.10.4 on page 103
- **omp_get_num_devices** routine, see Section 3.2.31 on page 264
- **omp_get_initial_device** routine, see Section 3.2.35 on page 267
- **omp_target_free** routine, see Section 3.5.2 on page 283

## 3.5.2 **`omp_target_free`**

**Summary**

The **`omp_target_free`** routine frees the device memory allocated by the
**`omp_target_alloc`** routine.

**Format**

```
void omp_target_free(void * device_ptr, int device_num);
```

**Constraints on Arguments**

A program that calls **`omp_target_free`** with a non-**`NULL`** pointer that does not have a value
returned from **`omp_target_alloc`** is non-conforming. The *device_num* must be greater than or
equal to zero and less than the result of **`omp_get_num_devices()`** or the result of a call to
**`omp_get_initial_device()`**.

**Effect**

The **`omp_target_free`** routine frees the memory in the device data environment associated
with *device_ptr*. If *device_ptr* is **`NULL`**, the operation is ignored.

Synchronization must be inserted to ensure that all accesses to *device_ptr* are completed before the
call to **`omp_target_free`**.

When called from within a **`target`** region the effect of this routine is unspecified.

**Cross References**

- **`target`** construct, see Section 2.10.4 on page 103
- **`omp_get_num_devices`** routine, see Section 3.2.31 on page 264
- **`omp_get_initial_device`** routine, see Section 3.2.35 on page 267
- **`omp_target_alloc`** routine, see Section 3.5.1 on page 282

## 3.5.3 **omp_target_is_present**

**Summary**

The **omp_target_is_present** routine tests whether a host pointer has corresponding storage on a given device.

**Format**

```
int omp_target_is_present(void * ptr, int device_num);
```

**Constraints on Arguments**

The value of *ptr* must be a valid host pointer or **NULL**. The *device_num* must be greater than or equal to zero and less than the result of **omp_get_num_devices()** or the result of a call to **omp_get_initial_device()**.

**Effect**

This routine returns *true* if the specified pointer would be found present on device *device_num* by a **map** clause; otherwise, it returns *false*.

When called from within a **target** region the effect of this routine is unspecified.

**Cross References**

- **target** construct, see Section 2.10.4 on page 103
- **map** clause, see Section 2.15.5.1 on page 216.
- **omp_get_num_devices** routine, see Section 3.2.31 on page 264
- **omp_get_initial_device** routine, see Section 3.2.35 on page 267

1    ## 3.5.4 `omp_target_memcpy`

2    **Summary**

3    The **omp_target_memcpy** routine copies memory between any combination of host and device
4    pointers.

5    **Format**

```
int omp_target_memcpy(void * dst, void * src, size_t length,
                      size_t dst_offset, size_t src_offset,
                      int dst_device_num, int src_device_num);
```

6    **Constraints on Arguments**

7    Each device must be compatible with the device pointer specified on the same side of the copy. The
8    *dst_device_num* and *src_device_num* must be greater than or equal to zero and less than the result
9    of **omp_get_num_devices()** or equal to the result of a call to
10   **omp_get_initial_device()**.

11   **Effect**

12   *length* bytes of memory at offset *src_offset* from *src* in the device data environment of device
13   *src_device_num* are copied to *dst* starting at offset *dst_offset* in the device data environment of
14   device *dst_device_num*. The return value is zero on success and non-zero on failure. The host
15   device and host device data environment can be referenced with the device number returned by
16   **omp_get_initial_device**. This routine contains a task scheduling point.

17   When called from within a **target** region the effect of this routine is unspecified.

18   **Cross References**

19   • **target** construct, see Section 2.10.4 on page 103

20   • **omp_get_initial_device** routine, see Section 3.2.35 on page 267

21   • **omp_target_alloc** routine, see Section 3.5.1 on page 282

## 3.5.5 `omp_target_memcpy_rect`

### Summary

The `omp_target_memcpy_rect` routine copies a rectangular subvolume from a
multi-dimensional array to another multi-dimensional array. The copies can use any combination of
host and device pointers.

### Format

```
int omp_target_memcpy_rect(
            void * dst, void * src,
            size_t element_size,
            int num_dims,
            const size_t* volume,
            const size_t* dst_offsets,
            const size_t* src_offsets,
            const size_t* dst_dimensions,
            const size_t* src_dimensions,
            int dst_device_num, int src_device_num);
```

### Constraints on Arguments

The length of the offset and dimension arrays must be at least the value of *num_dims*. The
`dst_device_num` and `src_device_num` must be greater than or equal to zero and less than
the result of `omp_get_num_devices()` or equal to the result of a call to
`omp_get_initial_device()`.

The value of *num_dims* must be between 1 and the implementation-defined limit, which must be at
least three.

1 **Effect**

2 This routine copies a rectangular subvolume of *src*, in the device data environment of device
3 *src_device_num*, to *dst*, in the device data environment of device *dst_device_num*. The volume is
4 specified in terms of the size of an element, number of dimensions, and constant arrays of length
5 *num_dims*. The maximum number of dimensions supported is at least three, support for higher
6 dimensionality is implementation defined. The volume array specifies the length, in number of
7 elements, to copy in each dimension from *src* to *dst*. The *dst_offsets* (*src_offsets*) parameter
8 specifies number of elements from the origin of *dst* (*src*) in elements. The *dst_dimensions*
9 (*src_dimensions*) parameter specifies the length of each dimension of *dst* (*src*)

10 The routine returns zero if successful. If both *dst* and *src* are **NULL** pointers, the routine returns the
11 number of dimensions supported by the implementation for the specified device numbers. The host
12 device and host device data environment can be referenced with the device number returned by
13 **omp_get_initial_device**. Otherwise, it returns a non-zero value. The routine contains a
14 task scheduling point.

15 When called from within a **target** region the effect of this routine is unspecified.

16 **Cross References**

17 • **target** construct, see Section 2.10.4 on page 103

18 • **omp_get_initial_device** routine, see Section 3.2.35 on page 267

19 • **omp_target_alloc** routine, see Section 3.5.1 on page 282

20 **3.5.6  omp_target_associate_ptr**

21 **Summary**

22 The **omp_target_associate_ptr** routine maps a device pointer, which may be returned
23 from **omp_target_alloc** or implementation-defined runtime routines, to a host pointer.

24 **Format**

```
int omp_target_associate_ptr(void * host_ptr, void * device_ptr,
                             size_t size, size_t device_offset,
                             int device_num);
```

**Constraints on Arguments**

The value of *device_ptr* value must be a valid pointer to device memory for the device denoted by the value of *device_num*. The *device_num* argument must be greater than or equal to zero and less than the result of **omp_get_num_devices()** or equal to the result of a call to **omp_get_initial_device()**.

**Effect**

The **omp_target_associate_ptr** routine associates a device pointer in the device data environment of device *device_num* with a host pointer such that when the host pointer appears in a subsequent **map** clause, the associated device pointer is used as the target for data motion associated with that host pointer. The *device_offset* parameter specifies what offset into *device_ptr* will be used as the base address for the device side of the mapping. The reference count of the resulting mapping will be infinite. After being successfully associated, the buffer pointed to by the device pointer is invalidated and accessing data directly through the device pointer results in unspecified behavior. The pointer can be retrieved for other uses by disassociating it. When called from within a **target** region the effect of this routine is unspecified.

The routine returns zero if successful. Otherwise it returns a non-zero value.

Only one device buffer can be associated with a given host pointer value and device number pair. Attempting to associate a second buffer will return non-zero. Associating the same pair of pointers on the same device with the same offset has no effect and returns zero. Associating pointers that share underlying storage will result in unspecified behavior. The **omp_target_is_present** function can be used to test whether a given host pointer has a corresponding variable in the device data environment.

**Cross References**

- **target** construct, see Section 2.10.4 on page 103
- **map** clause, see Section 2.15.5.1 on page 216.
- **omp_target_alloc** routine, see Section 3.5.1 on page 282
- **omp_target_disassociate_ptr** routine, see Section 3.5.6 on page 287

## 3.5.7 `omp_target_disassociate_ptr`

**Summary**

The **`omp_target_disassociate_ptr`** removes the associated pointer for a given device from a host pointer.

**Format**

```
int omp_target_disassociate_ptr(void * ptr, int device_num);
```

**Constraints on Arguments**

The *device_num* must be greater than or equal to zero and less than the result of **`omp_get_num_devices()`** or equal to the result of a call to **`omp_get_initial_device()`**.

**Effect**

The **`omp_target_disassociate_ptr`** removes the associated device data on device *device_num* from the presence table for host pointer *ptr*. A call to this routine on a pointer that is not **`NULL`** and does not have associated data on the given device results in unspecified behavior. The reference count of the mapping is reduced to zero, regardless of its current value.

When called from within a **`target`** region the effect of this routine is unspecified.

After a call to **`omp_target_disassociate_ptr`**, the contents of the device buffer are invalidated.

**Cross References**

- **`target`** construct, see Section 2.10.4 on page 103
- **`omp_target_associate_ptr`** routine, see Section 3.5.6 on page 287

———————————— C / C++ ————————————

<br>

# 2 Environment Variables

3  This chapter describes the OpenMP environment variables that specify the settings of the ICVs that
4  affect the execution of OpenMP programs (see Section 2.3 on page 36). The names of the
5  environment variables must be upper case. The values assigned to the environment variables are
6  case insensitive and may have leading and trailing white space. Modifications to the environment
7  variables after the program has started, even if modified by the program itself, are ignored by the
8  OpenMP implementation. However, the settings of some of the ICVs can be modified during the
9  execution of the OpenMP program by the use of the appropriate directive clauses or OpenMP API
10  routines.

11  The environment variables are as follows:

12  - **OMP_SCHEDULE** sets the *run-sched-var* ICV that specifies the runtime schedule type and chunk
13    size. It can be set to any of the valid OpenMP schedule types.

14  - **OMP_NUM_THREADS** sets the *nthreads-var* ICV that specifies the number of threads to use for
15    parallel regions.

16  - **OMP_DYNAMIC** sets the *dyn-var* ICV that specifies the dynamic adjustment of threads to use for
17    **parallel** regions.

18  - **OMP_PROC_BIND** sets the *bind-var* ICV that controls the OpenMP thread affinity policy.

19  - **OMP_PLACES** sets the *place-partition-var* ICV that defines the OpenMP places that are
20    available to the execution environment.

21  - **OMP_NESTED** sets the *nest-var* ICV that enables or disables nested parallelism.

22  - **OMP_STACKSIZE** sets the *stacksize-var* ICV that specifies the size of the stack for threads
23    created by the OpenMP implementation.

24  - **OMP_WAIT_POLICY** sets the *wait-policy-var* ICV that controls the desired behavior of waiting
25    threads.

26  - **OMP_MAX_ACTIVE_LEVELS** sets the *max-active-levels-var* ICV that controls the maximum
27    number of nested active **parallel** regions.

- **OMP_THREAD_LIMIT** sets the *thread-limit-var* ICV that controls the maximum number of threads participating in a contention group.

- **OMP_CANCELLATION** sets the *cancel-var* ICV that enables or disables cancellation.

- **OMP_DISPLAY_ENV** instructs the runtime to display the OpenMP version number and the initial values of the ICVs, once, during initialization of the runtime.

- **OMP_DEFAULT_DEVICE** sets the *default-device-var* ICV that controls the default device number.

- **OMP_MAX_TASK_PRIORITY** sets the *max-task-priority-var* ICV that specifies the maximum value that can be specified in the **priority** clause of the **task** construct.

The examples in this chapter only demonstrate how these variables might be set in Unix C shell (csh) environments. In Korn shell (ksh) and DOS environments the actions are similar, as follows:

- csh:

```
setenv OMP_SCHEDULE "dynamic"
```

- ksh:

```
export OMP_SCHEDULE="dynamic"
```

- DOS:

```
set OMP_SCHEDULE=dynamic
```

# 4.1 `OMP_SCHEDULE`

The **`OMP_SCHEDULE`** environment variable controls the schedule type and chunk size of all loop directives that have the schedule type **`runtime`**, by setting the value of the *run-sched-var* ICV.

The value of this environment variable takes the form:

*type*[, *chunk*]

where

- *type* is one of **`static`**, **`dynamic`**, **`guided`**, or **`auto`**

- *chunk* is an optional positive integer that specifies the chunk size

If *chunk* is present, there may be white space on either side of the "**,** ". See Section 2.7.1 on page 56 for a detailed description of the schedule types.

The behavior of the program is implementation defined if the value of **`OMP_SCHEDULE`** does not conform to the above format.

Implementation specific schedules cannot be specified in **`OMP_SCHEDULE`**. They can only be specified by calling **`omp_set_schedule`**, described in Section 3.2.12 on page 243.

Examples:

```
setenv OMP_SCHEDULE "guided,4"
setenv OMP_SCHEDULE "dynamic"
```

**Cross References**

- *run-sched-var* ICV, see Section 2.3 on page 36.

- Loop construct, see Section 2.7.1 on page 56.

- Parallel loop construct, see Section 2.11.1 on page 124.

- **`omp_set_schedule`** routine, see Section 3.2.12 on page 243.

- **`omp_get_schedule`** routine, see Section 3.2.13 on page 245.

# 4.2 OMP_NUM_THREADS

The **OMP_NUM_THREADS** environment variable sets the number of threads to use for **parallel** regions by setting the initial value of the *nthreads-var* ICV. See Section 2.3 on page 36 for a comprehensive set of rules about the interaction between the **OMP_NUM_THREADS** environment variable, the **num_threads** clause, the **omp_set_num_threads** library routine and dynamic adjustment of threads, and Section 2.5.1 on page 50 for a complete algorithm that describes how the number of threads for a **parallel** region is determined.

The value of this environment variable must be a list of positive integer values. The values of the list set the number of threads to use for **parallel** regions at the corresponding nested levels.

The behavior of the program is implementation defined if any value of the list specified in the **OMP_NUM_THREADS** environment variable leads to a number of threads which is greater than an implementation can support, or if any value is not a positive integer.

Example:

```
setenv OMP_NUM_THREADS 4,3,2
```

### Cross References

- *nthreads-var* ICV, see Section 2.3 on page 36.
- **num_threads** clause, Section 2.5 on page 46.
- **omp_set_num_threads** routine, see Section 3.2.1 on page 231.
- **omp_get_num_threads** routine, see Section 3.2.2 on page 232.
- **omp_get_max_threads** routine, see Section 3.2.3 on page 233.
- **omp_get_team_size** routine, see Section 3.2.19 on page 251.

# 4.3 OMP_DYNAMIC

The **OMP_DYNAMIC** environment variable controls dynamic adjustment of the number of threads to use for executing **parallel** regions by setting the initial value of the *dyn-var* ICV. The value of this environment variable must be **true** or **false**. If the environment variable is set to **true**, the OpenMP implementation may adjust the number of threads to use for executing **parallel** regions in order to optimize the use of system resources. If the environment variable is set to **false**, the dynamic adjustment of the number of threads is disabled. The behavior of the program is implementation defined if the value of **OMP_DYNAMIC** is neither **true** nor **false**.

Example:

```
setenv OMP_DYNAMIC true
```

### Cross References

- *dyn-var* ICV, see Section 2.3 on page 36.
- **omp_set_dynamic** routine, see Section 3.2.7 on page 237.
- **omp_get_dynamic** routine, see Section 3.2.8 on page 239.

# 4.4 OMP_PROC_BIND

The **OMP_PROC_BIND** environment variable sets the initial value of the *bind-var* ICV. The value of this environment variable is either **true**, **false**, or a comma separated list of **master**, **close**, or **spread**. The values of the list set the thread affinity policy to be used for parallel regions at the corresponding nested level.

If the environment variable is set to **false**, the execution environment may move OpenMP threads between OpenMP places, thread affinity is disabled, and **proc_bind** clauses on **parallel** constructs are ignored.

Otherwise, the execution environment should not move OpenMP threads between OpenMP places, thread affinity is enabled, and the initial thread is bound to the first place in the OpenMP place list.

The behavior of the program is implementation defined if the value in the **OMP_PROC_BIND** environment variable is not **true**, **false**, or a comma separated list of **master**, **close**, or **spread**. The behavior is also implementation defined if an initial thread cannot be bound to the first place in the OpenMP place list.

1    Examples:

```
setenv OMP_PROC_BIND false
setenv OMP_PROC_BIND "spread, spread, close"
```

2    **Cross References**

3    • *bind-var* ICV, see Section 2.3 on page 36.

4    • **proc_bind** clause, see Section 2.5.2 on page 52.

5    • **omp_get_proc_bind** routine, see Section 3.2.22 on page 254.

# 6   **4.5  OMP_PLACES**

7    A list of places can be specified in the **OMP_PLACES** environment variable. The
8    *place-partition-var* ICV obtains its initial value from the **OMP_PLACES** value, and makes the list
9    available to the execution environment. The value of **OMP_PLACES** can be one of two types of
10   values: either an abstract name describing a set of places or an explicit list of places described by
11   non-negative numbers.

12   The **OMP_PLACES** environment variable can be defined using an explicit ordered list of
13   comma-separated places. A place is defined by an unordered set of comma-separated non-negative
14   numbers enclosed by braces. The meaning of the numbers and how the numbering is done are
15   implementation defined. Generally, the numbers represent the smallest unit of execution exposed by
16   the execution environment, typically a hardware thread.

17   Intervals may also be used to define places. Intervals can be specified using the *<lower-bound>* :
18   *<length>* : *<stride>* notation to represent the following list of numbers: "*<lower-bound>*,
19   *<lower-bound>* + *<stride>*, ..., *<lower-bound>* + (*<length>*- 1)*<stride>*." When *<stride>* is
20   omitted, a unit stride is assumed. Intervals can specify numbers within a place as well as sequences
21   of places.

22   An exclusion operator "**!**" can also be used to exclude the number or place immediately following
23   the operator.

24   Alternatively, the abstract names listed in Table 4.1 should be understood by the execution and
25   runtime environment. The precise definitions of the abstract names are implementation defined. An
26   implementation may also add abstract names as appropriate for the target platform.

27   The abstract name may be appended by a positive number in parentheses to denote the length of the
28   place list to be created, that is *abstract_name(num-places)*. When requesting fewer places than

available on the system, the determination of which resources of type *abstract_name* are to be included in the place list is implementation defined. When requesting more resources than available, the length of the place list is implementation defined.

**TABLE 4.1:** Defined Abstract Names for **OMP_PLACES**

| Abstract Name | Meaning |
|---|---|
| **threads** | Each place corresponds to a single hardware thread on the target machine. |
| **cores** | Each place corresponds to a single core (having one or more hardware threads) on the target machine. |
| **sockets** | Each place corresponds to a single socket (consisting of one or more cores) on the target machine. |

The behavior of the program is implementation defined when the execution environment cannot map a numerical value (either explicitly defined or implicitly derived from an interval) within the **OMP_PLACES** list to a processor on the target platform, or if it maps to an unavailable processor. The behavior is also implementation defined when the **OMP_PLACES** environment variable is defined using an abstract name.

The following grammar describes the values accepted for the **OMP_PLACES** environment variable.

$$
\begin{aligned}
\langle \text{list} \rangle &\models \langle \text{p-list} \rangle \mid \langle \text{aname} \rangle \\
\langle \text{p-list} \rangle &\models \langle \text{p-interval} \rangle \mid \langle \text{p-list} \rangle , \langle \text{p-interval} \rangle \\
\langle \text{p-interval} \rangle &\models \langle \text{place} \rangle : \langle \text{len} \rangle : \langle \text{stride} \rangle \mid \langle \text{place} \rangle : \langle \text{len} \rangle \mid \langle \text{place} \rangle \mid !\langle \text{place} \rangle \\
\langle \text{place} \rangle &\models \{\langle \text{res-list} \rangle\} \\
\langle \text{res-list} \rangle &\models \langle \text{res-interval} \rangle \mid \langle \text{res-list} \rangle , \langle \text{res-interval} \rangle \\
\langle \text{res-interval} \rangle &\models \langle \text{res} \rangle : \langle \text{num-places} \rangle : \langle \text{stride} \rangle \mid \langle \text{res} \rangle : \langle \text{num-places} \rangle \mid \langle \text{res} \rangle \mid !\langle \text{res} \rangle \\
\langle \text{aname} \rangle &\models \langle \text{word} \rangle (\langle \text{num-places} \rangle) \mid \langle \text{word} \rangle \\
\langle \text{word} \rangle &\models \text{sockets} \mid \text{cores} \mid \text{threads} \mid \text{<implementation-defined abstract name>} \\
\langle \text{res} \rangle &\models \textit{non-negative integer} \\
\langle \text{num-places} \rangle &\models \textit{positive integer} \\
\langle \text{stride} \rangle &\models \textit{integer} \\
\langle \text{len} \rangle &\models \textit{positive integer}
\end{aligned}
$$

1    Examples:

```
setenv OMP_PLACES threads
setenv OMP_PLACES "threads(4)"
setenv OMP_PLACES "{0,1,2,3},{4,5,6,7},{8,9,10,11},{12,13,14,15}"
setenv OMP_PLACES "{0:4},{4:4},{8:4},{12:4}"
setenv OMP_PLACES "{0:4}:4:4"
```

2    where each of the last three definitions corresponds to the same 4 places including the smallest
3    units of execution exposed by the execution environment numbered, in turn, 0 to 3, 4 to 7, 8 to 11,
4    and 12 to 15.

5    **Cross References**

6    • *place-partition-var*, Section 2.3 on page 36.

7    • Controlling OpenMP thread affinity, Section 2.5.2 on page 52.

8    • **omp_get_num_places** routine, see Section 3.2.23 on page 256.

9    • **omp_get_place_num_procs** routine, see Section 3.2.24 on page 257.

10    • **omp_get_place_proc_ids** routine, see Section 3.2.25 on page 258.

11    • **omp_get_place_num** routine, see Section 3.2.26 on page 259.

12    • **omp_get_partition_num_places** routine, see Section 3.2.27 on page 260.

13    • **omp_get_partition_place_nums** routine, see Section 3.2.28 on page 261.

## 14    **4.6    OMP_NESTED**

15    The **OMP_NESTED** environment variable controls nested parallelism by setting the initial value of
16    the *nest-var* ICV. The value of this environment variable must be **true** or **false**. If the
17    environment variable is set to **true**, nested parallelism is enabled; if set to **false**, nested
18    parallelism is disabled. The behavior of the program is implementation defined if the value of
19    **OMP_NESTED** is neither **true** nor **false**.

20    Example:

```
setenv OMP_NESTED false
```

**Cross References**

- *nest-var* ICV, see Section 2.3 on page 36.

- **omp_set_nested** routine, see Section 3.2.10 on page 240.

- **omp_get_team_size** routine, see Section 3.2.19 on page 251.

# 4.7 **OMP_STACKSIZE**

The **OMP_STACKSIZE** environment variable controls the size of the stack for threads created by the OpenMP implementation, by setting the value of the *stacksize-var* ICV. The environment variable does not control the size of the stack for an initial thread.

The value of this environment variable takes the form:

*size* | *size***B** | *size***K** | *size***M** | *size***G**

where:

- *size* is a positive integer that specifies the size of the stack for threads that are created by the OpenMP implementation.

- **B**, **K**, **M**, and **G** are letters that specify whether the given size is in Bytes, Kilobytes (1024 Bytes), Megabytes (1024 Kilobytes), or Gigabytes (1024 Megabytes), respectively. If one of these letters is present, there may be white space between *size* and the letter.

If only *size* is specified and none of **B**, **K**, **M**, or **G** is specified, then *size* is assumed to be in Kilobytes.

The behavior of the program is implementation defined if **OMP_STACKSIZE** does not conform to the above format, or if the implementation cannot provide a stack with the requested size.

Examples:

```
setenv OMP_STACKSIZE 2000500B
setenv OMP_STACKSIZE "3000 k "
setenv OMP_STACKSIZE 10M
setenv OMP_STACKSIZE " 10 M "
setenv OMP_STACKSIZE "20 m "
setenv OMP_STACKSIZE " 1G"
setenv OMP_STACKSIZE 20000
```

1 **Cross References**

2 • *stacksize-var* ICV, see Section 2.3 on page 36.

# 4.8 `OMP_WAIT_POLICY`

4 The `OMP_WAIT_POLICY` environment variable provides a hint to an OpenMP implementation
5 about the desired behavior of waiting threads by setting the *wait-policy-var* ICV. A compliant
6 OpenMP implementation may or may not abide by the setting of the environment variable.

7 The value of this environment variable takes the form:

8 `ACTIVE | PASSIVE`

9 The `ACTIVE` value specifies that waiting threads should mostly be active, consuming processor
10 cycles, while waiting. An OpenMP implementation may, for example, make waiting threads spin.

11 The `PASSIVE` value specifies that waiting threads should mostly be passive, not consuming
12 processor cycles, while waiting. For example, an OpenMP implementation may make waiting
13 threads yield the processor to other threads or go to sleep.

14 The details of the `ACTIVE` and `PASSIVE` behaviors are implementation defined.

15 Examples:

```
setenv OMP_WAIT_POLICY ACTIVE
setenv OMP_WAIT_POLICY active
setenv OMP_WAIT_POLICY PASSIVE
setenv OMP_WAIT_POLICY passive
```

16 **Cross References**

17 • *wait-policy-var* ICV, see Section 2.3 on page 36.

## 4.9 `OMP_MAX_ACTIVE_LEVELS`

The **OMP_MAX_ACTIVE_LEVELS** environment variable controls the maximum number of nested active **parallel** regions by setting the initial value of the *max-active-levels-var* ICV.

The value of this environment variable must be a non-negative integer. The behavior of the program is implementation defined if the requested value of **OMP_MAX_ACTIVE_LEVELS** is greater than the maximum number of nested active parallel levels an implementation can support, or if the value is not a non-negative integer.

**Cross References**

- *max-active-levels-var* ICV, see Section 2.3 on page 36.
- **omp_set_max_active_levels** routine, see Section 3.2.15 on page 246.
- **omp_get_max_active_levels** routine, see Section 3.2.16 on page 248.

## 4.10 `OMP_THREAD_LIMIT`

The **OMP_THREAD_LIMIT** environment variable sets the maximum number of OpenMP threads to use in a contention group by setting the *thread-limit-var* ICV.

The value of this environment variable must be a positive integer. The behavior of the program is implementation defined if the requested value of **OMP_THREAD_LIMIT** is greater than the number of threads an implementation can support, or if the value is not a positive integer.

**Cross References**

- *thread-limit-var* ICV, see Section 2.3 on page 36.
- **omp_get_thread_limit** routine, see Section 3.2.14 on page 246.

## 4.11 `OMP_CANCELLATION`

The **OMP_CANCELLATION** environment variable sets the initial value of the *cancel-var* ICV.

The value of this environment variable must be **true** or **false**. If set to **true**, the effects of the **cancel** construct and of cancellation points are enabled and cancellation is activated. If set to **false**, cancellation is disabled and the **cancel** construct and cancellation points are effectively ignored.

**Cross References**

- *cancel-var*, see Section 2.3.1 on page 36.
- **cancel** construct, see Section 2.14.1 on page 172.
- **cancellation point** construct, see Section 2.14.2 on page 176.
- **omp_get_cancellation** routine, see Section 3.2.9 on page 240.

# 4.12 **OMP_DISPLAY_ENV**

The **OMP_DISPLAY_ENV** environment variable instructs the runtime to display the OpenMP version number and the value of the ICVs associated with the environment variables described in Chapter 4, as *name* = *value* pairs. The runtime displays this information once, after processing the environment variables and before any user calls to change the ICV values by runtime routines defined in Chapter 3.

The value of the **OMP_DISPLAY_ENV** environment variable may be set to one of these values:

**TRUE | FALSE | VERBOSE**

The **TRUE** value instructs the runtime to display the OpenMP version number defined by the **_OPENMP** version macro (or the **openmp_version** Fortran parameter) value and the initial ICV values for the environment variables listed in Chapter 4. The **VERBOSE** value indicates that the runtime may also display the values of runtime variables that may be modified by vendor-specific environment variables. The runtime does not display any information when the **OMP_DISPLAY_ENV** environment variable is **FALSE** or undefined. For all values of the environment variable other than **TRUE**, **FALSE**, and **VERBOSE**, the displayed information is unspecified.

The display begins with "OPENMP DISPLAY ENVIRONMENT BEGIN", followed by the **_OPENMP** version macro (or the **openmp_version** Fortran parameter) value and ICV values, in the format *NAME* '=' *VALUE*. *NAME* corresponds to the macro or environment variable name, optionally prepended by a bracketed *device-type*. *VALUE* corresponds to the value of the macro or ICV associated with this environment variable. Values should be enclosed in single quotes. The display is terminated with "OPENMP DISPLAY ENVIRONMENT END".

1     Example:

```
% setenv OMP_DISPLAY_ENV TRUE
```

2     The above example causes an OpenMP implementation to generate output of the following form:

```
OPENMP DISPLAY ENVIRONMENT BEGIN
  _OPENMP='201511'
  [host] OMP_SCHEDULE='GUIDED,4'
  [host] OMP_NUM_THREADS='4,3,2'
  [device] OMP_NUM_THREADS='2'
  [host,device] OMP_DYNAMIC='TRUE'
  [host] OMP_PLACES='0:4,4:4,8:4,12:4'
  ...
OPENMP DISPLAY ENVIRONMENT END
```

# 3  **4.13**  **OMP_DEFAULT_DEVICE**

4     The **OMP_DEFAULT_DEVICE** environment variable sets the device number to use in device
5     constructs by setting the initial value of the *default-device-var* ICV.

6     The value of this environment variable must be a non-negative integer value.

7     **Cross References**

8     • *default-device-var* ICV, see Section 2.3 on page 36.

9     • device constructs, Section 2.10 on page 95.

## 1 4.14   `OMP_MAX_TASK_PRIORITY`

2   The **`OMP_MAX_TASK_PRIORITY`** environment variable controls the use of task priorities by
3   setting the initial value of the *max-task-priority-var* ICV. The value of this environment variable
4   must be a non-negative integer.

5   Example:

```
% setenv OMP_MAX_TASK_PRIORITY 20
```

6   **Cross References**

7   • *max-task-priority-var* ICV, see Section 2.3 on page 36.

8   • Tasking Constructs, see Section 2.9 on page 83.

9   • **`omp_get_max_task_priority`** routine, see Section 3.2.36 on page 268.

**APPENDIX A**

# Stubs for Runtime Library Routines

---

This section provides stubs for the runtime library routines defined in the OpenMP API. The stubs
are provided to enable portability to platforms that do not support the OpenMP API. On these
platforms, OpenMP programs must be linked with a library containing these stub routines. The stub
routines assume that the directives in the OpenMP program are ignored. As such, they emulate
serial semantics executing on the host.

Note that the lock variable that appears in the lock routines must be accessed exclusively through
these routines. It should not be initialized or otherwise modified in the user program.

In an actual implementation the lock variable might be used to hold the address of an allocated
memory block, but here it is used to hold an integer value. Users should not make assumptions
about mechanisms used by OpenMP implementations to implement locks based on the scheme
used by the stub procedures.

Fortran

Note – In order to be able to compile the Fortran stubs file, the include file **omp_lib.h** was split
into two files: **omp_lib_kinds.h** and **omp_lib.h** and the **omp_lib_kinds.h** file included
where needed. There is no requirement for the implementation to provide separate files.

Fortran

# A.1   C/C++ Stub Routines

```c
#include <stdio.h>
#include <stdlib.h>
#include "omp.h"

void omp_set_num_threads(int num_threads)
{
}

int omp_get_num_threads(void)
{
    return 1;
}

int omp_get_max_threads(void)
{
    return 1;
}

int omp_get_thread_num(void)
{
    return 0;
}

int omp_get_num_procs(void)
{
    return 1;
}

int omp_in_parallel(void)
{
    return 0;
}

void omp_set_dynamic(int dynamic_threads)
{
}

int omp_get_dynamic(void)
{
    return 0;
}

int omp_get_cancellation(void)
{
    return 0;
```

```
1              }
2
3              void omp_set_nested(int nested)
4              {
5              }
6
7              int omp_get_nested(void)
8              {
9                  return 0;
10             }
11
12             void omp_set_schedule(omp_sched_t kind, int chunk_size)
13             {
14             }
15
16             void omp_get_schedule(omp_sched_t *kind, int *chunk_size)
17             {
18                 *kind = omp_sched_static;
19                 *chunk_size = 0;
20             }
21
22             int omp_get_thread_limit(void)
23             {
24                 return 1;
25             }
26
27             void omp_set_max_active_levels(int max_active_levels)
28             {
29             }
30
31             int omp_get_max_active_levels(void)
32             {
33                 return 0;
34             }
35
36             int omp_get_level(void)
37             {
38                 return 0;
39             }
40
41             int omp_get_ancestor_thread_num(int level)
42             {
43                 if (level == 0)
44                 {
45                     return 0;
46                 }
47                 else
```

```
      {
          return -1;
      }
  }

  int omp_get_team_size(int level)
  {
      if (level == 0)
      {
          return 1;
      }
      else
      {
          return -1;
      }
  }

  int omp_get_active_level(void)
  {
      return 0;
  }

  int omp_in_final(void)
  {
      return 1;
  }

  omp_proc_bind_t omp_get_proc_bind(void)
  {
      return omp_proc_bind_false;
  }

  int omp_get_num_places(void)
  {
      return 0;
  }

  int omp_get_place_num_procs(int place_num)
  {
      return 0;
  }

  void omp_get_place_proc_ids(int place_num, int *ids)
  {
  }

  int omp_get_place_num(void)
```

```
{
    return -1;
}

int omp_get_partition_num_places(void)
{
    return 0;
}

void omp_get_partition_place_nums(int *place_nums)
{
}

void omp_set_default_device(int device_num)
{
}

int omp_get_default_device(void)
{
    return 0;
}

int omp_get_num_devices(void)
{
    return 0;
}

int omp_get_num_teams(void)
{
    return 1;
}

int omp_get_team_num(void)
{
    return 0;
}

int omp_is_initial_device(void)
{
    return 1;
}

int omp_get_initial_device(void)
{
    return -10;
}
```

```
1    int omp_get_max_task_priority(void)
2    {
3        return 0;
4    }
5
6    struct __omp_lock
7    {
8        int lock;
9    };
10
11   enum { UNLOCKED = -1, INIT, LOCKED };
12
13   void omp_init_lock(omp_lock_t *arg)
14   {
15       struct __omp_lock *lock = (struct __omp_lock *)arg;
16       lock->lock = UNLOCKED;
17   }
18
19   void omp_init_lock_with_hint(omp_lock_t *arg, omp_lock_hint_t hint)
20   {
21       omp_init_lock(arg);
22   }
23
24   void omp_destroy_lock(omp_lock_t *arg)
25   {
26       struct __omp_lock *lock = (struct __omp_lock *)arg;
27       lock->lock = INIT;
28   }
29
30   void omp_set_lock(omp_lock_t *arg)
31   {
32       struct __omp_lock *lock = (struct __omp_lock *)arg;
33       if (lock->lock == UNLOCKED)
34       {
35           lock->lock = LOCKED;
36       }
37       else if (lock->lock == LOCKED)
38       {
39           fprintf(stderr, "error: deadlock in using lock variable\n");
40           exit(1);
41       }
42
43       else
44       {
45           fprintf(stderr, "error: lock not initialized\n");
46           exit(1);
47       }
```

```
1              }
2
3          void omp_unset_lock(omp_lock_t *arg)
4          {
5              struct __omp_lock *lock = (struct __omp_lock *)arg;
6              if (lock->lock == LOCKED)
7              {
8                  lock->lock = UNLOCKED;
9              }
10             else if (lock->lock == UNLOCKED)
11             {
12                 fprintf(stderr, "error: lock not set\n");
13                 exit(1);
14             }
15             else
16             {
17                 fprintf(stderr, "error: lock not initialized\n");
18                 exit(1);
19             }
20         }
21
22         int omp_test_lock(omp_lock_t *arg)
23         {
24             struct __omp_lock *lock = (struct __omp_lock *)arg;
25             if (lock->lock == UNLOCKED)
26             {
27                 lock->lock = LOCKED;
28                 return 1;
29             }
30             else if (lock->lock == LOCKED)
31             {
32                 return 0;
33             }
34             else
35             {
36                 fprintf(stderr, "error: lock not initialized\ n");
37                 exit(1);
38             }
39         }
40
41         struct __omp_nest_lock
42         {
43             short owner;
44             short count;
45         };
46
47         enum { NOOWNER = -1, MASTER = 0 };
```

```
1
2        void omp_init_nest_lock(omp_nest_lock_t *arg)
3        {
4            struct __omp_nest_lock *nlock=(struct __omp_nest_lock *)arg;
5            nlock->owner = NOOWNER;
6            nlock->count = 0;
7        }
8
9        void omp_init_nest_lock_with_hint(omp_nest_lock_t *arg,
10                                          omp_lock_hint_t hint)
11        {
12            omp_init_nest_lock(arg);
13        }
14
15        void omp_destroy_nest_lock(omp_nest_lock_t *arg)
16        {
17            struct __omp_nest_lock *nlock=(struct __omp_nest_lock *)arg;
18            nlock->owner = NOOWNER;
19            nlock->count = UNLOCKED;
20        }
21
22        void omp_set_nest_lock(omp_nest_lock_t *arg)
23        {
24            struct __omp_nest_lock *nlock=(struct __omp_nest_lock *)arg;
25            if (nlock->owner == MASTER && nlock->count >= 1)
26            {
27                nlock->count++;
28            }
29            else if (nlock->owner == NOOWNER && nlock->count == 0)
30            {
31                nlock->owner = MASTER;
32                nlock->count = 1;
33            }
34            else
35            {
36                fprintf(stderr, "error: lock corrupted or not initialized\n");
37                exit(1);
38            }
39        }
40
41        void omp_unset_nest_lock(omp_nest_lock_t *arg)
42        {
43            struct __omp_nest_lock *nlock=(struct __omp_nest_lock *)arg;
44            if (nlock->owner == MASTER && nlock->count >= 1)
45            {
46                nlock->count--;
47                if (nlock->count == 0)
```

```
  1                     {
  2                         nlock->owner = NOOWNER;
  3                     }
  4                 }
  5                 else if (nlock->owner == NOOWNER && nlock->count == 0)
  6                 {
  7                     fprintf(stderr, "error: lock not set\n");
  8                     exit(1);
  9                 }
 10                 else
 11                 {
 12                     fprintf(stderr, "error: lock corrupted or not initialized\n");
 13                     exit(1);
 14                 }
 15             }
 16
 17             int omp_test_nest_lock(omp_nest_lock_t *arg)
 18             {
 19                 struct __omp_nest_lock *nlock=(struct __omp_nest_lock *)arg;
 20                 omp_set_nest_lock(arg);
 21                 return nlock->count;
 22             }
 23
 24             double omp_get_wtime(void)
 25             {
 26             /* This function does not provide a working
 27              * wallclock timer. Replace it with a version
 28              * customized for the target machine.
 29              */
 30                 return 0.0;
 31             }
 32
 33             double omp_get_wtick(void)
 34             {
 35             /* This function does not provide a working
 36              * clock tick function. Replace it with
 37              * a version customized for the target machine.
 38              */
 39                 return 365. * 86400.;
 40             }
 41
 42             void * omp_target_alloc(size_t size, int device_num)
 43             {
 44                 if (device_num != -10)
 45                   return NULL;
 46                 return malloc(size)
 47             }
```

```
1
2          void omp_target_free(void *device_ptr, int device_num)
3          {
4              free(device_ptr);
5          }
6
7          int omp_target_is_present(void *ptr, int device_num)
8          {
9              return 1;
10         }
11
12         int omp_target_memcpy(void *dst, void *src, size_t length,
13                               size_t dst_offset, size_t src_offset,
14                               int dst_device, int src_device)
15         {
16             // only the default device is valid in a stub
17             if (dst_device != -10 || src_device != -10
18                     || ! dst || ! src )
19                 return EINVAL;
20             memcpy((char *)dst + dst_offset,
21                     (char *)src + src_offset,
22                     length);
23             return 0;
24         }
25
26         int omp_target_memcpy_rect(
27             void *dst, void *src,
28             size_t element_size,
29             int num_dims,
30             const size_t *volume,
31             const size_t *dst_offsets,
32             const size_t *src_offsets,
33             const size_t *dst_dimensions,
34             const size_t *src_dimensions,
35             int dst_device_num, int src_device_num)
36         {
37             int ret=0;
38             // Both null, return number of dimensions supported,
39             // this stub supports an arbitrary number
40             if (dst == NULL && src == NULL) return INT_MAX;
41
42             if (!volume || !dst_offsets || !src_offsets
43                     || !dst_dimensions || !src_dimensions
44                     || num_dims < 1 ) {
45                 ret = EINVAL;
46                 goto done;
47             }
```

```
 1                    if (num_dims == 1) {
 2                        ret = omp_target_memcpy(dst, src,
 3                                                element_size * volume[0],
 4                                                dst_offsets[0] * element_size,
 5                                                src_offsets[0] * element_size,
 6                                                dst_device_num, src_device_num);
 7                        if(ret) goto done;
 8                    } else {
 9                        size_t dst_slice_size = element_size;
10                        size_t src_slice_size = element_size;
11                        for (int i=1; i < num_dims; i++) {
12                            dst_slice_size *= dst_dimensions[i];
13                            src_slice_size *= src_dimensions[i];
14                        }
15                        size_t dst_off = dst_offsets[0] * dst_slice_size;
16                        size_t src_off = src_offsets[0] * src_slice_size;
17                        for (size_t i=0; i < volume[0]; i++) {
18                            ret = omp_target_memcpy_rect(
19                                        (char *)dst + dst_off + dst_slice_size*i,
20                                        (char *)src + src_off + src_slice_size*i,
21                                        element_size,
22                                        num_dims - 1,
23                                        volume + 1,
24                                        dst_offsets + 1,
25                                        src_offsets + 1,
26                                        dst_dimensions + 1,
27                                        src_dimensions + 1,
28                                        dst_device_num,
29                                        src_device_num);
30                            if (ret) goto done;
31                        }
32                    }
33                done:
34                    return ret;
35                }
36
37                int omp_target_associate_ptr(void *host_ptr, void *device_ptr,
38                                             size_t size, size_t device_offset,
39                                             int device_num)
40                {
41                    // No association is possible because all host pointers
42                    // are considered present
43                    return EINVAL;
44                }
45
46                int omp_target_disassociate_ptr(void *ptr, int device_num)
47                {
```

```
1              return EINVAL;
2          }
3
```

# A.2   Fortran Stub Routines

```fortran
      subroutine omp_set_num_threads(num_threads)
        integer num_threads
        return
      end subroutine

      integer function omp_get_num_threads()
        omp_get_num_threads = 1
        return
      end function

      integer function omp_get_max_threads()
        omp_get_max_threads = 1
        return
      end function

      integer function omp_get_thread_num()
        omp_get_thread_num = 0
        return
      end function

      integer function omp_get_num_procs()
        omp_get_num_procs = 1
        return
      end function

      logical function omp_in_parallel()
        omp_in_parallel = .false.
        return
      end function

      subroutine omp_set_dynamic(dynamic_threads)
        logical dynamic_threads
        return
      end subroutine

      logical function omp_get_dynamic()
        omp_get_dynamic = .false.
        return
      end function

      logical function omp_get_cancellation()
        omp_get_cancellation = .false.
        return
      end function
```

```fortran
          subroutine omp_set_nested(nested)
            logical nested
            return
          end subroutine

          logical function omp_get_nested()
            omp_get_nested = .false.
            return
          end function

          subroutine omp_set_schedule(kind, chunk_size)
            include 'omp_lib_kinds.h'
            integer (kind=omp_sched_kind) kind
            integer chunk_size
            return
          end subroutine

          subroutine omp_get_schedule(kind, chunk_size)
            include 'omp_lib_kinds.h'
            integer (kind=omp_sched_kind) kind
            integer chunk_size
            kind = omp_sched_static
            chunk_size = 0
            return
          end subroutine

          integer function omp_get_thread_limit()
            omp_get_thread_limit = 1
            return
          end function

          subroutine omp_set_max_active_levels(max_level)
            integer max_level
          end subroutine

          integer function omp_get_max_active_levels()
            omp_get_max_active_levels = 0
            return
          end function

          integer function omp_get_level()
            omp_get_level = 0
            return
          end function

          integer function omp_get_ancestor_thread_num(level)
            integer level
```

```fortran
      if ( level .eq. 0 ) then
         omp_get_ancestor_thread_num = 0
      else
         omp_get_ancestor_thread_num = -1
      end if
      return
    end function

    integer function omp_get_team_size(level)
      integer level
      if ( level .eq. 0 ) then
         omp_get_team_size = 1
      else
         omp_get_team_size = -1
      end if
      return
    end function

    integer function omp_get_active_level()
      omp_get_active_level = 0
      return
    end function

    logical function omp_in_final()
      omp_in_final = .true.
      return
    end function

    function omp_get_proc_bind()
      include 'omp_lib_kinds.h'
      integer (kind=omp_proc_bind_kind) omp_get_proc_bind
      omp_get_proc_bind = omp_proc_bind_false
    end function

    integer function omp_get_num_places()
      return 0
    end function

    integer function omp_get_place_num_procs(place_num)
      integer place_num
      return 0
    end function

    subroutine omp_get_place_proc_ids(place_num, ids)
      integer place_num
      integer ids(*)
      return
```

```fortran
      end subroutine

      integer function omp_get_place_num()
        return -1
      end function

      integer function omp_get_partition_num_places()
        return 0
      end function

      subroutine omp_get_partition_place_nums(place_nums)
        integer place_nums(*)
        return
      end subroutine

      subroutine omp_set_default_device(device_num)
        integer device_num
        return
      end subroutine

      integer function omp_get_default_device()
        omp_get_default_device = 0
        return
      end function

      integer function omp_get_num_devices()
        omp_get_num_devices = 0
        return
      end function

      integer function omp_get_num_teams()
        omp_get_num_teams = 1
        return
      end function

      integer function omp_get_team_num()
        omp_get_team_num = 0
        return
      end function

      logical function omp_is_initial_device()
        omp_is_initial_device = .true.
        return
      end function

      integer function omp_get_initial_device()
        omp_get_initial_device = -10
```

```fortran
1                          return
2                     end function
3
4                     integer function omp_get_max_task_priority()
5                       omp_get_max_task_priority = 0
6                       return
7                     end function
8
9                     subroutine omp_init_lock(lock)
10                      ! lock is 0 if the simple lock is not initialized
11                      !         -1 if the simple lock is initialized but not set
12                      !          1 if the simple lock is set
13                      include 'omp_lib_kinds.h'
14                      integer(kind=omp_lock_kind) lock
15
16                      lock = -1
17                      return
18                    end subroutine
19
20                    subroutine omp_init_lock_with_hint(lock, hint)
21                      include 'omp_lib_kinds.h'
22                      integer(kind=omp_lock_kind) lock
23                      integer(kind=omp_lock_hint_kind) hint
24
25                      call omp_init_lock(lock)
26                      return
27                    end subroutine
28
29                    subroutine omp_destroy_lock(lock)
30                      include 'omp_lib_kinds.h'
31                      integer(kind=omp_lock_kind) lock
32
33                      lock = 0
34                      return
35                    end subroutine
36
37                    subroutine omp_set_lock(lock)
38                      include 'omp_lib_kinds.h'
39                      integer(kind=omp_lock_kind) lock
40
41                      if (lock .eq. -1) then
42                        lock = 1
43                      elseif (lock .eq. 1) then
44                        print *, 'error: deadlock in using lock variable'
45                        stop
46                      else
47                        print *, 'error: lock not initialized'
```

```fortran
            stop
          endif
          return
        end subroutine

        subroutine omp_unset_lock(lock)
          include 'omp_lib_kinds.h'
          integer(kind=omp_lock_kind) lock

          if (lock .eq. 1) then
            lock = -1
          elseif (lock .eq. -1) then
            print *, 'error: lock not set'
            stop
          else
            print *, 'error: lock not initialized'
            stop
          endif
          return
        end subroutine

        logical function omp_test_lock(lock)
          include 'omp_lib_kinds.h'
          integer(kind=omp_lock_kind) lock

          if (lock .eq. -1) then
            lock = 1
            omp_test_lock = .true.
          elseif (lock .eq. 1) then
            omp_test_lock = .false.
          else
            print *, 'error: lock not initialized'
            stop
          endif

          return
        end function

        subroutine omp_init_nest_lock(nlock)
          ! nlock is
          ! 0 if the nestable lock is not initialized
          ! -1 if the nestable lock is initialized but not set
          ! 1 if the nestable lock is set
          ! no use count is maintained
          include 'omp_lib_kinds.h'
          integer(kind=omp_nest_lock_kind) nlock
```

```fortran
1            nlock = -1
2
3            return
4          end subroutine
5
6          subroutine omp_init_nest_lock_with_hint(nlock, hint)
7            include 'omp_lib_kinds.h'
8            integer(kind=omp_nest_lock_kind) nlock
9            integer(kind=omp_lock_hint_kind) hint
10
11           call omp_init_nest_lock(nlock)
12           return
13         end subroutine
14
15         subroutine omp_destroy_nest_lock(nlock)
16           include 'omp_lib_kinds.h'
17           integer(kind=omp_nest_lock_kind) nlock
18
19           nlock = 0
20
21           return
22         end subroutine
23
24         subroutine omp_set_nest_lock(nlock)
25           include 'omp_lib_kinds.h'
26           integer(kind=omp_nest_lock_kind) nlock
27
28           if (nlock .eq. -1) then
29             nlock = 1
30           elseif (nlock .eq. 0) then
31             print *, 'error: nested lock not initialized'
32             stop
33           else
34             print *, 'error: deadlock using nested lock variable'
35             stop
36           endif
37
38           return
39         end subroutine
40
41         subroutine omp_unset_nest_lock(nlock)
42           include 'omp_lib_kinds.h'
43           integer(kind=omp_nest_lock_kind) nlock
44
45           if (nlock .eq. 1) then
46             nlock = -1
47           elseif (nlock .eq. 0) then
```

```fortran
          print *, 'error: nested lock not initialized'
          stop
        else
          print *, 'error: nested lock not set'
          stop
        endif

        return
      end subroutine

      integer function omp_test_nest_lock(nlock)
        include 'omp_lib_kinds.h'
        integer(kind=omp_nest_lock_kind) nlock

        if (nlock .eq. -1) then
          nlock = 1
          omp_test_nest_lock = 1
        elseif (nlock .eq. 1) then
          omp_test_nest_lock = 0
        else
          print *, 'error: nested lock not initialized'
          stop
        endif

        return
      end function

      double precision function omp_get_wtime()
        ! this function does not provide a working
        ! wall clock timer. replace it with a version
        ! customized for the target machine.

        omp_get_wtime = 0.0d0

        return
      end function

      double precision function omp_get_wtick()
        ! this function does not provide a working
        ! clock tick function. replace it with
        ! a version customized for the target machine.
        double precision one_year
        parameter (one_year=365.d0*86400.d0)

        omp_get_wtick = one_year

        return
```

1       **end function**

1 *This page intentionally left blank*

# ² **Interface Declarations**

---

3   This appendix gives examples of the C/C++ header file, the Fortran **include** file and Fortran
4   **module** that shall be provided by implementations as specified in Chapter 3. It also includes an
5   example of a Fortran 90 generic interface for a library routine. This is a non-normative section,
6   implementation files may differ.

# ₁ B.1   Example of the `omp.h` Header File

```
2          #ifndef _OMP_H_DEF
3          #define _OMP_H_DEF
4
5          /*
6           * define the lock data types
7           */
8          typedef void *omp_lock_t;
9
10         typedef void *omp_nest_lock_t;
11
12         /*
13          * define the lock hints
14          */
15         typedef enum omp_lock_hint_t
16         {
17          omp_lock_hint_none = 0,
18          omp_lock_hint_uncontended = 1,
19          omp_lock_hint_contended = 2,
20          omp_lock_hint_nonspeculative = 4,
21          omp_lock_hint_speculative = 8
22         /* , Add vendor specific constants for lock hints here,
23             starting from the most-significant bit. */
24         } omp_lock_hint_t;
25
26         /*
27          * define the schedule kinds
28          */
29         typedef enum omp_sched_t
30         {
31          omp_sched_static = 1,
32          omp_sched_dynamic = 2,
33          omp_sched_guided = 3,
34          omp_sched_auto = 4
35         /* , Add vendor specific schedule constants here */
36         } omp_sched_t;
37
38         /*
39          * define the proc bind values
40          */
41         typedef enum omp_proc_bind_t
42         {
43          omp_proc_bind_false = 0,
44          omp_proc_bind_true = 1,
45          omp_proc_bind_master = 2,
46          omp_proc_bind_close = 3,
```

```
1              omp_proc_bind_spread = 4
2            } omp_proc_bind_t;
3
4            /*
5             * exported OpenMP functions
6             */
7            #ifdef __cplusplus
8            extern "C"
9            {
10           #endif
11
12           extern void omp_set_num_threads(int num_threads);
13           extern int omp_get_num_threads(void);
14           extern int omp_get_max_threads(void);
15           extern int omp_get_thread_num(void);
16           extern int omp_get_num_procs(void);
17           extern int omp_in_parallel(void);
18           extern void omp_set_dynamic(int dynamic_threads);
19           extern int omp_get_dynamic(void);
20           extern int omp_get_cancellation(void);
21           extern void omp_set_nested(int nested);
22           extern int omp_get_nested(void);
23           extern void omp_set_schedule(omp_sched_t kind, int chunk_size);
24           extern void omp_get_schedule(omp_sched_t *kind, int *chunk_size);
25           extern int omp_get_thread_limit(void);
26           extern void omp_set_max_active_levels(int max_active_levels);
27           extern int omp_get_max_active_levels(void);
28           extern int omp_get_level(void);
29           extern int omp_get_ancestor_thread_num(int level);
30           extern int omp_get_team_size(int level);
31           extern int omp_get_active_level(void);
32           extern int omp_in_final(void);
33           extern omp_proc_bind_t omp_get_proc_bind(void);
34           extern int omp_get_num_places(void);
35           extern int omp_get_place_num_procs(int place_num);
36           extern void omp_get_place_proc_ids(int place_num, int *ids);
37           extern int omp_get_place_num(void);
38           extern int omp_get_partition_num_places(void);
39           extern void omp_get_partition_place_nums(int *place_nums);
40           extern void omp_set_default_device(int device_num);
41           extern int omp_get_default_device(void);
42           extern int omp_get_num_devices(void);
43           extern int omp_get_num_teams(void);
44           extern int omp_get_team_num(void);
45           extern int omp_is_initial_device(void);
46           extern int omp_get_initial_device(void);
47           extern int omp_get_max_task_priority(void);
```

```
1
2              extern void omp_init_lock(omp_lock_t *lock);
3              extern void omp_init_lock_with_hint(omp_lock_t *lock,
4                                                  omp_lock_hint_t hint);
5              extern void omp_destroy_lock(omp_lock_t *lock);
6              extern void omp_set_lock(omp_lock_t *lock);
7              extern void omp_unset_lock(omp_lock_t *lock);
8              extern int omp_test_lock(omp_lock_t *lock);
9
10             extern void omp_init_nest_lock(omp_nest_lock_t *lock);
11             extern void omp_init_nest_lock_with_hint(omp_nest_lock_t *lock,
12                                                      omp_lock_hint_t hint);
13             extern void omp_destroy_nest_lock(omp_nest_lock_t *lock);
14             extern void omp_set_nest_lock(omp_nest_lock_t *lock);
15             extern void omp_unset_nest_lock(omp_nest_lock_t *lock);
16             extern int omp_test_nest_lock(omp_nest_lock_t *lock);
17
18             extern double omp_get_wtime(void);
19             extern double omp_get_wtick(void);
20
21             extern void * omp_target_alloc(size_t size, int device_num);
22             extern void omp_target_free(void * device_ptr, int device_num);
23             extern int omp_target_is_present(void * ptr, int device_num);
24             extern int omp_target_memcpy(void *dst, void *src, size_t length,
25                                          size_t dst_offset, size_t src_offset,
26                                          int dst_device_num, int src_device_num);
27             extern int omp_target_memcpy_rect(
28                 void *dst, void *src,
29                 size_t element_size,
30                 int num_dims,
31                 const size_t *volume,
32                 const size_t *dst_offsets,
33                 const size_t *src_offsets,
34                 const size_t *dst_dimensions,
35                 const size_t *src_dimensions,
36                 int dst_device_num, int src_device_num);
37             extern int omp_target_associate_ptr(void * host_ptr,
38                                                 void * device_ptr,
39                                                 size_t size,
40                                                 size_t device_offset,
41                                                 int device_num);
42             extern int omp_target_disassociate_ptr(void * ptr,
43                                                    int device_num);
44
45             #ifdef __cplusplus
46             }
47             #endif
```

1
2      **#endif**

# B.2 Example of an Interface Declaration `include` File

*omp_lib_kinds.h:*

```fortran
      integer omp_lock_kind
            integer omp_nest_lock_kind
            integer omp_lock_hint_kind
      ! this selects an integer that is large enough to hold a 64 bit integer
            parameter ( omp_lock_kind = selected_int_kind( 10 ) )
            parameter ( omp_nest_lock_kind = selected_int_kind( 10 ) )
            parameter ( omp_lock_hint_kind = selected_int_kind( 10 ) )

            integer omp_sched_kind
      ! this selects an integer that is large enough to hold a 32 bit integer
            parameter ( omp_sched_kind = selected_int_kind( 8 ) )
            integer ( omp_sched_kind ) omp_sched_static
            parameter ( omp_sched_static = 1 )
            integer ( omp_sched_kind ) omp_sched_dynamic
            parameter ( omp_sched_dynamic = 2 )
            integer ( omp_sched_kind ) omp_sched_guided
            parameter ( omp_sched_guided = 3 )
            integer ( omp_sched_kind ) omp_sched_auto
            parameter ( omp_sched_auto = 4 )

            integer omp_proc_bind_kind
            parameter ( omp_proc_bind_kind = selected_int_kind( 8 ) )
            integer ( omp_proc_bind_kind ) omp_proc_bind_false
            parameter ( omp_proc_bind_false = 0 )
            integer ( omp_proc_bind_kind ) omp_proc_bind_true
            parameter ( omp_proc_bind_true = 1 )
            integer ( omp_proc_bind_kind ) omp_proc_bind_master
            parameter ( omp_proc_bind_master = 2 )
            integer ( omp_proc_bind_kind ) omp_proc_bind_close
            parameter ( omp_proc_bind_close = 3 )
            integer ( omp_proc_bind_kind ) omp_proc_bind_spread
            parameter ( omp_proc_bind_spread = 4 )

            integer ( omp_lock_hint_kind ) omp_lock_hint_none
            parameter ( omp_lock_hint_none = 0 )
            integer ( omp_lock_hint_kind ) omp_lock_hint_uncontended
            parameter ( omp_lock_hint_uncontended = 1 )
            integer ( omp_lock_hint_kind ) omp_lock_hint_contended
            parameter ( omp_lock_hint_contended = 2 )
            integer ( omp_lock_hint_kind ) omp_lock_hint_nonspeculative
            parameter ( omp_lock_hint_nonspeculative = 4 )
```

```fortran
      integer ( omp_lock_hint_kind ) omp_lock_hint_speculative
      parameter ( omp_lock_hint_speculative = 8 )
```

*omp_lib.h:*

```fortran
! default integer type assumed below
! default logical type assumed below
! OpenMP API v4.5

      include 'omp_lib_kinds.h'
      integer openmp_version
      parameter ( openmp_version = 201511 )

      external omp_set_num_threads
      external omp_get_num_threads
      integer omp_get_num_threads
      external omp_get_max_threads
      integer omp_get_max_threads
      external omp_get_thread_num
      integer omp_get_thread_num
      external omp_get_num_procs
      integer omp_get_num_procs
      external omp_in_parallel
      logical omp_in_parallel
      external omp_set_dynamic
      external omp_get_dynamic
      logical omp_get_dynamic
      external omp_get_cancellation
      logical omp_get_cancellation
      external omp_set_nested
      external omp_get_nested
      logical omp_get_nested
      external omp_set_schedule
      external omp_get_schedule
      external omp_get_thread_limit
      integer omp_get_thread_limit
      external omp_set_max_active_levels
      external omp_get_max_active_levels
      integer omp_get_max_active_levels
      external omp_get_level
      integer omp_get_level
      external omp_get_ancestor_thread_num
      integer omp_get_ancestor_thread_num
      external omp_get_team_size
      integer omp_get_team_size
      external omp_get_active_level
      integer omp_get_active_level
      external omp_set_default_device
```

```
1              external omp_get_default_device
2              integer omp_get_default_device
3              external omp_get_num_devices
4              integer omp_get_num_devices
5              external omp_get_num_teams
6              integer omp_get_num_teams
7              external omp_get_team_num
8              integer omp_get_team_num
9              external omp_is_initial_device
10             logical omp_is_initial_device
11             external omp_get_initial_device
12             integer omp_get_initial_device
13             external omp_get_max_task_priority
14             integer omp_get_max_task_priority
15
16             external omp_in_final
17             logical omp_in_final
18
19             integer ( omp_proc_bind_kind ) omp_get_proc_bind
20             external omp_get_proc_bind
21             integer omp_get_num_places
22             external omp_get_num_places
23             integer omp_get_place_num_procs
24             external omp_get_place_num_procs
25             external omp_get_place_proc_ids
26             integer omp_get_place_num
27             external omp_get_place_num
28             integer omp_get_partition_num_places
29             external omp_get_partition_num_places
30             external omp_get_partition_place_nums
31
32             external omp_init_lock
33             external omp_init_lock_with_hint
34             external omp_destroy_lock
35             external omp_set_lock
36             external omp_unset_lock
37             external omp_test_lock
38             logical omp_test_lock
39
40             external omp_init_nest_lock
41             external omp_init_nest_lock_with_hint
42             external omp_destroy_nest_lock
43             external omp_set_nest_lock
44             external omp_unset_nest_lock
45             external omp_test_nest_lock
46             integer omp_test_nest_lock
47
```

```
1              external omp_get_wtick
2              double precision omp_get_wtick
3              external omp_get_wtime
4              double precision omp_get_wtime
```

# B.3 Example of a Fortran Interface Declaration `module`

```
!       the "!" of this comment starts in column 1
!23456

       module omp_lib_kinds
       integer, parameter :: omp_lock_kind = selected_int_kind( 10 )
       integer, parameter :: omp_nest_lock_kind = selected_int_kind( 10 )
       integer, parameter :: omp_lock_hint_kind = selected_int_kind( 10 )
       integer (kind=omp_lock_hint_kind), parameter ::
     &   omp_lock_hint_none = 0
       integer (kind=omp_lock_hint_kind), parameter ::
     &   omp_lock_hint_uncontended = 1
       integer (kind=omp_lock_hint_kind), parameter ::
     &   omp_lock_hint_contended = 2
       integer (kind=omp_lock_hint_kind), parameter ::
     &   omp_lock_hint_nonspeculative = 4
       integer (kind=omp_lock_hint_kind), parameter ::
     &   omp_lock_hint_speculative = 8

       integer, parameter :: omp_sched_kind = selected_int_kind( 8 )
       integer(kind=omp_sched_kind), parameter ::
     &   omp_sched_static = 1
       integer(kind=omp_sched_kind), parameter ::
     &   omp_sched_dynamic = 2
       integer(kind=omp_sched_kind), parameter ::
     &   omp_sched_guided = 3
       integer(kind=omp_sched_kind), parameter ::
     &   omp_sched_auto = 4

       integer, parameter :: omp_proc_bind_kind = selected_int_kind( 8 )
       integer (kind=omp_proc_bind_kind), parameter ::
     &   omp_proc_bind_false = 0
       integer (kind=omp_proc_bind_kind), parameter ::
     &   omp_proc_bind_true = 1
       integer (kind=omp_proc_bind_kind), parameter ::
     &   omp_proc_bind_master = 2
       integer (kind=omp_proc_bind_kind), parameter ::
     &   omp_proc_bind_close = 3
       integer (kind=omp_proc_bind_kind), parameter ::
     &   omp_proc_bind_spread = 4
        end module omp_lib_kinds

       module omp_lib
```

```fortran
                       use omp_lib_kinds

!                                                 OpenMP API v4.5
                integer, parameter :: openmp_version = 201511

              interface

                subroutine omp_set_num_threads (num_threads)
                 integer, intent(in) :: num_threads
                end subroutine omp_set_num_threads

                function omp_get_num_threads ()
                 integer :: omp_get_num_threads
                end function omp_get_num_threads

                function omp_get_max_threads ()
                 integer :: omp_get_max_threads
                end function omp_get_max_threads

                function omp_get_thread_num ()
                 integer :: omp_get_thread_num
                end function omp_get_thread_num

                function omp_get_num_procs ()
                 integer :: omp_get_num_procs
                end function omp_get_num_procs

                function omp_in_parallel ()
                 logical :: omp_in_parallel
                end function omp_in_parallel

                subroutine omp_set_dynamic (dynamic_threads)
                 logical, intent(in) ::dynamic_threads
                end subroutine omp_set_dynamic

                function omp_get_dynamic ()
                 logical :: omp_get_dynamic
                end function omp_get_dynamic

                function omp_get_cancellation ()
                 logical :: omp_get_cancellation
                end function omp_get_cancellation

                subroutine omp_set_nested (nested)
                 logical, intent(in) :: nested
                end subroutine omp_set_nested
```

```fortran
1                        function omp_get_nested ()
2                          logical :: omp_get_nested
3                        end function omp_get_nested
4
5                        subroutine omp_set_schedule (kind, chunk_size)
6                          use omp_lib_kinds
7                          integer(kind=omp_sched_kind), intent(in) :: kind
8                          integer, intent(in) :: chunk_size
9                        end subroutine omp_set_schedule
10
11                       subroutine omp_get_schedule (kind, chunk_size)
12                         use omp_lib_kinds
13                         integer(kind=omp_sched_kind), intent(out) :: kind
14                         integer, intent(out)::chunk_size
15                       end subroutine omp_get_schedule
16
17                       function omp_get_thread_limit ()
18                         integer :: omp_get_thread_limit
19                       end function omp_get_thread_limit
20
21                       subroutine omp_set_max_active_levels (max_levels)
22                         integer, intent(in) :: max_levels
23                       end subroutine omp_set_max_active_levels
24
25                       function omp_get_max_active_levels ()
26                         integer :: omp_get_max_active_levels
27                       end function omp_get_max_active_levels
28
29                       function omp_get_level()
30                         integer :: omp_get_level
31                       end function omp_get_level
32
33                       function omp_get_ancestor_thread_num (level)
34                         integer, intent(in) :: level
35                         integer :: omp_get_ancestor_thread_num
36                       end function omp_get_ancestor_thread_num
37
38                       function omp_get_team_size (level)
39                         integer, intent(in) :: level
40                         integer :: omp_get_team_size
41                       end function omp_get_team_size
42
43                       function omp_get_active_level ()
44                         integer :: omp_get_active_level
45                       end function omp_get_active_level
46
47                       function omp_in_final ()
```

```
1                        logical :: omp_in_final
2                        end function omp_in_final
3
4                        function omp_get_proc_bind ()
5                         use omp_lib_kinds
6                         integer(kind=omp_proc_bind_kind) :: omp_get_proc_bind
7                         omp_get_proc_bind = omp_proc_bind_false
8                        end function omp_get_proc_bind
9
10                       function omp_get_num_places ()
11                       integer :: omp_get_num_places
12                       end function omp_get_num_places
13
14                       function omp_get_place_num_procs (place_num)
15                       integer, intent(in) :: place_num
16                       integer :: omp_get_place_num_procs
17                       end function omp_get_place_num_procs
18
19                       subroutine omp_get_place_proc_ids (place_num, ids)
20                       integer, intent(in) :: place_num
21                       integer, intent(out) :: ids(*)
22                       end subroutine omp_get_place_proc_ids
23
24                       function omp_get_place_num ()
25                       integer :: omp_get_place_num
26                       end function omp_get_place_num
27
28                       function omp_get_partition_num_places ()
29                       integer :: omp_get_partition_num_places
30                       end function omp_get_partition_num_places
31
32                       subroutine omp_get_partition_place_nums (place_nums)
33                       integer, intent(out) :: place_nums(*)
34                       end subroutine omp_get_partition_place_nums
35
36                       subroutine omp_set_default_device (device_num)
37                        integer :: device_num
38                       end subroutine omp_set_default_device
39
40                       function omp_get_default_device ()
41                        integer :: omp_get_default_device
42                       end function omp_get_default_device
43
44                       function omp_get_num_devices ()
45                        integer :: omp_get_num_devices
46                       end function omp_get_num_devices
47
```

```fortran
1                            function omp_get_num_teams ()
2                             integer :: omp_get_num_teams
3                            end function omp_get_num_teams
4
5                            function omp_get_team_num ()
6                             integer :: omp_get_team_num
7                            end function omp_get_team_num
8
9                            function omp_is_initial_device ()
10                            logical :: omp_is_initial_device
11                           end function omp_is_initial_device
12
13                           function omp_get_initial_device ()
14                            integer :: omp_get_initial_device
15                           end function omp_get_initial_device
16
17                           function omp_get_max_task_priority ()
18                            integer :: omp_get_max_task_priority
19                           end function omp_get_max_task_priority
20
21                           subroutine omp_init_lock (svar)
22                            use omp_lib_kinds
23                            integer(kind=omp_lock_kind), intent(out) :: svar
24                           end subroutine omp_init_lock
25
26                           subroutine omp_init_lock_with_hint (svar, hint)
27                            use omp_lib_kinds
28                            integer(kind=omp_lock_kind), intent(out) :: svar
29                            integer(kind=omp_lock_hint_kind), intent(in) :: hint
30                           end subroutine omp_init_lock_with_hint
31
32                           subroutine omp_destroy_lock (svar)
33                            use omp_lib_kinds
34                            integer(kind=omp_lock_kind), intent(inout) :: svar
35                           end subroutine omp_destroy_lock
36
37                           subroutine omp_set_lock (svar)
38                            use omp_lib_kinds
39                            integer(kind=omp_lock_kind), intent(inout) :: svar
40                           end subroutine omp_set_lock
41
42                           subroutine omp_unset_lock (svar)
43                            use omp_lib_kinds
44                            integer(kind=omp_lock_kind), intent(inout) :: svar
45                           end subroutine omp_unset_lock
46
47                           function omp_test_lock (svar)
```

```fortran
                        use omp_lib_kinds
                        logical :: omp_test_lock
                        integer(kind=omp_lock_kind), intent(inout) :: svar
                       end function omp_test_lock

                       subroutine omp_init_nest_lock (nvar)
                        use omp_lib_kinds
                        integer(kind=omp_nest_lock_kind), intent(out) :: nvar
                       end subroutine omp_init_nest_lock

                       subroutine omp_init_nest_lock_with_hint (nvar, hint)
                        use omp_lib_kinds
                        integer(kind=omp_nest_lock_kind), intent(out) :: nvar
                        integer(kind=omp_lock_hint_kind), intent(in) :: hint
                       end subroutine omp_init_nest_lock_with_hint

                       subroutine omp_destroy_nest_lock (nvar)
                        use omp_lib_kinds
                        integer(kind=omp_nest_lock_kind), intent(inout) :: nvar
                       end subroutine omp_destroy_nest_lock

                       subroutine omp_set_nest_lock (nvar)
                        use omp_lib_kinds
                        integer(kind=omp_nest_lock_kind), intent(inout) :: nvar
                       end subroutine omp_set_nest_lock

                       subroutine omp_unset_nest_lock (nvar)
                        use omp_lib_kinds
                        integer(kind=omp_nest_lock_kind), intent(inout) :: nvar
                       end subroutine omp_unset_nest_lock

                       function omp_test_nest_lock (nvar)
                        use omp_lib_kinds
                        integer :: omp_test_nest_lock
                        integer(kind=omp_nest_lock_kind), intent(inout) :: nvar
                       end function omp_test_nest_lock

                       function omp_get_wtick ()
                        double precision :: omp_get_wtick
                       end function omp_get_wtick

                       function omp_get_wtime ()
                        double precision :: omp_get_wtime
                       end function omp_get_wtime

                       end interface
```

1              **end module omp_lib**

3 Any of the OpenMP runtime library routines that take an argument may be extended with a generic
4 interface so arguments of different **KIND** type can be accommodated.

5 The **OMP_SET_NUM_THREADS** interface could be specified in the **omp_lib** module as follows:

```fortran
interface omp_set_num_threads

    subroutine omp_set_num_threads_4(num_threads)
      use omp_lib_kinds
      integer(4), intent(in) :: num_threads
    end subroutine omp_set_num_threads_4

    subroutine omp_set_num_threads_8(num_threads)
      use omp_lib_kinds
      integer(8), intent(in) :: num_threads
    end subroutine omp_set_num_threads_8

end interface omp_set_num_threads
```

2 # OpenMP Implementation-Defined
3 # Behaviors

4 This appendix summarizes the behaviors that are described as implementation defined in this API.
5 Each behavior is cross-referenced back to its description in the main specification. An
6 implementation is required to define and document its behavior in these cases.

7 • **Processor**: a hardware unit that is implementation defined (see Section 1.2.1 on page 2).

8 • **Device**: an implementation defined logical execution engine (see Section 1.2.1 on page 2).

9 • **Device address**: an address in a *device data environment* (see Section 1.2.6 on page 11).

10 • **Memory model**: the minimum size at which a memory update may also read and write back
11 adjacent variables that are part of another variable (as array or structure elements) is
12 implementation defined but is no larger than required by the base language (see Section 1.4.1 on
13 page 17).

14 • **Memory model**: Implementations are allowed to relax the ordering imposed by implicit flush
15 operations when the result is only visible to programs using non-sequentially consistent atomic
16 directives (see Section 1.4.4 on page 20).

17 • **Internal control variables**: the initial values of *dyn-var*, *nthreads-var*, *run-sched-var*,
18 *def-sched-var*, *bind-var*, *stacksize-var*, *wait-policy-var*, *thread-limit-var*, *max-active-levels-var*,
19 *place-partition-var*, and *default-device-var* are implementation defined. The method for
20 initializing a target device's internal control variable is implementation defined (see Section 2.3.2
21 on page 37).

22 • **Dynamic adjustment of threads**: providing the ability to dynamically adjust the number of
23 threads is implementation defined . Implementations are allowed to deliver fewer threads (but at
24 least one) than indicated in Algorithm 2-1 even if dynamic adjustment is disabled (see
25 Section 2.5.1 on page 50).

- **Thread affinity**: For the **close** thread affinity policy, if $T > P$ and $P$ does not divide $T$ evenly, the exact number of threads in a particular place is implementation defined. For the **spread** thread affinity, if $T > P$ and $P$ does not divide $T$ evenly, the exact number of threads in a particular subpartition is implementation defined. The determination of whether the affinity request can be fulfilled is implementation defined. If not, the number of threads in the team and their mapping to places become implementation defined (see Section 2.5.2 on page 52).

- **Loop directive**: the integer type (or kind, for Fortran) used to compute the iteration count of a collapsed loop is implementation defined. The effect of the **schedule(runtime)** clause when the *run-sched-var* ICV is set to **auto** is implementation defined. The *simd_width* used when a **simd** schedule modifier is specified is implementation defined (see Section 2.7.1 on page 56).

- **sections construct**: the method of scheduling the structured blocks among threads in the team is implementation defined (see Section 2.7.2 on page 65).

- **single construct**: the method of choosing a thread to execute the structured block is implementation defined (see Section 2.7.3 on page 67)

- **simd construct**: the integer type (or kind, for Fortran) used to compute the iteration count for the collapsed loop is implementation defined. The number of iterations that are executed concurrently at any given time is implementation defined. If the *alignment* parameter is not specified in the **aligned** clause, the default alignments for the SIMD instructions are implementation defined (see Section 2.8.1 on page 72).

- **declare simd construct**: if the parameter of the **simdlen** clause is not a constant positive integer expression, the number of concurrent arguments for the function is implementation defined. If the *alignment* parameter of the **aligned** clause is not specified, the default alignments for SIMD instructions are implementation defined (see Section 2.8.2 on page 76).

- **taskloop construct**: The number of loop iterations assigned to a task created from a **taskloop** construct is implementation defined, unless the **grainsize** or **num_tasks** clauses are specified. The integer type (or kind, for Fortran) used to compute the iteration count for the collapsed loop is implementation defined (see Section 2.9.2 on page 87).

- **is_device_ptr clause**: Support for pointers created outside of the OpenMP device data management routines is implementation defined (see Section 2.10.4 on page 103).

- **teams construct**: the number of teams that are created is implementation defined but less than or equal to the value of the **num_teams** clause if specified. The maximum number of threads participating in the contention group that each team initiates is implementation defined but less than or equal to the value of the **thread_limit** clause if specified (see Section 2.10.7 on page 114).

- **distribute construct**: the integer type (or kind, for Fortran) used to compute the iteration count for the collapsed loop is implementation defined (see Section 2.10.8 on page 117).

- **distribute construct**: If no **dist_schedule** clause is specified then the schedule for the

1    **distribute** construct is implementation defined (see Section 2.10.8 on page 117).

2    • **critical construct**: the effect of using a **hint** clause is implementation defined (see
3    Section 2.13.2 on page 149).

4    • **atomic construct**: a compliant implementation may enforce exclusive access between
5    **atomic** regions that update different storage locations. The circumstances under which this
6    occurs are implementation defined. If the storage location designated by $x$ is not size-aligned
7    (that is, if the byte alignment of $x$ is not a multiple of the size of $x$), then the behavior of the
8    atomic region is implementation defined (see Section 2.13.6 on page 155).

──────────── Fortran ────────────

9    • **threadprivate directive**: if the conditions for values of data in the threadprivate objects of
10   threads (other than an initial thread) to persist between two consecutive active parallel regions do
11   not all hold, the allocation status of an allocatable variable in the second region is
12   implementation defined (see Section 2.15.2 on page 183).

13   • **shared clause**: passing a shared variable to a non-intrinsic procedure may result in the value of
14   the shared variable being copied into temporary storage before the procedure reference, and back
15   out of the temporary storage into the actual argument storage after the procedure reference.
16   Situations where this occurs other than those specified are implementation defined (see
17   Section 2.15.3.2 on page 190).

18   • **Runtime library definitions**: it is implementation defined whether the include file **omp_lib.h**
19   or the module **omp_lib** (or both) is provided. It is implementation defined whether any of the
20   OpenMP runtime library routines that take an argument are extended with a generic interface so
21   arguments of different **KIND** type can be accommodated (see Section 3.1 on page 230).

──────────── Fortran ────────────

22   • **omp_set_num_threads routine**: if the argument is not a positive integer the behavior is
23   implementation defined (see Section 3.2.1 on page 231).

24   • **omp_set_schedule routine**: for implementation specific schedule types, the values and
25   associated meanings of the second argument are implementation defined. (see Section 3.2.12 on
26   page 243).

27   • **omp_set_max_active_levels routine**: when called from within any explicit **parallel**
28   region the binding thread set (and binding region, if required) for the
29   **omp_set_max_active_levels** region is implementation defined and the behavior is
30   implementation defined. If the argument is not a non-negative integer then the behavior is
31   implementation defined (see Section 3.2.15 on page 246).

32   • **omp_get_max_active_levels routine**: when called from within any explicit **parallel**
33   region the binding thread set (and binding region, if required) for the
34   **omp_get_max_active_levels** region is implementation defined (see Section 3.2.16 on
35   page 248).

- **omp_get_place_proc_ids routine**: the meaning of the nonnegative numerical identifiers returned by the **omp_get_place_proc_ids** routine is implementation defined (see Section 3.2.25 on page 258).

- **omp_get_initial_device routine**: the value of the device number is implementation defined (see Section 3.2.35 on page 267).

- **omp_init_lock_with_hint** and **omp_init_nest_lock_with_hint routines**: if hints are stored with a lock variable, the effect of the hints on the locks are implementation defined (see Section 3.3.2 on page 273).

- **omp_target_memcpy_rect routine**: the maximum number of dimensions supported is implementation defined, but must be at least three (see Section 3.5.5 on page 286).

- **OMP_SCHEDULE environment variable**: if the value does not conform to the specified format then the result is implementation defined (see Section 4.1 on page 292).

- **OMP_NUM_THREADS environment variable**: if any value of the list specified in the **OMP_NUM_THREADS** environment variable leads to a number of threads that is greater than the implementation can support, or if any value is not a positive integer, then the result is implementation defined (see Section 4.2 on page 293).

- **OMP_PROC_BIND environment variable**: if the value is not **true**, **false**, or a comma separated list of **master**, **close**, or **spread**, the behavior is implementation defined. The behavior is also implementation defined if an initial thread cannot be bound to the first place in the OpenMP place list (see Section 4.4 on page 294).

- **OMP_DYNAMIC environment variable**: if the value is neither **true** nor **false** the behavior is implementation defined (see Section 4.3 on page 294).

- **OMP_NESTED environment variable**: if the value is neither **true** nor **false** the behavior is implementation defined (see Section 4.6 on page 297).

- **OMP_STACKSIZE environment variable**: if the value does not conform to the specified format or the implementation cannot provide a stack of the specified size then the behavior is implementation defined (see Section 4.7 on page 298).

- **OMP_WAIT_POLICY environment variable**: the details of the **ACTIVE** and **PASSIVE** behaviors are implementation defined (see Section 4.8 on page 299).

- **OMP_MAX_ACTIVE_LEVELS environment variable**: if the value is not a non-negative integer or is greater than the number of parallel levels an implementation can support then the behavior is implementation defined (see Section 4.9 on page 300).

- **OMP_THREAD_LIMIT environment variable**: if the requested value is greater than the number of threads an implementation can support, or if the value is not a positive integer, the behavior of the program is implementation defined (see Section 4.10 on page 300).

- **OMP_PLACES environment variable**: the meaning of the numbers specified in the environment variable and how the numbering is done are implementation defined. The precise definitions of

<table>
<tr><td>1</td><td>the abstract names are implementation defined. An implementation may add</td></tr>
<tr><td>2</td><td>implementation-defined abstract names as appropriate for the target platform. When creating a</td></tr>
<tr><td>3</td><td>place list of n elements by appending the number $n$ to an abstract name, the determination of</td></tr>
<tr><td>4</td><td>which resources to include in the place list is implementation defined. When requesting more</td></tr>
<tr><td>5</td><td>resources than available, the length of the place list is also implementation defined. The behavior</td></tr>
<tr><td>6</td><td>of the program is implementation defined when the execution environment cannot map a</td></tr>
<tr><td>7</td><td>numerical value (either explicitly defined or implicitly derived from an interval) within the</td></tr>
<tr><td>8</td><td>**OMP_PLACES** list to a processor on the target platform, or if it maps to an unavailable processor.</td></tr>
<tr><td>9</td><td>The behavior is also implementation defined when the **OMP_PLACES** environment variable is</td></tr>
<tr><td>10</td><td>defined using an abstract name (see Section 4.5 on page 295).</td></tr>
</table>

# <sub>2</sub> **Features History**

---

<sub>3</sub> This appendix summarizes the major changes between recent versions of the OpenMP API since
<sub>4</sub> version 2.5.

# <sub>5</sub> **D.1  Version 4.0 to 4.5 Differences**

<sub>6</sub>  • Support for several features of Fortran 2003 was added (see Section 1.6 on page 21 for features
<sub>7</sub>    that are still not supported).

<sub>8</sub>  • A parameter was added to the **ordered** clause of the loop construct (see Section 2.7.1 on
<sub>9</sub>    page 56) and clauses were added to the **ordered** construct (see Section 2.13.8 on page 166) to
<sub>10</sub>    support doacross loop nests and use of the **simd** construct on loops with loop-carried backward
<sub>11</sub>    dependences.

<sub>12</sub>  • The **linear** clause was added to the loop construct (see Section 2.7.1 on page 56).

<sub>13</sub>  • The **simdlen** clause was added to the **simd** construct (see Section 2.8.1 on page 72) to support
<sub>14</sub>    specification of the exact number of iterations desired per SIMD chunk.

<sub>15</sub>  • The **priority** clause was added to the **task** construct (see Section 2.9.1 on page 83) to
<sub>16</sub>    support hints that specify the relative execution priority of explicit tasks. The
<sub>17</sub>    **omp_get_max_task_priority** routine was added to return the maximum supported
<sub>18</sub>    priority value (see Section 3.2.36 on page 268) and the **OMP_MAX_TASK_PRIORITY**
<sub>19</sub>    environment variable was added to control the maximum priority value allowed (see
<sub>20</sub>    Section 4.14 on page 303).

<sub>21</sub>  • Taskloop constructs (see Section 2.9.2 on page 87 and Section 2.9.3 on page 91) were added to
<sub>22</sub>    support nestable parallel loops that create OpenMP tasks.

- To support interaction with native device implementations, the **use_device_ptr** clause was added to the **target data** construct (see Section 2.10.1 on page 95) and the **is_device_ptr** clause was added to the **target** construct (see Section 2.10.4 on page 103).

- The **nowait** and **depend** clauses were added to the **target** construct (see Section 2.10.4 on page 103) to improve support for asynchronous execution of **target** regions.

- The **private**, **firstprivate** and **defaultmap** clauses were added to the **target** construct (see Section 2.10.4 on page 103).

- The **declare target** directive was extended to allow mapping of global variables to be deferred to specific device executions and to allow an *extended-list* to be specified in C/C++ (see Section 2.10.6 on page 110).

- To support unstructured data mapping for devices, the **target enter data** (see Section 2.10.2 on page 97) and **target exit data** (see Section 2.10.3 on page 100) constructs were added and the **map** clause (see Section 2.15.5.1 on page 216) was updated.

- To support a more complete set of device construct shortcuts, the **target parallel** (see Section 2.11.5 on page 129), target parallel loop (see Section 2.11.6 on page 131), target parallel loop SIMD (see Section 2.11.7 on page 132), and **target simd** (see Section 2.11.8 on page 134), combined constructs were added.

- The **if** clause was extended to take a *directive-name-modifier* that allows it to apply to combined constructs (see Section 2.12 on page 147).

- The **hint** clause was addded to the **critical** construct (see Section 2.13.2 on page 149).

- The **source** and **sink** dependence types were added to the **depend** clause (see Section 2.13.9 on page 169) to support doacross loop nests.

- The implicit data-sharing attribute for scalar variables in **target** regions was changed to **firstprivate** (see Section 2.15.1.1 on page 179).

- Use of some C++ reference types was allowed in some data sharing attribute clauses (see Section 2.15.3 on page 188).

- Semantics for reductions on C/C++ array sections were added and restrictions on the use of arrays and pointers in reductions were removed (see Section 2.15.3.6 on page 201).

- The **ref**, **val**, and **uval** modifiers were added to the **linear** clause (see Section 2.15.3.7 on page 207).

- Support was added to the map clauses to handle structure elements (see Section 2.15.5.1 on page 216).

- Query functions for OpenMP thread affinity were added (see Section 3.2.23 on page 256 to Section 3.2.28 on page 261).

- The lock API was extended with lock routines that support storing a hint with a lock to select a desired lock implementation for a lock's intended usage by the application code (see Section 3.3.2 on page 273).

- Device memory routines were added to allow explicit allocation, deallocation, memory transfers and memory associations (see Section 3.5 on page 282).

- C/C++ Grammar (previously Appendix B) was moved to a separate document.

# D.2 Version 3.1 to 4.0 Differences

- Various changes throughout the specification were made to provide initial support of Fortran 2003 (see Section 1.6 on page 21).

- C/C++ array syntax was extended to support array sections (see Section 2.4 on page 44).

- The **proc_bind** clause (see Section 2.5.2 on page 52), the **OMP_PLACES** environment variable (see Section 4.5 on page 295), and the **omp_get_proc_bind** runtime routine (see Section 3.2.22 on page 254) were added to support thread affinity policies.

- SIMD constructs were added to support SIMD parallelism (see Section 2.8 on page 72).

- Device constructs (see Section 2.10 on page 95), the **OMP_DEFAULT_DEVICE** environment variable (see Section 4.13 on page 302), the **omp_set_default_device**, **omp_get_default_device**, **omp_get_num_devices**, **omp_get_num_teams**, **omp_get_team_num**, and **omp_is_initial_device** routines were added to support execution on devices.

- Implementation defined task scheduling points for untied tasks were removed (see Section 2.9.5 on page 94).

- The **depend** clause (see Section 2.13.9 on page 169) was added to support task dependences.

- The **taskgroup** construct (see Section 2.13.5 on page 153) was added to support more flexible deep task synchronization.

- The **reduction** clause (see Section 2.15.3.6 on page 201) was extended and the **declare reduction** construct (see Section 2.16 on page 220) was added to support user defined reductions.

- The **atomic** construct (see Section 2.13.6 on page 155) was extended to support atomic swap with the **capture** clause, to allow new atomic update and capture forms, and to support sequentially consistent atomic operations with a new **seq_cst** clause.

1. • The **cancel** construct (see Section 2.14.1 on page 172), the **cancellation point**
2. construct (see Section 2.14.2 on page 176), the **omp_get_cancellation** runtime routine
3. (see Section 3.2.9 on page 240) and the **OMP_CANCELLATION** environment variable (see
4. Section 4.11 on page 300) were added to support the concept of cancellation.

5. • The **OMP_DISPLAY_ENV** environment variable (see Section 4.12 on page 301) was added to
6. display the value of ICVs associated with the OpenMP environment variables.

7. • Examples (previously Appendix A) were moved to a separate document.

# D.3 Version 3.0 to 3.1 Differences

9. • The **final** and **mergeable** clauses (see Section 2.9.1 on page 83) were added to the **task**
10. construct to support optimization of task data environments.

11. • The **taskyield** construct (see Section 2.9.4 on page 93) was added to allow user-defined task
12. scheduling points.

13. • The **atomic** construct (see Section 2.13.6 on page 155) was extended to include **read**, **write**,
14. and **capture** forms, and an **update** clause was added to apply the already existing form of the
15. **atomic** construct.

16. • Data environment restrictions were changed to allow **intent(in)** and **const**-qualified types
17. for the **firstprivate** clause (see Section 2.15.3.4 on page 196).

18. • Data environment restrictions were changed to allow Fortran pointers in **firstprivate** (see
19. Section 2.15.3.4 on page 196) and **lastprivate** (see Section 2.15.3.5 on page 199).

20. • New reduction operators **min** and **max** were added for C and C++

21. • The nesting restrictions in Section 2.17 on page 227 were clarified to disallow closely-nested
22. OpenMP regions within an **atomic** region. This allows an **atomic** region to be consistently
23. defined with other OpenMP regions so that they include all code in the atomic construct.

24. • The **omp_in_final** runtime library routine (see Section 3.2.21 on page 253) was added to
25. support specialization of final task regions.

26. • The *nthreads-var* ICV has been modified to be a list of the number of threads to use at each
27. nested parallel region level. The value of this ICV is still set with the **OMP_NUM_THREADS**
28. environment variable (see Section 4.2 on page 293), but the algorithm for determining the
29. number of threads used in a parallel region has been modified to handle a list (see Section 2.5.1
30. on page 50).

- The *bind-var* ICV has been added, which controls whether or not threads are bound to processors (see Section 2.3.1 on page 36). The value of this ICV can be set with the **OMP_PROC_BIND** environment variable (see Section 4.4 on page 294).

- Descriptions of examples (previously Appendix A) were expanded and clarified.

- Replaced incorrect use of **omp_integer_kind** in Fortran interfaces (see Section B.3 on page 335 and Section B.4 on page 342) with **selected_int_kind(8)**.

# D.4 Version 2.5 to 3.0 Differences

The concept of tasks has been added to the OpenMP execution model (see Section 1.2.5 on page 9 and Section 1.3 on page 14).

- The **task** construct (see Section 2.9 on page 83) has been added, which provides a mechanism for creating tasks explicitly.

- The **taskwait** construct (see Section 2.13.4 on page 153) has been added, which causes a task to wait for all its child tasks to complete.

- The OpenMP memory model now covers atomicity of memory accesses (see Section 1.4.1 on page 17). The description of the behavior of **volatile** in terms of **flush** was removed.

- In Version 2.5, there was a single copy of the *nest-var*, *dyn-var*, *nthreads-var* and *run-sched-var* internal control variables (ICVs) for the whole program. In Version 3.0, there is one copy of these ICVs per task (see Section 2.3 on page 36). As a result, the **omp_set_num_threads**, **omp_set_nested** and **omp_set_dynamic** runtime library routines now have specified effects when called from inside a **parallel** region (see Section 3.2.1 on page 231, Section 3.2.7 on page 237 and Section 3.2.10 on page 240).

- The definition of active **parallel** region has been changed: in Version 3.0 a **parallel** region is active if it is executed by a team consisting of more than one thread (see Section 1.2.2 on page 2).

- The rules for determining the number of threads used in a **parallel** region have been modified (see Section 2.5.1 on page 50).

- In Version 3.0, the assignment of iterations to threads in a loop construct with a **static** schedule kind is deterministic (see Section 2.7.1 on page 56).

- In Version 3.0, a loop construct may be associated with more than one perfectly nested loop. The number of associated loops may be controlled by the **collapse** clause (see Section 2.7.1 on page 56).

- Random access iterators, and variables of unsigned integer type, may now be used as loop iterators in loops associated with a loop construct (see Section 2.7.1 on page 56).

- The schedule kind **auto** has been added, which gives the implementation the freedom to choose any possible mapping of iterations in a loop construct to threads in the team (see Section 2.7.1 on page 56).

- Fortran assumed-size arrays now have predetermined data-sharing attributes (see Section 2.15.1.1 on page 179).

- In Fortran, **firstprivate** is now permitted as an argument to the **default** clause (see Section 2.15.3.1 on page 189).

- For list items in the **private** clause, implementations are no longer permitted to use the storage of the original list item to hold the new list item on the master thread. If no attempt is made to reference the original list item inside the **parallel** region, its value is well defined on exit from the **parallel** region (see Section 2.15.3.3 on page 192).

- In Version 3.0, Fortran allocatable arrays may appear in **private**, **firstprivate**, **lastprivate**, **reduction**, **copyin** and **copyprivate** clauses. (see Section 2.15.2 on page 183, Section 2.15.3.3 on page 192, Section 2.15.3.4 on page 196, Section 2.15.3.5 on page 199, Section 2.15.3.6 on page 201, Section 2.15.4.1 on page 211 and Section 2.15.4.2 on page 213).

- In Version 3.0, static class members variables may appear in a **threadprivate** directive (see Section 2.15.2 on page 183).

- Version 3.0 makes clear where, and with which arguments, constructors and destructors of private and threadprivate class type variables are called (see Section 2.15.2 on page 183, Section 2.15.3.3 on page 192, Section 2.15.3.4 on page 196, Section 2.15.4.1 on page 211 and Section 2.15.4.2 on page 213).

- The runtime library routines **omp_set_schedule** and **omp_get_schedule** have been added; these routines respectively set and retrieve the value of the *run-sched-var* ICV (see Section 3.2.12 on page 243 and Section 3.2.13 on page 245).

- The *thread-limit-var* ICV has been added, which controls the maximum number of threads participating in the OpenMP program. The value of this ICV can be set with the **OMP_THREAD_LIMIT** environment variable and retrieved with the **omp_get_thread_limit** runtime library routine (see Section 2.3.1 on page 36, Section 3.2.14 on page 246 and Section 4.10 on page 300).

- The *max-active-levels-var* ICV has been added, which controls the number of nested active **parallel** regions. The value of this ICV can be set with the **OMP_MAX_ACTIVE_LEVELS** environment variable and the **omp_set_max_active_levels** runtime library routine, and it can be retrieved with the omp_get_max_active_levels runtime library routine (see Section 2.3.1 on page 36, Section 3.2.15 on page 246, Section 3.2.16 on page 248 and Section 4.9 on page 300).

- The *stacksize-var* ICV has been added, which controls the stack size for threads that the OpenMP implementation creates. The value of this ICV can be set with the **OMP_STACKSIZE** environment variable (see Section 2.3.1 on page 36 and Section 4.7 on page 298).

- The *wait-policy-var* ICV has been added, which controls the desired behavior of waiting threads. The value of this ICV can be set with the **OMP_WAIT_POLICY** environment variable (see Section 2.3.1 on page 36 and Section 4.8 on page 299).

- The **omp_get_level** runtime library routine has been added, which returns the number of nested **parallel** regions enclosing the task that contains the call (see Section 3.2.17 on page 249).

- The **omp_get_ancestor_thread_num** runtime library routine has been added, which returns, for a given nested level of the current thread, the thread number of the ancestor (see Section 3.2.18 on page 250).

- The **omp_get_team_size** runtime library routine has been added, which returns, for a given nested level of the current thread, the size of the thread team to which the ancestor belongs (see Section 3.2.19 on page 251).

- The **omp_get_active_level** runtime library routine has been added, which returns the number of nested, active **parallel** regions enclosing the task that contains the call (see Section 3.2.20 on page 252).

- In Version 3.0, locks are owned by tasks, not by threads (see Section 3.3 on page 270).

# Index